
LambLisp

Lisp for Real-Time Control
01 Red Fox Alpha



1	Getting Started With Real-time <i>LambLisp</i>	3
1.1	License	3
1.2	Prerequisites	3
1.3	Quick Start: Try LambLisp on Linux	4
1.4	Using the REPL	4
1.5	Deploy to ESP32	5
1.6	What Happens at Startup	9
1.7	What's in the repository?	9
2	<i>Lisp</i> in a nutshell	11
3	Why Lisp?	15
4	Why <i>LambLisp</i>?	17
4.1	Real-time control	17
4.2	Multi-platform	17
4.3	Widely recognized <i>Scheme</i> language specification	17
4.4	Reuse of existing C++/Arduino libraries	18
4.5	Adaptive real-time garbage collection	18
4.6	Open API for new <i>Lisp</i> primitives	18
4.7	Incremental over-the-air updates without reboot	18
4.8	First-class hierarchical dictionaries	19
4.9	Incremental, just-in-time compilation	19
4.10	Lexical scoping	19
4.11	Tail recursion and tail-calls	20
4.12	Procedures are first-class data types	21
4.13	Runtime collaboration with C++ code	21
4.14	Key feature summary	22
5	LambLisp Architecture Internals	23
5.1	Background	23
5.2	The <i>dictionary</i> type	24
5.3	Hierarchical dictionaries	25
5.4	Execution environment is a <i>hierarchical dictionary</i>	25
5.5	Objects are wrappers around <i>hierarchical dictionaries</i>	27
5.6	Interpreter & compiler organization	28
5.7	LambLisp Virtual Machine	32
5.8	Memory Management	35
5.9	Real-time Guarantees	41

5.10	Scalability	44
5.11	Language Internals and C++ API	51
5.12	C++ Integration	58
6	LambLisp Extensions	67
6.1	Language Extensions	67
6.2	Dictionary (Environment) Operations	68
6.3	Structured Data	69
6.4	Arbitrary-Precision Integers (Bignums)	73
6.5	Cryptography	73
6.6	Bytecode Virtual Machine	74
6.7	Native Code Generator (NCG)	74
6.8	Diagnostics and Development Tools	75
6.9	Arduino / Embedded I/O (ESP32 and compatible)	75
6.10	I2C (Wire)	76
6.11	WiFi	76
6.12	Over-the-Air Updates	78
6.13	Runtime and GC Tuning	80
6.14	LLIP — LambLisp Interaction Protocol	82
6.15	LCD1602 Display	93
6.16	Ultrasonic Sonar (HC-SR04)	94
6.17	PCA9685 PWM Controller	94
6.18	WS2812 Addressable LED Strip (NeoPixel)	94
6.19	SPI Bus	95
6.20	OneWire (1-Wire Bus)	95
6.21	Freenove 4WD Car Kit — Peripheral Summary	96
6.22	Camera and On-Board Display (esp32-s3-eye)	97
6.23	Bundled Scheme Libraries	99
6.24	Behaviors: Cooperative Scheduling	103
6.25	ESP32 System	104
6.26	LambLisp Modbus Interface	106
6.27	LambLisp PROFIBUS and PROFINET Interface	118
6.28	Bundled Demos	130
7	Worked Examples	135
7.1	Worked Example: Autonomous Floor Mapping with the Freenove 4WD Car	136
8	LambLisp Compatibility Matrix	151
8.1	Legend	151

8.2	Coverage Summary	151
8.3	Detailed Compatibility	153
9	LambLisp Frequently Asked Questions	167
9.1	Getting Started	167
9.2	Concepts	168
9.3	Language and Conformance	170
9.4	Integration and Operation	171
10	Other Scheme Implementations	173
11	Topic Documentation	175
11.1	Hardware & Platform Builtins (xmop3)	175
12	Class Documentation	227
12.1	Cell Class Reference	227
12.2	Lamb Class Reference	259
12.3	LambPreSettings Struct Reference	298
12.4	LambStdioClass Class Reference	299
12.5	LL_File Class Reference	299
12.6	LL_File_System Class Reference	300
12.7	NcgErrJump Struct Reference	300
13	Acknowledgements	303
14	Glossary	305
15	List of Figures	313
16	Index	315



Figure 1 Eponymous Lamb

LambLisp by Copyright 2026 [Frobenius Norm LLC](#).

LambLisp is a **real-time** implementation of the **Scheme** dialect of the **LISP** programming language. *Scheme* is governed by standards known as *Scheme R5RS* and *Scheme R7RS*. **LambLisp** is designed primarily for *Scheme R5RS*, with some features from *Scheme R7RS* and **Common Lisp** that aid in development of **real-time embedded control systems**.

LambLisp provides a robust and scalable Hardware Abstraction Layer (HAL). The **LambLisp** HAL adds many features focused on embedded controls, and enables convenient interoperability with existing device drivers or other C/C++ libraries.

LambLisp runs across a three-tier platform stack:

- **ESP32 / ESP32-S3** (32-bit Xtensa LX7) — real-time embedded boards running physical hardware. The Espressif ESP32-S3-DevKitC-1 is the primary test board. The Freenove 4WD Car Kit (classic ESP32) is an autonomous vehicle demo exercising I2C, PWM, analog/digital I/O, and real-time task queues.
- **Linux x86-64** — edge server or development workstation tier. Full **LambLisp** with CUDA bindings for desktop GPU development. Acts as a supervisory controller for one or more embedded boards via the LLIP protocol.
- **Linux AArch64 (ARM64) / Nvidia Jetson Orin Nano** — edge AI tier. Full **LambLisp** with CUDA bindings for on-board neural inference and control. Demonstrated on the Nvidia Jetson Orin Nano development kit.

All three tiers communicate peer-to-peer via **LLIP** (**LambLisp Interaction Protocol**), enabling a scalable supervisory Lisp application that spans embedded sensors, local GPU inference, and cloud compute within a single unified language environment.

For more information:

- Read the LambLisp manual online at
https://wawhite3.github.io/LambLispRT-v01-linux_x86_64/html/index.html
- Find the LambLisp manual in PDF at
https://wawhite3.github.io/LambLispRT-v01-linux_x86_64/LambLisp.pdf
- Download the full LambLisp repositories from
<https://github.com/wawhite3/?tab=repositories&q=LambLispRT>
- Visit the developer Frobenius Norm at <https://frobeniusnorm.com>.

Note

The code repository is renewed frequently, with history deleted. Always check back to see if a new repo is available, not just an update to the previous repo.

Chapter 1

Getting Started With Real-time *LambLisp*

1.1 License

LambLisp is copyright 2026 by Frobenius Norm LLC, a Massachusetts corporation.

LambLisp is a commercial product and requires a license when used to generate revenue, to promote other products or services, or for other commercial activity. Licensing provides *LambLisp* customization, hardware application development assistance, post-deployment support, and a warranty.

LambLisp may be used for non-commercial purposes, but comes with no support and no warranty of any kind.

Warning

LambLisp is in **ALPHA** state. Do not deploy it on any application where life or property may be endangered.

1.2 Prerequisites

- Linux development machine (x86_64 or aarch64).
- Familiarity with embedded cross-compilation and basic Linux shell use.
- *platformio* (`pio`) — the build system used for all *LambLisp* targets.

To install *platformio*:

```
curl -fsSL https://raw.githubusercontent.com/platformio/platformio-core-installer/master/get-platformio.py | python3  
ln -s ~/.platformio/penv/bin/pio ~/.local/bin/pio
```

Clone the repository (replace `vxx` with the current version number):

```
git clone https://github.com/wawhite3/LambLispRT-vxx ~/src/LambLispRT-vxx  
cd ~/src/LambLispRT-vxx
```

1.3 Quick Start: Try LambLisp on Linux

The Linux build is the fastest way to explore *LambLisp* — no hardware required.

```
pio run -e linux_x86_64
```

The executable is placed in `.pio/build/linux_x86_64/program`. Run it from the repo root so it can find the `data/` directory, which contains `setup.scm` and the standard library files:

```
cd ~/src/LambLispRT-vxx
.pio/build/linux_x86_64/program
```

You will see startup messages followed by the *LambLisp* REPL prompt:

```
LL=>
```

Try a few expressions:

```
LL=> (+ 1 2)
3
LL=> (define (greet name) (string-append "Hello, " name "!"))
LL=> (greet "world")
"Hello, world!"
LL=> (load "data/PCA9685.scm" 1)
```

Press `Ctrl-C` to exit.

1.4 Using the REPL




Figure 1.1 LambLisp REPL Cycle

The `LL=>` prompt accepts any Scheme expression. Input is line-buffered — the REPL evaluates each line as it arrives.

Loading files:

On Linux, the working directory is the repo root, so prefix paths with `data/`:

```
LL=> (load "data/PCA9685.scm" 1) ; verbose load
LL=> (load "data/PCA9685.scm" 0) ; quiet load
```

On ESP32, the working directory is `data/` on SPIFFS, so use the bare filename:

```
LL=> (load "PCA9685.scm" 1)
```

Error messages:

When an error occurs, *LambLisp* prints a diagnostic and returns to the REPL — the session is not lost. The GC and VM state remain valid; you can continue evaluating expressions immediately.

Exiting:

- Linux: press `Ctrl-C`.
- ESP32: reset the board or power-cycle.

1.5 Deploy to ESP32

1.5.1 Supported Platforms and Boards

LambLisp ships with pre-configured PlatformIO environments for all supported hardware. The table below lists every environment in `platformio.ini` (generated from `w3_pio/boards.json`).

Environment	Architecture	Flash / PSRAM	Notes
esp32-s3-devkitc-1	ESP32-S3	4 MB / 8 MB	Primary test board; onboard JTAG debugger
esp32s3-n8r2	ESP32-S3	8 MB / 2 MB	Higher flash, less PSRAM
Freenove-4WD-Car-Kit-ESP32	ESP32 (classic)	4 MB / 8 MB	Robotics demo platform
linux_x86_64	Linux x86-64	—	Full-featured POSIX build, no CUDA
linux_x86_64_cuda	Linux x86-64	—	POSIX build with CUDA support
linux_x86_64_profile	Linux x86-64	—	gprof profiling build (<code>-O1 -pg</code>)
linux_aarch64	Linux AArch64	—	Cross-compile for Jetson / ARM servers
linux_aarch64_cuda	Linux AArch64	—	Jetson native build with CUDA

ESP32 Hardware

- **ESP32-S3-DevKitC-1** — the primary test board. It has two USB ports: the first (UART) is used for firmware upload and the serial REPL; the second (USB) connects to the onboard JTAG debugger. Connect both to your development computer. The JTAG port works with VS Code, codelite, and other debugger GUIs.
- **ESP32-S3 N8R2** — a compact ESP32-S3 module with 8 MB flash and 2 MB PSRAM. Suited to applications that need more code storage but can work with a smaller heap.
- **Freenove 4WD Car Kit (ESP32)** — an alternative that exercises LambLisp's full control capabilities: digital & analog I/O, I2C motor control, strip LEDs, timers, and task queues. This board uses the classic ESP32 (not S3) and has no onboard JTAG debugger.

Note

It is useful to keep both an ESP32-DevKit and your target hardware available. When a difficult bug appears, the DevKit's JTAG debugger can save hours. Arrange pin definitions and error handling so the application runs safely on the DevKit even when target hardware is absent.

Linux / Jetson Builds

LambLisp runs natively on Linux without any microcontroller hardware. The Linux builds share the same Scheme interpreter, GC, and standard library as the ESP32 builds; only the hardware-specific modules (Arduino GPIO, I2C, WS2812, etc.) are omitted.

To build and run the Linux x86-64 version:

```
pio run -e linux_x86_64
.pio/build/linux_x86_64/program
```

For the CUDA-enabled build (requires an NVIDIA GPU and CUDA toolkit):

```
pio run -e linux_x86_64_cuda
.pio/build/linux_x86_64_cuda/program
```

For profiling with gprof:

```
pio run -e linux_x86_64_profile
.pio/build/linux_x86_64_profile/program
gprof .pio/build/linux_x86_64_profile/program gmon.out | less
```

The AArch64 environments cross-compile on Linux x86-64 using the AArch64 GCC toolchain, or build natively on the Jetson Orin Nano itself:

```
pio run -e linux_aarch64_cuda    # on Jetson: builds and runs natively
```

CUDA Integration

CUDA support is compiled in on the `linux_x86_64_cuda` and `linux_aarch64_cuda` environments. The C++ layer registers thin bindings to the CUDA runtime API, keeping their **native names** so they map one-to-one onto the familiar `cuda*` calls. At startup, `setup.scm` automatically loads `Cuda.scm`, which adds device-introspection and demonstration helpers on top of those bindings.

Runtime bindings (registered by the CUDA build):

Procedure	Description
<code>cudaGetDeviceCount</code>	Number of CUDA-capable devices visible to the process.
<code>cudaGetDeviceProperties</code>	Device properties (name, total memory, compute capability, ...) for a device index.
<code>cudaMalloc</code> / <code>cudaFree</code>	Allocate / free device (GPU) memory.
<code>cudaMallocHost</code>	Allocate page-locked (pinned) host memory.
<code>cudaMemcpy</code>	Copy between host and device; direction constants <code>cudaMemcpyHostToDevice</code> , <code>cudaMemcpyDeviceToHost</code> , <code>cudaMemcpyDeviceToDevice</code> .
<code>cudaMemset</code>	Fill device memory with a byte value.
<code>cudaStreamCreate</code> / <code>cudaStreamSynchronize</code>	Create a stream / block until its queued work completes.
<code>cudaDeviceSynchronize</code>	Block until all device work completes.
<code>cudaKernel-vectorAdd</code>	Built-in element-wise vector-add kernel (example / benchmark).
<code>cudaLaunch-kernel1</code>	Launch a registered device kernel.

Helpers from `Cuda.scm`:

Procedure	Description
<code>have-Cuda</code>	#t when the CUDA bindings are present in the current build.
<code>cuda-printprops</code> / <code>cuda-printprops-brief</code>	Print full / one-line device properties.
<code>cudaStart</code>	Report device count and properties at startup.
<code>cudaTest1</code> / <code>cudaTest2</code>	Self-test / demo kernels (vector-add round-trip).

An AMD **HIP/ROCm** backend (`linux_x86_64_hip`) mirrors the same binding names, so code written against the `cuda*` API runs unchanged on AMD GPUs.

`Cuda.scm` is always loaded by `setup.scm` on CUDA builds. On non-CUDA builds the file is absent and the load is skipped silently. The `linux_x86_64_cuda` environment is convenient for developing and testing CUDA-dependent code on a desktop GPU before deploying to the Jetson.




Figure 1.2 ESP32-S3 Hardware Context

1.5.2 Build and Flash

Install the ESP32 platform for *platformio* (first time only):

```
pio pkg install --platform espressif32
```

Build, upload the filesystem (data/ directory), and upload the firmware:

```
pio run -e esp32-s3-devkitc-1
pio run -e esp32-s3-devkitc-1 --target uploadfs
pio run -e esp32-s3-devkitc-1 --target upload
```

For the Freenove 4WD Car Kit, substitute *Freenove-4WD-Car-Kit-ESP32* for the environment name. All supported environments are listed in *platformio.ini*.

1.5.3 Connect to the REPL

After flashing, open a serial monitor at 115200 baud:

```
pio device monitor -e esp32-s3-devkitc-1
```

Or use any terminal emulator (*screen*, *minicom*, etc.) on the appropriate */dev/ttyUSB** or */dev/ttyACM** device. You will see the startup log followed by the *LL=>* prompt.

1.6 What Happens at Startup

When *LambLisp* starts, `main.cpp` (provided in the repo) does the following:

1. Creates a *Lamb* virtual machine instance.
2. Installs all native C++ operators (hardware I/O, math, etc.).
3. Loads `setup.scm` from the `data/` directory.
4. Calls the Scheme procedure `loop` defined in `setup.scm` on every iteration of the Arduino `loop()` cycle.

The provided `setup.scm` loads the standard *LambLisp* library files and defines a `loop` procedure suited to the demo hardware. For a new application, edit `setup.scm` (or the files it loads) to define the behavior you need. The minimum `setup.scm` for a custom application is:

```
(define (loop)
  ;; your control code here
)
```

The file name `setup.scm` is fixed. The procedure name `loop` follows the Arduino convention but is not enforced by the VM; any number of *LambLisp* procedures can be called from C++ `loop()` at different points.

1.7 What's in the repository?

LambLisp is delivered as a finished executable, but more importantly, as a set of library components that allow scalability by adding more native functions to *LambLisp*. There is a separate repository for each supported architecture.

The root of the repository contains these files and directories:

File or Dir	Description
README.md	A mini-README
LambLisp.bin	Pre-built application binary
LambLisp.pdf	Detailed documentation in PDF form
html	Detailed documentation in HTML form
src	C++ main program (<code>main.cpp</code>) and device driver code
data	Contains:
	- <code>setup.scm</code>
	- <i>LambLisp</i> features - object system, timers, etc
	- <i>LambLisp</i> I/O interfaces
lib	Contains:
	- <code>liblambLisp.a</code> - <i>LambLispRT</i> library archive
	- source code - selected examples for extensibility

File or Dir	Description
platformio.ini	Sample configurations for ESP32 and Linux
w3_pio	Sample low-level board config files for platformio

There are also architecture-specific copies of the generically named LambLisp.bin and liblamblisp.a, which may assist those deploying multiple architectures.

Chapter 2

Lisp in a nutshell

Elements of <i>Lisp</i>
Spaces are separators between expressions.
If an expression starts with (, then it is a list and requires a matching).
Otherwise, the next group of characters, up to a space, is a symbol .
In both cases the expression is called a symbolic expression , or S-expression .
A <i>list</i> contains 0 or more S-expressions between its parenthesis, as in () (foo) (1 (2 three) 4).
There is an environment that associates symbols with values.
There is a small foundational set of procedures that accept arguments and return new S-expressions.
An evaluator continuously reads <i>S-expressions</i> , evaluates them, and returns the result.
If the evaluator encounters a list, the first item in the list must be a <i>procedure</i> and it will be applied to the arguments provided.

Note

To **apply** a procedure to a list of arguments, the *formal parameters* of the *procedure* are paired with the actual arguments provided in the application. Then those pairs are added as bindings to the environment of the *lambda expression*. Finally, the body of the *procedure* is evaluated in that augmented environment.

In the *Lisp* world, any item not a list (or more accurately a *pair*) is called an **atom**. Numbers, strings, and vectors all are *atoms*. Perhaps surprisingly, the foundational *Lisp* does not require these, as they can be implemented using purely symbolic programming. There is an important distinction here between the *mathematical foundations of Lisp*, the *programming language called Lisp*, and *implementations of Lisp*.

Many areas of mathematics have only thrived after a notation was introduced to facilitate reasoning about it. Think of *sine*, *cosine*, the *square root* symbol, integral and differential calculus, matrix algebra, etc.

Likewise, *Lisp* established a small set of symbols and their behavior when encountered in a program. In the days of pen and paper, Greek letters would have been used, as they were in the **lambda calculus** from which *Lisp* inherits its intellectual grounding and the **lambda** operator. In the teletype era, only the typewriter symbols were available.

Lisp was created to perform *symbolic reasoning*. *Lisp's* elements of reasoning about symbolic expressions:

<i>Lisp</i> symbolic expressions	Description of operation
symbols	A sequence of non-spaces, such as <code>x</code> , <code>2day</code> , and <code>next<>thursday</code> . Non-space characters may appear anywhere in a <i>symbol</i> .
<code>#f</code> and <code>#t</code>	<code>#f</code> represents false , while <code>#t</code> represents true . Anything not <i>false</i> is considered <i>true</i> .
<code>()</code> or <code>NIL</code> or <code>nil</code>	A singleton that is an <i>atom</i> and a <i>list</i> , but not a <i>pair</i> . Used as a sentinel to terminate lists.
lists	A set of parenthesis enclosing 0 or more <i>symbols</i> or <i>lists</i> , as in <code>(1-element-list)</code> and <code>(nested (nested () list) lists)</code>
environment	An association between symbols and values. Often implemented as a list of pairs: <code>((x 12) (2day 42))</code> .
<code>if</code>	Execute test and choose from alternate code paths.
<code>eq?</code>	Accepts 2 arguments and returns false if they are not the same S-expression.
<code>define</code>	Update or install new <code>(symbol value)</code> pair in the current environment.
<code>set!</code>	Update existing <code>(symbol value)</code> pair in the current environment.
<code>lambda</code>	Creates a procedure ; i.e., a set of formal parameters, a body of <i>Lisp</i> code, and the environment in which it was defined.
<code>eval</code>	S-expression evaluation ; if a list, the <i>procedure</i> at the front of the list is applied to the remaining arguments.
<code>apply</code>	Execute the body of a <i>lambda</i> , in its original environment enhanced with the arguments to <code>apply</code> paired to the procedure's formal parameters.
<code>quote</code>	Prevent evaluation of something that would normally be evaluated.
<code>macro</code>	A procedure that accepts its arguments unevaluated, i.e., as source code, and returns new source code to be executed in its place.

Some terms should be defined in the *Lisp* context; for example *operators*, *functions*, *procedures*, *macros* are related concepts, but not identical.

Important axiomatic behaviors
The word operator is used for any S-expression found in the first position of a list at evaluation time.
An <i>operator</i> must resolve to a <i>procedure</i> before it can be <i>applied</i> , i.e., the operator position may be occupied by a symbol or other expression.
The word function is ambiguous and may refer to a <i>lambda</i> expression or a <i>procedure</i> .
A lambda expression defines a function's parameters and body, such as <code>(lambda (p1 p2 p3) (do-something p1 p2 p3))</code>
The first argument of a <i>lambda expression</i> is the list of <i>formal parameters</i> , and the remaining arguments (the <i>body</i>) are S-expressions to be evaluated at runtime.
A procedure is a <i>pair</i> consisting of a <i>lambda expression</i> and an <i>environment</i> in which the procedure was created.
A macro is a procedure whose arguments are not evaluated before the code body is applied; afterwards the result of applying the code body is further evaluated.

Note

When a *lambda expression* is executed, it returns a *procedure*. A *procedure* records the formal arguments and the body of the *lambda expression*, and also the *environment* in which the *lambda expression* was evaluated.

This by itself is an incredibly powerful set of mathematics, but note that it has no numbers or strings. That is because, from the point of view of symbolic expressions, numbers and strings are just symbols. It is easy to write a function to add 2 digits when the digits are just 0 and 1. From this point of view numbers are not fundamental or axiomatic. Instead they are a result of agreement on the meaning of symbols and their combination. The same can be said of strings. When the problem is presented in terms of symbols and their relationships, programs become amenable to proofs of correctness by familiar methods such as substitution and induction.

Adding numbers, strings, and vectors to *Lisp*'s capabilities provides computing facilities for real-world problem types. Adding these types means adding representations for them, and adding operators that can manipulate them. *Lisp* follows the usual rules for representing numbers and strings, and so below we present a list of typical symbolic expressions:

```
"This is a string atom, below is a number atom"
1
"Below is a list of number atoms"
(1 2 3)
"Below is a list of mixed types, atoms and a nested list"
(42 "another string" (nested list))
"Below is an arithmetic list expression; it must be evaluated to produce the result 15"
(+ 1 2 (* 3 4))
"Below is a symbol atom; when evaluated the result is the symbol's value in the current environment"
sym
```

Numbers, strings, and symbols are *atoms*. When a number or string is encountered by the evaluator it is not further evaluated, or put another way, it evaluates to itself. When a symbol is encountered by the evaluator, the symbol is located in the current environment, and the evaluator returns the value associated with the symbol.

Note that the expression (1 2 3) is a valid S-expression (a list), but the evaluator will reject it because the first element of the list (1) is not a procedure. The evaluator must be able to *apply* the first element, which therefore must be a *procedure*.

Note

An important feature of *Lisp* is the equivalence of code and data. In some sense, everything is data, until it is submitted to the evaluator function `eval`, at which point it is executed as code.

For example, the job of the *Lisp reader* is to accept *Lisp* input from an implementation-defined *port*, and submit that input to the evaluator. Until it gets to the evaluator, it's just data.

Lisp was developed by discovery, rather than predefinition, and each successful experiment led to another language "standard feature". Eventually this led to overstuffing, as witnessed by Guy Steele's 1000-page tome *Common Lisp*. With a specification that size, many partially compatible subsets emerged, losing some of the *Common* in *Common Lisp*.

In reaction to the size of Common Lisp, an effort was made to extract the most useful and essential behaviors of Lisp into a simpler package, which became *Scheme*. Just as effectively as Guy Steele had recorded *Common Lisp* for posterity, he was also a key figure in factoring out the fundamental behaviors exhibited by a Lisp machine. Those basic behaviors, in combination, could be used to produce all the other more complex behavior. The result is the set of functions found in *Scheme*.

The *Scheme* variant of *Lisp* has many features familiar to programmers in other languages, such as *lexical scoping* and *dynamic typing*, as well as advanced features such as tail recursion, first-class procedures, and continuation-passing style.

Chapter 3

Why Lisp?

LISP is the original language of artificial intelligence, and is the second-oldest programming language still in use (after Fortran). It was designed for *symbolic reasoning* rather than computation. From the theoretical base beginning about 1960, *LISP* quickly grew into a practical language with several well-supported variants. *LISP* and LISP programs were always uppercase; the 8-bit byte had not yet been standardized and text was often packed into 6-bit chunks with no room for lowercase.

Today, approaching the mid-21st century, *Lisp* is active in many variants: SKILL from Cadence Design, Auto↵ Desk Lisp, ChezScheme from Cisco, and others. Like *LambLisp*, these are often tailored to solve problems in a specific domain, and descend from either *Scheme* or the earlier *Common Lisp*. *LambLisp* is optimized for intelligent real-time control of physical processes.

The *LambLisp* bibliography contains a mini-history of automated symbolic reasoning since the Renaissance. It contains work beginning with Isaac Newton and Bertrand Russell, through to Alonzo Church, John Mc↵ Carthy, and Edsger Dijkstra, as well as many less famous and more recent authors. In the 21st century development of automated symbolic reasoning continues, with recent publications on already well-trodden topics such as real-time garbage collection and macro expansion. New approaches to artificial intelligence have created renewed interest in systems (like *Lisp* and its descendents) that offer provable correctness and traceable results.

Any papers cited in the *LambLisp* documentation can be found in the bibliography at <https://github.com/wawhite3/LambLisp-Bibliography>

Additionally, these books provide excellent exposure to many of the issues encountered when developing a *Lisp* programming system:

- Structure and Interpretation of Computer Programs (Abelson and Sussman). Known as SICP, it was a standard Computer Science textbook for many years.
- Lisp in Small Pieces (Queinnec). Another excellent textbook, which drew lessons from SICP and examines more topics and more detail.
- The Art of Computer Programming (Knuth). Written in the pre-C era, TAOCP is an encyclopedia of algorithms that remain foundational to modern computer science.

Lisp invented, aggregated, or evolved all the important techniques used in programming languages today, including interactive computing, just-in-time compilation, on-the-fly code updates, macro expansion, polymorphism, object systems, lexical and dynamic scoping, tail recursion, garbage collection, and exception handling.

Lisp offers tremendous scalability. At the *micro* end, *Lisp* makes an ideal *assembly language*, with a simple prefix notation that can transform directly into machine instructions. At the *macro* end, new functions can be defined and combined dynamically in a way not possible in C++ or Python. These new functions can be the results of substantial analysis performed elsewhere (call it **AI**), putting it at a scale billions of times larger than the assembly language application.

Lisp is **interactive** and *Lisp* programs can be modified on-the-fly, while they are running. Programs can evolve. In other languages, predefined parameters may be tuned during operation, but the code does not change. In *Lisp*, the code may continually evolve, and provides *Lisp* with the power of induction and of continuous improvement.

The concept of interactive computing applies not only to interaction with people. **Machine-to-machine interaction** is an inherent behavior of the interactive nature of *Lisp*. In *LambLisp*-based control systems new algorithms can be developed off-system, and then downloaded without interrupting the controlled process. Today, those off-system resources can be **AI assets** in the cloud, and *LambLisp* can be the target of **AI-generated code**.

An important part of the attraction of *Lisp* is the simple syntax, based on just a few rules that are compact and fast on any processor, even in 1960. This led to a variety of *Lisp* dialects; two of the most important are *Common Lisp* and *Scheme*. *Common Lisp* is in the "large Lisp" category, although there are many partially-compatible subset implementations.

Scheme is in the "medium-size Lisp" category, and was a purposeful attempt to distill the fundamental behaviors of *Common Lisp* into a simpler set.

In the lifetime of *Lisp*, many processors have come and gone. Today ARM-based processors have scaled to be competitive in many control applications. There is a wide variety of boards and modules available, and plenty of inexpensive parts running on the commonly available I2C bus. Today's 32-bit and 64-bit microprocessors and edge SoCs present an ideal platform for *Lisp* real-time control.

Chapter 4

Why *LambLisp*?

There are some features that make *LambLisp* special in the *Lisp* menagerie.

4.1 Real-time control

LambLisp provides bounded loop times, sub-microsecond interrupt response, and non-blocking asynchronous scheduling through eight real-time guarantees spanning GC latency, loop control, ISR paths, and deferred tasks.

See: [Real-time Guarantees](#) · [Garbage Collection](#) · [Memory and GC Tuning](#)

4.2 Multi-platform

LambLisp is a multi-platform library: the same application Scheme code runs unchanged on ESP32, Jetson, and Linux. The platform-specific layer is thin — Arduino-style I/O on ESP32, POSIX I/O on Linux, CUDA bindings on Jetson. All language semantics, the GC, and the REPL behave identically. Develop interactively on Linux; deploy to ESP32 by copying the `.scm` file.

See: [Multi-Platform Support](#)

4.3 Widely recognized *Scheme* language specification

LambLisp implements the **Scheme R5RS** core in full. The one deliberate departure is the omission of `call/cc` (and `dynamic-wind`, which depends on it); all other R5RS features are supported. R7RS-small coverage is extensive — see the Compatibility Matrix. *Scheme* is widely taught and there is plenty of information online. There are also many desktop versions of *Scheme* available for study.

See: [Compatibility Matrix](#)

4.4 Reuse of existing C++/Arduino libraries

LambLisp can incorporate existing C/C++ libraries unchanged; existing Arduino drivers require no modification. Six integration patterns — stateless procedures, stateful GC-managed objects, device facades, syntax forms, proprietary distribution, and data marshaling — cover the full range of embedding needs.

See: [C++ Integration](#)

4.5 Adaptive real-time garbage collection

Lisp behavior depends on a continuous supply of memory, which must be reclaimed when no longer in use, and then later reissued. The reclamation process is called **garbage collection** or just **GC**.

Until the 1990s, *Lisp* could not be applied to real-time control applications due to unpredictable lapses in control during GC. While GC was occurring, the controlled application was left unsupervised. This problem was solved by Taichi Yuasa, based on earlier work by Edsger Dijkstra. Yuasa provided a provably correct solution that divided GC into increments, and provided formulas for the key metrics in a correct implementation.

In *LambLisp* real-time guarantees are provided by Yuasa's GC formula. *LambLisp* continuously calculates the GC start threshold (Yuasa's M metric) and allocates additional memory when required (Yuasa's N metric).

LambLisp realizes additional performance gains with GC idle-time look-ahead. By setting a value for *idle threshold*, idle time can be applied to garbage collection; later during peak loads GC passes can be omitted altogether. This effectively pushes some workload out of the peak periods and into idle periods, reducing peak loop time.

These idle-time collections will be larger than the normal GC quantum. A small additional benefit is therefore obtained because GC loop overhead is spread over a larger collection than the normal GC quantum.

See: [Garbage Collection · Memory and GC Tuning](#)

4.6 Open API for new *Lisp* primitives

Any C++ function with the three-argument `mop3` signature becomes a native Scheme procedure with no FFI overhead — the same mechanism used to implement all *LambLisp* built-ins.

See: [C++ Integration](#)

4.7 Incremental over-the-air updates without reboot

LambLisp supports incremental hot-reload without a reboot — new Scheme code takes effect immediately in the running control loop, with no flash erase and no control gap.

See: [Code Lifecycle: Incremental Hot-Reload](#)

4.8 First-class hierarchical dictionaries

A part of the success of Python has been its pervasive use of hash tables. Hash tables are also used extensively in *LambLisp*, supporting the representation of environments, dictionaries, and objects. *LambLisp* dictionaries are *hierarchical*, providing a parent/child relationship between hash tables. The internal execution environment of *LambLisp* is a dictionary.




Figure 4.1 Dictionary: Hash Table Layout (djb2)

See: [The dictionary type](#) · [Hierarchical dictionaries](#) · [Execution environment](#)

4.9 Incremental, just-in-time compilation

Much of *LambLisp* behavior is implemented in the *runtime system*, a collection of compiled routines that perform the basic operations defined in the *Scheme* specification. The process of compiling begins with "expanding" those symbols whose values are functions. In this case, "expanding" means to replace the symbol with its value (a *procedure*) wherever appropriate. This can be done case-by-case; when a *LambLisp* program is fully transformed this way, execution becomes a series of linked C++ function calls. The downside to this is that debug information is lost.

See: [Scalable Execution Pipeline](#)

4.10 Lexical scoping

Lexical scoping is what most programmers are already familiar with. C/C++ and Python programmers will discuss *local variables* and *global variables*; these are typical in lexical scoping. Some *Lisps* (but not *Scheme* or *LambLisp*) have *dynamic environments* and *special variables*, which do not follow the familiar lexical convention. In *Scheme* and *LambLisp* there is only *lexical scoping*.

See: [Execution environment is a hierarchical dictionary](#)

4.11 Tail recursion and tail-calls

Tail calls allow many problems to be solved recursively without a growing stack, allowing compact, inductive code in place of iteration.





Figure 4.2 Tail Recursion: Stack Growth vs. Constant Stack

Continuation-passing style is a programming idiom on which the next step in a computation is always an extra argument to every procedure. That "next step" is called a *continuation*. When a procedure has determined its result, instead of returning, it tail-calls the continuation. This is an elegant way to solve many problems by induction.


Tail-Call Trampoline — $O(1)$ Stack vs $O(N)$ Recursive Eval

Naive Recursive eval
stack depth = $O(N)$



↓
stack grows...
...N more frames...
C stack: $O(N)$

Trampoline eval
stack depth = $O(1)$



tail call returns thunk



re-enter (no stackgrowth)

C stack: $O(1)$

*tail_state_set() signals tail position
ll_gc_protect2(sexpr, env_exec, ...) keeps roots live*

Figure 4.3 Trampoline: $O(1)$ Tail Calls

See: [Interpreter & compiler organization](#)

4.12 Procedures are first-class data types

Procedures are created as the result of evaluating a **lambda expression**. They get created, stored, and destroyed, just as any other value. They can be passed as arguments to other procedures. In LambLisp the **nlambda** and **macro** expressions are specialized *procedures*.

See: [lambda](#), [nlambda](#), and [macro](#)

4.13 Runtime collaboration with C++ code

C++ objects with persistent state can be placed under LambLisp GC control and passed as first-class Scheme values; they are automatically deleted when they become unreachable.

See: [C++ Integration](#)

4.14 Key feature summary

LambLisp Features
R5RS core fully implemented; <code>call/cc</code> and <code>dynamic-wind</code> omitted by design; extensive R7RS-small coverage.
Additional features from Scheme R7RS and Common Lisp specifications.
Runs on 32-bit ESP32/ESP32-S3, 64-bit x86-64, and 64-bit AArch64 (Jetson).
Full Linux builds — essential edge and supervisory tier, not a port.
High-level implementation of language primitives for compact high performance.
Bytecode compiler and native code generator (JIT) for x86-64, AArch64, and Xtensa LX7.
Arduino-compatible interfaces for digital & analog I/O, I2C, WiFi, etc.
Easy addition of manufacturer-provided C/C++ hardware drivers.
Easy integration of application-specific C++ functions.
Leverage existing Arduino skill set, gcc-based tool chain, C runtime.
Adaptive incremental garbage collector for fast, uniform loop times.
Selectable run-time packages (math, ethernet etc) to control program size.
Virtual memory capability from remote or local storage.
Incremental code updates over-the-air without pausing control or rebooting.

Chapter 5

LambLisp Architecture Internals

Terms introduced in this chapter — `cell`, `sexpr`, `dict`, `env`, `GC phase`, `thunk`, `mop3`, `NCG` — are defined in the Glossary at the end of this document.

5.1 Background

5.1.1 Brief history of Lisp and Scheme

During the 1960s, a great many of the desired and desirable behaviors of programming languages were still being worked out. While Fortran was used heavily for numeric computation, *lexical scoping* (which C programmers experience as "local variables") arose in the Algol language, and later "message passing" systems (primarily *SmallTalk*) resulted in the "tail recursion" technique and interactive computing.

LISP provided a significant abstraction of the computation process. By operating primarily on addresses rather than values, *Lisp* provided a method for creating new and arbitrary data structures. A great deal of research and discussion was expended on functions and their abstract properties (fexprs, property lists, `nlambdas` and the rest). With these tools and techniques, *Lisp* provided a facility for reasoning, not just calculating.

The experimental nature of *Lisp* led to a large number of overlapping solutions to overlapping problems. This led to the development of **Common Lisp**, a 1000-page compendium of *Lisp* functions that had proven useful.

At the same time, there was a realization that "more and bigger" *Lisp* was not necessarily a good thing, producing many incompatible partial implementations. An alternative mindset emerged: a small set of language primitives would be easier to standardize, provided that set was sufficient to construct the richer functions found in *Common Lisp*.

These and other ideas — especially *tail recursion* and *continuation-passing style* — came together in *Scheme*, designed in the 1970s to integrate the theoretical advances and practical understanding accumulated over the previous 15 years.

Scheme went through several iterations and the usual technical-committee process, resulting in two main specifications today: R5RS and R7RS. There has always been tension between *Scheme* minimalists, who prefer a small set of primitives, and utilitarians, who prefer a language useful for real-world problems. This led to the R7RS specification being divided into R7RS-small and R7RS-large, with R7RS-small already being significantly larger than R5RS.

5.2 The *dictionary* type

In computer science, a *dictionary* is a data structure and set of operations for managing (key value) pairs.

LambLisp function	Dictionary operations
dict-set!	install or replace a (key value) pair
dict-ref?	return a (key value) pair based on its key
dict-ref	return a value based on its key

The name *dictionary* applies to the features and functions performed, and not to any particular implementation. A sequential search or a database query are both useful implementations of a dictionary for different purposes, with different performance expectations.

LambLisp uses several advanced techniques to ensure high performance dictionaries in embedded control.

Large or long-lived dictionaries are implemented as *hash tables*. When a dictionary is large or long-lived, great performance gains can be obtained by distributing the (key value) pairs throughout an array. To perform a lookup, compute a *hash value* for the key, and use that as an array index. The (key value) pair, if present, will be in the list of items at that index. The hash function, used to create a hash value from a key, produces a fixed-length number from the contents of the key. The output of the hash function must be the same whenever the same key is presented as input. The hash values produced should be widely distributed among the hash table indexes.

Small dictionaries are implemented as association lists, or **alists**. When a dictionary is small or short-lived, the one-time cost of creating a hash table is unjustified.

In hashed dictionaries, at each element of the hashed array there is a (very short) *alist*, so *alist lookup* is a foundational operation and should be as fast as possible.

LambLisp optimizes symbol lookup by computing hash values once when a symbol is created, and storing the hash value with the symbol. The most common symbol operation (test for equality of symbols) reduces to a fast pointer comparison when using a dictionary.

The dictionary type described above is a *flat dictionary*. That is, there is just one vector, and a (key value) pair is either present or not.

5.3 Hierarchical dictionaries

LambLisp implements *hierarchical dictionaries*. A LambLisp *hierarchical dictionary* consists of a list of *frames*, where each frame is a *flat dictionary*. Hereafter, any reference to *dictionaries* should be understood to refer to LambLisp *hierarchical dictionaries*. In cases where a *flat dictionary* is required, it will be called out specifically.

The fundamental operation in a dictionary is **lookup**. A *lookup* must be done to support every other dictionary operation: create, modify, delete etc. The *lookup* operation consists of providing a key K, and querying the dictionary for a (key value) pair having a key matching the key K provided.

When performing a dictionary lookup, each frame is checked in turn for a matching key. If a key is not present in the top frame of the dictionary, it may be present in succeeding (parent) frames. The *dictionary* is a first-class type in LambLisp, available to the Lisp application.

The dictionary type is a fairly general purpose data structure. Keys and values may be any type. Each frame may be an alist or vector, with the choice made on frame size and expected lifetime. Symbol keys are computed once and stored with the symbol, speeding up dictionary operations where the keys are symbols (as in environments and commonly as in object member identifiers).

The primary method of creating dictionaries is with the function `alist->dict`, which traverses an association list and returns a dictionary. The (key value) pairs of the alist are distributed into a hash table, which becomes a frame in the dictionary. It is also possible to specify a parent dictionary when creating a dictionary, or NIL if there is no parent (as in a *flat dictionary*). Because dictionaries are lists (of frames), they can be appended or examined with *car* and *cdr*, just as with other lists.

In addition to being available in the *Lisp* application, *hierarchical dictionaries* are central to several other aspects of LambLisp.

5.4 Execution environment is a *hierarchical dictionary*

The execution environment is a hierarchical dictionary. In the *Scheme* specification, there are 2 environments mentioned, the *syntactic environment* and the *interaction environment*. The syntactic environment contains all those definitions that are found in the *Scheme* specification. The other is called the *interaction environment*. This is where application definitions are stored.

In LambLisp, the execution environment (which is a dictionary) is initialized with 2 frames. The base frame consists of the *syntactic environment* described in the *Scheme* spec, and the second (top) frame is the *interaction environment* also described there. Top-level definitions are installed in the interaction environment.

When a function is called, its formal parameters are paired with the associated runtime values, and that set of pairs is added as a new frame on top of the current execution environment. This naturally and elegantly implements lexical scoping.




Figure 5.1 Environment Chain and Closure Capture

5.5 Objects are wrappers around *hierarchical dictionaries*

Hierarchical dictionaries also form the basis for LambLisp's Object System (LOBS). Just as the dictionary presents an efficient model for the execution environment, it also provides an excellent base for an object-based approach.

In the object view each frame in a dictionary represents an instance of a class, while the parent frames in the dictionary represent the parent object instance. This leads naturally to inheritance behavior, including multiple inheritance.

In the LambLisp Object System, there is no need to separate class definition from class instantiation. An association list is used to create a new dictionary, and the same alist can be used to produce new instances. In effect, the alist *is* the class definition.

This leads to some interesting freedoms as compared with the C++ or Python approach to class definition and instantiation.

LOBS features
Objects are created from an association list and (optionally) parent objects.
Objects inherit from their parent objects.
Multiple objects may inherit from the same parent object.
Objects may initialize from the same alist, but have different parent objects.
Objects may initialize from different alists, but may share a common parent object.
Objects may have new fields added/deleted any time, as any other dictionary.

To use LOBS as a C++-style object system, do the following:

1. Initialize all objects of the same type from the same association list.
2. Create and initialize new parent objects for each new child object.

To round out the object model, LambLisp provides `dict->obj`. You have correctly perceived that this will convert a dictionary into an "object". This provides a minimal object-style accessor for a dictionary, such that `(some-object 'foo)` will return the value of the field `foo` in the object `some-object`. and `(some-object 'foo 42)` will set the corresponding field on `some-object`.

The simple *set/get* object model suffices for many cases. It is possible to add more complex behaviors by assigning procedures to dictionary values, and/or by enclosing dictionaries in a wrapper procedure.

5.6 Interpreter & compiler organization

Language implementations depend on *primitives*; these are the language elements that are assumed to exist when a programmer is programming in that language. For a high-level language, the *if-then-else* sequence is a common language primitive.

Low-level languages, such as assembly language, also have a set of primitives. Each primitive does less work when used, but by coding directly in assembly language optimization can be implemented that cannot at the higher level, and the result can be many times faster than a higher-level language. At some level, even a compiled high-level language is still interpreting a lower-level machine.

Many available Lisp implementations are based on underlying *abstract machines*. Two of the most important such abstract machines are the **SECD machine** (Landin, 1964) and the **CEK machine** (Felleisen and Friedman, 1987). *TinyScheme*, for example, uses an *SECD machine* as its execution target.

LambLisp implements interpretation of Lisp at a high level; each of the functions in the *Scheme R5RS/R7RS* specifications is implemented in C++. They all share the same C++ signature:

```
Sexpr_t any_function(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

The `lamb` parameter is the `Lamb` virtual machine instance which is to execute the operation. The parameter `sexpr` is the *symbolic expression* to be operated upon by the function, and `env_exec` is the environment in which to execute the function. The return type is also a *symbolic expression*. Note that `env_exec` is known to be a dictionary, and that its keys are all symbols.

This organization has several implications.

First, the language primitives are coded directly in C++, rather than interpreted at run time through a lower-level bytecode system. This makes compilation of the language primitives easy; simply replace the symbol with its value at read time, using a macro facility.

Second, the *Scheme* specification describes only a small set of primitives, with much of the language being "derived expressions". These derived expressions are not language primitives, but instead they either get expanded once at read time, or interpreted at run time. Their performance depends entirely on the expanded versions.

In the highest-performing cases, macros expand into machine code, utilizing the CPU registers directly to carry out *Lisp* operations. That high performance comes with costs. In an embedded system, that type of compiler would be included onboard, taking up space even though used only occasionally. Alternatively, an offboard compiler could be used. That makes some features difficult to implement, such as on-the-fly incremental code updates.

The derived expressions in the *Scheme* specification form a "compiler" that transforms *Scheme* source code into expanded code that relies on relatively few low-level primitives. Some of those lower-level primitives, such as *if*, *define*, and *set!*, are fundamental to the application of *lambda calculus* to computing, while others (such as mathematical functions) are necessary to solve practical problems.

The simple compiler has a disadvantage: the code it generates is simpler to execute than the source, but not simple enough to be fast unless it is designed to produce assembly-level native code. For example, an expansion of `cond` results in many sequential `if` functions to be executed. In an optimizing compiler, a series of low-level test-and-jump sequences will be coded inline, one for each `if`. This will be very fast.

Short of that, executing `cond` requires a pass through the main loop for each test clause. A system that expanded `cond`, but required sequential execution of `if` (in the *Lisp* context), would be relatively slower.

LambLisp has implemented nearly all the *Scheme* language elements as primitives, written in C++. This reduces the number of steps to be interpreted between source code intake, transformation into an *abstract syntax tree* (AST), and its execution in C++. In the design of *LambLisp*, a `cond` in *Lisp* source code corresponds directly to a C++ function `cond(...)`. Encountering `cond` in an expression does not result in a chain of small bits of `if` execution. Instead, the C++ `cond` is called, either after its symbol is resolved in interpreted mode, or called directly through a pointer in compiled mode. All the clauses of the `cond` are executed in a loop at C++ speed, and not in a loop at the *Lisp*-language level or in a bytecode interpreter.




Figure 5.2 Performance Characteristics of Several Implementation Strategies

The chart above shows where each strategy lands on a log scale of relative execution speed. *LambLisp*'s measured results confirm that bytecode and NCG tiers outperform the AST interpreter by 2–100× depending on the benchmark: loop-intensive code (*iota*, *sieve*) benefits most; recursive code (*fib*, *tak*) shows only a modest gain because each AST eval step invokes a full C++ primitive and performs more useful work per dispatch than a bytecode fetch-decode cycle.

The resource chart below shows the cost side of the same picture.




Figure 5.3 Runtime RAM by Implementation Strategy

The AST, Bytecode, and NCG tiers of *LambLisp* cluster tightly at ~2–3 MB — the same GC heap fits all three. The dashed line marks the boundary where an offboard compiler becomes required; above it, the runtime image alone exceeds the total addressable PSRAM of an ESP32-S3. Chez Scheme's runtime stands as far to the right on the resource axis as it does on the performance axis, confirming that the ~250× speed advantage comes with a ~100× RAM cost — a trade-off that does not exist for embedded targets.

Why the compiled tier does not run on embedded hardware

The optimizing compilers — Chez Scheme, SBCL, Chicken, Bigloo, Racket — deliver their performance by making assumptions that are false on an ESP32:

Requirement	Desktop	ESP32 / LambLisp
External toolchain	Required (compiler + linker)	None — source loaded on device
RAM	10 MB – 1 GB	~2 MB PSRAM
Operating system	Full POSIX (mmap, fork, exec)	FreeRTOS (no VM, no fork)
Writable+executable memory	Available	Constrained; IRAM alias required
Garbage collector pauses	Tolerated (ms to tens of ms)	Must not exceed loop period (10 ms)
Hot reload without reboot	Not supported	Core requirement

Compilation step. Chez and SBCL compile Scheme to native code on the development machine, producing a binary that requires linking against a full runtime. The compiler itself requires hundreds of megabytes. Chicken compiles Scheme to C and then to native code via GCC — the device never sees source code. *LambLisp* reads `.scm` files directly from SPIFFS at runtime; no offboard tool is ever involved.

Memory. The Chez runtime image alone exceeds the entire addressable PSRAM of an ESP32-S3. *LambLisp* fits in approximately 300 KB of Flash plus a configurable heap, making it a Scheme runtime designed for the target, not a port of a desktop system.

Real-time GC. Desktop collectors are designed for large heaps where a multi-millisecond stop-the-world pause is acceptable. A 10 ms GC pause violates a 10 ms control loop deadline. LambLisp's Yuasa incremental collector bounds the pause to a small constant per allocation, regardless of heap size.

Hot reload. Compiled systems require edit → compile → link → flash → reboot — tens of seconds of downtime during which the controlled process is uncontrolled. LambLisp can receive new code over LLIP and execute (`load "/new-behavior.scm"`) inside the running loop.

The performance gap between LambLisp and an optimizing compiler (~250× for recursive code) is the price of eliminating all these dependencies. For applications that need that last 250×, the NCG tier is the path forward.

The two execution paths available in *LambLisp* — direct AST interpretation and optional bytecode compilation — are shown in detail in the pipeline diagram below. Both paths share the same GC heap, environment chain, and interaction environment; the bytecode compiler is an optional acceleration step triggered explicitly by the application.




Figure 5.4 Bytecode Compilation Pipeline

The compilation process begins by reading the source code and creating an *Abstract Syntax Tree* (AST). The AST is an internal representation of the program, with each element of the source code being an element of the AST.

To represent the AST, *LambLisp* uses a "direct-connect" technique, in which high-level Lisp operations (implemented in C++) are executed through pointers directly embedded into Lisp objects. These objects correspond directly to high-level features in the language specification (`define`, `lambda`, etc). All of the LambLisp language primitives are implemented this way, and it is possible to add additional LambLisp functions implemented in C++. They are treated as any other first-class function.

5.7 LambLisp Virtual Machine

5.7.1 Control Applications as Virtual Machines.

It will help us to think of the running application as layers of virtual machines. In fact, there is a stack of virtual machines, starting at the bottom with a single transistor, and finishing at the top with whatever runtime inputs the system is designed to respond to.

Those inputs themselves arrive from another, external virtual machine, which has its own internal organization and its own set of inputs, computation, and outputs.

It is also useful to identify the boundary between the *hardware virtual machines* and *software virtual machines*. The illustration below shows a few layers representing the virtual machines involved.




Figure 5.5 A CPU shown as a stack of Virtual Machines

The item labeled "Programs" above is its own virtual machine, consisting of multiple software layers. The layers below that are typically considered "hardware".

At the bottom is the very lowest virtual machine - a field effect transistor. These are operated such that the gate either allows or blocks electron flow through the substrate. When we ascribe meaning to the flow (either a 1 or a 0), the physical flows become meaningful symbols that people can use to communicate.

Transistors are assembled into several types of micro-machines that, when used together, become what we recognize as a "CPU". Symbols of additional complexity are created at every level: numbers, sums and products of numbers, comparison of numbers, interpreting numbers as instructions, etc.




Figure 5.6 Typical Control Application Organization

The illustration above shows a typical control application breakdown. In this view, the code entities form a set with a *partial ordering* relationship, so any of the code entities in the diagram can rely on any of the other entities at an equal or lower level. They can also rely on any entity to the left.

In many loop-based control applications, there are also a few exceptional dependencies, whereby an entity may call another to its right or above. These often occur at system startup. For example, if there is one entity TIME that is responsible for knowing the time, it should not fail under any circumstances, even if not initialized and is unsure of the time. It must provide an answer of some kind.

Another example is when several different types of measurements are required to meet thresholds before the next process step can proceed. These measurements are often averaged over time, and a simultaneous value is required for all them to ensure process synchronization. Each software entity may have some work to do at synchronization time. Each may query the application control layer to obtain information about the larger environment it is operating in, such as whether the current loop is the synchronization loop, and adapt its behavior as needed.

5.7.2 LambLisp Block Diagram

LambLisp Block Diagram — Interpreter and Bytecode Paths




Figure 5.7 LambLisp Block Diagram

When source code is read in, the *reader* front end of the interpreter tokenizes the text, parses the resulting code, and constructs an *abstract syntax tree* (AST). The AST is the executable version of the code, an ordered collection containing all the parts and their relationships. Because the resulting executable is a series of direct links, the AST can be traversed rapidly without a lookup table, and because the tree end nodes are high-level functions (rather than low-level bytecodes), the result is a fast interpreter.

Once the interpreter is available, and correctly executes the AST, the macro facility of *Lisp* is used to transform high-level nodes in the AST into a collection of lower-level nodes that perform the same computation. In a simple *Lisp*, consisting only of *if*, *eq?*, *define*, *lambda*, *set!* and *cons*, every correct program will be factored into those primitives.

When macros are expanded at *read time*, then the effect is one of incremental compilation. After each macro is expanded into a series of pointers to native C++ code, the result is a compiled program consisting of a tree of native code pointers to be traversed. In addition to allowing for faster execution, this tree-structured architecture is what allows *Lisp* programs to be modified on-the-fly.

One clear example of the performance advantage is in the *Scheme cond* function, which acts as a sequential if-then-elseif. Using the minimal set of primitives, *cond* would expand into a series of *if* functions, each of which would need to be processed by the *Lisp* eval loop. In *LambLisp*, *cond* is a primitive implemented in C++, and does not require its own set of passes through the *eval loop* of *Lisp*. This provides improved performance.

5.8 Memory Management

5.8.1 Overview

Memory reuse was recognized as a challenge right from the early days of programmable computers. Most computational results are intermediate; once calculated, they are quickly combined with other results and no longer required. *Garbage collection* is the pejorative term for activities related to the identification and reclamation of unused memory. Every programming language must have a strategy for reusing the memory space that no longer holds useful results.

There are 4 main ways to allocate memory: 1) static allocation, 2) stack allocation, 3) heap allocation, and 4) "manually".

Static allocation is done once before the program starts or immediately after. The memory remains available for use by the application for the life of the program. If this memory is to be used for multiple purposes, the processes involved in reuse must be done "manually" by the running application.

Stack allocation is what C programmers experience as "local variables". When a function is called, a new block is allocated from the top of the execution stack. These blocks are interwoven with the return address of the calling functions. When a function returns, the local storage block is popped off and the calling function's context is now at top of stack. In embedded control systems, the available stack space is usually small, often just 4k or 8k bytes.

The memory not allocated statically or to the stack is referred to as the "heap". C and C++ programmers must manage this space manually, using the `alloc()/free()` or `new/delete` functions, and employ bespoke tracking mechanisms to distinguish used space from unused. In C++ the bespoke aspect can be somewhat mitigated by requiring exclusive use of constructors/destructors for memory allocation, but still it is a problem for which the programmer must provide a solution.

Many other languages allocate large blocks once from the heap, and automatically manage the interior allocation of those blocks according to the requirements of the language. Among these languages are Python, Java, JavaScript, Lua, and of course Lisp.

5.8.2 Cell Memory Model and List Structures

At the lowest level, LambLisp depends on a *cell* structure in 3 consecutive words of computer memory. There are several assumptions about the computer words:

- each word consists of multiple bytes, and the bytes are individually addressable
- each word is capable of holding an address
- each word is capable of holding a signed integer of useful size
- 2 consecutive words of memory may be used to hold a floating-point value.

Of the 3 words in a **Cell**, the first word is called the *tag* and contains information about the cell, such as the cell type and its garbage collection state. If the tag indicates that the cell is a pair type, then the Lisp *car* and *cdr* are represented in the following 2 words of the cell. If the tag indicates an atom (such as a symbol, integer, string etc), then the remaining parts of the **Cell** are used to encode the details of the atom.

Nearly all *Lisp* implementations use this *tagged* arrangement, sometimes with variations such as putting the tag, *car*, and *cdr* in separate parallel arrays, or storing some of the tag bits in unused bits of the *car* and *cdr* fields. These variations are generally not useful on today's microprocessors for 2 reasons: 1) Data structure fields are aligned on word boundaries by the C++ compiler. Defeating this feature results in reduced performance. 2) Separate, parallel arrays will reduce the efficiency of the hardware processor's onboard cache.

Instead of a barrier, these features provide a new opportunity for embedded real-time control with *LambLisp*.

First, *LambLisp* devotes an entire byte for type codes. At present only 5 of the 8 type bits are used. It has proven beneficial to performance to add new types that correspond to the objects handled by the *LambLisp* interpreter. For example, the tail recursion implementation results in many *thunks*, which are code expressions paired with an environment to run them in. As it turns out, there are 2 kinds of thunks: a single S-expression, or a list of S-expressions. Once freed from the imperative to reduce the number of types, adding just a few types provides immense leverage in the *LambLisp* Virtual Machine. Also, using the entire byte means no masking operation is required, speeding up the very common type-testing operation.

Second, *LambLisp* devotes the second byte to **Cell** state information. Garbage collection relies on a state-based assessment of each **Cell**, and cells have 5 GC states, requiring 3 bits. That leaves 5 bits unused, as well as 3 additional unused values for the GC state. These are available for debugging purposes.

Third, in word 0 of each **Cell** only bytes 0 and 1 are required for the tag. In *LambLisp*, the extra bytes in the tag word (word 0 of each **Cell**) are used to store *immediate* types, such as short strings or bytevectors. Using these types for short strings or byte vectors eliminates the need for heap operations. These can be especially useful for small-packet interactions with onboard devices via SPI or I2C.




Figure 5.8 LambLisp VM Cell Structure

At the *Lisp* level, only the *car* and *cdr* are visible; the tag word is not directly available. However, the *Lisp* language provides *predicates* for identifying the type of an S-expression, either a pair (using the *pair?* predicate), or a type of atom (using predicates *integer?*, *real?*, etc). This pattern is encouraged in applications programs, where new data types may be created and predicates provide confirmation of an object's type.

The following illustrates an example of a list structure constructed from the cells described above:




Figure 5.9 Typical List Structure

5.8.3 Garbage Collection

The literature on garbage collection is broad and deep, with several important behavioral features identified among the various possible algorithms, primarily:

Garbage collector distinguishing features
copying vs. not copying
compacting vs. not compacting
support for circular structures
need for auxiliary memory during GC
need for computational pause during GC
average throughput vs. peak pause time
concurrent access to shared memory

In a real-time control system, the discriminating factor for implementing a garbage collector is that it can perform its function incrementally. Time spent on GC is time spent not actively monitoring the controlled process. LambLisp implements a tunable control parameter (the *GC time quantum*), that limits the time spent on each GC increment, ensuring that the cell allocator will not become starved of free cells.

LambLisp's memory reuse strategy relies on these key papers:

Key GC Papers	
Knuth 1963	In "The Art of Programming", Algorithm 2.3.5b describes the fundamental GC stack-based marking strategy
Dijkstra 1976	Describes the "tricolor abstraction" and proves the correctness of incremental garbage collection
Yuasa 1990	Provides a method for determining optimum memory size and free reserve during incremental GC

LambLisp implements the stack-based idle-mark-sweep garbage collection cycle described by Knuth, the tricolor abstraction described by Dijkstra, and actively adapts GC performance based on analysis of the kind described by Yuasa.

Yuasa developed 2 parameters, M (GC start threshold), and N (minimum required cell population). The garbage collector begins in the *idle* state. *Marking* begins when the free cell reserve falls below the GC start threshold, and proceeds incrementally. When marking is complete, the *sweep* phase is begun and also proceeds incrementally. Each increment is less than or equal to the configured GC time quantum. Once sweeping has finished, the GC state is *idle* again until the next M threshold is reached.





Figure 5.10 LambLisp VM Cell Life Cycle

Dijkstra Tricolor Mark — Yuasa Snapshot Write Barrier



Concrete Write Barrier Example




Figure 5.11 Tricolor Marking and Yuasa Write Barrier




Figure 5.12 Incremental GC Timeline

The GC strategy followed by LambLisp is as follows:

- Each cell produced must be matched by an incremental amount of GC, and on the average, that GC must free up at least one cell per cons.
- LambLisp performs an incremental GC pass at *cons* time.
- LambLisp measures the time required for mark and sweep operations.
- LambLisp provides a tunable parameter, the *GC time quantum*, which is the number of microseconds allocated to garbage collection at *cons* time.
- The time quantum is converted into mark and sweep quanta, which are the number of cells that can be marked/swept in the configured time quantum. The numeric quanta have a minimum value of 2.
- The Yuasa parameters are recomputed at the end of every sweep. The parameter M (minimum free cell reserve) is continuously adjusted as required, and the parameter N (minimum required number of cells) is also checked; if N is greater than the available memory then more memory is requested from the system heap.

5.9 Real-time Guarantees

A real-time system guarantees that outputs are updated within a bounded time after inputs change. LambLisp uses loop-based control: the application's `loop` procedure runs repeatedly, reading sensors and driving actuators each cycle. Deadlines are *hard* (missing one is a failure) or *soft* (the system degrades gracefully); LambLisp is soft real-time at the loop level but achieves hard latency in individual ISR operations.

Group	Dimension	Guarantee	Mechanism
GC	Incremental GC	No stop-the-world pause	Yuasa tricolor mark/sweep
GC	GC time quantum	Bounded time per cons	Tunable <code>gc-quantum</code> parameter
GC	Idle-time GC	GC absorbs loop slack	Target loop time + idle pass
GC	Adaptive sizing	GC never starves allocator	Yuasa M/N analysis per cycle
Loop	Target loop time	Bounded total loop duration	<code>loop-target-ms</code> parameter
Loop	Hard vs soft distinction	Design guidance	Factoring hard deadlines out of the loop
Interrupt	ISR latency	200–400 ns GPIO response	IRAM_ATTR + direct register writes
Async	Deferred task queue	Non-blocking callbacks	Time-triggered thunks, no <code>sleep</code>

5.9.1 GC Latency — No Stop-the-World Pauses

The central GC guarantee is that no single allocation ever triggers an arbitrarily long pause. LambLisp's garbage collector is **incremental**: each `cons` operation performs a small, bounded quantum of mark or sweep work before returning. The pause per allocation equals the GC time quantum — a tunable parameter — not the total collection time.

The algorithm is Yuasa's incremental tricolor mark/sweep. Three colors track GC state: *white* (not yet reached), *grey* (reached, children not yet scanned), *black* (fully scanned). The Yuasa write barrier — applied at every mutation during the mark phase — ensures that no live object can be missed and no object is freed while still reachable.

The **GC time quantum** (`gc-quantum`) bounds how many cells are marked or swept at each `cons`. Setting it lower reduces worst-case pause time; setting it higher increases GC throughput. The quantum is expressed in microseconds; LambLisp converts it to a cell count calibrated against measured mark/sweep rates.

Idle-time GC pushes additional collection work into loop slack. When a loop finishes early relative to the target loop time, LambLisp calculates the remaining time budget and runs an extra GC pass using the full slack as its quantum. This moves GC effort from peak periods to idle periods, reducing worst-case loop time during heavy allocation.

Adaptive sizing implements Yuasa's M/N analysis after every sweep. M (the minimum free reserve) and N (the minimum cell population) are recomputed from the current allocation rate. If N exceeds the available pool, additional cell blocks are requested from the system heap up to the configured maximum. The GC never runs dry; the allocator never stalls.

See: [Garbage Collection · Memory and GC Tuning](#)

5.9.2 Loop Time — Bounded Control Cycle

The `loop-target-ms` parameter sets the target duration for a single control loop iteration. LambLisp measures elapsed time at loop-end and uses any remaining budget for idle-time GC (described above) and idle-time printing.

Setting `loop-target-ms` correctly requires measuring the worst-case loop duration under load — peak sensor reads, maximum GC pressure, and any blocking I/O — and adding a margin. LambLisp exposes loop timing statistics accessible from the REPL or over LLIP from a supervisory host.

Hard vs soft factoring: LambLisp is soft real-time at the Scheme level. For operations with a hard deadline — a pulse that must be generated within 1 μ s of a trigger — the correct pattern is to handle the deadline in C++ (an `IRAM_ATTR` ISR) and communicate the result to the Scheme loop via a shared flag. The Scheme loop reads the flag every iteration and acts on it; the hard deadline is met in C++, the policy response is handled in Scheme.

See: [Memory and GC Tuning](#)

5.9.3 Interrupt Latency — Sub-microsecond ISR

On ESP32, code placed in on-chip SRAM with `IRAM_ATTR` executes from fast SRAM rather than slower PSRAM. Combining IRAM placement with direct hardware register writes (bypassing the Arduino `digitalWrite` abstraction) yields GPIO response latencies of **200–400 ns** at 240 MHz.

```
void IRAM_ATTR fast_isr() {
  GPIO.out_wts.val = (1UL < LED_PIN); // ~200 ns from interrupt assertion
}
```

The ISR sets a flag or writes a hardware register. It does not call any LambLisp API — the VM is not reentrant from an ISR context. The Scheme loop polls the shared flag each iteration and responds; response latency is bounded by the loop time, not the ISR latency.

```
volatile bool g_trigger = false;

void IRAM_ATTR edge_isr() { g_trigger = true; }

// C++ mop3 called from the Scheme loop each iteration:
Sexpr_t fn_trigger_check(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec) {
  if (!g_trigger) return HASHF;
  g_trigger = false;
  return HASHT;
}

;;; In the control loop - polled every iteration, bounded by loop-target-ms:
(define (loop)
  (when (trigger-check)
    (handle-trigger))
  (loop))
```

5.9.4 Asynchronous Work — Deferred Task Queue

Blocking the Scheme loop with `(sleep ms)` wastes CPU and violates loop-time guarantees. The **deferred task queue** replaces `sleep` with time-triggered thunks that the loop services without blocking.

LambLisp maintains a time-ordered queue of `(deadline . thunk)` pairs. Each loop iteration dequeues and calls all thunks whose deadline has elapsed. Between calls the loop continues normally; no polling or busy-waiting occurs.


```
;;; Schedule a thunk 500 ms from now - the loop keeps running in between:
(task-after 500 (lambda () (set! led-state (not led-state))))

;;; Blink indefinitely - each invocation reschedules itself:
(define (blink-loop)
  (set! led-state (not led-state))
  (digitalWrite LED_PIN (if led-state HIGH LOW))
  (task-after 250 blink-loop))

(blink-loop)
```

The `TimedQueue` container in `data/Lobs.scm` implements this pattern at the Scheme level. Motor durations, LED animations, and buzzer tones are expressed as timed thunks rather than blocking `delay` calls, keeping the control loop responsive.

See: [Memory and GC Tuning](#)

5.10 Scalability

LambLisp is scalable across eight dimensions, organized into three groups. *Reach* dimensions determine where code runs and how nodes connect. *Speed* dimensions determine how fast it runs. *Footprint* dimensions determine how much memory and flash it consumes. Each dimension can be tuned independently; they compose.

Group	Dimension	From	To	Mechanism
Reach	Platform tier	ESP32: 240 MHz, 2 MB PSRAM	Linux x86-64 / Jetson: 8 GB RAM	Same source, same GC, same REPL
Reach	Network reach	Single node	Multi-node LLIP mesh	LLIP TCP/TLS: eval + filesystem
Reach	AI / machine-to-machine	Fixed compiled algorithm	Runtime AI-generated Scheme	llip-eval + load without reboot
Speed	Execution tier	AST interpreter	NCG native: 2–10× over BC, up to 742× over interpreter on ESP32	procedure->bytecode, ncg-compile
Speed	C++ extension	Pure Scheme	Any C++ library at native speed	mop3 interface
Footprint	Feature footprint	Full R5RS/R7RS binary	Stripped minimal build	Build-time module selection
Footprint	Memory / GC	4k cells	All available PSRAM	GC pool + Yuasa adaptive tuning
Footprint	Code lifecycle	Full-image OTA (reboot)	Incremental hot-reload (no reboot)	(load) in running loop via LLIP

5.10.1 Group A — Reach

Multi-Platform Support

LambLisp is a multi-platform library. Application Scheme code written for one target runs on all three without modification. The same .scm source file, the same garbage collector, and the same LL=> REPL are present whether the host is an ESP32-S3 at 240 MHz with 2 MB of PSRAM, a Jetson Orin Nano at 1.5 GHz with 8 GB of RAM, or an x86-64 Linux workstation. Platform selection is a build-time flag in platformio.ini; no source changes are needed to move between targets.

What is portable — identical on all three targets:

Category	Examples
Core language	Closures, tail calls, define, let, letrec, quasiquote
Macro system	syntax-rules, macro, nlambdas, macroexpand
Standard library	R5RS + R7RS-small: lists, strings, vectors, chars, numbers

Category	Examples
Exception system	with-exception-handler, raise, guard, error
Compilation	procedure->bytecode, ncg-compile (all three NCG backends)
GC	Yuasa incremental tricolor, same parameters
REPL + LLIP	LL=> prompt, (load), remote eval server
Source files	.scm files loaded with (load "file.scm")

What is platform-specific — the thin HAL:

Category	ESP32	Linux / Jetson
Digital I/O	digitalRead, digitalWrite, analogRead	not available
I2C / SPI / WiFi	Wire.begin, SPI.begin, WiFi.begin, etc.	not available
GPU / CUDA	not available	cudaMalloc, inference bindings
Loop driver	(loop) called by Arduino loop()	top-level runs directly

Platform detection — the string `lamb-board` identifies the current target at runtime. Three boolean predicates derived from it are defined in `setup.scm`:

```
linux?    ;; #t on linux_x86_64 and linux_aarch64
esp32?    ;; #t on all ESP32 / ESP32-S3 board variants
jetson?   ;; #t on linux_aarch64 (Jetson Orin Nano)
```

Use them for the small amount of platform-specific initialization that belongs in application code:

```
;;; Portable core - runs identically everywhere:
(define (read-temperature) (i2c-read-byte sensor-addr))

;;; Platform-specific startup - thin and explicit:
(when esp32?
  (Wire.begin SDA-PIN SCL-PIN)
  (Wire.setClock 400000))

(when linux?
  (define (i2c-read-byte addr) (linux-i2c-read "/dev/i2c-1" addr)))
```

Develop on Linux, deploy to ESP32 — the recommended workflow:

1. Write and test application logic interactively at the LL=> REPL on Linux. The full language, GC, and standard library are available.
2. Wrap hardware calls in `(when esp32? ...)` guards or stub them for Linux.
3. Copy the .scm files to the ESP32 flash filesystem via LLIP; evaluate `(load)`. No recompilation, no reflashing of application code.

See: [Deploy to ESP32 · LLIP](#)

Network Reach

A single LambLisp node on an ESP32 can be reached by any other LambLisp node — a Jetson, a Linux workstation, or another ESP32 — using **LLIP** (LambLisp Interaction Protocol), a line-oriented S-expression protocol over TCP.

LLIP exposes two capabilities over a single connection: remote filesystem operations (upload, download, directory listing) and remote evaluation. The `llip-eval` call sends any Scheme expression to the remote node and returns its result as a Scheme value. A supervisory controller can query live sensor state, reconfigure parameters, and deploy new behavior without interrupting the controlled process.

Multiple ESP32 boards can be addressed from a single Jetson or Linux host, forming a mesh where the supervisory node coordinates behavior across the fleet. All nodes speak the same language; no translation layer is needed.

```
;; From a Jetson supervisory node, reach two ESP32 boards:
(define left-wheel (llip-connect "10.0.0.11" 5555))
(define right-wheel (llip-connect "10.0.0.12" 5555))

(llip-eval left-wheel '(Motor.speed))      ; query live RPM
(llip-eval right-wheel '(set! pid-kp 0.85)) ; tune PID gain live
```

See: [LLIP — LambLisp Interaction Protocol](#)

AI and Machine-to-Machine

Because LambLisp code is data — S-expressions that can be constructed, inspected, and transmitted as strings — an AI system or supervisory process can generate new Scheme code and deploy it to a running embedded controller without reboot or recompilation.

The workflow: an off-system AI analyzes telemetry (sensor traces, timing logs, error rates), generates a revised algorithm as a Scheme expression or `.scm` file, uploads it over LLIP, and evaluates `(load "/↵new-policy.scm")` in the running loop. The new definition replaces the old one in the interaction environment immediately; the GC reclaims the old code. No flash erase, no reboot, no control gap.

This is the embedded embodiment of Lisp's historical strength: interactive, live, self-modifying programs. The "off-system resource" is no longer a human typing at a REPL — it is an AI agent operating at machine speed.

```
;; Supervisory AI node - generate and deploy a new control policy:
(define board (llip-connect "10.0.0.11" 5555))

(llip-write board "/policy.scm" (ai-generate-policy telemetry)) ; upload
(llip-eval board '(load "/policy.scm"))                        ; hot-load
(llip-eval board '(policy-version))                            ; → "2026-04-28-r3"
```

5.10.2 Group B — Speed

Execution Tiers: AST → Bytecode → Native Code

LambLisp evaluates Scheme directly from the parsed S-expression (AST) by default, with no mandatory compilation step. This gives fast startup and interactive REPL response without any warm-up phase.

For procedures that run frequently, two optional compilation stages are available:

- **Bytecode** (`procedure->bytecode`, `compile-environment!`): compiles a procedure or an entire environment to a compact bytecode representation. The bytecode executor is a tight dispatch loop that eliminates the overhead of tree-walking the AST on every call. Measured speedup on the trampoline VM: 1.7–3.3×.
- **Native Code Generator** (`ncg-compile`, `ncg-compile-environment!`): compiles bytecode further to native machine instructions for x86-64, AArch64, and Xtensa LX6/LX7 (ESP32/ESP32-S3). The NCG eliminates the dispatch loop and produces code that runs directly on the CPU with no interpreter overhead. On ESP32 LX6 at 240 MHz (2026-05-20), measured speedup over bytecode: `tak` 7×, `ackermann` 6×, `tree` 7×, `fib-rec` 10×, `qsort` 5×; over interpreter: `tak` 324×, `sieve` 341×, `fib-memo` 742×, `qsort` 219×. NCG→NCG recursive calls bypass argument-list allocation entirely, which explains the large gains on recursion-heavy benchmarks. Several inline fast paths accelerate the most common operations: `vector-ref/vector-set!` skip the shim for `T_SVEC_HEAP + T_INT32` index; NCG→NCG call dispatch bypasses the three-level shim chain; `cons` grabs from the free list without calling `gc_pass`; 9 primops (`not`, `eq?`, `zero?`, `positive?`, `negative?`, `set-car!`, `set-cdr!`, `integer?`, `boolean?`) and composed accessors (`caar`, `cadr`, `caddr`, `cddr`, etc.) are expanded inline by the bytecode compiler to single-opcode instructions.

The three levels are complementary. Interactive development stays at the AST level; time-critical routines can be promoted to bytecode or native code with a single function call. NCG is available and recommended on all three targets (ESP32, Jetson, Linux).

C++ Extension

LambLisp has a clean, simple interface to C++ code. All the LambLisp native functions have the same signature. As part of the control application, additional native features can be added and will run at full C++ speed. These C++ additions may be useful in several circumstances:

- Implement a low-level hardware abstraction and make it available to Lisp.
- Provide a Lisp interface to a specialized computation library.
- Improve performance at key algorithmic bottlenecks.

For example, LambLisp implements the *digitalWrite* function this way:

```
Sexpr_t mop3_cio_digitalWrite(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
{
    Sexpr_t sx_pin = lamb.car(sexpr);           //Get the pin # in S-expression form
    Sexpr_t val     = lamb.cadr(sexpr);         //Get the value in S-expression form
    Int_t pin       = sx_pin->mustbe_Int_t();    //throw error if an INT not provided

    digitalWrite(pin, val != HASHF);           //in Scheme, any value "not false" is "true"
    return val;                                //return whatever value the pin was set to
}
```

To use this C++ function from within LambLisp to set pin 38 HIGH and then LOW:

```
(digitalWrite 38 #t)
(digitalWrite 38 #f)
```

See: [C++ Integration](#)

5.10.3 Group C — Footprint

Feature Footprint

At build time, LambLisp allows fine-grained control of the system functions that are included in the application image.

For example, LambLisp supports all the math functions described in the *Scheme R5RS/R7RS* specifications, and by default all are included in the build. Functions that are not needed can be commented out in the source code and will not be linked into the control application.

This same approach can be used for other groups of function, such as ports and strings. These groups of functions are completely defined in the spec to provide a full solution in their respective domains, but in practice each application uses only a subset of the available capability.

As another example, few embedded control applications will need the full set of case-independent string operators, and they are easy to add in Lisp later if needed. These functions are candidates for implementation in Lisp and loading at runtime, rather than including them in every binary application image.

Memory and GC Tuning

LambLisp has an adaptable real-time garbage collection implementation that can be scaled in several ways. At every *cons* operation, the GC has an opportunity to run, and may mark or sweep cells at that time.

GC parameter	Function
GC time quantum	Amount of time to spend on GC for each cons
Cell block size	# cells per block
Max cell blocks	maximum # of blocks to allocate
GC % threshold	Increase memory until threshold reached (or max)

The GC time quantum determines how many mark or sweep operations to do during each increment. A smaller number results in shorter GC interruptions, but lower total throughput.

The cell block size is not critical, but the minimum system setup needs 4k cells, and so 8k is a natural default for this parameter, allowing space for a useful *Lisp* application.

The maximum number of cell blocks can be chosen to tune the amount of time spent in GC. The amount of time spent marking depends on the number of cells used by the *LambLisp* program, and not the total number of cells available. Conversely, when sweeping, the entire cell population must be swept.

With more cells in the population, more cells are swept and reclaimed with each GC mark/sweep cycle, because the fixed cost of marking the used cells is amortized over the entire cell population. More memory results in less time spent in garbage collection, but with diminishing returns as more memory is added with constant program size.

LambLisp implements several adaptive tuning mechanisms in the GC implementation.

First, *LambLisp* provides an incremental garbage collector, based on the work of Taichi Yuasa. This eliminates extended pauses for garbage collection. Yuasa began with a stack-based, mark & sweep, stop-the-world garbage collector, and transformed it into an incremental adaptation. In the incremental version, each time a new memory cell is issued, some quantum of garbage collection must also be done. The requirement is to keep the reclamation of memory ahead of the issuance, so for each memory cell issued, at least 1 must be reclaimed.

Yuasa determined several parameters required to optimize incremental garbage collection.

- The size of the *Lisp* program.
- How many marks or sweeps to perform at each new memory issuance.
- The minimal amount of free memory before GC begins.
- The size of memory required to support the ongoing issuance rate.

The amount of memory and the amount of free cell reserve is adjusted dynamically to preserve realtime behavior, based on a Yuasa analysis after every marking phase.

Another principal concern is the amount of CPU time devoted to garbage collection. In *LambLisp*, the percentage of time used in GC is one factor used to determine whether to expand the cell population. If the percentage time in GC is above the configured threshold, then an additional cell block is added. As long as the threshold is exceeded, new blocks will be added up to the configured maximum.

The Yuasa analysis determines how many cells total are required to maintain realtime production, and how many must be kept in free reserve to issue while GC collects some new free cells. The Yuasa parameters are recalculated during every garbage collection.

To allow for maximum throughput while maintaining low GC pause time, *LambLisp* provides a *target loop time* parameter. When each loop has finished, but before returning to the C++ `main()` that called it, the *LambLisp* virtual machine will check if the target loop time has been reached yet. If not, then there are potentially some idle tasks that could be done, that might perhaps usefully fill the time.

The most important of these is garbage collection. If there is loop time left before the target time limit, and if the garbage collection process is active (marking or sweeping) and not idle, *LambLisp* will calculate a new, temporary GC time quantum based on the time remaining, and run a GC incremental pass with that time quantum. This tends to put much of the GC effort into "idle time", so the GC subsystem moves more quickly into its own *idle* phase, where no GC is done at all until the Yuasa *M* parameter is breached.

Interestingly, the printing of messages on *Serial* (the Arduino serial port, or equivalent on other platforms) also takes time, despite the UARTs and software buffering involved. Two messages printed during the same loop might easily run the loop time up to 10 milliseconds, regardless of what the loop time is without the messages. *LambLisp* provides a means to make the printing occur only if idle time is available. This is useful for informational but noncritical messages. This feature has the added advantage of not computing the message contents if time is not available; it is not just skipping printing, but the entire message evaluation.

The *idle time* feature is available to *LambLisp* applications, as well as in the underlying C++ interface.

See: [Garbage Collection](#)

Code Lifecycle: Incremental Hot-Reload

In C++-based systems, half of the persistent storage must be devoted to over-the-air updates. An entire new program image is downloaded before the existing image is marked obsolete. After that, the storage for the old image is still required for future OTA updates, but otherwise is not used again. A reboot is required to start the new image.

With *LambLisp*, updates are incremental. New Scheme definitions are loaded into the running interaction environment; the new code takes effect immediately and the GC reclaims the old definitions. No flash erase, no reboot, no control gap.

LambLisp can receive new code over LLIP and execute `(load "/new-behavior.scm")` inside the running loop — the controlled process keeps running throughout. Because programs are ASTs and ASTs are data structures, startup code used once can be purged from memory and reloaded on demand, algorithms can be chosen and downloaded at runtime, and new behaviors can be introduced in the field without interrupting the controlled process.

```
;; Deploy a new behavior to a running board - no reboot:  
(llip-write board "/new-behavior.scm" new-code)  
(llip-eval board '(load "/new-behavior.scm"))
```

See: [Execution environment · LLIP — LambLisp Interaction Protocol](#)

5.11 Language Internals and C++ API

5.11.1 lambda, nlambda, and macro




Figure 5.13 Argument Evaluation: lambda vs. nlambda vs. macro

From the beginning, *Lisp* functions have been divided into **special functions** and other functions, which have no special designation. They are *special* primarily in that they do not interact with their parameters in the same way as non-special functions. The main differences between these *special functions* and the non-special variety are these:

- Normal lambda functions have their arguments evaluated before the function is applied.
- Special functions receive some or all of their arguments unevaluated.
- Special functions may also be special in many other ways that are function-specific.

While the notion of "function-specific features of special functions" may seem vague, it is a result of the evolutionary nature of early *Lisp*. It was the first computer programming language that understood itself, and that meta-understanding required and produced a lot of exploration.

The topic of *special functions* became inextricably bound up with the topic of **macros**. In all the many *Lisp* implementations over the years, macros have the largest variety of implementations of any other feature in the language.

A major shortcoming of *Scheme R5RS* is the inclusion of `syntax-rules` as the sole language element for a macro capability. Prior to *R5RS*, `syntax-rules` was optional. It is not required to implement the rest of *Scheme*, when the simpler **nlambda** construct will do. There is no lower-level macro feature in the language

itself that could support `syntax-rules`, even though there are several to choose from in previous *Lisps*. To use `syntax-rules` for creating macros first requires an implementation of another (non-Lisp) pattern-matching language. Therefore, every *Scheme* implementation needs to invent or borrow some other lower-level macro facility to support implementation of `syntax-rules`. This puts the cart before the horse, and goes against the original *Scheme* minimalist values.

As a result, there is no common best practice on the implementation of `syntax-rules`. And it gets worse: `syntax-rules` has now been superseded by `syntax-case`, which has the same drawbacks, while presenting new backward-compatibility challenges.

In *Common Lisp*, a macro system is described that has its own set of quirks, such as not being able to use `apply` with a macro.

lambda

A **lambda expression** takes a formal parameter list and a code body and returns a **procedure** – a first-class value that encapsulates three things: the parameter list, the body, and the **environment at the point of creation**. That captured environment is called a *closure*.

In *LambLisp* a procedure is represented as a pair: one part holds the lambda expression, the other holds the environment dict that was current when `lambda` was evaluated. When the procedure is later applied, a new frame is pushed on top of the captured environment, the arguments are bound into that frame, and the body executes.

Because the environment is captured **by reference**, not by value, any top-level definitions added after the closure was created are visible to it – they exist in the same dict chain.

```
(define (make-adder n)
  (lambda (x) (+ x n)))    ;; captures n from make-adder's frame

(define add5 (make-adder 5))
(add5 3)                  ;;=> 8
```

When applying a lambda procedure, every argument is **evaluated** before the procedure body runs.

nlambda

In *LambLisp*, the foundation for implementing *special forms* is **nlambda**, a concept inherited from **InterLisp**. The *nlambda* operator returns a value of type **nprocedure**. An *nprocedure* has the same internal structure as a lambda procedure – a body and a captured environment – but it differs in one critical way:

Note

The key distinction between `lambda` and `nlambda` is that **nprocedure arguments are not evaluated**. The unevaluated argument expressions are bound directly to the formal parameters. This allows the procedure body to receive and manipulate source code as data.

nlambda supports two formal parameter styles. The standard list form binds each argument expression to its corresponding parameter:

```
;;; list form: each argument expression bound to one parameter
(define show-args (nlambda (a b) (list a b)))
(show-args (+ 1 1) foo)    ;;=> ((+ 1 1) foo)  -- no evaluation
```

The **bare-symbol** form binds the *entire* argument list to a single name. This is the pattern used to implement Scheme syntax forms; in *LambLisp*, primitives like `when`, `cond`, and `let` are implemented as C++ `llOp` functions with the same unevaluated-argument contract:

```
;;; bare-symbol form: entire argument list bound to one name
(define show-all (nlambda args args))
(show-all x (+ 1 2) 'hello)    ;;=> (x (+ 1 2) (quote hello))
```

The *nlambda* feature, combined with `quasiquote`, is sufficient to implement all of *Scheme*'s syntax forms. It is also the foundation on which the *LambLisp* implementation of `syntax-rules` is built.

macro

LambLisp provides the **macro** operator, which returns a **macro transformer**. Like *nlambda*, a macro transformer receives its arguments unevaluated. The critical difference is what happens to the return value:

- An **nlambda** body computes a value and returns it directly to the caller.
- A **macro** body computes a value (*the expansion*) and that expansion is then **evaluated**, with the result returned to the caller.

Applying the *macro transformer* to arguments is called *expanding* the macro; evaluating the expansion is called *executing* it. A typical macro body uses quasiquote to construct the expansion:

```
(define my-when
  (macro (test . body)          ;; test and body received unevaluated
    `(if ,test (begin ,@body))) ;; expansion is an if-expression

(my-when (> x 0)
  (display x)
  (newline))
;; expands to: (if (> x 0) (begin (display x) (newline)))
;; then that if-expression is evaluated
```

Read-time vs. evaluation-time expansion. The *LambLisp* reader calls `macroexpand` as it builds each S-expression. When it encounters a symbol in operator position that is already bound to a macro, the expansion is performed immediately and the *expanded form* replaces the original in the input structure. In this case the macro runs once, at load time, and only the expansion is stored.

If a macro is defined *after* the code that uses it has already been read, the evaluator encounters the macro binding at run time and expands on each call. This produces correct results but at a performance cost. The practical rule is: **define macros before loading the files that use them.**

In *LambLisp* macros are first-class values. They can be stored in variables, passed as arguments, and are garbage-collected like any other pair.

Hygiene

The *nlambda* and *macro* forms described above are **unhygienic**: a name introduced in the macro expansion may accidentally capture a variable from the caller's lexical scope, or be captured by one. For example, a macro that introduces a helper binding named `tmp` will conflict if the caller also has a `tmp` in scope.

LambLisp provides *syntax-rules* (R5RS/R7RS) as a **hygienic** alternative. *syntax-rules* macros are pattern-based and use *gensym* internally to rename introduced bindings, ensuring that macro-introduced names never collide with names at the call site. Use *nlambda*/*macro* when hygiene is not required (syntax forms, DSLs, performance-critical paths); use *syntax-rules* when writing portable, reusable macro libraries.

Note

lambda, **nlambda**, and **macro** all create subtypes of procedures, represented uniformly as pairs of (formals . body) and a captured environment. The table below summarizes the differences:

Form	Arguments	Return value	Hygienic
lambda	evaluated	value	n/ $\leftrightarrow$ a
nlambda	unevaluated	value	no
macro	unevaluated	eval(expansion)	no
syntax-rules	unevaluated	eval(expansion)	yes

5.11.2 Hash table and hashing details

The earlier sections describe the dictionary and hierarchical dictionary types. This section covers implementation details relevant to performance and to using non-symbol keys.

Hash table sizing and collisions

When multiple keys hash to the same index it is called a *collision*. Collisions are common but not usually a problem because the alist chains at each index are short. Hash tables are typically sized at 1–3× the expected number of entries, leaving many empty slots and keeping chains short.

Hash table size is rounded up to the next power of 2. This allows the C++ bitwise-and (&) operator to extract the index from the full hash value, which is significantly faster than the modulo (%) operator.

The oblist

Symbol hashes are computed from the symbol's text when the symbol is first created, and stored with the symbol record in a hash table called *oblist*. The *oblist* is a hash table used as a **set** — it is not a dictionary. Its purpose is to map each symbol's text to a unique address; once interned, two symbols are equal if and only if they share the same *oblist* address. This reduces every symbol-equality test to a single pointer comparison.

Objects and non-symbol keys

Dictionary keys need not be symbols. The hashing rules for non-symbol types are:

- The system hash function is djb2.
- Symbol hashes are precomputed once at creation time.
- Hashes for immutable atoms (integers, reals, etc.) are computed from the atom's value.
- All other hashes use the S-expression address — effectively applying the *eq?* predicate for key matching.

Unlike Python, which requires dictionary keys to be immutable types (tuples rather than lists), LambLisp hashes mutable objects by address, so any type may be used as a key. Dictionaries themselves may be keys or values in other dictionaries.

Because objects are dictionaries and dictionaries are lists of frames, child objects can be created from parent objects on the fly using *append* — no predeclared class hierarchy is required.

5.11.3 Lightweight object system

LambLisp's dynamic nature — first-class closures, mutable hierarchical dicts, and runtime environment manipulation — makes several patterns that are awkward in C++ clean and natural in embedded Lisp code.

`dict->obj` converts any dict into a read/write closure accessor:

```
(define (dict->obj dict-in)
  (let* ((dict-out (dict-add-alist-frame dict-in '(__dict__) (__self__)))
        (dict-rdwr (lambda (key . args)
                      (if (null? args)
                          (dict-ref dict-out key)
                          (dict-rebind! dict-out key (car args))))))
    (dict-rebind! dict-out '__dict__ dict-out)
    (dict-rebind! dict-out '__self__ dict-rdwr)
    dict-rdwr))
```

- `(obj 'field)` — read a field.
- `(obj 'field val)` — write a field.
- `(obj '__dict__)` — obtain the underlying dict for direct dict operations.
- `(obj '__self__)` — obtain the closure itself as a first-class value.
- Prototype inheritance: pass a parent dict as `dict-in`; the new frame sits on top and shadowed keys in the parent remain accessible through the chain.

Performance: LambLisp dicts are hash tables, so `dict-ref` is $O(1)$. The cost is one hash lookup vs one array index for `define-record-type` — negligible in practice.

To add a type tag, the recommended pattern is:

```
(define (typed-obj type dict-in)
  (let ((obj (dict->obj dict-in)))
    (obj '__type__ type)
    obj))

(define (lobs-type? obj type)
  (and (procedure? obj)
       (equal? (obj '__type__) type)))
```

Per-type predicate: `(define (MyRecord? x) (lobs-type? x 'MyRecord))`

The object system and containers described here are implemented in `data/Lobs.scm`.

5.11.4 Containers

`data/Lobs.scm` provides three closure-based containers:

Queue — FIFO backed by a list pair (head/tail pointers).

- `(Queue)` — create a new empty queue.
- `(q item ...)` — enqueue one or more items; returns the updated queue list.
- `(q)` — dequeue the front item; returns `#f` if empty.

Enqueue and dequeue are both $O(1)$ amortized.

Stack — LIFO backed by a list.

- `(Stack)` — create a new empty stack.
- `(s item ...)` — push one or more items; returns the updated stack list.
- `(s)` — pop the top item; returns `#f` if empty.

Push and pop are $O(1)$.

TimedQueue — scheduled task queue for motors, LEDs, buzzer, and other timed actuators. Each entry is `(duration . thunk)`. The main loop calls `(tq)` on every tick; entries whose deadline has elapsed are dequeued and their thunk is called.

- `(TimedQueue)` — create a new timed queue.
- `(tq ms thunk)` — schedule `thunk` to run after `ms` milliseconds.
- `(tq)` — service the queue; runs at most one expired entry per call.

The timed queue is the core scheduling mechanism for the Freenove 4WD Car demo: motor durations, LED sequences, and buzzer tones are all expressed as timed thunks rather than blocking `delay` calls, keeping the control loop responsive.

5.12 C++ Integration

LambLisp is designed so that C++ and Scheme are peers, not separate layers. Any C++ function with the three-argument `mop3` signature becomes a native Scheme procedure with no foreign-function overhead — the same mechanism used to implement all LambLisp built-ins. Existing Arduino libraries, hardware drivers, and C++ classes require no modification. Six dimensions of C++ integration are available; they compose freely.

Group	Dimension	From	To	Mechanism
Binding	Stateless procedure	C++ function	Scheme proc	<code>mk_Mop3_probst_t + dict_bind_bang</code>
Binding	Stateful object	C++ class instance	GC-managed Scheme value	<code>mk_cppobj + deleter</code>
Binding	Syntax form	C++ function	Scheme <code>nlambda</code>	<code>mk_Mop3_nprobst_t</code>
Patterns	Device patterns	Arduino driver unchanged	Callable from Scheme	Singleton / I2C / large-driver recipes
Distribution	Proprietary distribution	C++ source private	Binary <code>.o</code> linked to LambLisp	<code>mop3</code> installer linked at build time
Data	Data marshaling	Scheme value	C++ scalar or buffer	<code>as_Int32_t</code> , <code>coerce_Real_t</code> , <code>any_bvec_get_info</code>

5.12.1 Stateless Procedures

All of the LambLisp language primitives conform to the same function signature:

```
Sexpr_t f(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec);
```

Where:

Parameter	Description
<code>Sexpr_t</code>	Return value is a symbolic expression type, which is a pointer to a LambLisp memory cell.
<code>lamb</code>	An instance of a <code>Lamb</code> virtual machine; in C++ produced by <code>Lamb lamb = new Lamb;</code>
<code>sexpr</code>	The symbolic expression to be evaluated
<code>env_exec</code>	The environment in which the evaluation should take place.

To make the C++ function available in LambLisp, it must be bound into the execution environment. Because the environment is a hierarchical dictionary, the binding can occur in any frame of the dictionary. To bind function `cpp_func` to the LambLisp function `Lamb-func`, use the `Lamb dict_bind_bang` function.

A few things to note about the `dict_bind_bang` function:

- Like most other functions in Scheme-based Lisps, a *bang* or exclamation point ! is used to highlight functions that cause mutations to existing data.

- `dict_bind_bang` takes 4 arguments, 2 of which are environments (i.e., dictionaries). One is the execution environment. All LambLisp native C++ functions require this argument. The other is the target environment (which is a dictionary) that is to be modified.

- The full signature is:

```
Sexpr_t dict_bind_bang(Sexpr_t dict, Sexpr_t key, Sexpr_t val, Sexpr_t env_exec);
```

Sample code for adding a C++ function to LambLisp:

```
//Define your native function somewhere, with the necessary signature.
Sexpr_t my_native_operator(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec);

Sexpr_t sym = lamb.mk_symbol("Lamb-func", env_exec); //Create a symbol in Lisp.
Sexpr_t proc = OBJ_UNDEF;
mop3_gc_protect(sym, {                               //Protect sym across the next allocation.
    proc = lamb.mk_Mop3_procst_t(my_native_operator, env_exec);
});

//Bind the symbol to the function in the top-level interaction environment.
lamb.dict_bind_bang(lamb.r5_interaction_environment(), sym, proc, env_exec);
```

As a concrete example, the ESP32-S3 devkit has an onboard LED accessible via the C++ `neopixel` function. Its native signature is:

```
void neonpixel(int pin, int red, int green, int blue);
```

The *mop3* wrapper and binding look like this:

```
Sexpr_t mop3_neopixelWrite(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
{
    int pin    = lamb.car(sexpr)->as_Int_t();
    int red    = lamb.cadr(sexpr)->as_Int_t();
    int green  = lamb.caddr(sexpr)->as_Int_t();
    int blue   = lamb.caddr(sexpr)->as_Int_t();
    neonpixel(pin, red, green, blue);
    return OBJ_VOID;
}

Lamb *lamb = new Lamb;
lamb->setup();

Sexpr_t sym = lamb->mk_symbol("neopixel", env_exec);
lamb->gc_root_push(sym);                               // protect sym across the next allocation
Sexpr_t proc = lamb->mk_Mop3_procst_t(mop3_neopixel, env_exec);
lamb->gc_root_pop();
lamb->dict_bind_bang(env_target, sym, proc, env_exec);
```

Then from *Lisp*:

```
(neopixel 38 128 0 0) ; pin 38, half red, no green, no blue
```

Note

Garbage collection can happen whenever a new cell is created. Any value not in the current environment or on the GC protect stack is subject to collection. If a native function allocates several cells, each must be pushed onto the GC protect stack before the next is created, and all must be popped before returning.




Figure 5.14 Adding a C++ Primitive to the Scheme Environment (mop3)

5.12.2 Stateful Objects

The neopixel example above wraps a *stateless* C function — each call is independent. When a C++ object carries persistent state (member variables, dynamically allocated buffers, open handles), use `mk_cppobj()` to place it under LambLisp GC control. The GC owns the wrapper cell. When the cell becomes unreachable, the GC automatically calls the registered *deleter*, which frees the C++ object. No manual `free()` or `delete` is needed from Scheme.




Figure 5.15 T_CPP_HEAP cell layout and GC lifecycle

API:

```
// Wrap a heap-allocated C++ object. deleter is called by the GC on sweep.
Sexpr_t mk_cppobj(void *obj, CPPDeleterPtr deleter, Sexpr_t env_exec);

// Retrieve the pointer inside a T_CPP_HEAP cell.
void * any_cppobj_get_ptr();

// typedef for the deleter signature
typedef void CPPDeleter(void *cpp_obj);
```

Example — PID controller:

A PID controller requires state (integral, prev_error) that persists across calls. Wrapping it in a T_CPP_HEAP cell makes each controller a first-class Scheme value: multiple controllers can coexist, be stored in lists, and are collected when no longer referenced.

```
struct PIDState {
  Real_t kp, ki, kd;
  Real_t integral, prev_error;
  Real_t out_min, out_max;
};

static void pid_deleter(void *p) { delete static_cast<PIDState *>(p); }

Sexpr_t PID_mop3_make(Lamb &lamb, Sexpr_t args, Sexpr_t env_exec)
{
  ME("::PID_mop3_make()");
  ll_try {
    auto *pid = new PIDState;
    pid->kp = lamb.car(args)->coerce_Real_t();
    pid->ki = lamb.cadr(args)->coerce_Real_t();
    pid->kd = lamb.caddr(args)->coerce_Real_t();
    pid->out_min = lamb.caddr(args)->coerce_Real_t();
    pid->out_max = lamb.caddr(args)->coerce_Real_t();
    pid->integral = 0;
    pid->prev_error = 0;
    return lamb.mk_cppobj(pid, pid_deleter, env_exec);
  }
  ll_catch();
}

Sexpr_t PID_mop3_update(Lamb &lamb, Sexpr_t args, Sexpr_t env_exec)
{
  ME("::PID_mop3_update()");
  ll_try {
    auto *pid = static_cast<PIDState *>(lamb.car(args)->any_cppobj_get_ptr());
    Real_t sp = lamb.cadr(args)->coerce_Real_t();
    Real_t meas = lamb.caddr(args)->coerce_Real_t();
    Real_t dt = lamb.caddr(args)->coerce_Real_t();
    Real_t err = sp - meas;
    pid->integral += err * dt;
    Real_t deriv = (err - pid->prev_error) / dt;
    pid->prev_error = err;
    Real_t out = pid->kp * err + pid->ki * pid->integral + pid->kd * deriv;
    if (out < pid->out_min) out = pid->out_min;
    if (out > pid->out_max) out = pid->out_max;
    return lamb.mk_real(out, env_exec);
  }
  ll_catch();
}

Sexpr_t PID_mop3_reset(Lamb &lamb, Sexpr_t args, Sexpr_t env_exec)
{
  ME("::PID_mop3_reset()");
  ll_try {
    auto *pid = static_cast<PIDState *>(lamb.car(args)->any_cppobj_get_ptr());
    pid->integral = 0;
    pid->prev_error = 0;
    return OBJ_VOID;
  }
  ll_catch();
}
```

From Scheme:

```
;;; Create a velocity PID: kp=1.5 ki=0.02 kd=0.05 out=[-100,100]
```

```
(define motor-pid (pid-make 1.5 0.02 0.05 -100.0 100.0))

;;; Control loop - called every dt seconds
(define (control-step setpoint encoder-reading dt)
  (pid-update! motor-pid setpoint encoder-reading dt))

;;; motor-pid is a normal Scheme value - store in a dict, pass to procedures, etc.
;;; When motor-pid goes out of scope, the GC calls pid_deleter automatically.
```

5.12.3 Device Patterns

LambLisp provides several examples demonstrating how to interface with a variety of devices. Here we review the different approaches used for different types of devices.

Low level input/output, such as `digitalRead()`, `analogRead()` and similar functions are represented directly in LambLisp, identical to their Arduino counterparts, if one discounts the shift in parenthesis. The Sonar example illustrates the techniques required.

The Sonar example also shows that some things are best done in C++; in this case busy-waiting for a pin to change state and measure its time in state with microsecond resolution. Although the system is blocked waiting for the pin, other approaches fail on the ESP32-Arduino. Using interrupts avoids blocking, but the interrupt latency is large and variable, leading to a useless measurement. Use of threads (pthreads) with a dedicated wait thread leads to the same problem.

Interestingly, the Sonar measurement falls in the same range as the loop time, and that is the source of the requirement for busy-waiting. If the loop time was much longer, the interrupt jitter would not matter. If the loop time was much shorter, we would loop and poll instead of waiting.

The Sonar example also exhibits the *singleton* pattern. There is only 1 Sonar, and it has C++ operations, so it makes sense to bundle it into a C++ class and to create a single instance of that class. The interface then becomes very thin, just a matter of marshalling the parameters in the *mop3* interfaces for each function required in Lisp.

I2C and SPI are common board-level communication protocols used to control embedded devices. SPI is point-to-point, while I2C is a command/response protocol with one command device on the bus at any time, and multiple response devices available to receive commands and report data.




Figure 5.16 Hardware Device Access Patterns

I2C and SPI are both implemented and included in LambLisp. The SPI bindings (`SPI.begin`, `SPI.beginTransaction` / `SPI.endTransaction`, `SPI.transfer`, `SPI.transfer-bytevector!`, `SPI.write-bytevector`, `SPI.end`) mirror the Arduino SPI API, following the same principles as the I2C (`Wire.*`) bindings.

Interfaces to I2C devices fall into 3 sizes: small, medium, and large.

A simple device like the PCF8574 can be operated directly from Lisp over I2C with minimal overhead; this device accepts only I2C write and read commands, which are available directly in Lisp. See the file PCF8574 for the minimal example using I2C low-level operations.

Devices like Wire (aka I2C) and WiFi are in the medium category. These libraries depend on predefined *singletons* of a predefined class. The LambLisp implementation is similar to Sonar, but a bit simpler because the Wire/WiFi class and instances are predefined. On the other hand, both Wire and WiFi have a lot more methods and parameters to deal with.

In the case of Wire and WiFi, C++ does a bit of work, checking and returning high-level errors rather than returning raw error codes to Lisp. Mostly though, the interface layer is concerned with marshalling parameters between C++ and Lisp.

The PCA9685 is a 16 channel PWN controller, originally advertised as an LED controller, but in fact it can be used anywhere a controllable square wave is needed. As it turns out, a variety of motors and linear actuators can be controlled this way.

The PCA9685 is simple in principle, but it allows multiple devices on the bus, and allows each of them to respond to a secondary address, which is programmable and may be shared between devices. This allows operating them in groups with a single command.

That is relevant because it makes the driver code much larger, more than most people would want to rework and take ownership. The PCA9685 LambLisp example shows how to incorporate a large driver without needing to understand it in detail or modify it, by interfacing only with the behaviors required for your application.

5.12.4 Syntax Forms

A `T_MOP3_PROC` receives its arguments already evaluated. When you need a C++ function that receives its arguments *unevaluated* — to implement a special form, a short-circuit operator, or a macro-like construct — register it as a `T_MOP3_NPROC` by setting `syntax=true` in the installer table.

The `sexpr` parameter then contains the raw, unevaluated argument list exactly as written at the call site. The C++ function is responsible for evaluating whatever it needs via `lamb.eval()`.

```
// An nlambda that evaluates only its second argument if condition is true:
Sexpr_t fn_when(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec) {
    Sexpr_t cond = lamb.eval(lamb.car(sexpr), env_exec);
    if (cond != HASHF)
        return lamb.eval(lamb.cadr(sexpr), env_exec);
    return OBJ_VOID;
}

// In the installer table:
{ fn_when, "when", true } // true → T_MOP3_NPROC
```

5.12.5 Proprietary Distribution

Because the mop3 binding mechanism is pure C++, a library vendor can distribute compiled object files (.o or a static .a) without revealing source code. The installer function — `__lamb_NAME_install_mop3` — is the only symbol the linker needs; the customer adds one call to it at startup and the full C++ API is available as native Scheme procedures.

The Scheme source files that *use* the library (.scm) are human-readable and can be shared freely; the C++ implementation remains private. This is the recommended pattern for commercial sensor and actuator drivers distributed as LambLisp extensions.

```
// Vendor delivers: libmysensor.a + mysensor_install.h
// Customer's main.cpp:
#include "mysensor_install.h"

void setup() {
  Lamb *lamb = new Lamb;
  lamb->setup();
  Sexpr_t env = lamb->r5_interaction_environment();
  __lamb_mysensor_install_mop3(*lamb, env, env); // one call exposes all APIs
}

; Customer's control.scm - no C++ visible:
(sensor-init 0x48)
(define (read-temperature) (sensor-read-celsius 0x48))
```

5.12.6 Data Marshaling

Every value crossing the Scheme/C++ boundary is a Sexpr_t cell pointer. The marshaling API converts between cells and C++ native types.

Direction	Scheme type	C++ call	C++ type
Scheme → C++	integer	<code>p->as_Int32_t()</code>	<code>int32↔_t</code>
Scheme → C++	integer	<code>p->as_Int64_t()</code>	<code>int64↔_t</code>
Scheme → C++	real	<code>p->coerce_Real_t()</code>	<code>double</code>
Scheme → C++	bytevector	<code>p->any_bvec_get_info(len, ptr)</code>	<code>size_t, ByteVec_t</code>
Scheme → C++	string	<code>p->str()</code>	<code>const char*</code>
C++ → Scheme	integer	<code>lamb.mk_integer(v, env)</code>	<code>int64↔_t</code>
C++ → Scheme	real	<code>lamb.mk_real(v, env)</code>	<code>double</code>
C++ → Scheme	bytevector	<code>lamb.mk_bytevector(n, env)</code>	<code>size↔_t</code>
C++ → Scheme	string	<code>lamb.mk_string(s, env)</code>	<code>const char*</code>
C++ → Scheme	boolean	<code>HASHT / HASHF</code>	—
C++ → Scheme	void	<code>OBJ_VOID</code>	—

Errors: throw a C++ exception inside a mop3 and LambLisp's `ll_catch()` intercepts it, converts it to a Scheme error object, and propagates it through `with-exception-handler / guard` as a normal R7RS error. No special C++ error-handling code is needed.

```
Sexpr_t fn_safe_divide(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec) {
  ll_try {
    int32_t a = lamb.car(sexpr)->as_Int32_t();
    int32_t b = lamb.cadr(sexpr)->as_Int32_t();
    if (b == 0) throw std::runtime_error("division by zero");
    return lamb.mk_integer(a / b, env_exec);
  }
  ll_catch();
}
```


Chapter 6

LambLisp Extensions

The features in this section are present in LambLisp but are outside the R5RS and R7RS-small specifications. They fall into three broad groups: language and control extensions to the Scheme core, typed numeric storage types for embedded numerical work, and hardware and I/O interfaces specific to the supported platforms (ESP32, Jetson, Linux). The C++ API for adding new native procedures is described separately in the *Language Internals and C++ API* section.

6.1 Language Extensions

Feature	Description
<code>nlambda</code>	Defines a procedure that receives its arguments <i>unevaluated</i> as a list; the InterLisp mechanism underlying all LambLisp syntax forms.
<code>macro</code>	First-class macro value: a transformer procedure applied at eval-time; the expansion is then evaluated. Distinct from <code>syntax-rules</code> hygienic macros.
<code>macro?</code>	Returns <code>#t</code> if its argument is a macro value.
<code>nproc?</code> / <code>nprocedure?</code>	Returns <code>#t</code> if its argument is an <code>nlambda</code> procedure.
<code>llop-proc?</code> / <code>llop-nproc?</code>	Returns <code>#t</code> for a C++ native procedure or C++ <code>nlambda</code> respectively.
<code>macroexpand</code>	Fully expands a macro call form, applying all transformations recursively.
<code>macroexpand1</code>	Expands a macro call form exactly one level.
<code>atom?</code>	<code>(atom? x)</code> — true if <code>x</code> is not a pair; the complement of <code>pair?</code> .
<code>ifnot</code>	<code>(ifnot test then)</code> — shorthand for <code>(if (not test) then)</code> .
<code>while</code>	<code>(while test body ...)</code> — loop that executes <code>body</code> while <code>test</code> is true; returns unspecified.
<code>until</code>	<code>(until test body ...)</code> — loop that executes <code>body</code> until <code>test</code> becomes true; returns unspecified.
<code>reverse!</code>	Destructive in-place list reversal; reuses the existing pairs.

Feature	Description
vector->alist / alist->vector	Convert between association lists and vectors.
current-environment	Returns the current lexical environment as a first-class dictionary object.
add1 / sub1	Increment or decrement a number by one.
pi	The floating-point value of pi.
gensym	Returns a fresh uninterned symbol guaranteed unique across all calls.
sizeof	(sizeof type-keyword) — returns the C++ sizeof of a named type; useful for FFI layout calculations.

Platform detection

Feature	Description
lamb-board	String identifying the current build target; set at compile time from the PlatformIO environment name (e.g., "linux_x86_64", "esp32-s3-devkitc-1").
linux?	#t on all Linux targets (x86-64 and AArch64 / Jetson).
esp32?	#t on all ESP32 and ESP32-S3 board variants.
jetson?	#t on linux_aarch64 (Jetson Orin Nano).

6.2 Dictionary (Environment) Operations

Feature	Description
dict-ref	Looks up a key in a hierarchical dictionary, searching parent frames.
dict-ref?	Like dict-ref but returns #f instead of signalling an error when the key is not found.
dict-bind!	Binds a new key–value pair in the innermost frame of a dictionary, walking the chain if the key already exists.
dict-rebind!	Updates an existing binding in a dictionary without adding a new frame entry.
dict-bind-alist!	Binds all key–value pairs from an association list into a dictionary frame.
dict-rebind-alist!	Updates multiple existing bindings from an association list.
dict-add-empty-frame	Appends a new empty hash-table frame on top of a dictionary chain.
dict-add-keyval-frame	Appends a new frame pre-populated with given keys and values.
dict-add-alist-frame	Appends a new frame pre-populated from an association list.
dict-append!	Appends one dictionary chain to another, linking their frames.
dict-keys	Returns a list of all keys visible in a dictionary (all frames).
dict-values	Returns a list of all values visible in a dictionary (all frames).
dict->alist	Converts a dictionary frame to an association list of (key . value) pairs.
dict->2list	Returns two parallel lists (keys values) for all bindings in a dictionary.

Feature	Description
<code>alist->dict</code>	Converts an association list to a single-frame dictionary.
<code>dict-analyze</code>	Prints diagnostic statistics about a dictionary (load factor, chain depth, collisions).
<code>list-analyze</code>	Prints length and structural statistics about a list.

6.3 Structured Data

LambLisp bundles three lightweight facilities for moving structured data in and out of an embedded target — JSON documents, packed binary frames, and typed numeric arrays — so a device can talk to cloud APIs, wire protocols, and numerical code without external libraries.

6.3.1 JSON

`json-read` parses a JSON string into LambLisp values; `json-write` serialises them back. JSON objects map to association lists with string keys, arrays to lists, `true/false` to `#t/#f`, and `null` to the symbol `null`; numbers become exact integers or inexact reals as written.

```
(json-read "{ \"pin\":4, \"on\":true}")    ;; => ((\"pin\" . 4) (\"on\" . #t))
(json-write (list 1 2.5 #f 'null))      ;; => "[1,2.5,false,null]"
```

6.3.2 CSV

The `csv` feature (`features/csv.scm`, auto-loaded where shipped) reads and writes comma-separated values with RFC 4180 quoting. A quoted field (one wrapped in double quotes) may contain the delimiter, a newline, or a quote written as `\"`. The reader works one character at a time, so a quoted field that spans several physical lines is read correctly. Field values are strings unless you supply a coercion function.

For reading, `csv-read` returns a list of records, where each record is a list of field strings. `csv-read->dicts` treats the first record as a header row and returns a list of dicts keyed by those field names. This is the equivalent of Python's `csv.DictReader`.

Procedure	Result
<code>(csv-read port [opts])</code>	list of records (each a list of field strings)
<code>(csv-read->dicts port [opts])</code>	list of dicts; the header row keys the rest
<code>(csv-string->rows str [opts])</code> / <code>(csv-string->dicts str [opts])</code>	same, from a string
<code>(read-csv path [opts])</code> / <code>(read-csv->dicts path [opts])</code>	same, from a file path
<code>(csv-parse-line str [opts])</code>	parse one record string → list of fields

```
(dict-ref (car (read-csv->dicts "sensors.csv")) "temperature")  ;; => "21.5"
(csv-string->rows "a,\"x,y\",c\n")                             ;; => (("a" "x,y" "c")) ; quoted comma preserved
```

For writing, the same data goes the other way (records or dicts to CSV text). A field is quoted if it contains the delimiter, a quote, or a newline, and any embedded quotes are doubled.

Procedure	Effect
<code>(rows->csv-string rows [opts]) / (dicts->csv-string dicts [opts])</code>	render to a string
<code>(csv-write port rows [opts]) / (csv-write-dicts port dicts [opts])</code>	write to a port

```
(rows->csv-string '(("id" "note") ("1" "has, comma")))
;; => "id,note\r\n1,\"has, comma\"\r\n"
```

The `opts` argument is an alist; all keys are optional:

Key	Default	Meaning
<code>delimiter</code>	<code>#\,</code>	field separator; e.g. <code>(integer->char 59)</code> for <code>;</code> , <code>(integer->char 9)</code> for tab
<code>quote</code>	double-quote (U+0022)	quote character
<code>trim</code>	<code>#f</code>	strip surrounding whitespace from unquoted fields (quoted fields kept verbatim)
<code>coerce</code>	identity	<code>(lambda (col value-string) -> value)</code> applied per cell in the <code>->dicts</code> readers
<code>columns</code>	first dict's keys	explicit column order/header for <code>csv-write-dicts</code>
<code>quote-all</code>	<code>#f</code>	writer quotes every field, not just those that need it
<code>newline</code>	<code>"\r\n"</code>	record terminator the writer emits (<code>"\n"</code> for unix)

```
;; coerce two numeric columns to numbers while reading:
(read-csv->dicts "sensors.csv"
  (list (cons 'coerce
    (lambda (col v)
      (if (member col '("temperature" "humidity")) (string->number v) v))))))
```

A few behavior notes. Blank lines are skipped. A data row with too few columns fills the missing ones with `"`; a row with too many drops the extras. Both cases are logged. For any rows of strings, reading back the output of `csv-write` reproduces the original rows. One case is not yet handled: a leading UTF-8 byte-order mark is not stripped automatically, so remove a BOM upstream if your source emits one.

6.3.3 Binary framing (struct)

`struct-pack` / `struct-unpack` (with `struct-pack-into!` for in-place fills and `struct-calcsz` for sizing) convert between LambLisp numbers and packed binary buffers for Modbus registers, sensor maps, and MQTT payloads. Layout is driven by a Python-struct-style format string with an explicit endianness prefix, so results never depend on host byte order.

```
(struct-unpack ">HhB" (struct-pack ">HhB" 40001 -120 7)) ;; => (40001 -120 7)
```

6.3.4 Typed numeric arrays (array-*)

Typed arrays hold homogeneous, unboxed numeric data. The elements are stored as raw machine numbers, contiguously, in C row-major order. A vector is just the rank-1 case: the same `array-*` operations work on vectors and on higher-rank arrays. The element types are `int32`, `int64`, `float32`, and `float64`. There are also two aliases that follow the host: `int` is the native integer (`int64` on 64-bit hosts, `int32` on ESP32) and

`real` is always `float32`. This set of operations replaces the older `tvec`, `intvector`, `realvector`, and `ndarray` systems.

These are distinct from R7RS `vector` (which holds arbitrary Scheme values) and from R7RS `bytevector` (which holds `u8` bytes).

There are two kinds of constructor. `make-<T>-vector` takes a single length `N` and builds a rank-1 array; passing a shape list to it is an error. `make-<T>-array` takes a shape (a list of dimensions, or the dimensions as separate arguments) and builds an `N`-dimensional array. Both take an optional `fill` argument, which defaults to zero.

```
(define v (make-float32-vector 4 1.5))      ;; length-4 float32 vector, filled 1.5
(array-ref v 2)                             ;; => 1.5
(define m (make-float32-array (list 2 3)))  ;; 2x3 float32 matrix, zero-filled
(array-set! m 1 2 12.0)                     ;;
(array-ref (array-transpose m) 2 1)         ;; => 12.0
```

Constructors

Feature	Description
(<code>make-int32-vector</code> <code>N</code> [<code>fill</code>]) <code>make-float64-vector</code> ...	Rank-1 array of the given element type and length <code>N</code> .
(<code>make-int-vector</code> <code>N</code> [<code>fill</code>])	Rank-1 native-integer array (<code>int64</code> on 64-bit hosts, <code>int32</code> on ESP32).
(<code>make-real-vector</code> <code>N</code> [<code>fill</code>])	Rank-1 <code>float32</code> array.
(<code>make-int32-array</code> <code>shape</code> [<code>fill</code>]) <code>make-float64-array</code> ...	N-D array; <code>shape</code> is a list of dims or spread dims.
(<code>make-int-array</code> <code>shape</code> [<code>fill</code>]) (<code>make-real-array</code> <code>shape</code> [<code>fill</code>]) /	N-D native-integer / <code>float32</code> array.

Operations (work on any rank)

Feature	Description
(<code>array-ref</code> <code>a</code> <code>i</code> <code>j</code> ...)	Reads one element; element type inferred from the array.
(<code>array-set!</code> <code>a</code> <code>i</code> <code>j</code> ... <code>v</code>)	Writes one element.
<code>array-type</code>	Element-type tag (compare against <code>array-type-int32</code> / <code>-int64</code> / <code>-float32</code> / <code>-float64</code>).
<code>array-rank</code>	Number of dimensions.
(<code>array-dim</code> <code>a</code> <code>k</code>)	Extent of dimension <code>k</code> (0-based).
<code>array-size</code>	Total element count.
<code>array-length</code>	Length of a rank-1 array.
<code>array-shape</code>	Shape as a Scheme list of integers.
<code>array-data</code>	The underlying raw data bytevector (row-major).
(<code>array-reshape</code> <code>a</code> <code>shape</code>)	Fresh array with a new shape over a copy of the data (element count must match).
<code>array-transpose</code>	Fresh array with reversed axes (for a 2-D array, the matrix transpose).

Feature	Description
(array-slice a dim start end)	Fresh sub-array over [start, end) along dimension dim.
array-copy	Fresh, independent copy.
(array-fill! a v)	Sets every element to v.
array-sum	Sum of all elements (integer for int arrays, real for float arrays).
array-randomize01!	Fills a float32 array with uniform random values in [0, 1).
array-randomize11!	Fills a float32 array with uniform random values in [-1, 1).
(array-dot a b)	Inner product of two rank-1 arrays of equal length and element type.
array?	#t if its argument is a typed array.
vector->sparsevec	Converts a Scheme vector to a list of (index . value) pairs for non-zero entries.
sparsevec->vector	Reconstructs a dense Scheme vector from a sparse-vector representation.

Array math (linear algebra)

Linear-algebra and element-wise operations over typed arrays of any rank (a rank-1 array is a vector). Each operation dispatches on the array's element type, and reports an error if the shapes or element types do not match. The layout is row-major (C order), meaning the last index varies fastest in memory, so loops run fastest when the last index is the innermost one. The built-in operations already do this; for example, the CPU path of `array-matmul!` uses i-k-j loop order. See [CUDA Integration](#) for moving an array's backing store (`array-data`) to and from GPU device memory.

Feature	Description
(array-matmul! C A B)	Matrix product $C = A \cdot B$ for 2-D arrays, written in place. A is $m \times k$, B is $k \times n$, C must be preallocated $m \times n$; all three share one element type. On CUDA builds the float32 / float64 cases delegate to cuBLAS <code>sgemm</code> / <code>dgemm</code> (using the $(A \cdot B)^T = B^T \cdot A^T$ swap so no host-side transpose is needed); otherwise a row-major CPU loop runs. Integer element types always use the CPU loop.
(array-scale! a s)	Multiplies every element of a by scalar s, in place.
(array-add! a b)	Element-wise $a[i] += b[i]$, in place (same shape and element type).
(array-mul! a b)	Element-wise $a[i] *= b[i]$, in place (same shape and element type).
(array-copy! dst src)	Copies src into dst, in place (same shape and element type). Distinct from the allocating <code>array-copy</code> .
array-contiguous?	Always #t. The (shape . data) representation is always contiguous (C row-major); strided zero-copy views are not part of the model. The procedure is kept for API completeness.

```
(define A (make-float32-array (list 2 3)))
(define B (make-float32-array (list 3 2)))
(define C (make-float32-array (list 2 2)))
;; ... fill A and B ...
(array-matmul! C A B)           ;; C = A·B, cuBLAS-accelerated on CUDA builds
(array-scale! C 0.5)           ;; halve every element in place
```

6.4 Arbitrary-Precision Integers (Bignums)

LambLisp integers are exact and have no fixed size limit. They start as machine words and grow to arbitrary precision as needed, so the same code works for a small loop counter and for a 2048-bit cryptographic key. The [cryptography](#) library is built on these integers.

A bignum is stored as a 32-bit header (bit 31 is the sign, bits 30:0 are the limb count) followed by N 32-bit limbs (LL_int32 words). Arithmetic runs in native 32-bit words, the register width of the ESP32-S3 and the other 32-bit MCUs LambLisp runs on, so a 2048-bit modular exponentiation behaves the same on an ESP32-S3 as on a 64-bit Linux host.

LambLisp does not pack small integers into spare pointer bits (by low-bit tagging or NaN-boxing), as some Schemes do. Each integer is its own typed cell holding its own value, and pointers keep their full width. An integer takes one of three types by size: T_INT32, T_INT64, or T_BIGNUM. Arithmetic promotes a value to the next wider type as it grows and narrows it again when it shrinks. An exact integer is never silently converted to an inexact float.

The ordinary operators (+, -, *, quotient, and so on) already promote through these three types, so most code never calls a bignum-* procedure directly. The procedures below exist for when you want to build or inspect a bignum explicitly.

Bignum primitive	Description
bignum-add / -sub / -mul / -div / -mod / -gcd	Core arithmetic on integers far larger than a machine word.
bignum-shift / bignum-cmp	Bit shift (×/÷ powers of two) and three-way compare.
to-bignum	Coerce any exact integer to a T_BIGNUM cell.
bignum-to-string / bignum-from-string	Decimal/radix text ⇔ bignum (number->string / string->number handle bignums too).
bignum?	#t if its argument is a T_BIGNUM cell.

6.5 Cryptography

LambLisp includes a small cryptography library for embedded key exchange and message authentication. It is built on the [arbitrary-precision integers](#) described above; the same bignum code runs a 2048-bit modular exponentiation on an ESP32-S3 and on a Linux host.

The low-level building blocks are hashing (sha256 over a bytevector), a cryptographic random source (ll-random-bytes and ll-random-bignum, which read the platform CSPRNG: esp_fill_random on ESP32, /dev/urandom on Linux), and modular exponentiation (modular-exponent). Diffie-Hellman and RSA are both built from modular-exponent.

The bignum, hashing, RNG, and modular-exponent operations are C++, compiled in when LL_BIGNUM is enabled. The higher-level protocols are written in Scheme. rsa.scm provides RSA (encrypt, decrypt, sign, verify), HMAC-SHA256, HKDF-SHA256, bytevector XOR, and bignum/bytevector conversion. dh_groups.↔ scm supplies the RFC 3526 2048-bit MODP group (dh-group-14-prime, dh-group-14-generator). Both files load at startup (from setup.scm) on targets whose filesystem image includes them: linux_x86_64, linux_↔ aarch64, esp32-s3-devkitc-1, esp32s3-n8r2, and the Freenove 4WD car kit.

Area	Provided by	Highlights
Hashing & randomness	C++ primitives	sha256 (bytevector → 32-byte digest), 11-random-bytes, 11-random-bignum
Modular arithmetic & Diffie-Hellman	C++ + dh_groups.scm	modular-exponent, dh-group-14-prime, dh-group-14-generator
RSA, HMAC & HKDF	rsa.scm	rsa-encrypt/decrypt/sign/↔ verify, hmac-sha256, hkdf-sha256, bytevector-xor

A Diffie-Hellman exchange over the built-in 2048-bit group. Both sides compute the same shared secret without transmitting a private key:

```
(define p dh-group-14-prime)
(define g dh-group-14-generator)
(define a (diffie-hellman-keygen p g)) ; Alice: (public . private)
(define b (diffie-hellman-keygen p g)) ; Bob:  (public . private)
;; each side combines its own private key with the peer's public key
(define s-alice (diffie-hellman-shared (car b) (cdr a) p))
(define s-bob   (diffie-hellman-shared (car a) (cdr b) p))
(equal? s-alice s-bob) ; => #t (shared secret agrees)
```

See the source comments in 11_vm_mop3_bignum.cpp, rsa.scm, and dh_groups.scm for the exact argument order and semantics of each procedure.

6.6 Bytecode Virtual Machine

Feature	Description
procedure->bytecode	Compiles a lambda or nlambda to a bytecode object (T_BYTECODE); returns the original procedure if the form is not compilable.
bytecode?	Returns #t if its argument is a bytecode-compiled procedure.
bytecode-disasm	Returns a list describing the bytecode instruction sequence of a compiled procedure; intended for debugging.
compile-environment!	Walks an environment dictionary and compiles every lambda/nlambda binding to bytecode in place.

6.7 Native Code Generator (NCG)

Feature	Description
ncg-compile	JIT-compiles a bytecode procedure to native machine code for the current platform (x86-64, AArch64, or Xtensa LX7).
ncg-compiled?	Returns #t if a bytecode procedure has been JIT-compiled to native code.
ncg-disasm	Logs the native code address and size of an NCG-compiled procedure.

Feature	Description
<code>ncg-compile-environment!</code>	Walks an environment dictionary and NCG-compiles every bytecode binding in place; returns the count compiled.

6.8 Diagnostics and Development Tools

Feature	Description
<code>lamb-debug</code>	Sets or queries the LambLisp diagnostic logging level.
<code>lamb-debug?</code>	Returns #t if diagnostic logging is currently enabled.
<code>lamb-breakpoint</code>	Triggers a debugger breakpoint or diagnostic halt; behaviour is platform-dependent.
<code>lamb-expand</code>	Expands a form using the current macro environment and prints the result; development aid.
<code>syslog</code>	Writes a message to the platform system log (serial console on ESP32, <code>syslog</code> on Linux).
<code>lamb-oblist</code>	Returns the interned symbol table (oblist) as a dictionary; useful for inspecting all known symbols.

6.9 Arduino / Embedded I/O (ESP32 and compatible)

Feature	Description
<code>digitalRead / digitalWrite</code>	Read or write a digital GPIO pin.
<code>pinMode</code>	Configure a GPIO pin as input, output, or input with pull-up.
<code>analogRead / analogWrite</code>	Read a 12-bit ADC value or write a PWM duty cycle.
<code>analogReadResolution / analogWriteResolution</code>	Configure ADC/PWM bit resolution.
<code>delay_ms / delay_us</code>	Busy-wait for a number of milliseconds or microseconds.
<code>millis / micros</code>	Return elapsed time since boot in milliseconds or microseconds.
<code>pulseIn / pulseInLong</code>	Measure the duration of a pulse on a GPIO pin.
<code>shiftIn / shiftOut</code>	Bit-bang serial transfer on arbitrary GPIO pins.
<code>tone / noTone</code>	Generate or stop a square-wave audio tone on a GPIO pin.
<code>attachInterrupt / detachInterrupt</code>	Register or remove an ISR callback for a GPIO pin edge.
<code>digitalPinToInterrupt</code>	Convert a GPIO pin number to the corresponding interrupt number.
<code>interrupts / noInterrupts</code>	Enable or disable all hardware interrupts.

Feature	Description
<code>bit / bitRead / bitSet / bitClear / bitWrite</code>	Bit-manipulation helpers operating on integer words.
<code>highByte / lowByte</code>	Extract the high or low byte of a 16-bit integer.
<code>HIGH / LOW / INPUT / INPUT_PULLUP / OUTPUT</code>	Arduino GPIO level and pin-mode constants.

6.10 I2C (Wire)

Feature	Description
<code>Wire.begin</code>	Initialize the I2C bus, optionally as a peripheral at a given address.
<code>Wire.beginTransmission</code> / <code>Wire.endTransmission</code>	Start or end an I2C write transaction to a given device address.
<code>Wire.requestFrom</code>	Request bytes from an I2C peripheral.
<code>Wire.write</code>	Queue bytes for transmission in the current I2C transaction.
<code>Wire.read / Wire.available</code>	Read a byte or check how many bytes are available from a peripheral.
<code>Wire.setClock / Wire.setPins / Wire.setTimeout</code>	Configure bus speed, GPIO pin assignment, or transaction timeout.

6.11 WiFi

Feature	Description
<code>WiFi.begin</code>	Connect to a WiFi network with an SSID and password.
<code>WiFi.disconnect</code>	Disconnect from the current WiFi network.
<code>WiFi.status</code>	Return the current WiFi connection status as a two-element list (<code>code string</code>).
<code>WiFi.localIP / WiFi.gatewayIP / WiFi.subnetMask</code>	Query the assigned IP addresses.
<code>WiFi.macAddress</code>	Return the device's MAC address as a string.
<code>WiFi.SSID / WiFi.RSSI / WiFi.BSSID</code>	Query the connected network name, signal strength, or access-point MAC.
<code>WiFi.scanNetworks</code>	Perform an active scan and return the number of visible networks.
<code>WiFi.config / WiFi.setDNS</code>	Set a static IP configuration or DNS server address.
<code>open-tcp-client-port / open-tls-client-port</code>	Open a TCP or TLS socket to a remote host; returns a Scheme port.

6.11.1 WiFi Station — Advanced

Beyond the basic connect/query verbs, the station API exposes reconnection control, blocking connect waits, and detailed network-info accessors. IP-valued accessors return dotted-decimal strings; MAC/BSSID accessors return colon-separated hex strings.

Procedure	Description
<code>WiFi.reconnect</code>	Reconnect to the last access point; returns <code>#t</code> on success.
<code>WiFi.waitForConnectResult</code>	Block until connected or timeout, then return the <code>wl_status_t</code> code. Optional timeout in ms (default 60000).
<code>WiFi.isConnected</code>	Return <code>#t</code> if currently associated with an access point.
<code>WiFi.setAutoReconnect</code>	Enable or disable automatic reconnection. Optional boolean arg; defaults to enable (<code>#t</code>).
<code>WiFi.broadcastIP</code>	Return the station's broadcast address.
<code>WiFi.networkID</code>	Return the network base (subnet) address.
<code>WiFi.dnsIP</code>	Return a configured DNS server address. Optional index selects primary (0, default) or secondary (1).
<code>WiFi.subnetCIDR</code>	Return the subnet prefix length as an integer (CIDR notation, e.g. 24).
<code>WiFi.psk</code>	Return the WPA/WPA2 pre-shared key as a string.
<code>WiFi.encryptionType</code>	Return the encryption type of the connected network as an integer. Optional interface index (default 0).
<code>WiFi.mode</code>	Set the radio mode: 0=off, 1=STA, 2=AP, 3=AP+STA. Returns <code>#t</code> on success.
<code>WiFi.getHostname</code> / <code>WiFi.setHostname</code>	Query or set the station mDNS hostname. <code>setHostname</code> returns <code>#t</code> on success.
<code>WiFi.setSleep</code>	Set power-save mode. Boolean (<code>#t</code> / <code>#f</code>) toggles modem sleep, or an integer selects 0=none, 1=modem, 2=light. Returns <code>#t</code> on success.
<code>WiFi.channel</code>	Return the current WiFi channel number.
<code>WiFi.hostByName</code>	Resolve a hostname to an IP-address string; returns <code>#f</code> on failure.

```
;; Connect, then block (up to 15 s) for the result instead of polling status.
(WiFi.begin "my-ssid" "my-password")
(WiFi.setAutoReconnect #t) ; auto-reconnect on drop
(let ((code (WiFi.waitForConnectResult 15000)))
  (if (WiFi.isConnected)
      (begin
        (display (list 'ip      (WiFi.localIP)
                      'cidr    (WiFi.subnetCIDR)
                      'dns      (WiFi.dnsIP)
                      'chan     (WiFi.channel)))
        (newline)
        (WiFi.hostByName "example.com")) ; => "93.184.216.34" or #f
      (display (list 'connect-failed code))))
```

6.11.2 WiFi Access Point (softAP)

The ESP32 can run its own access point instead of (or alongside) joining one, letting other devices connect directly to the board. This is how a companion computer — for example a Jetson Orin Nano — can reach the board over a private, self-contained network with no external router (the "Jetson control over SoftAP" arrangement).

Procedure	Description
<code>WiFi.softAP</code>	Start the soft access point. Args: <code>ssid</code> , optional password, <code>channel</code> (default 1), <code>hidden</code> (0/1, default 0), <code>max-connections</code> (default 4). Returns <code>#t</code> on success.
<code>WiFi.softAPConfig</code>	Set the soft-AP IP configuration. Args: <code>ip</code> , <code>gateway</code> , <code>subnet</code> (dotted-decimal strings). Returns <code>#t</code> on success.
<code>WiFi.softAPdisconnect</code>	Shut down the soft-AP. Optional boolean also powers the radio off. Returns <code>#t</code> on success.
<code>WiFi.softAPgetStationNum</code>	Return the number of stations currently associated with the soft-AP.
<code>WiFi.softAPIP</code>	Return the soft-AP's own IP address.
<code>WiFi.softAPBroadcastIP</code>	Return the soft-AP broadcast address.
<code>WiFi.softAPSSID</code>	Return the soft-AP SSID as a string.
<code>WiFi.softAPmacAddress</code>	Return the soft-AP MAC address as a colon-separated hex string.
<code>WiFi.softAPgetHostname</code> / <code>WiFi.softAPsetHostname</code>	Query or set the soft-AP mDNS hostname. <code>set</code> returns <code>#t</code> on success.

```
;; Bring up a private access point for a directly-attached companion computer.
(WiFi.mode 2) ; AP-only mode
(WiFi.softAPConfig "192.168.5.1" ; our AP address
  "192.168.5.1" ; gateway (self)
  "255.255.255.0") ; /24 subnet
(WiFi.softAP "lamb-net" "s3cr3t-key" 6 0 4) ; ch 6, visible, up to 4 clients

(display (list 'ap-ip (WiFi.softAPIP) ; => "192.168.5.1"
  'ap-ssid (WiFi.softAPSSID) ; => "lamb-net"
  'ap-mac (WiFi.softAPmacAddress)))
(newline)

;; Poll how many clients (e.g. the Jetson) have joined.
(display (list 'stations (WiFi.softAPgetStationNum)))
(newline)
```

6.12 Over-the-Air Updates

LambLisp can [hot-reload Scheme code without a reboot](#); new code takes effect immediately in the running control loop. The over-the-air update procedures cover the other case: replacing the VM binary itself, which does require a reboot.

The device carries a small, separate factory application (the *loader*) alongside the LambLisp VM. A running VM cannot rewrite its own flash image, so an update is done as a handshake. The VM records a request

in the NVS namespace `ota-p125` and reboots into `factory`, which does the actual fetch-and-flash and then hands control to the newly installed VM. Each procedure returns `#t` on success and `#f` on failure; on the Linux host build they are safe no-ops that return `#f` (or `#t` for `ota-boot-ok!`, which is vacuously true off-device).

Procedure	Description
<code>ota-request-install!</code>	Request an INSTALL on the next boot: the loader fetches and writes a new VM image from the configured URL. Returns <code>#t</code> if the request was recorded.
<code>ota-request-recover!</code>	Request a RECOVER on the next boot: the loader re-fetches a known-good image (recovery from a bad install).
<code>ota-request-console!</code>	Request a CONSOLE fallback on the next boot: the loader drops to an operator BLE/UART console instead of updating.
<code>ota-set-network! ssid pass url</code>	Store the Wi-Fi SSID, password, and image URL the loader will use for the network fetch. All three arguments are strings.
<code>ota-reboot-to-loader</code>	Point the OTA boot selector at the <code>factory</code> loader and restart immediately. Does not return on the device.
<code>ota-boot-ok!</code>	Confirm the currently running image is healthy: cancel the native rollback and clear the pending request plus the loader's install-retry counter.

6.12.1 The loader wire contract (NVS)

The VM and the `factory` loader are separate applications that never run at the same time, so they communicate through a small record in NVS (non-volatile storage). The procedures above are wrappers over this record; it is documented here so you can inspect or drive an update from either side.

- **Namespace:** `ota-p125`
- **Keys:**

Key	Written by	Meaning
<code>pending_update</code>	<code>ota-request-install! / -recover! / -console!</code>	1 = install, 2 = recover, 3 = console. The loader reads it on the next boot to decide what to do, then clears it.
<code>wifi_ssid / wifi_pass</code>	<code>ota-set-network!</code>	Credentials the loader brings Wi-Fi up with for the network fetch.
<code>img_url</code>	<code>ota-set-network!</code>	The URL the loader fetches the new VM image from.
<code>install_tries</code>	the loader (cleared by <code>ota-boot-ok!</code>)	The loader's retry counter; it stops re-installing after a few failed attempts.

`ota-boot-ok!` clears `pending_update` and `install_tries` together. Clearing that pair is what tells the loader the new image booted successfully.

6.12.2 The boot-confirm / rollback model

If an update rebooted into a broken VM, there would be no way to recover it remotely. To prevent that, the update flow requires the new image to confirm itself after it boots:

- After the loader installs a new image and reboots into it, the image starts in a *pending-verify* state. If it never confirms itself, the platform's OTA rollback (and the loader's own retry guard) revert to the previous working image on the next boot.
- The new VM confirms itself by calling `ota-boot-ok!` early in a good boot. That call cancels the rollback and clears both `pending_update` and the loader's `install_tries` retry counter.
- Call `ota-boot-ok!` only after you are satisfied the new image is actually working, for example once your control loop has come up cleanly, rather than at the very start of the program. Calling it too early would defeat the rollback.

6.12.3 Example

Configuring the network, requesting an install, and handing off to the loader:

```
;; 1. Tell the loader where to fetch the new image.
(ota-set-network! "my-ssid" "my-passphrase"
  "https://example.com/firmware/LambLisp.bin")

;; 2. Mark the next boot as an INSTALL.
(ota-request-install!)

;; 3. Reboot into the factory loader (does not return).
(ota-reboot-to-loader)
```

And the confirm-after-boot pattern, run inside the freshly installed VM once it is satisfied it is healthy:

```
;; ... startup and self-checks have completed successfully ...
(ota-boot-ok!) ;; cancel rollback, clear the pending request + retry counter
```

If the new VM crashes or hangs before reaching `ota-boot-ok!`, the loader reverts to the previous image on the next boot — the install cannot brick the device.

6.13 Runtime and GC Tuning

The [adaptive real-time garbage collector](#) tunes itself continuously from a Yuasa analysis, but LambLisp also exposes its timing, system, and GC/memory knobs directly to Scheme. These procedures let an application measure loop time, query free memory, and — where the automatic policy is not what you want — read and override the GC's tuning parameters. This section names the actual procedures behind the memory-tuning prose above.

Most procedures live in the `Platform.namespaces`. A few of the most common carry bare aliases as well (`micros`, `millis`, `delay-ms`, `gc!`). By convention a name ending in `!` is a **setter** (it takes an argument and returns an unspecified value); the same name without `!` is the matching **getter** (it takes no argument and returns the current value).

6.13.1 Timing and system

Procedure	Description
<code>Platform.micros / micros</code>	Microseconds since boot, as an integer.
<code>Platform.millis / millis</code>	Milliseconds since boot, as an integer.
<code>delay-ms n</code>	Busy-wait for <code>n</code> milliseconds using the platform delay.
<code>Platform.loop-elapsed-us</code>	Microseconds elapsed since the start of the current loop iteration.
<code>Platform.loop-elapsed-ms</code>	Milliseconds elapsed since the start of the current loop iteration.
<code>Platform.reboot</code>	Reboot the device (no-op on POSIX).
<code>lamblisp-exit</code>	Flush and exit the process (POSIX) or reboot (embedded).
<code>Platform.repl-quiet! flag</code>	Enable (<code>flag non-#f</code>) or disable line-mode quiet console: suppress async log output and per-char prompt redisplay. Used by the LLIP tunnel.
<code>Platform.random-integer</code>	A cryptographically random integer.
<code>Platform.random-real</code>	A cryptographically random float in $[0, 1)$.
<code>Platform.random-bytevector</code>	<code>(Platform.random-bytevector n) → a bytevector of <code>n</code> cryptographically random bytes.</code>
<code>Platform.free-heap</code>	Bytes of free heap memory on the platform.
<code>Platform.free-stack</code>	Bytes of free stack space on the platform.
<code>Platform.verbosity / Platform.verbosity! n</code>	Get / set the logging verbosity level (0 = silent, higher = more logging).

6.13.2 GC and memory tuning

Procedure	Description
<code>gc!</code>	Run one complete GC cycle (mark + sweep); useful to reset the GC phase before a benchmark.
<code>Platform.gc-idle-task! us</code>	Run the incremental GC idle task for up to <code>us</code> microseconds.
<code>Platform.cell-block-size</code>	Number of cells in the first (fixed) memory block.
<code>Platform.n-cell-blocks</code>	Current number of allocated memory blocks.
<code>Platform.nfree</code>	Current number of free cells in the GC heap.
<code>Platform.max-cell-blocks</code> / <code>Platform.↔ max-cell-blocks! n</code>	Get / set the maximum number of GC blocks to allocate before the memory-cap warning fires.
<code>Platform.extension-block-size</code> / <code>Platform.↔ extension-block-size! n</code>	Get / set the number of cells in each future extension (non-first) block.
<code>Platform.expand-to-n-blocks! n</code>	Grow the heap to at least <code>n</code> blocks now, allocating extension blocks as needed.

Procedure	Description
<code>Platform.gc-budget-ns</code> / <code>Platform.gc-budget-ns! ns</code>	Get / set the GC time budget per loop iteration, in nanoseconds.
<code>Platform.gcloud-target-pct</code> / <code>Platform.gcloud-target-pct! pct</code>	Get / set the target GC load as a percentage (0–100) of loop time; when exceeded, the heap grows toward the max.
<code>Platform.ncg-frame-pool-n</code>	Number of native-code-generator (NCG) frames currently in the free-list pool.
<code>Platform.ncg-frame-pool-init! n</code>	Add <code>n</code> NCG frames to the free-list pool (useful after pool exhaustion).

6.13.3 Example

Reading the current GC-load target and giving the collector a larger per-loop time budget:

```
(Platform.gcloud-target-pct)    ;; => e.g. 30  (grow memory once GC exceeds 30% of loop time)
(Platform.gc-budget-ns)       ;; => current per-loop GC budget in ns

(Platform.gc-budget-ns! 500000) ;; allow up to 0.5 ms of GC per loop iteration
(Platform.gcloud-target-pct! 20) ;; keep GC under ~20% of loop time; expand memory sooner

(Platform.free-heap)           ;; check headroom after tuning
(Platform.n-cell-blocks)       ;; how many blocks are currently allocated
```

To force a full collection before timing a benchmark (so no partial GC cycle skews the result):

```
(gc!)                          ;; complete mark + sweep now
(let ((t0 (micros)))
  (run-benchmark)
  (- (micros) t0))             ;; elapsed microseconds
```

6.14 LLIP — LambLisp Interaction Protocol

LLIP (LambLisp Interaction Protocol) is a line-oriented S-expression protocol over TCP or TLS that connects LambLisp instances. It serves two purposes: remote filesystem management (upload, download, directory operations) and remote evaluation (query state, invoke procedures, push new code). A single protocol covers both, using the same wire format and connection lifecycle.

6.14.1 Wire Format

Every message — request or response — is a single S-expression in `write` format followed by a newline. The server reads one line per message using LambLisp's `read` procedure, so strings, symbols, numbers, lists, and booleans all round-trip correctly without escaping or framing overhead. There is no length prefix, no binary encoding, and no session multiplexing — one request at a time per connection.

6.14.2 Connection Lifecycle




Figure 6.1 LLIP Connection Lifecycle

6.14.3 Operations

Request	Response	Description
<code>(ping)</code>	<code>(ok)</code>	Liveness check.
<code>(ls path)</code>	<code>(ok ("f1" ...))</code>	List directory; returns list of filename strings.
<code>(stat path)</code>	<code>(ok size)</code>	File size in bytes.
<code>(read path)</code>	<code>(ok "content")</code>	Full file content as string.
<code>(read path offset n)</code>	<code>(ok "chunk")</code>	Partial read; n bytes from offset.
<code>(write path "content")</code>	<code>(ok)</code>	Overwrite or create file.
<code>(append path "content")</code>	<code>(ok)</code>	Append to existing file.
<code>(delete path)</code>	<code>(ok)</code>	Delete file.
<code>(rename from to)</code>	<code>(ok)</code>	Rename file; both paths must be absolute.
<code>(load path)</code>	<code>(ok result)</code>	Load and eval file; return last expression value.
<code>(eval expr)</code>	<code>(ok result)</code>	Eval arbitrary expression in server's interaction environment.
<code>(bye)</code>	<code>(ok)</code>	Close connection.

Errors are returned as (error "message") and raised as Scheme errors on the client side, so guard can catch them.

6.14.4 The eval Operation

(eval expr) evaluates any Scheme expression in the server's interaction environment and returns the result as a Scheme value. This means the client can query live state, invoke procedures, reconfigure parameters, or upload and execute new behavior — all without restarting the board:

```
(llip-eval c '(Motor.speed))           ; query live value
(llip-eval c '(Motor.motion 1000 0.5 0.5 0.5 0.5)) ; enqueue a motion command
(llip-eval c '(set! loop-target-ms 5))   ; reconfigure a parameter live
(llip-eval c '(load "/new-behavior.scm")) ; deploy new code without reboot
(llip-eval c '(dict-keys (current-environment))) ; introspect all bound names
```

6.14.5 Client API (llip-client.scm)

Two calling modes are provided.

Synchronous — blocks until the result arrives. Simple, suitable for scripting and one-shot operations:

```
(define c (llip-connect-tcp "192.168.1.42" 8080 "token"))
(llip-eval c '(+ 1 2))           ; => 3
(define remote-fib (llip-make-remote c 'fib))
(remote-fib 30)                  ; => 832040
(llip-disconnect c)
```

Asynchronous — sends without blocking and polls for the result from the main loop. Suitable for managing multiple boards or remote GPUs concurrently:

```
(define fire-inference
  (llip-make-remote-async gpu-conn 'run-model
    (lambda (result) (apply-actuation board-conn result))))

(fire-inference sensor-vec)      ; returns immediately
;; in supervisor loop:
(llip-poll-async gpu-conn)       ; delivers result to callback when ready
```

Chunked file transfer helpers (llip-upload, llip-download) copy files between local and remote filesystems in fixed 512-byte chunks, keeping peak heap usage bounded on memory-constrained boards.

6.14.6 Server (llip-server.scm)

The server runs on the embedded board as a non-blocking subsystem polled from the main loop:

```
(define srv (llip-make-tcp-server 8080 "secret-token")) ; TLS variant: (llip-make-server port cert-pem key-pem token)
;; in loop:
(llip-server-poll srv)
```

llip-server-poll accepts one pending connection and processes one pending request per call, keeping per-iteration cost bounded. Path safety is enforced: all paths must be absolute, must not contain .., and must not contain null bytes.

6.14.7 Transports: TCP/TLS vs Serial

LLIP is transport-agnostic: the same S-expression-per-line codec runs over TCP, TLS, or a raw serial line. The difference is *who supplies reliability*.

Over **TCP/TLS** the transport already guarantees ordered, error-checked delivery, so the LLIP layer is just the line codec — write one S-expression, flush, read one line back. Frame helpers `llip-frame-write!` and `llip-frame-read` (in `llip-transport.scm`) implement exactly this: `read-line + parse-from-string` rather than `read` directly off the socket, because the reader's peek can mistake a split TCP segment for EOF mid-datum.

Over a **raw serial line** there is no transport reliability at all — a UART can drop, flip, or reorder bytes silently. LLIP therefore layers its own **stop-and-wait ARQ** on top of the serial REPL (`llip-serial.scm`): every data frame carries a Fletcher-16 checksum and a sequence number, and the sender retransmits on NAK or time-out. This is the reliability LLIP gets for free over TCP but must provide itself over serial.

Property	TCP / TLS	Serial
Reliability	native (TCP checksums + ARQ)	app-layer stop-and-wait ARQ (Fletcher-16 + seq)
Framing	one S-expr per line	one S-expr per line (same codec)
Integrity	TCP checksum	per-frame Fletcher-16 (<i>ck</i>)
Session setup	connect / accept	reset → boot → quiet mode
Best for	networked boards, supervisor mesh	a board on a USB cable, no network

6.14.8 Serial Transport and ARQ (`llip-serial.scm`)

The serial receiver is not a blocking loop. The frame verbs (`op` / `da` / `c1`) are plain top-level procedures that the ordinary REPL — itself non-blocking, assembling a line across main-loop ticks — evaluates one line at a time. Nothing blocks: the main loop keeps ticking sonar, LEDs, buzzer, and motors between frames.

Reliability is **stop-and-wait ARQ**. Each chunk frame is acknowledged before the next is sent; a bad checksum or wrong sequence number is answered with a NAK so the host retransmits.

Verb	Reply	Description
(<code>llt</code>)	—	Enter quiet mode: silence per-char echo and the async log flood so frames are not garbled.
(<code>op "path"</code>)	ACK	Open + truncate the destination file; reset sequence and counters.
(<code>da SEQ CRC "chunk"</code>)	ACK SEQ / NAK <i>exp</i>	Write one chunk if the Fletcher-16 CRC checks and SEQ is the expected next one; otherwise NAK <i><expected></i> . A duplicate (lost ACK) is re-ACKed.
(<code>c1</code>)	DONE <i>N</i> <i>FC</i>	Close the file; report total chars <i>N</i> and the aggregate checksum <i>FC</i> for an end-to-end check.
(<code>en</code>)	—	Leave quiet mode, back to interactive.

Verb	Reply	Description
<i>any other S-expr</i>	its own output	Evaluated normally; in quiet mode only its own (display ...) streams — this is how a host runs a command.

The checksum procedure `ck` computes Fletcher-16 over a string's codepoints and must match the host tool's `fletcher16` byte-for-byte:

```
(define (ck s)
  (let ((a 0) (b 0) (L (string-length s)))
    (let lp ((i 0))
      (if (< i L)
          (begin
             (set! a (modulo (+ a (char->integer (string-ref s i))) 255))
             (set! b (modulo (+ b a) 255))
             (lp (+ i 1)))
          (+ (* b 256) a))))))
```

Quiet mode is what makes literal sentinels safe on a serial line: with no per-char echo and no Input: <form> log, a host sends one frame line, reads one ACK line, and moves on without its own markers being echoed back into the stream.

6.14.9 Device-Initiated Serving (`llip-transport.scm`)

LLIP is symmetric — *either* end may initiate an exchange. For the robot stack the roles are reversed from the filesystem case: the Jetson orchestrator is the TCP **server**, and the 4WD car is the TCP **client** (it duty-cycles its WiFi for power, so the client drives connect/disconnect). `llip-transport.scm` provides both sides:

Procedure	Description
<code>(llip-frame-write! port frame)</code>	Write one frame as a flat S-expression + newline, then flush.
<code>(llip-frame-read port)</code>	Read one line and parse it to a frame; #f on EOF or parse failure.
<code>(llip-accept server-port)</code>	Poll the non-blocking <code>server-accept</code> (20 ms) until a client connects.
<code>(llip-conn-transport conn)</code>	Wrap a connection as the <code>(cons send! recv)</code> pair the orchestrator expects.
<code>(llip-serve-once server-port goal . opts)</code>	Accept one 4WD session, run the orchestrator loop over it, close (the close tells the 4WD to drop WiFi); returns <code>(done N)</code> or <code>(halt R)</code> .
<code>(llip-4wd-serve host port run-fn)</code>	Client side: connect, then loop receiving move-programs, running <code>run-fn</code> , and sending telemetry until the orchestrator closes the link.

The transport is a plain `(cons send! recv)` pair, so a stub transport swaps out for real TCP with no change to the orchestrator:

```
;;; Jetson side: serve one car session toward a goal.
(define server (open-tcp-server-port 5555))
(llip-serve-once server "drive to the dock" 4) ; ttl = 4 s

;;; 4WD side: run received move-programs and report telemetry.
(llip-4wd-serve "10.0.0.1" 5555 move-run-frame!)
```

6.14.10 Command-Line Client: `tools/llip`

`tools/llip` is the host-side LLIP client and listener — a single-file Python 3 tool (only `--serial` needs `pyserial`). Because LLIP is symmetric, it can both drive a device and serve a device that initiates.

Verb	Description
<code>send SRC [DEST]</code>	Push a file into the device (chunked, ARQ-verified on serial).
<code>recv REMOTE [LOCAL]</code>	Pull a file from the device (needs the device-side pull primitive; not in firmware yet).
<code>eval "FORM"</code>	Evaluate a form and stream its output up to a done marker.
<code>repl</code>	Interactive line REPL relayed over the link (Ctrl-D to exit).
<code>listen DIR</code>	Serve a device that <i>initiates</i> — this host becomes the frame receiver, mirroring the device's <code>op/da/c1</code> handler, so a device can push files or drive eval whenever it connects.

Transport selection. `--serial PORT [--baud N]` uses a serial line (default 115200) and supplies its own Fletcher-16 stop-and-wait ARQ; `--tcp HOST:PORT` uses TCP, whose native reliability means no app-layer ARQ and no reset/boot/quiet dance. With neither flag, a single `/dev/ttyUSB*` or `/dev/ttyACM*` is auto-detected. On serial the tool performs a clean reset + boot wait, then enters quiet mode (`,quiet`) so frames are not garbled, and restores the interactive console (`,verbose`) on close.

```
# Push a behavior file to the only attached board (serial auto-detected), then run it.
tools/llip send new-behavior.scm /new-behavior.scm
tools/llip eval '(load "/new-behavior.scm")'

# Same over TCP to a running LLIP TCP server on the board.
tools/llip --tcp 10.0.0.11:5555 eval '(Motor.speed)'

# Serve a device that initiates a file push into ./inbound (":PORT" binds all interfaces).
tools/llip --tcp :6000 listen ./inbound
```

The dev-side `w3_ai_scripts/llip_tunnel.py` provides the same serial ARQ tunnel as a Python library (`open_session / push_file / ev / close_session`) for scripted bring-up and the test runners; `tools/llip` is the shipped, general-purpose client.

6.14.11 HTTP / LLM Layer (`llip-http.scm`)

For AI-driven control, LambLisp talks to a local `llama-server` (`llama.cpp`) in-process over `open-tcp-client-port` — no subprocess, no Python: "only LL in the stack." JSON exists *only* here, as the thin envelope `llama.cpp`'s `/completion` API requires; everything the orchestrator passes around is LLIP S-expressions, and the LLM's "content" is itself a grammar-constrained LLIP move-program.

Procedure	Description
<code>(llip-http-post-json host port path body)</code>	POST a JSON body (HTTP/1.1 + <code>Connection: close</code>) and return the response body string, or <code>#f</code> on connect failure.

Procedure	Description
<code>(llip-llm-complete prompt grammar [host [port]])</code>	The LLM call: send a prompt + a GBNF grammar, return the grammar-constrained move-program string (the JSON content field), or #f. Host/port default to 127.0.0.1:8080.

```
(llip-llm-complete
  "go forward for two seconds then stop"
  llip-grammar) ; => "((forward (s 2)) (stop))"
```

6.14.12 AI Orchestration and the Robot Stack

The robot stack is a three-tier control loop built entirely in LambLisp. An LLM proposes a move-program, the grammar makes invalid output impossible, the 4WD re-validates and executes it through a bounded behaviour queue, and a reactive safety layer can veto any move at any instant. Intent flows down; safety is guaranteed at the bottom.

Deliberative tier — the orchestrator (`llip-orchestrator.scm`). Owns all I/O and state; the LLM stays a pure text function. The loop is: goal → plan (LLM) → move-program frame → send → recv telemetry → done / replan / halt.

Procedure	Description
<code>llip-grammar</code>	The GBNF grammar for the move DSL; mirrors <code>w3↔_ai_scripts/llip_move.gbnf</code> . Constraining generation to this grammar makes plans guaranteed-valid.
<code>(llip-build-prompt instruction)</code>	Wrap an instruction in the qwen ChatML system/↔user/assistant prompt.
<code>(llip-plan instruction [host [port]])</code>	Deliberative tier: call the LLM, read the result, and return a validated move list, or #f.
<code>(llip-make-move-program seq ttl moves)</code>	Build a <code>(move-program (seq ...) (ttl ...) (moves ...))</code> frame.
<code>(llip-orchestrate goal transport . opts)</code>	The orchestration loop over a <code>(cons send! recv) transport</code> ; returns <code>(done SEQ)</code> or <code>(halt REASON)</code> . Auto-recovers obstacle aborts by re-planning; deadman / planner-failure / unknown status → halt.

Executive tier — the executor (`llip-executor.scm`). Turns a validated move list into motor action on the shared cooperative motion queue. Safety by construction: it *interprets*, never *evals* — a fixed case dispatch (verb → binder) is the entire restricted environment, so no move-program can express anything but the seven verbs.

Procedure	Description
<code>(move-run-frame! frame)</code>	Execute a full move-program frame: feed the deadman, record seq, then run its moves.
<code>(move-run-program! bq moves)</code>	Re-validate the whole program (defence in depth — the 4WD never trusts the link) and enqueue each move on behaviour queue bq.
<code>(move-exec! bq m)</code>	The dispatch that <i>is</i> the sandbox: verb → binder, nothing else reachable; unknown verb is an error.
<code>(move-tank! l r)</code>	Tank drive: left/right side speeds → the four motors.
<code>(mv-forward ...)</code> <code>(mv-backward ...)</code> <code>(mv-turn-left ...)</code> <code>(mv-turn-right ...)</code> <code>(mv-wait ...)</code> <code>(mv-stop ...)</code> <code>(mv-dock ...)</code>	The seven per-verb binders; each enqueues one timed once-for task so a long move holds its slot without blocking the main loop.
<code>(move-estop!)</code>	Reactive e-stop: abort the motion queue and cut the motors.

Reactive tier — safety (`llip-reactive.scm`). Lowest in the stack, highest in priority. Safety lives on the 4WD, not the planner: the robot stays safe even when the plan is wrong, stale, or the link drops. `(reactive-tick!)` runs every main-loop iteration *before* the motor tick, so a reflex or the deadman can veto the executor before it advances another move.

Procedure	Description
<code>(reactive-tick!)</code>	Run the reactive vetoes; call before <code>Motor.loop</code> each iteration.
<code>(reactive-veto! signal reason)</code>	Preempt the plan: e-stop, raise a signal, and report an aborted frame up so the planner can re-plan.
<code>(reactive-deadman!)</code>	Halt if no valid frame arrived within <code>deadman-ms</code> (default 1500); latched, clears on recovery.
<code>(llip-feed-deadman!)</code>	Reset the heartbeat — called whenever an LLIP frame arrives.

The obstacle reflex is edge-triggered with hysteresis (trip below 0.15 m, re-arm above 0.25 m) so a recovery maneuver can run without being e-stopped every tick.

Signals (`llip-signals.scm`). `(robot-signal! event)` maps a robot event to the LED ring: base modes (`idle` / `running` / `dock`) persist; momentary signals (`bad-move` / `obstacle` / `low-battery`) are bounded blinks that auto-restore the base; deadman is the one sticky, highest-priority signal and the only one that paints the whole ring — safety wins.

Framing and validation (`llip-frame.scm`). Pure data, no I/O, so it loads and tests on any LambLisp (Linux or ESP32):

Procedure	Description
<code>(llip-move-program? frame)</code>	Is frame a well-formed <code>(move-program (seq N) (ttl T) (moves ...))</code> ?

Procedure	Description
<code>(llip-program-valid? moves)</code>	Is <code>moves</code> a valid move list — every move legal, length ≤ 32 ?
<code>(llip-make-telemetry seq move-index status reason sensors)</code>	Build a telemetry ("up") frame; <code>status</code> $\in$ running/done/aborted/halted.

The move vocabulary is seven verbs (forward backward turn-left turn-right wait stop dock) with unit-tagged magnitudes (`s m cm deg`); stop and dock are nullary. This vocabulary is defined once and mirrored by both the GBNF grammar (`w3_ai_scripts/llip_move.gbnf`) and the JSON schema (`w3_ai_scripts/llip_move_schema.json`), so the LLM, the validator, and the executor all agree on exactly what a legal move is.

6.14.13 Security

TLS is the production transport; plain TCP is available for development on a trusted LAN. The auth token is a shared secret that grants full filesystem and eval access. The `llip-safe-path?` predicate on the server prevents directory traversal.

6.14.14 Role in the Supervisory Architecture

LLIP is the uniform communication layer across all deployment tiers — embedded boards, edge supervisor (Jetson), and cloud GPU. The same client API is used regardless of whether the remote is an ESP32 on a LAN or a cloud LambLisp instance over WAN. Latency differences are absorbed by the async API: fast LAN connections resolve in the same loop tick, slow WAN connections resolve in a later tick, and the supervisor loop handles both transparently.

6.14.15 Worked Example: Supervisory Control of the Freenove 4WD Car

The following is drawn from `data/demo-supervisor.scm`. The Linux supervisor connects to the car over plain TCP, wraps each car operation as a thin Scheme procedure, and drives the car from the supervisor's own Scheme environment without any change to the car's on-board code:

```
;;; 1. Connect to the car from the Linux supervisor.
(define car (llip-connect-tcp "192.168.1.100" 8080 "demo-car"))

;;; 2. Wrap remote calls as local procedures.
;;; llip-eval sends a quoted expression to the car and returns its result.
(define (car-move! dir ms) (llip-eval car `(robot-move! ',dir ,ms)))
(define (car-stop) (llip-eval car `(robot-stop)))
(define (car-turn! deg) (llip-eval car `(robot-turn! ,deg)))
(define (car-sweep) (llip-eval car `(robot-sweep)))
(define (car-set-leds! pat) (llip-eval car `(set-leds! ',pat)))

;;; 3. Drive the car from the supervisor -- same syntax as on-board code.
(car-set-leds! '(red green blue red green blue red green blue red green blue))
(car-move! 'forward 1000) ; move forward for 1 s
(car-turn! 90) ; turn 90 degrees
(car-sweep) ; trigger sonar sweep; returns scan result list
(car-stop)

;;; 4. Query live state.
```



```
(llip-eval car '(get-sonar)) ; => distance in cm
(llip-eval car '(get-battery)) ; => battery %

;;; 5. Deploy new behavior without rebooting the car.
(llip-upload car "new-behavior.scm" "/new-behavior.scm")
(llip-eval car '(load "/new-behavior.scm"))

;;; 6. Disconnect when done.
(llip-disconnect car)
```

The car runs `llip-server.scm` on its main loop and never needs to know anything about the supervisor — it just handles `eval` requests as they arrive. The supervisor can inspect live telemetry, tune parameters, upload new `.scm` files, and hot-reload them, all while the car remains under continuous control.

6.14.16 Digital Twin

A *digital twin* is a LambLisp process on the edge or cloud whose Scheme environment mirrors the live state of a physical board. Telemetry is pulled from the board via LLIP; a Scheme model evolves in lockstep. The result is a live, queryable, scriptable replica of the physical device. It can answer questions the board cannot answer without interrupting its control loop:

- Replay the last N seconds of sensor history.
- Simulate a fault condition and observe how the control law responds.
- Validate that new code produces the same trajectory on the model before deploying it to hardware.
- Run a faster-than-realtime forecast of future system state.

The twin is a Scheme program that holds the device model as ordinary variables and procedures. A sync procedure runs periodically, reading live values from the board and writing them into the twin's environment:

```
;;; twin-state: last-known board state as a closure-backed dict object.
(define twin-state
  (dict->obj (list->dict '((speed . 0.0) (heading . 0.0) (battery . 100.0)
                        (sonar . ()) (loop-ms . 0)))))

;;; Pull one tick of live telemetry from the board into the twin.
(define (twin-sync! board)
  (twin-state 'speed (llip-eval board '(Motor.speed)))
  (twin-state 'heading (llip-eval board '(get-heading)))
  (twin-state 'battery (llip-eval board '(get-battery)))
  (twin-state 'sonar (llip-eval board '(get-sonar)))
  (twin-state 'loop-ms (llip-eval board '(loop-elapsed-ms))))

;;; Simplified kinematic model: predict (dx dy) over interval dt.
(define (twin-predict-position dt)
  (let ((v (twin-state 'speed))
        (hdg (twin-state 'heading)))
    (list (* v dt (cos hdg))
          (* v dt (sin hdg)))))
```

Validation before deployment. Before uploading new behavior to the board, run it against the twin. The twin loads the candidate file and evaluates the same sequence of inputs the board would receive, then compares outputs against an expected trajectory:

```
(define (validate-candidate! path)
  (llip-eval twin-conn `(load ,path))
  (let loop ((steps 100) (ok #t))
    (if (= steps 0)
        ok
        (let* ((predicted (llip-eval twin-conn '(step-model)))
               (expected (expected-trajectory steps)))
```

```
(loop (- steps 1) (and ok (close-enough? predicted expected))))))

(if (validate-candidate! "new-behavior.scm")
    (begin
      (llip-upload board "new-behavior.scm" "/new-behavior.scm")
      (llip-eval board '(load "/new-behavior.scm")))
    (display "Validation failed -- not deploying\n"))
```

Anomaly replay. The twin records a rolling window of telemetry snapshots. When an anomaly is detected the window is frozen and replayed offline:

```
(define telemetry-log '())
(define log-window 200) ; keep last 200 ticks

(define (twin-sync-and-log! board)
  (twin-sync! board)
  (set! telemetry-log
    (take (cons (twin-snapshot) telemetry-log) log-window)))

;;; On anomaly: print the frozen history.
(define (analyse-anomaly)
  (for-each (lambda (snap)
    (display (format "t=~a speed=~a sonar=~a\n"
      (snap 'loop-ms)
      (snap 'speed)
      (snap 'sonar))))
    (reverse telemetry-log)))
```

The twin runs entirely in Scheme on the edge machine; the on-board code requires no changes.




Figure 6.2 LambLisp Digital Twin

6.14.17 Unified Distributed Environment

When a supervisory controller manages hardware on multiple nodes — an ESP32 running sensors and actuators, a Jetson running a GPU inference engine, a Linux host running vision — the application logic should not have to know which function lives where. `llip-make-remote` projects a named procedure from a remote node into the local Scheme namespace. The result is an ordinary procedure: calling it sends the arguments over LLIP and returns the result. The distribution is invisible to the rest of the program.

```
;;; One-time setup: project remote bindings into the local namespace.
(define board (llip-connect "192.168.1.42" 7700))
(define gpu (llip-connect "192.168.1.43" 7700))

;;; Board bindings - sensor reads and motor commands execute on the ESP32.
(define get-sonar (llip-make-remote board 'get-sonar))
(define robot-move! (llip-make-remote board 'robot-move!))
```

```
(define Motor.speed (llip-make-remote board 'Motor.speed))

;;; GPU bindings - inference executes on the Jetson.
(define run-inference (llip-make-remote gpu 'run-model))
(define get-embeddings (llip-make-remote gpu 'get-embeddings))

;;; Local procedures run here on the supervisor.
(define (get-camera-frame) ...)
(define (choose-action embedding sonar) ...)
```

With the namespace unified, the perception–action loop reads as if everything were local:

```
(define (perception-action-loop)
  (let* ((frame (get-camera-frame)) ; local
        (embedding (get-embeddings frame)) ; Jetson GPU
        (dist (get-sonar)) ; ESP32 sensor
        (action (choose-action embedding dist)) ; local
        (speed (cdr (assq 'speed action))))
    (Motor.speed speed) ; ESP32 actuator
    (robot-move! (assq 'direction action)) ; ESP32 actuator
    (perception-action-loop)))
```

Moving a binding to a different node requires changing exactly one `llip-make-remote` call at setup; the loop body is unchanged. Adding a new node — say a second camera on a Raspberry Pi — is adding new bindings at setup.




Figure 6.3 Unified Distributed Environment

6.15 LCD1602 Display

Feature	Description
<code>LCD1602.begin</code>	Initialize an LCD1602 character display over I2C.
<code>LCD1602.clear</code> / <code>LCD1602.home</code>	Clear the display or move the cursor to position (0, 0).
<code>LCD1602.setCursor</code>	Move the cursor to a specific column and row.
<code>LCD1602.print</code> / <code>LCD1602.write</code>	Print a string or write a raw character code to the display.

Feature	Description
LCD1602.display / LCD1602.backlight / LCD1602.cursor	Control display-on, backlight-on, and cursor-visibility.

6.16 Ultrasonic Sonar (HC-SR04)

Feature	Description
Sonar.begin	Initialize a sonar sensor on given trigger and echo GPIO pins.
Sonar.start	Trigger a new distance measurement (non-blocking; ISR-driven).
Sonar.result	Return the most recent echo duration in microseconds.
Sonar.ping	Trigger and synchronously wait for one measurement; returns echo time.
Sonar.poke	Schedule a sonar ping and register a callback for the result.
Sonar.us->distance	Convert an echo time in microseconds to a distance in centimetres.

6.17 PCA9685 PWM Controller

Feature	Description
PCA9685.begin	Initialize a PCA9685 16-channel PWM controller over I2C.
PCA9685.setupSingleDevice	Configure one PCA9685 device on the bus.
PCA9685.setSingleDeviceToFrequency	Set the PWM output frequency for a device.
PCA9685.getSingleDeviceFrequency	Read back the configured PWM frequency.
PCA9685.setChannelDutyCycle	Set one channel's duty cycle as a percentage.
PCA9685.setChannelPulseWidth	Set one channel's pulse width in microseconds.
PCA9685.setChannelServoPulseDuration	Set one channel's pulse width for standard servo control.
PCA9685.setChannelOnAndOffTime	Set raw on/off tick values for one channel.

6.18 WS2812 Addressable LED Strip (NeoPixel)

Feature	Description
WS2812.begin	Initialize an addressable LED strip on a given data GPIO pin.
WS2812.setBrightness	Set the global brightness scaling factor (0–255).
WS2812.setLedColorData	Set the RGB color of one LED by index.
WS2812.show	Transmit buffered color data to the strip.
WS2812.Wheel	Convert a 0–255 position to an RGB colour traversing the hue wheel.

6.19 SPI Bus

The SPI procedures mirror the Arduino `SPIClass` API for full-duplex master-mode communication. The bus deliberately does **not** own the chip-select (CS) line — the caller drives CS explicitly with `digitalWrite` (from `CommonIO`), exactly as Adafruit's `SPIDevice` does — so a single bus can address any number of devices. A transaction (`beginTransaction ... endTransaction`) latches the clock rate and SPI mode (0–3) for the duration of one device exchange. On the Linux host build every procedure is a compile-clean no-op, so code loads and smoke-tests unchanged.

Procedure	Description
<code>SPI.begin</code>	Start the hardware SPI peripheral. Call with no arguments for the board defaults, or (<code>SPI.begin sck miso mosi</code>) to assign the clock, MISO, and MOSI GPIO pins (CS is left to the caller).
<code>SPI.end</code>	Release the SPI peripheral.
<code>SPI.beginTransaction</code>	Open a transaction: (<code>SPI.beginTransaction clock-hz [mode]</code>) sets the clock rate in Hz and the SPI mode 0–3 (clock polarity/phase; default 0), MSB-first.
<code>SPI.endTransaction</code>	Close the current transaction.
<code>SPI.transfer</code>	Full-duplex single-byte exchange: (<code>SPI.transfer byte</code>) clocks one byte out on MOSI and returns the byte simultaneously read on MISO.
<code>SPI.transfer-bytevector!</code>	In-place full-duplex block transfer: (<code>SPI.transfer-bytevector! bv</code>) sends each byte of <code>bv</code> and overwrites it with the byte read back, returning the same bytevector.
<code>SPI.write-bytevector</code>	Write-only block send: (<code>SPI.write-bytevector bv</code>) clocks every byte of <code>bv</code> out on MOSI, discarding MISO, and returns the number of bytes written.

```
;; Talk to an SPI device whose chip-select is on GPIO 10.
(define cs 10)
(pinMode cs 3)                ; 3 = OUTPUT
(digitalWrite cs 1)           ; CS idle high

(SPI.begin 12 13 11)           ; sck=12 miso=13 mosi=11
(SPI.beginTransaction 1000000 0) ; 1 MHz, mode 0

(digitalWrite cs 0)           ; select the device
(SPI.transfer #x9F)           ; send a "read-ID" command byte
(define id (bytevector 0 0 0))
(SPI.transfer-bytevector! id)  ; read 3 ID bytes back in place
(digitalWrite cs 1)           ; deselect

(SPI.endTransaction)
id                             ; => the device's 3-byte ID
```

6.20 OneWire (1-Wire Bus)

The OneWire procedures implement the Maxim/Dallas 1-Wire protocol directly in software, with microsecond-precise, IRAM-resident bit slots timed under a critical section (interrupts masked) for reliable signalling. Every procedure takes the GPIO **pin** as its first argument, so a single build can drive several independent 1-Wire buses. The pin is configured as open-drain with the internal pull-up enabled, but an external 4.7 kΩ pull-up to 3V3 is still required by the bus. On the Linux host build every procedure is a compile-clean no-op.

Procedure	Description
<code>onewire-reset</code>	<code>(onewire-reset pin)</code> issues a reset pulse and returns <code>#t</code> if at least one device answered with a presence pulse, else <code>#f</code> .
<code>onewire-write-byte</code>	<code>(onewire-write-byte pin b)</code> writes one byte, LSB first.
<code>onewire-read-byte</code>	<code>(onewire-read-byte pin)</code> reads and returns one byte, LSB first.
<code>onewire-write-bit</code>	<code>(onewire-write-bit pin b)</code> writes a single bit (0 or 1).
<code>onewire-read-bit</code>	<code>(onewire-read-bit pin)</code> reads and returns a single bit (0 or 1).
<code>onewire-search</code>	<code>(onewire-search pin)</code> runs the Maxim ROM-search algorithm and returns a list of 8-byte ROM bytevectors, one per device found (family code / LSB first).

```
;; Read a DS18B20 temperature sensor on GPIO 4 (single device on the bus).
(define ow 4)
```

```
(onewire-reset ow)
(onewire-write-byte ow #xCC)          ; SKIP ROM (address the only device)
(onewire-write-byte ow #x44)          ; CONVERT T (start a conversion)
(delay-ms 750)                        ; 12-bit conversion takes ~750 ms
```

```
(onewire-reset ow)
(onewire-write-byte ow #xCC)          ; SKIP ROM
(onewire-write-byte ow #xBE)          ; READ SCRATCHPAD
(define lsb (onewire-read-byte ow))
(define msb (onewire-read-byte ow))
(define raw (+ lsb (* 256 msb)))      ; signed 1/16 °C units
(/ raw 16.0)                         ; => temperature in °C
```

6.21 Freenove 4WD Car Kit — Peripheral Summary

The Freenove 4WD Car Kit (ESP32) demo uses the following hardware. Implemented devices have full LambLisp Scheme API bindings. Devices present on the board but not yet implemented in LambLisp are noted explicitly.

Device	Implementation status
4x DC motors (via PCA9685 I2C, address 0x5F)	Implemented — <code>PCA9685.*</code> procedures
2x pan/tilt servos (PCA9685, channels 0–1)	Implemented — <code>PCA9685.setChannelServo↔PulseDuration</code>
12x WS2812 RGB LEDs (GPIO 32)	Implemented — <code>WS2812.*</code> procedures
HC-SR04 ultrasonic sonar (TRIG=12, ECHO=15)	Implemented — <code>Sonar.*</code> procedures
Photosensitive light sensor (ADC GPIO 33)	Implemented — <code>analogRead</code>
Battery voltage divider (ADC GPIO 32)	Implemented — <code>analogRead</code>
Buzzer (GPIO 2, LEDC channel 0)	Implemented — <code>tone / noTone / ledc*</code> procedures
3x line-tracking IR sensors (PCF8574 I2C, address 0x20)	Implemented — <code>Wire.*</code> low-level I2C
IR receiver (GPIO 0)	<i>(not implemented)</i> — pin defined, no IR decode driver

Device	Implementation status
OV2640 camera (DVP interface)	LambLisp camera driver implemented (camera-init, camera-capture, camera->lcd, ...) — currently targeted at the esp32-s3-eye; not yet wired to this board's camera

6.22 Camera and On-Board Display (esp32-s3-eye)

LambLisp drives the OV2640 camera and the built-in ST7789 240×240 IPS LCD on the **esp32-s3-eye** dev board (8 MB PSRAM). The binding follows a strict division of labour: **C++ owns the pixels, LambLisp owns the policy**. The frame hot-paths — grab a frame, byte-swap and blit it to the panel, compute a mean-luma scalar — stay in C++ (ll_xmop3_EyeCam.cpp), while Scheme (scm/features/eye-cam.scm) drives the demo choreography and the eye's own main loop. The C++ layer needs only the ESP-IDF components already bundled in the Arduino-ESP32 core (esp_camera for the OV2640 DVP sensor, esp_lcd for the ST7789 over SPI3) — no external Arduino libraries.

All framebuffers and blits are **RGB565**. The OV2640 emits RGB565 in **big-endian** (byte-swapped) byte order, and esp_lcd ships memory to the panel in that same order, so a camera frame blits straight through with no swap — but a colour built in native order (e.g. an lcd-box fill) is byte-swapped once on its way to the wire. These procedures ship on the esp32-s3-eye target only (they are compiled under LL_CAMERA); on other boards the symbols are not registered.

6.22.1 Camera / LCD primitives (C++)

Registered by the EyeCam-install-mop3 installer. RGB565 colours are plain integers (0–65535).

Procedure	Description
(lcd-init) → #t	Bring up SPI3, the ST7789 panel, and the backlight. Idempotent; also reserves the internal-DMA band buffer used by every later draw.
(lcd-backlight [percent]) → void	Set backlight brightness 0–100 (default 100). 0 blanks the screen without tearing down the panel; useful to isolate a backlight fault from a draw fault.
(lcd-clear [rgb565]) → void	Clear the whole screen to one colour (default black).
(lcd-box x y w h rgb565) → void	Draw a filled axis-aligned rectangle in one RGB565 colour.
(lcd-rgb565! bvec x y w h) → void	Blit a big-endian RGB565 region from a bytevector to the panel at (x, y). Silently skips if the bytevector is shorter than w*h*2 bytes.
(lcd-testpat) → #t	Blit a known 240×240 test pattern (colour bars, grayscale ramp, fine stripes) through the same path as camera->lcd — pin-points display-path corruption with no camera involved.

Procedure	Description
<code>(camera-init [mode]) → #t/#f</code>	Bring up the OV2640. <code>mode</code> is 'rgb565 (default, drives the LCD) or 'jpeg (for camera-capture). Re-calling in the same format is a no-op that keeps the reserved DMA line buffer.
<code>(camera->lcd) → #t/#f</code>	Grab one RGB565 frame and blit it to the LCD — the live-preview hot path. <code>#f</code> if the camera or LCD is not up, or no frame is available.
<code>(camera-preview [n]) → n</code>	Blit <code>n</code> frames to the LCD in a tight blocking loop (default 300), then return. Freezes on the last frame; the eye-cam app-loop is the non-blocking alternative.
<code>(camera-thumb [n] [x0 y0]) → #t/#f</code>	Grab one RGB565 frame and print an <code>n×n</code> thumbnail over serial as hex RGB565 rows (<code>n</code> default 24, clamped 4–48). With <code>x0 y0</code> , prints a raw <code>n×n</code> crop at that origin instead of a subsample. For off-device visual verification without the LCD. Output is gated by the quiet REPL — send <code>,verbose</code> first.
<code>(camera-set key val) → void</code>	Tune the OV2640 sensor. <code>key</code> is one of <code>framesize</code> , <code>brightness</code> , <code>contrast</code> , <code>hmirror</code> , <code>vflip</code> , <code>special-effect</code> ; <code>val</code> is an integer.
<code>(camera-capture) → bytevector</code>	Return one JPEG frame as a bytevector (requires <code>(camera-init 'jpeg)</code>); <code>#f</code> if no frame is available.
<code>(frame-brightness) → integer</code>	Cheap mean luma (Rec. 601 weights) over a downsampled RGB565 frame, 0–255 — a scalar for trigger-policy demos ("brighter than <code>N</code> → fire"). Returns <code>-1</code> if the camera is down or not in RGB565 mode.

6.22.2 eye-cam.scm live-preview layer

Portable one-liners and the eye's application loop, built on top of the C++ primitives. This file also defines the RGB565 colour constants `red`, `green`, `blue`, `black`, and `white`.

Procedure	Description
<code>(eye-orient!)</code>	Correct the sensor orientation for the lab mount (<code>vflip + hmirror</code> , i.e. a 180° flip). Adjust here if the physical mount changes.
<code>(eye-preview)</code>	The one-liner demo: <code>lcd-init</code> , <code>camera-init</code> , <code>eye-orient!</code> , then a blocking <code>camera-preview</code> .
<code>(eye-demo)</code>	Flash a red/green/blue/white/black colour splash on the LCD to prove the panel is alive, then hand over to the live camera preview.
<code>eye-live?</code>	Boolean flag (initially <code>#f</code>); when true the app-loop blits a fresh camera frame every tick.
<code>(eye-live-start) → 'live</code>	Init the LCD + camera once, apply <code>eye-orient!</code> , and set <code>eye-live?</code> so the app-loop begins live blitting.
<code>(eye-live-stop) → 'stopped</code>	Clear <code>eye-live?</code> — freezes on the current frame while the loop keeps running (no blit).

Procedure	Description
(eye-make-loop)	The eye's main-loop <i>maker</i> . The eye target setup installs it via (set! app-loop eye-make-loop); setup.scm calls it once to build the loop that replaces the shared robot loop. The loop auto-starts the live preview on the first tick (safely after boot-time flash reads finish), blits when eye-live?, and donates spare time to the incremental GC.

Bring up the panel and camera and run a non-blocking live preview (the app-loop keeps the panel live while the REPL stays responsive):

```
(eye-live-start)      ; => live  -- LCD + camera up, blitting each loop tick
;; ... panel now shows the live camera feed ...
(eye-live-stop)       ; => stopped -- freeze on the current frame
```

Capture a thumbnail over serial for off-device verification (no LCD needed):

```
,verbose              ; un-gate log output so the thumbnail prints
(camera-init)          ; => #t   -- RGB565 mode
(eye-orient!)          ; correct the sensor orientation
(camera-thumb 32)      ; => #t   -- prints a 32x32 hex RGB565 thumbnail
```

6.23 Bundled Scheme Libraries

The low-level device primitives documented above (PCA9685.*, WS2812.*, Sonar.*, Wire.*, tone, ...) are C++ *mops* — thin, fast, and unopinionated. On top of them LambLisp ships a set of .scm **convenience layers** that turn those primitives into a friendly, board-aware API: named servos, a four-motor drive with an I2C write-cache, a fuel-gauge LED ring, buzzer patterns, drift-free timers, and I2C bus scanners. These files live under scm/core/ and are selected into each target's filesystem image by its manifest (scm/← targets/<env>.manifest), so a given board ships only the layers relevant to it.

Every layer **degrades gracefully**: if the underlying mop or the board's pin map is absent, the public procedures are rebound to a no-op (Lambda0) rather than erroring, so the same application code loads on a board that lacks the peripheral. The flagship Freenove 4WD Car worked example ([below](#)) is built almost entirely from these APIs.

6.23.1 Servos and Motors (PCA9685.scm)

Wraps the raw PCA9685.* PWM mops into normalized servos (input range −1..+1) and a four-wheel drive whose commands are queued as timed *behaviours* (see [Behaviors](#)) and de-duplicated against a per-channel write-cache so the I2C bus is touched only when a command actually changes.

Procedure	Description
pca9685-info	Object accessor for board config: (pca9685-info 'Xservo), 'MotorPins, 'servo-frequency, etc.
make-servo	(make-servo chan pwmin pwmax) → a (lambda (−1..1) ...) mapping the normalized range onto a channel's pulse-width span.

Procedure	Description
<code>xservo / yservo</code>	Ready-made pan / tilt servos (channels 0 / 1); call with $-1..+1$. <code>servo2...servo7</code> also defined.
<code>Motor.setpins</code>	Low-level: drive one motor from a forward/back channel pair at a signed fraction.
<code>Motor.set</code>	Set motor n (0–3) to $-1..+1$; writes I2C only when the value changes.
<code>Motor.set4</code>	Set all four motors at once: (<code>Motor.set4 m0 m1 m2 m3</code>).
<code>Motor.setAll</code>	Set all four motors to one speed.
<code>Motor.motion</code>	Queue a timed move: hold (<code>m0 m1 m2 m3</code>) for duration <code>ms</code> , then run the next queued motion.
<code>Motor.loop</code>	Tick the motion queue once per main loop; auto-stops the motors when the queue drains.
<code>Motor.estop!</code>	Emergency stop: flush the motion queue and cut all four motors immediately.
<code>Motor.last</code>	Vector of the last value written to each motor channel (the change-cache).

```
(xservo 0) (yservo 0)           ; center the pan/tilt head
(Motor.motion 1000 0.5 0.5 0.5 0.5) ; forward at half speed for 1 s
(Motor.motion 1000 -0.5 -0.5 -0.5 -0.5); then reverse for 1 s
;; ... then, once per main-loop iteration:
(Motor.loop)                   ; advances the queue; stops when idle
```

6.23.2 Addressable LEDs (`WS2812.scm`)

A ring/strip helper over the `WS2812.*` mops: whole-strip and per-LED colour setters, a symbolic colour palette, nautical navigation lights, and non-blocking blink behaviours driven by a shared LED queue (`led-bq`).

Procedure	Description
<code>WS2812.setAll</code>	Set every LED to one <code>r g b</code> colour and show.
<code>WS2812.off</code>	Turn the whole strip off.
<code>WS2812.set</code>	Set one LED by index to <code>r g b</code> and show.
<code>WS2812.setColor</code>	Set one LED to a named colour: (<code>WS2812.setColor 0 'rgb-red</code>).
<code>WS2812.test</code>	Queued startup self-test: flash all LEDs red → green → blue → off.
<code>WS2812.navLights</code>	Port red, starboard green, fore/aft white.
<code>rgb-colors</code>	Palette accessor: (<code>rgb-colors 'rgb-cyan</code>) → (0 128 128). Colours: black, red, green, blue, cyan, yellow, magenta, white.
<code>ws2812-info</code>	Config accessor: 'Nleds, 'ledchan, 'ledtype, colour-order type codes.
<code>Led.blinkOnce</code>	Queue one blink of LED <code>nled</code> : on color for on-time <code>ms</code> , off off-time <code>ms</code> .
<code>Led.blinkAll</code>	Queue one blink of the whole ring.
<code>Led.loop</code>	Tick the LED queue once per main loop.

```
(WS2812.navLights)           ; running lights
(Led.blinkAll 'rgb-red 200 200) ; queue a 200/200 ms red blink
;; ... once per main-loop iteration:
(Led.loop)
```

6.23.3 LED Ring Service (Leds.scm)

(Ships only on the Freenove 4WD Car target.) A higher-level, robot-agnostic service layered on WS2812.scm: it splits the ring into an illumination arc and a signal arc, runs a **priority arbiter** (transient signals preempt a persistent base mode, then restore it), and paints the ring as a **battery fuel gauge**.

Procedure	Description
led-illuminate!	Turn the forward white illumination arc on/off (#f clears it).
led-emit!	(led-emit! pri task) — play a transient signal task if its priority $\geq$ the current one, then restore the base mode.
led-clear!	Hard-clear the arbiter: drop any signal, reset priority, restore base mode.
led-blink	Build a blink-pattern task for a zone: $n > 0$ blinks then done, $n = 0$ blinks forever.
battery-check	Read the battery ADC, paint the ring green/yellow/red as a charge gauge, log volts, return raw ADC.
battery-lit	Compute how many LEDs to light for a raw ADC reading (the gauge math).

```
(battery-check)      ; ring shows remaining charge; logs "ADC ... N/12 LEDs ... ~ V"
(led-illuminate! #t) ; front arc white for the camera
```

6.23.4 Buzzer (Buzzer.scm)

Non-blocking tone patterns over the tone mop, queued on a dedicated behaviour queue (Buzzer.bq).

Procedure	Description
Buzzer.tone	The underlying tone primitive (rebound to no-op if the board has no buzzer pin).
Buzzer.on / Buzzer.off	Start a continuous tone at freq Hz / silence it.
Buzzer.click	A single short 8 ms acknowledgement click (queued).
Buzzer.onoff	Queue tone-on for on-time ms, then off for off-time ms.
Buzzer.offon	Queue off then on (the inverse phase order).
Buzzer.arf	A rising chirp "bark" (60→4000 Hz). Built on Buzzer.chirp.
Buzzer.loop	Tick the buzzer queue once per main loop.

```
(Buzzer.click)      ; short acknowledgement beep
(Buzzer.onoff 2000 100 100) ; queue a 2 kHz 100/100 ms beep
;; ... once per main-loop iteration:
(Buzzer.loop)
```

6.23.5 Timers (Timers.scm)

Closure-based, drift-free timing helpers. Each constructor returns a procedure you poll; none of them block.

Procedure	Description
Stopwatch_ms / Stopwatch_us	Return a thunk giving elapsed ms / μ s since the stopwatch was created.
AutoTimer_ms	Countdown timer: returns remaining ms while running, #f when expired (auto-restarts). Starts expired.
KitchenTimer_ms / KitchenTimer_us	Return an object with 'remaining, 'set!, 'expired?, 'running? messages.
Oneshot	Return a thunk that yields #t the first call and #f forever after.
uptime_ms	Milliseconds since boot (a shared stopwatch).
make-idler	Rate-limiter: returns a thunk that yields #t when its interval has elapsed <i>and</i> the loop has spare time.

```
(define t (Stopwatch_ms))
(do-some-work)
(info "took ~a ms\n" (t))

(define beat (AutoTimer_ms 1000))
;; in the loop: (unless (beat) (heartbeat)) ; fires ~once/second
```

6.23.6 I2C GPIO Expander (PCF8574.scm)

Driver for the PCF8574 8-bit I2C port expander (on the 4WD car it reads the line-tracking IR sensors). PCF8574 builds a per-address object with 'read8 / 'write8 messages.

Procedure	Description
PCF8574	(PCF8574 addr) → an object with 'read8 (returns last input byte) and 'write8 messages.
pcf8574	A ready-made instance at the default address 0x20.
PCF8574-I0	Low-level raw read8 / write8 taking an explicit address (over Wire).

```
(define ir pcf8574)
((dict-ref ir 'read8)) ; read the 8-bit input byte (e.g. IR line sensors)
```

6.23.7 Light Sensor (LightSensor.scm)

Procedure	Description
light-sensor	Read the photosensor ADC, normalized to 0.0–1.0. No-op returning nothing if the board has no pin-photo.

```
(light-sensor) ; => e.g. 0.42 (brighter = larger)
```

6.23.8 I2C Bus Utilities (I2C.scm)

Convenience helpers layered on the native Wire module for bringing up and scanning the bus.

Procedure	Description
<code>I2C.begin</code>	Set the SDA/SCL pins and begin the bus as master: <code>(I2C.begin sda scl addr)</code> .
<code>I2C.ping</code>	Probe one address; returns the transmission error code (<code>#f/0</code> = present).
<code>I2C.inventory</code>	Scan addresses 0x00–0x7F, log expected vs. unexpected devices, return the count found.

```
(I2C.begin pin-sda pin-scl 0) ; 0 = we are bus master
(I2C.inventory)               ; => N ; logs each device found
```

6.23.9 Sonar Convenience Wrappers (`Sonar.scm`)

The low-level `Sonar.*` mops (`Sonar.begin`, `Sonar.start`, `Sonar.result`, `Sonar.us->distance`, ...) are documented under [Ultrasonic Sonar](#). This `.scm` layer adds board-aware setup and a **continuous, non-blocking** ranging loop that keeps the latest distance in metres.

Procedure	Description
<code>Sonar.pins</code>	The <code>(trig echo)</code> pin pair from the board map (or <code>(0 0)</code>).
<code>Sonar.range-m</code>	Rated maximum range in metres (reported when nothing is in range).
<code>Sonar.latest</code>	The most recent measured distance in metres (updated by <code>Sonar.loop</code>).
<code>Sonar.setup</code>	Initialize the sensor from <code>Sonar.pins</code> and arm the first ping.
<code>Sonar.loop</code>	Send-or-expect tick: resolve the in-flight ping, update <code>Sonar.latest</code> , launch the next. Call once per main loop. Never blocks.

```
(Sonar.setup)
;; ... once per main-loop iteration:
(Sonar.loop)
(when (< Sonar.latest 0.30) (Motor.estop!)) ; obstacle reflex
```

6.24 Behaviors: Cooperative Scheduling

`scm/core/Behaviors.scm` is the cooperative, non-blocking scheduling DSL that the device layers above are built on. It exists because a `(loop)`-driven robot must never block: every timed move, blink, or beep has to interleave with sensor reads, comms, and reflexes.

The whole protocol is a **step-thunk**: a zero-argument procedure, called once per main-loop tick, that returns `#t` when it is finished (pop it) or `#f` to keep running. A *behaviour queue* runs the head step-thunk each tick and advances when it returns `#t`; you drive one queue per subsystem with `('tick)` once per main loop. Motor moves, the LED ring, the buzzer, and the LLIP executor are all just consumers of this one module.

6.24.1 Task Constructors

Procedure	Description
while-do	(while-do pred action) — run action every tick while pred; done when pred is false. The basis of all the rest.
until-do	Run action every tick until pred becomes true.
once	Do action exactly once, then pop.
behavior-hold	A task that does nothing and never finishes (pure placeholder).
for-ms	Run action every tick for ms milliseconds.
once-for	Do action once, then hold the slot for ms ms before popping (e.g. one motor move held for a duration).
after-ms	Wait ms ms, then do action once and pop (deferred action / start delay).
wait-ms	Hold the slot for ms ms doing nothing (inter-move gap).
wait	The canonical drift-free timer task: hold the queue head for ms ms, then pop. Deadline latches on the first tick at the head.
forever	Run action every tick, forever (ended only by queue abort).
cycle	Cycle through (list duration-ms thunk) phases forever — a whole blink/animation as one task.
seq	Compose sub-tasks into one task, run in order; done when all are consumed. Lets a multi-step sequence occupy a single queue slot.

6.24.2 The Queue

Procedure / Message	Description
make-behavior-queue	Return a message-dispatch queue procedure.
(bq 'push task)	Enqueue a step-thunk.
(bq 'tick)	Run one slice; call once per main loop.
(bq 'abort)	Reactive e-stop: drop the running task and flush everything behind it.
(bq 'idle?)	#t when nothing is running and the queue is empty.
behavior-queue-seq	Push several tasks in order: (behavior-queue-seq bq t1 t2 ...).

```
(define bq (make-behavior-queue))
;; A start-delayed, timed motion sequence, all non-blocking:
(behavior-queue-seq bq
  (after-ms 500 (lambda () (Motor.setAll 0.5))) ; wait 0.5 s, then go
  (wait 1000) ; ...run for 1 s
  (once (lambda () (Motor.setAll 0)))) ; then stop
;; drive it every main-loop iteration:
(bq 'tick)
;; ... and to bail out at any moment:
(bq 'abort)
```

6.25 ESP32 System

Feature	Description
millis / micros	Elapsed time since boot (millisecond / microsecond resolution).
esp32-efuse-mac-get-default	Read the factory-programmed 48-bit MAC address from eFuse.
esp32-efuse-pkg-ver	Read the ESP32 package version from eFuse.
esp32-efuse-read-block esp32-efuse-write-block	/ Read or write raw eFuse data blocks.

6.25.1 Memory Tuning: Stack Size, DMA Memory, and Pre-Reserve

On the ESP32 the scarce resource is **internal DRAM** (~320–350 KB on the S3), *not* PSRAM. The LambLisp cell heap lives in PSRAM (megabytes), so it is rarely the constraint; the tasks that fail on a tight board are the ones that need **internal DRAM**, and specifically its **DMA-capable** subset. Three levers interact — get them wrong and a board with plenty of PSRAM still can't read its own filesystem.

1. loopTask stack size. The application runs on the Arduino `loopTask`, whose stack is carved from internal DRAM. LambLisp overrides the framework's weak `getArduinoLoopTaskStackSize()` (in `main.cpp`) — e.g. 48 KB. Bigger is not free: every KB of stack is a KB of internal DRAM denied to DMA. Two things bound how small it can go:

- `LL_EVAL_STACK_BUDGET` (`ll_platform_generic.h`) — this guard raises a *catchable* "recursion too deep" error once `eval()` grows this far past its outermost frame. **Keep it below the loopTask stack size** (e.g. 40 KB budget under a 48 KB stack, ~8 KB headroom). If the budget exceeds the real stack, deep recursion overflows the C stack — a hard SIGSEGV/panic — instead of a clean error.
- **The reader is *not* guarded.** `read_list` recurses once per paren-nesting level *and* runs a read-time `macroexpand` at each level, so a single deeply-nested top-level form can overflow the stack while it is being *read* (array.scm's transpose nested ~12 deep and blew the 96 KB stack while other files were fine). Fix by keeping .scm forms shallow — factor deep helpers into flat top-level definitions rather than one monster expression.

2. Flash reads need DMA memory. `esp_flash_read` (and therefore every LittleFS read) needs a `MALLOC_CAP_DMA`-capable buffer in **internal** SRAM. PSRAM is *not* DMA-capable for flash, so a read into a PSRAM destination allocates an internal bounce buffer. Consequence: a read can fail with `ESP_ERR_NO_MEM` (**err 257** in the `esp_littlefs`: `failed to read ... err 257` message) even with **megabytes of PSRAM free** — what matters is the *largest contiguous free MALLOC_CAP_DMA block*, which can be a few KB while total internal free still looks healthy (100 KB+). This is the failure class behind the eye and N8R2 (2 MB PSRAM): once the DMA-capable block is fragmented or starved, the filesystem is effectively unreadable and boot-time (`load ...`) silently produces nothing.

3. The pre-reserve technique — and its trap. To guarantee a peripheral or the filesystem always has DMA memory, **reserve a contiguous MALLOC_CAP_DMA block early** — at install/setup time, before the heap and other subsystems fragment internal SRAM — and hold it for the code that needs it (or free it just-in-time around the operation). `ll_xmop3_EyeCam.cpp` does this for the camera/LCD band buffer.

The trap: an over-eager early reserve *starves* the very reads it was meant to protect. In that case the camera's early DMA-hole reserve dropped the largest free DMA block to ~16 KB, so **boot-time LittleFS reads returned garbage**; the fix was to gate that reserve *off* (`EYE_CAM_EARLY_DMA_RESERVE 0`), leaving ~32 KB DMA free through boot. So the technique cuts both ways:

- Reserve **enough** DMA for what genuinely needs it (camera frame/LCD band buffers).
- But **not so much** that boot-time flash reads have no DMA block left.
- On memory-tight boards (2 MB-PSRAM N8R2), that balance is the whole game — shrinking the loop↔ Task stack frees *general* internal DRAM but does not by itself defragment or enlarge the DMA-capable block, so a small reserved-and-held DMA buffer for the FS-read path is usually what actually fixes it.

Rule of thumb. Keep the loopTask stack as small as the reader/eval recursion allows (shallow `.scm` forms let you go smaller), keep `LL_EVAL_STACK_BUDGET` a notch under it, and protect the filesystem's DMA needs with a small, early, held reserve rather than hoping enough DMA-capable memory survives boot.

6.26 LambLisp Modbus Interface

LambLisp includes a full Modbus binding for both controller (master) and peripheral (slave) roles over TCP and RTU transports. This chapter covers the API for both roles, platform-specific build requirements, and worked examples for common industrial IoT scenarios.




Figure 6.4 Typical Modbus Field Architecture

6.26.1 1. Overview

Modbus is a serial communication protocol published in 1979 and now ubiquitous in industrial automation. It defines a request/response model between a **controller** (the device that initiates requests) and one or more **peripherals** (devices that respond to requests). The protocol is simple enough to be embedded in PLCs, variable-frequency drives (VFDs), energy meters, temperature controllers, and thousands of other field devices — many of which still ship with Modbus as their primary communication interface.

LambLisp supports two roles and two transports:

Role	Description
Controller	LambLisp initiates read/write requests to peripheral devices (PLCs, VFDs, meters, sensors)
Peripheral	LambLisp exposes a register bank to an external SCADA, HMI, or PLC controller

Transport	Description
TCP	Modbus framing over TCP/IP, port 502. Used over Ethernet and WiFi.
RTU	Binary framing over RS-485/RS-232 serial. Used for wired field buses.

The same function calls work regardless of transport; only the connect call differs.

6.26.2 2. LambLisp as Modbus Controller

In controller mode, LambLisp initiates requests to one or more peripheral devices and reads back the results. The binding uses a **non-blocking request/poll model** designed for real-time embedded loops.

2.1 Connecting

```
;; TCP: connect to a device at 192.168.1.10, Modbus unit ID 1
(define conn (modbus-tcp-connect "192.168.1.10" 502 1))

;; RTU: connect over UART2, 9600 baud, unit ID 3
(define conn (modbus-rtu-connect 2 9600 3))

;; Close when done
(modbus-close conn)
```

(modbus-tcp-connect host port unit-id) — Opens a TCP connection to host:port targeting peripheral unit-id. Returns a connection object. Raises an error if the connection cannot be established.

(modbus-rtu-connect uart-num baud unit-id) — Opens an RTU connection on the specified UART at the given baud rate targeting peripheral unit-id. On ESP32, configure the DE/RE pin with modbus-rtu-set-de-pin before opening (see [Platform Notes](#)).

(modbus-close conn) — Closes the connection and releases associated resources.

2.2 The Non-Blocking Request/Poll Model

Modbus devices — especially RTU peripherals — can take 50–500 ms to respond. On an embedded system running a real-time loop at 10–100 Hz, blocking that entire loop to wait for a sensor response is unacceptable. The LambLisp Modbus binding therefore uses a **two-phase model**: modbus-request queues the request and returns immediately; modbus-poll checks whether the response has arrived.




Figure 6.5 Non-blocking request/poll flow

```
;; Phase 1: issue request - returns a token immediately, never blocks
(define tok (modbus-request conn 'read-holding 40 2))

;; Phase 2: poll on each loop iteration
(define (check-result tok)
  (let ((result (modbus-poll tok)))
    (if result
        (begin (display "got: ") (display result) (newline))
        (display "not yet\n")))))
```

`modbus-poll` returns `#f` when the response has not yet arrived, and the result vector when it has. The calling loop simply does other work between polls and checks again on the next iteration. The loop never blocks.

`(modbus-request conn function-code start-address count)` — Queues a request and returns a token. `function-code` is a symbol: `read-coils`, `read-discrete-inputs`, `read-holding`, `read-input`, `write-coil`, `write-register`, `write-coils`, `write-registers`. `start-address` is the 0-based register address. `count` is the number of registers (for reads) or omitted for single-register writes.

`(modbus-poll token)` — Returns `#f` if the response has not arrived, or the result on completion. The result for read operations is a vector of integers or booleans; see FC-specific sections below. Raises an error if the device returned an exception code.

`(modbus-set-timeout conn ms)` — Sets the response timeout in milliseconds. Default is 1000 ms. If `modbus-poll` is called after the timeout has elapsed without a response, it raises an error.

2.3 Function Codes

All reads and writes use `modbus-request` with a symbol for the function code (see the [Quick Reference](#) for the full table). Reads return a vector; writes return `#t` on success.

```
;; Read 4 holding registers from address 0
(define tok (modbus-request conn 'read-holding 0 4))
(let ((regs (modbus-poll tok)))
  (display (/ (vector-ref regs 0) 10.0)) ; e.g. VFD frequency scaled *10
  (display " Hz\n"))

;; Write a single register; poll returns #t
(define tok2 (modbus-request conn 'write-register 9 1500))
(modbus-poll tok2) ; -> #t
```

For multi-value writes pass a vector as the value argument:

```
(modbus-request conn 'write-registers 10 (vector 1500 0))
(modbus-request conn 'write-coils      0 (vector #t #f #t #t))
```

2.4 Float32 Helpers

Many industrial devices store 32-bit floats in two consecutive 16-bit registers. Word order varies by manufacturer: `'abcd` (most Schneider/ABB/Siemens equipment) or `'cdab` (some Modicon PLCs and meters). Check your device manual.

`(modbus-request-float32 conn addr word-order)` — FC03 read of 2 registers. `(modbus-poll-float32 token)` — Returns the reassembled inexact real, or `#f`.

```
(define tok (modbus-request-float32 conn 100 'cdab))
(let loop ()
  (let ((temp (modbus-poll-float32 tok)))
    (or temp (begin (do-other-work) (loop))))))
```

2.5 Error Handling

`modbus-poll` raises when the device returns a Modbus exception, the request times out, or the connection is lost. Use `guard` to handle errors without crashing the loop:

```
(guard (e (#t (display (error-object-message e)) (newline) #f))  
  (let loop ()  
    (or (modbus-poll tok) (begin (do-other-work) (loop))))))
```

Exception codes appear in the message string, e.g. "Modbus exception 02: illegal data address".

6.26.3 3. LambLisp as Modbus Peripheral

In peripheral mode, LambLisp maintains a register bank that an external controller (SCADA system, HMI panel, or PLC) can read and write using standard Modbus function codes. The Scheme code updates the register bank from sensor readings, and optionally receives callbacks when the controller writes new setpoints or commands.




Figure 6.6 Modbus register types

3.1 Opening a Peripheral

```
;; TCP peripheral: listen on port 502, present as unit ID 1
(define peripheral (modbus-tcp-peripheral-open 502 1))
```

```
;; RTU peripheral: UART 2, 9600 baud, unit ID 5
(define peripheral (modbus-rtu-peripheral-open 2 9600 5))
```

(modbus-tcp-peripheral-open port unit-id) — Opens a TCP listening socket on port. On Linux, port 502 requires root or a capability grant; using port 5502 is common during development. Returns a peripheral object.

(modbus-rtu-peripheral-open uart-num baud unit-id) — Opens the specified UART as a Modbus RTU peripheral. Returns a peripheral object.

(modbus-peripheral-close peripheral) — Closes the peripheral and releases the port/UART.

3.2 The Four Register Tables

Modbus defines four data tables. The table below shows FC read/write codes and access permissions from both sides.

Table	FC Read	FC Write	Controller access	Scheme access
Coils	FC01	FC05, FC15	Read + Write	Read + Write
Discrete Inputs	FC02	—	Read only	Read + Write
Input Registers	FC04	—	Read only	Read + Write
Holding Registers	FC03	FC06, FC16	Read + Write	Read + Write

3.3 Register Accessors

All addresses are 0-based.

```

;;; --- Holding registers ---
(modbus-peripheral-hreg! peripheral address value) ; write from Scheme
(modbus-peripheral-hreg peripheral address)        ; read from Scheme

;;; --- Input registers ---
(modbus-peripheral-ireg! peripheral address value) ; write from Scheme
(modbus-peripheral-ireg peripheral address)        ; read from Scheme

;;; --- Coils ---
(modbus-peripheral-coil! peripheral address bool)  ; write from Scheme
(modbus-peripheral-coil peripheral address)        ; read from Scheme -> #t or #f

;;; --- Discrete inputs ---
(modbus-peripheral-dinput! peripheral address bool) ; write from Scheme
(modbus-peripheral-dinput peripheral address)      ; read from Scheme -> #t or #f

```

Values for registers must be exact integers in [0, 65535]. Coil and discrete input values must be booleans.

3.4 The Serve Loop

`modbus-peripheral-poll` processes any pending incoming requests from the controller. It is non-blocking: if no request is pending it returns `#f` immediately; if a request was processed it returns `#t`. Call it on each iteration of the main loop.

```

(define (main-loop)
  ; Update process data in the register bank
  (modbus-peripheral-ireg! peripheral 0
    (inexact->exact (round (* (read-temperature) 10))))
  (modbus-peripheral-ireg! peripheral 1
    (inexact->exact (round (* (read-pressure) 100))))
  ; Serve any pending controller requests
  (modbus-peripheral-poll peripheral)
  ; ... other real-time work ...
  (main-loop))

```

3.5 Write Callbacks

To be notified when the controller writes to coils or holding registers, register a callback with `modbus-peripheral-on-write!`. The callback receives the table symbol, the start address, and a vector of written values.

```
(modbus-peripheral-on-write! peripheral
  (lambda (table start-addr values)
    (cond
      ((eq? table 'holding)
        (when (= start-addr 10)
          ;; Controller wrote a new speed setpoint to HR40011
          (set-motor-speed! (/ (vector-ref values 0) 10.0))))
      ((eq? table 'coils)
        (when (= start-addr 0)
          ;; Controller toggled output relay 1
          (set-relay! 0 (vector-ref values 0)))))))
```

The callback is invoked synchronously inside `modbus-peripheral-poll` after the register bank has been updated and the response has been sent to the controller.

6.26.4 4. Platform Notes

4.1 Linux (x86_64 and aarch64)

The Linux backend uses **libmodbus** (<https://libmodbus.org>).

Install the development package:

```
sudo apt-get install libmodbus-dev
```

Enable Modbus in the build by defining `LL_MODBUS 1` in the board configuration header (e.g. `ll_board_↵
linux_x86.h`):

```
#define LL_MODBUS 1
```

The build system links `-lmodbus` automatically when `LL_MODBUS` is set.

For TCP peripheral mode on port 502, either run as root or grant the capability:

```
sudo setcap cap_net_bind_service=+ep ./lamblisp
```

Alternatively, use port 5502 during development and configure the SCADA simulator to connect to that port.

4.2 ESP32-S3

The ESP32 backend uses **no external library**. Modbus frames are built and parsed directly using `WiFi↵
Client` (TCP) and `HardwareSerial` (RTU) from the Arduino framework. No `lib_deps` entry is required for the Modbus environment.

Enable Modbus in the board header:

```
#define LL_MODBUS 1
```

RTU wiring for RS-485: RS-485 transceivers (e.g. MAX485, SN65HVD) require a direction-enable pin (DE/RE). Configure it before issuing any RTU requests:

```
(modbus-rtu-set-de-pin conn 4) ; example: GPIO 4 drives DE/RE on ESP32-S3 DevKit
```

`modbus-rtu-set-de-pin` sets the GPIO pin asserted HIGH before transmit and released LOW after `flush()`. Connect DE and RE of the transceiver together and wire to the specified GPIO. The default pin (`LL_MODBUS↵
_RTU_DE`) is `-1` (disabled), suitable for RS-232 or auto-direction transceivers.

Recommended UART-to-RS-485 pin assignments for ESP32-S3 DevKit:

Signal	GPIO
UART2 TX	17
UART2 RX	18
DE/RE	4

WiFi must be initialized before using Modbus TCP. Call `(wifi-connect ssid pass)` (or equivalent board setup) before `modbus-tcp-connect`.

TCP peripheral on ESP32: `modbus-tcp-peripheral-open` allocates a `WiFiServer` and calls `begin()`. `modbus-peripheral-poll` calls `server->available()` once per iteration to accept new clients. No Free↵
RTOS task is created.

4.3 Build Flag Summary

Flag	Effect
LL_MODBUS 1	Enables the Modbus binding; links libmodbus (Linux only)
LL_MODBUS 0 or undefined	Modbus functions are not available; calls raise "undefined"

6.26.5 5. Quick Reference

Controller Functions

Function	Description
<code>(modbus-tcp-connect host port unit)</code>	Connect to a TCP Modbus peripheral
<code>(modbus-rtu-connect uart baud unit)</code>	Connect to an RTU Modbus peripheral
<code>(modbus-close conn)</code>	Close connection
<code>(modbus-set-timeout conn ms)</code>	Set response timeout in milliseconds
<code>(modbus-request conn fc addr count)</code>	Issue a Modbus request; returns token
<code>(modbus-poll token)</code>	Check for response; returns result or #f
<code>(modbus-request-float32 conn addr order)</code>	Issue FC03 read of 2-register float
<code>(modbus-poll-float32 token)</code>	Poll float32 response; returns inexact or #f
<code>(modbus-rtu-set-de-pin conn gpio)</code>	Set RS-485 direction-enable pin (ESP32)

Controller Function Codes (symbols for `modbus-request`)

Symbol	FC	Direction	Result
<code>read-coils</code>	FC01	Read	Vector of booleans
<code>read-discrete-inputs</code>	FC02	Read	Vector of booleans
<code>read-holding</code>	FC03	Read	Vector of integers
<code>read-input</code>	FC04	Read	Vector of integers
<code>write-coil</code>	FC05	Write	#t on success
<code>write-register</code>	FC06	Write	#t on success
<code>write-coils</code>	FC15	Write	#t on success
<code>write-registers</code>	FC16	Write	#t on success

Peripheral Functions

Function	Description
<code>(modbus-tcp-peripheral-open port unit)</code>	Open TCP peripheral on port
<code>(modbus-rtu-peripheral-open uart baud unit)</code>	Open RTU peripheral
<code>(modbus-peripheral-close peripheral)</code>	Close peripheral
<code>(modbus-peripheral-poll peripheral)</code>	Process pending controller requests
<code>(modbus-peripheral-on-write! peripheral proc)</code>	Register write callback
<code>(modbus-peripheral-hreg! peripheral addr val)</code>	Write holding register from Scheme

Function	Description
<code>(modbus-peripheral-hreg peripheral addr)</code>	Read holding register from Scheme
<code>(modbus-peripheral-ireg! peripheral addr val)</code>	Write input register from Scheme
<code>(modbus-peripheral-ireg peripheral addr)</code>	Read input register from Scheme
<code>(modbus-peripheral-coil! peripheral addr bool)</code>	Write coil from Scheme
<code>(modbus-peripheral-coil peripheral addr)</code>	Read coil from Scheme
<code>(modbus-peripheral-dinput! peripheral addr bool)</code>	Write discrete input from Scheme
<code>(modbus-peripheral-dinput peripheral addr)</code>	Read discrete input from Scheme

6.27 LambLisp PROFIBUS and PROFINET Interface

LambLisp implements the two fieldbuses of the Siemens automation world: **PROFIBUS DP** over RS-485, and **PROFINET** over Ethernet. This chapter covers both, the API for each role, the build flags, and what you need to build locally.

PROFIBUS DP -- LambLisp as Master or as Slave




Figure 6.7 PROFIBUS DP: LambLisp as master or as slave

6.27.1 1. What You Get

Capability	Enable with	Also needs
PROFIBUS DP master (class 1, cyclic)	LL_PROFIBUS=1	an RS-485 transceiver
PROFIBUS DP slave	LL_PROFIBUS=1	as above; baud-capped, see §3
PROFINET IO-Device	LL_PROFINET=1	the pnd daemon, built locally per §5.1


PROFIBUS is self-contained: set the flag, wire up a transceiver, and it works.

PROFINET comes in two pieces. The `profinet-*` bindings are compiled into your LambLisp binary by `LL↔_PROFINET=1`; the protocol stack runs beside it as a separate daemon, `pnd`, which you build from the source in `w3_profinet/`. Section 5.1 has the commands.

6.27.2 2. PROFIBUS DP as a Master

A DP master polls each slave in turn, exchanging a fixed-size process image every cycle. LambLisp performs the whole handshake for you: a slave must be found, parameterised, configured, checked for diagnostics, and only then does cyclic data exchange begin.

How a Master Brings a Slave Into Cyclic Exchange



Each arrow is one request/response. The master must keep the cycle running: if it goes quiet, every slave watchdog fires.

Figure 6.8 The state machine the master drives, per slave

2.1 Opening the Bus

```

(define bus (profibus-open "/dev/ttyUSB0" 500000)) ; POSIX: a tty path
(define bus (profibus-open 1 500000))             ; ESP32: a UART index
(profibus-set-address bus 1)                       ; this station's FDL address
  
```

PROFIBUS uses **8E1** framing — 1 start bit, 8 data bits, **even parity**, 1 stop bit — not the 8N1 that Modbus RTU uses. LambLisp configures this for you; you only supply the baud rate.

Standard rates are 9600, 19200, 45450, 93750, 187500, 500000, and 1.5/3/6/12 Mbit/s. Higher rates mean shorter cable: 1200 m at 9.6 kbit/s, 100 m at 12 Mbit/s.

2.2 Declaring Slaves

Each slave needs its address, its **ident number** (from the vendor's GSD file), and its input/output lengths:

```

(define io (profibus-add-slave bus 3 #x0801 4 2)) ; addr 3, 4 bytes in, 2 out
(profibus-set-prm io (bytevector #x88 #x0A #x0A #x0B #x08 #x01 #x00))
(profibus-set-cfg io (bytevector #x13 #x21))      ; 4 bytes in, 2 bytes out
(profibus-set-watchdog io 1000)                   ; ms; 0 disables
  
```

The **configuration identifier bytes** passed to `profibus-set-cfg` must match the slave's GSD exactly, or it will refuse to leave `wait-cfg`. This is the single most common commissioning failure. The encoding is:

Bit	Meaning
0–3	length – 1
4	inputs present
5	outputs present
6	word (set) or byte (clear)
7	consistency over the whole module

So #x10 is one input byte, #x20 one output byte, #x30 one of each, #x13 four input bytes, #x50 one input word.

2.3 The Service Loop

```
(profibus-start bus)

(let loop ()
  (profibus-service bus)           ; advances the bus; never blocks

  (let drain ()                   ; 1. handle events
    (let ((ev (profibus-poll-event bus)))
      (when ev
        (case (car ev)
          ((slave-up) (log "slave ~a online" (cadr ev)))
          ((slave-down) (log "slave ~a DOWN: ~a" (cadr ev) (caddr ev)))
          ((diag) (log "slave ~a diagnostics ~a" (cadr ev) (caddr ev))))
        (drain))))))

  (let* ((in (profibus-inputs io)) ; 2. read a consistent snapshot
         (level (bytevector-u16-ref in 0)))
    (profibus-outputs-set! io      ; 3. stage the next outputs
      (u16->bytevector (control-law level))))
    (loop))
```

Three properties of this loop are worth understanding:

- `profibus-service` **never blocks**. Call it as often as you like; it does whatever work is possible and returns.
- `profibus-inputs` **returns a consistent snapshot**. The image is double-buffered, so you never see half of one cycle and half of the next.
- `profibus-outputs-set!` **stages, it does not transmit**. The staged image is published at a cycle boundary, so a partially-written image is never sent.

2.4 Events, Not Exceptions

A slave dropping off the bus is **not** an error. It is an event:

```
(slave-up 3)           ; reached cyclic exchange
(slave-down 3 timeout) ; also: prm, cfg, ident, fcs
(diag 3 #u8(...))      ; diagnostics available
(cfg-reject 3 #u8(...)) ; your cfg bytes, so you can see what was refused
```

Errors are raised only for programming mistakes — a bad handle, an undeclared slave, an unsafe baud rate. A control loop should not need guard around every cycle just to survive a station failing.

2.5 Diagnostics and Global Control

```
(profibus-diag io)           ; -> bytevector, or #f
(profibus-global-control bus #x02 0) ; CLEAR: zero all outputs, group 0
(profibus-stats bus)         ; -> (cycles retries timeouts fcs min mean max)
```

`profibus-stats` reports retries and timeouts separately from errors. A bus that silently retries is a bus that is about to fail in the field, so watch these counters rather than waiting for a station to drop.

6.27.3 3. PROFIBUS DP as a Slave

LambLisp can also *be* a DP slave, answering a third-party master:

```
(define self (profibus-slave-open "/dev/ttyUSB0" 19200 3 #x0801 2 1))
(profibus-slave-set-cfg self (bytevector #x11 #x20)) ; what we will accept
(profibus-slave-inputs-set! self (bytevector #xDE #xAD))

(let loop ()
  (profibus-slave-service self)
  (let ((out (profibus-slave-outputs self)))
    (apply-actuators out))
  (loop))
```

The baud ceiling, and why it exists

profibus-slave-open **refuses** a baud rate above 500 kbit/s, raising pb-baud-unsafe.

This is deliberate. A DP slave must begin its reply within T_{SDR} , a minimum of 11 bit times. That is 1.15 ms at 9.6 kbit/s but only **0.92 μ s at 12 Mbit/s** — far below what any software responder can guarantee. Devices that run at 12 Mbit/s use a dedicated ASIC for exactly this reason.

Baud	T_{SDR} minimum	Software slave
9.6 kbit/s	1.15 ms	comfortable
93.75 kbit/s	117 μ s	fine
500 kbit/s	22 μ s	the supported ceiling
1.5 Mbit/s	7.3 μ s	needs hardware turnaround
12 Mbit/s	0.92 μ s	ASIC only

Refusing loudly at open time is better than a segment that works on the bench and fails intermittently in a plant.

6.27.4 4. RS-485 Wiring

The standard violet DB9:

Pin	Signal	Note
3	B (RxD/TxD-P)	data +
8	A (RxD/TxD-N)	data –
5	DGND	terminator reference
6	VP (+5 V)	powers the terminator network
4	RTS	direction, for repeaters
1, 9	shield	to the connector housing

Both ends of a segment need the active terminator (390 Ω to VP, 220 Ω across A/B, 390 Ω to DGND); proper PROFIBUS connectors have it switchable.

A segment with unpowered terminators appears to work at low baud and fails intermittently at high baud. Rule this out before debugging software timing.


On ESP32, LambLisp uses the UART's hardware RS-485 half-duplex mode so the transceiver turnaround is handled in hardware. If your board cannot do this, set a direction pin with `(profibus-set-de-pin bus gpio)` and expect a lower baud ceiling. `(profibus-hardware-turnaround? bus)` tells you which you got.

For a plant segment use an isolated transceiver. For bench work a MAX3485 is fine; do not daisy-chain through a breadboard at 500 kbit/s.

6.27.5 5. PROFINET

PROFINET is Ethernet-based and needs a full protocol stack: DCP for discovery and commissioning, DCE/↔ RPC for connection establishment, and cyclic real-time frames.

PROFINET -- Two Processes, One Boundary



The daemon is a separate program, so its GPLv3 licence does not reach your application.

LambLisp keeps the process image and the policy; the daemon keeps the wire deadline.

The daemon is NOT part of the LambLisp distribution -- see the chapter text.

Figure 6.9 PROFINET runs as two processes

5.1 Building the daemon

The PROFINET stack runs as a separate program, `pnd`, and LambLisp talks to it over a local socket. The daemon holds the cyclic deadline, so a garbage collection in your Scheme code cannot make the device miss a bus cycle.

The daemon is shipped in source code form and needs to be built locally. Its source is in `w3_profinet/`, and the build needs the `rt-labs p-net` and `osal` sources, which you clone into `extern/`:

```

cd w3_profinet
mkdir -p extern && cd extern
git clone --depth 1 https://github.com/rtlabs-com/p-net.git
git clone --depth 1 https://github.com/rtlabs-com/osal.git
cd ..
./build.sh

```

`build.sh` is the entire build — the public `p-net` drop carries no CMake files, so the script compiles the stack, `osal`, and the LambLisp adaptation layer directly. It produces:

Output	What it is
<code>build/pnd</code>	the daemon
<code>build/↔ libprofinet.a</code>	the stack, if you would rather link your own

Output	What it is
build/linktest	calls <code>pnet_init()</code> ; run it to prove the build works

Run the daemon on the interface carrying PROFINET traffic. It needs `CAP_NET_RAW` to open an `AF_PACKET` socket and `CAP_NET_ADMIN` to honour a DCP Set-IP, so give it the capability or start it as root:

```
sudo setcap cap_net_raw,cap_net_admin+ep build/pnd
build/pnd --if pn0 --station lamb-io-1 &
```

`--socket PATH` moves the control socket (default is shared with the `profinet-*` bindings), and `--vendor N` / `--device N` set the identity reported in DCP. Start `pnd` before the LambLisp program that calls `profinet-open`.

5.2 The API

The bindings *are* in your LambLisp binary, and mirror the PROFIBUS shape:

```
(define dev (profinet-device-open "/tmp/lamblisp-profinet.sock" #x0493 #x0002))
(profinet-add-module dev 1 #x00000010)
(profinet-add-submodule dev 1 1 #x00000001 4 4) ; 4 bytes in, 4 out
(profinet-start dev "lamb-io-1") ; our NameOfStation

(let loop ()
  (let drain ()
    (let ((ev (profinet-poll-event dev)))
      (when ev
        (case (car ev)
          ((dcp-set-name) (save-station-name! (cadr ev))) ; commissioning!
          ((dcp-set-ip) (save-ip-config! (cadr ev)))
          ((connect) (log "controller connected"))
          ((param-end) (log "parameterised"))
          ((dcp-signal) (flash-led!))) ; "where are you?"
        (drain))))
    (profinet-inputs-set! dev 1 1 (read-sensors))
    (when (eq? (profinet-iocs dev 1 1) 'good)
      (apply-actuators (profinet-outputs dev 1 1)))
    (loop)))
```

`(dcp-set-name "...")` arriving is the moment an engineering tool commissions your device. Persisting that name is your application's job — which is exactly why it is surfaced to Scheme rather than handled silently.

5.3 Device model

PROFINET describes a device as **slots** holding **modules**, each with **subslots** holding **submodules** that carry the actual I/O. Your declaration must match the **GSDML** file you give the engineering tool, or the controller will refuse the connection at `Connect`. Generate the GSDML from the same definition you use in code; hand-maintaining both is a reliable way to lose an afternoon.

6.27.6 6. Certification and Naming

Both protocols are trademarks of PROFIBUS & PROFINET International (PI), and the right to use the names and marks on a product is tied to certification.

Item	Requirement	Cost
Vendor ID	one per company, from PI	free
Device Ident Number	required for every DP device except a master class 2	registration fee
GSD / GSDML file	one per device, matching the firmware	your effort
Certification	mandatory for PROFINET devices, per Conformance Class, at a PI Test Lab	lab + certificate fees

Two practical notes. A **DP master class 2** (an engineering or diagnostic tool) needs no Device Ident Number, which makes it the least encumbered thing to build first. And using a pre-certified communication module reduces your certification effort substantially but does **not** remove it — PI certifies devices, not chips.

6.27.7 7. Build Flags

Features are compiled in per target. Your package ships a `platformio.ini` already locked to the target you downloaded, so to change the feature set, edit its `build_flags` directly and rebuild.

Flag	Enables
LL_PROFIBUS=1	profibus-* and profibus-slave-*
LL_PROFINET=1	profinet-* (needs the daemon at run time)

Add them to the `build_flags` of your target's environment:

```
build_flags =  
-O3  
-DNDEBUG  
-DLL_PROFIBUS=1  
-DLL_PROFINET=1
```

Verify a feature actually reached your binary:

```
(defined? 'profibus-open)      ; => #t when LL_PROFIBUS was set
```

6.27.8 8. Quick Reference

PROFIBUS DP master

Function	Purpose
<code>(profibus-open port baud)</code>	open the bus; returns a handle
<code>(profibus-close bus)</code>	close it
<code>(profibus-set-address bus addr)</code>	this station's FDL address, 0–125
<code>(profibus-set-timing bus tsl retries)</code>	slot time (ms) and retry count
<code>(profibus-set-de-pin bus gpio)</code>	RS-485 direction pin, if not hardware
<code>(profibus-hardware-turnaround? bus)</code>	#t if the UART drives direction
<code>(profibus-add-slave bus addr ident in out)</code>	declare a slave
<code>(profibus-set-prm slave bv)</code>	parameterisation bytes
<code>(profibus-set-cfg slave bv)</code>	configuration identifiers
<code>(profibus-set-watchdog slave ms)</code>	slave watchdog; 0 disables
<code>(profibus-start bus) / (profibus-stop bus)</code>	begin/end cyclic exchange
<code>(profibus-service bus)</code>	advance the bus; non-blocking
<code>(profibus-state slave)</code>	offline wait-prm wait-cfg check-diag data-exch
<code>(profibus-inputs slave)</code>	consistent input snapshot
<code>(profibus-outputs-set! slave bv)</code>	stage outputs
<code>(profibus-diag slave)</code>	diagnostics bytevector, or #f
<code>(profibus-global-control bus cmd group)</code>	sync/freeze/clear broadcast
<code>(profibus-poll-event bus)</code>	next event, or #f
<code>(profibus-stats bus)</code>	cycles, retries, timeouts, FCS errors, timing

PROFIBUS DP slave

Function	Purpose
<code>(profibus-slave-open port baud addr ident in out)</code>	become a slave
<code>(profibus-slave-set-cfg self bv)</code>	the configuration we accept
<code>(profibus-slave-inputs-set! self bv)</code>	data we provide to the master
<code>(profibus-slave-outputs self)</code>	data received from the master
<code>(profibus-slave-set-diag self bv)</code>	extended diagnostics
<code>(profibus-slave-service self)</code>	answer pending telegrams
<code>(profibus-slave-state self)</code>	current state
<code>(profibus-slave-poll-event self)</code>	next event, or #f

PROFINET IO-Device

Function	Purpose
<code>(profinet-device-open path vendor device)</code>	connect to the daemon
<code>(profinet-add-module dev slot ident)</code>	declare a module
<code>(profinet-add-submodule dev slot sub ident in out)</code>	declare a submodule
<code>(profinet-start dev name)</code>	start with a NameOfStation
<code>(profinet-inputs-set! dev slot sub bv)</code>	provider data to the controller
<code>(profinet-outputs dev slot sub)</code>	consumer data from the controller
<code>(profinet-iops-set! dev slot sub good?)</code>	our provider status
<code>(profinet-iocs dev slot sub)</code>	good or bad, from the controller
<code>(profinet-state dev)</code>	offline dcp-only connected operate
<code>(profinet-connected? dev)</code>	#t when cyclic exchange is running
<code>(profinet-poll-event dev)</code>	next event, or #f
<code>(profinet-alarm dev slot sub type bv)</code>	send a process alarm
<code>(profinet-stats dev)</code>	cycles, missed, cycle counter, jitter, alarms

Error tags

Raised as the first element of the error message, so you can dispatch without string matching.

Tag	Meaning
<code>pb-baud-unsafe</code>	baud above the software responder ceiling
<code>pb-timeout</code>	no reply within the slot time after retries
<code>pb-fcs</code>	checksum or delimiter error
<code>pb-state</code>	operation invalid in this state
<code>pb-cfg-reject</code>	slave refused the configuration
<code>pb-ident</code>	slave's ident number does not match
<code>pn-transport</code>	lost the PROFINET daemon
<code>pn-not-connected</code>	I/O access before a controller connected
<code>pn-submodule</code>	slot/subslot not declared

6.28 Bundled Demos

Alongside the reusable [convenience layers](#), LambLisp ships a handful of complete **example applications** under `scm/demos/`. Where the libraries give you an API, the demos show that API assembled into something that runs end-to-end: a supervisory-control server for the Freenove 4WD car, a capability-limited eval sandbox, an occupancy-grid telemetry mapper, and two tiny character/data sketches. Each is selected into a target's filesystem image by its manifest (`scm/targets/<env>.manifest`); all five demos ship on the Linux hosts (`linux_x86_64`, `linux_x86_64_profile`, `linux_x86_64_cuda`, `linux_x86_64_hip`, `linux_aarch64`) and none of the ESP32 targets — they are host-side orchestration and teaching code, not on-board firmware.

The supervisor, sandbox, and telemetry demos are three tiers of one system and are used together (they underpin the [Worked Examples](#)); the two sketches stand alone.

6.28.1 Supervisory Control (demo-supervisor.scm)

The demo supervisor is a three-tier orchestrator: a browser talks to a Python web bridge (`w3_ai_scripts/↔demo-web.py`), which connects over LLIP to this supervisor, which in turn drives the Freenove 4WD car over a second LLIP link. It runs a non-blocking main loop that accepts bridge clients, polls telemetry, sweeps the sonar into the map, dispatches sandboxed viewer expressions to the car, and broadcasts results and map/telemetry updates to every client. Its load order and the required companion files (`llip-client.scm`, `llip-server.scm`, `demo-sandbox.scm`, `demo-telemetry.scm`) are documented in the file header.

`demo-start` is the single entry point; it opens the car link and bridge server (both configured by the `demo-car-*` / `demo-bridge-*` defines at the top of the file) and enters the loop. The **car wrappers** are thin conveniences over LLIP eval that you can also call directly from a REPL attached to the supervisor:

Procedure	Description
<code>demo-start</code>	Entry point: connect to the car, start the bridge LLIP server, run the main loop.
<code>car-eval</code>	<code>(car-eval expr)</code> — evaluate an arbitrary form on the car over LLIP; the primitive the others build on.
<code>car-move!</code>	<code>(car-move! dir ms)</code> — drive in direction <code>dir</code> for <code>ms</code> milliseconds.
<code>car-turn!</code>	<code>(car-turn! deg)</code> — turn in place by <code>deg</code> degrees.
<code>car-sweep</code>	<code>(car-sweep)</code> — run a sonar sweep; returns <code>(angle . dist)</code> readings for the map.
<code>car-set-leds!</code>	<code>(car-set-leds! pat)</code> — set the LED strip to pattern <code>pat</code> .

Also defined: `car-stop`, `car-leds-off`.

```
(demo-start)           ; connect, serve the bridge, and run the loop
;; ...or drive the car directly over the supervisor's LLIP link:
(car-move! 'forward 500) ; forward for half a second
(car-turn! 90)           ; quarter turn in place
(car-sweep)             ; => ((angle . dist) ...)
```

6.28.2 Eval Sandbox (demo-sandbox.scm)

The sandbox is what makes it safe to let an anonymous web viewer type Scheme that runs on your robot. Every submitted expression is parsed and checked before it is forwarded to the car, and each client is rate-limited. The safety model has three layers:

- **Whitelist.** Only calls whose head is a symbol in `sandbox-whitelist` (a fixed set of robot verbs and sensor getters — `robot-move!`, `get-sonar`, `set-leds!`, ...) are permitted, and the check recurses into nested argument calls, so every call anywhere in the form must be whitelisted. Bare atoms are allowed as data.
- **Blacklist.** Heads in `sandbox-forbidden-heads` (`define`, `set!`, `load`, `eval`, `apply`, `file/TCP open`, `system`, `shell`, ...) are rejected outright — this closes the obvious escapes before the whitelist is even consulted.
- **Rate limit.** Each `client-id` may submit at most `sandbox-rate-limit` (default 3) expressions per `sandbox-rate-window-ms` (default 2000 ms) window, tracked in an alist.

Procedure	Description
<code>sandbox-eval</code>	<code>(sandbox-eval expr-string client-id car-conn now-ms)</code> — rate-check, parse, safety-check, then eval on the car. Returns <code>(ok result-string)</code> or <code>(error msg-string)</code> .
<code>sandbox-safe?</code>	<code>(sandbox-safe? expr)</code> — #t if a parsed form passes the whitelist/blacklist rules; useful on its own.

```
(sandbox-safe? '(robot-move! 'forward 500)) ; => #t
(sandbox-safe? '(load "evil.scm"))          ; => #f (forbidden head)
(sandbox-safe? '(car (robot-pose)))         ; => #f (car not whitelisted)
```

6.28.3 Telemetry & Floor Mapping (`demo-telemetry.scm`)

This layer turns raw sensor polls into a live occupancy-grid map. Telemetry is kept as a flat alist (`sonar`, `battery`, `light`, `heading`, `pose-x/y`, `loop-ms`, `ncells`); the map is a 100×100 flat vector of cells at 10 cm resolution, each 0 unknown, 1 free, or 2 occupied. Sonar sweeps are ray-cast from the car's current pose to mark free space and wall hits, and the grid can be rendered as ASCII for a web view.

Procedure	Description
<code>telem-poll!</code>	<code>(telem-poll! car-conn now-ms)</code> — one LLIP round-trip that reads all sensors and updates the telemetry alist. Errors are swallowed (the car may be momentarily busy).
<code>grid-update-sweep!</code>	<code>(grid-update-sweep! sweep-result)</code> — fold a sweep (<code>(angle . dist)</code> pairs) into the occupancy grid from the current pose/heading.
<code>grid->ascii</code>	<code>(grid->ascii)</code> — render the grid as an ASCII string (<code>? unknown</code> , <code>. free</code> , <code># wall</code>).
<code>telem->json</code>	<code>(telem->json)</code> — render current telemetry as a JSON-like string for the web bridge.

```
(telem-poll! conn (elapsed-ms)) ; refresh sensor readings
(grid-update-sweep! (car-sweep)) ; fold a sonar sweep into the map
(display (grid->ascii))          ; print the map
```

6.28.4 Musical Notes (`Music.scm`)

A minimal note-table sketch for buzzer tunes: it defines the concert-pitch reference `a4` (440, the frequency of A4) and `notes`, the twelve chromatic note names. It is a starting point for buzzer melodies rather than a finished player.

Definition	Description
<code>a4</code>	440 — reference frequency of A in octave 4.
<code>notes</code>	List of the twelve chromatic note-name symbols.

6.28.5 Personality (`Personality.scm`)

A one-line character sketch: `Personality` names the device's identity (here `'Freenove-4WD-Car`). It is the smallest possible example of a per-target behavior/identity file that other demos or applications can key off of.

Definition	Description
Personality	The device's identity symbol, e.g. 'Freenove-4WD-Car.

Chapter 7

Worked Examples

7.1 Worked Example: Autonomous Floor Mapping with the Freenove 4WD Car

This chapter walks through a complete three-tier LambLisp application that uses the Freenove 4WD Car to map the free space in a room using its HC-SR04 ultrasonic sonar. The example is self-contained and intended to be run as-is; every code snippet is a fragment of a working file in `data/`.

The three tiers are:

Tier	Host	Role
Car	ESP32 (on the 4WD car)	Real-time sensor/actuator loop; autonomous gap-filling; map accumulation
Supervisor	Jetson Orin Nano	Mission orchestration; waypoint injection; map aggregation and display
Strategy	Cloud LLIP server (Linux)	Coverage-path planning; global map storage; mission replay

The tiers communicate exclusively over LLIP. The car runs as an LLIP server and never initiates a connection; the Jetson connects to the car and to the cloud server.




Figure 7.1 Three-Tier Architecture

7.1.1 1. Overview

The **floor-mapping mission** proceeds in three phases:

1. **Plan** — the cloud server uses the room outline (provided by the user) to compute a boustrophedon coverage path: a set of parallel strip waypoints that guarantee complete coverage.
2. **Execute** — the Jetson feeds waypoints to the car one at a time via LLIP. Between waypoints the car drives autonomously, taking a sonar reading at each step and storing (x y dist-cm heading-deg) tuples in an on-board map list. If the sonar detects an unexpected obstacle the car halts and waits for the Jetson to issue a new waypoint.
3. **Collect** — the Jetson periodically pulls the accumulated map list from the car and relays it to the cloud server, which merges it into a persistent grid and can serve it back to any connected display client.

A fourth, continuous thread runs on the Jetson: the **digital twin**, which mirrors live car state (speed, heading, battery, sonar) so that the supervisor can make decisions without polling the car in the middle of a move.

7.1.2 2. Hardware Setup

The only hardware modification required beyond the stock Freenove 4WD Car Kit is ensuring that the HC-SR04 sonar module is mounted facing forward on the pan servo bracket. The servo can be commanded to sweep left/right to widen the field of view; for the floor-mapping example the servo is fixed at 90° (forward-facing).

Room coordinate system:

```

      +Y (north)
      |
-----+----- +X (east)
      |
  car starts at (0,0)
  facing +X (east = 0 deg)

```

All distances and positions are in centimetres. Headings are in degrees, measured counter-clockwise from east, matching standard mathematical convention.

Required LambLisp modules (loaded by `setup.scm` or `(load ...)` on demand):

File	Purpose
<code>llip-server.scm</code>	LLIP server (already loaded on car)
<code>robot.scm</code>	Motor helpers: <code>robot-move!</code> , <code>robot-turn!</code> , <code>robot-stop</code>
<code>sonar.scm</code>	HC-SR04 wrapper: <code>sonar-read-cm</code>
<code>floor-map.scm</code>	Map accumulator (new, defined below)
<code>llip-client.scm</code>	LLIP client (Jetson and cloud side)
<code>coverage.scm</code>	Boustrophedon planner (cloud side)

7.1.3 3. On-Board Car Code

3.1 Map Accumulator (floor-map.scm)

The map is a plain Scheme list of readings. Each reading is a four-element list (x y dist-cm heading-deg). The car's current pose (x y heading) is maintained as mutable top-level variables so the Jetson can query them at any time via LLIP.

```
;;; floor-map.scm -- on-board map accumulator for the 4WD car
;;; Copyright 2026 by Frobenius Norm LLC

;;; Pose state -- updated by move! and turn! wrappers below.
(define car-x 0.0) ; centimetres from origin
(define car-y 0.0)
(define car-heading 0.0) ; degrees CCW from east

;;; Accumulated readings: list of (x y dist-cm heading-deg).
(define floor-map '())

;;; Record one sonar reading at the current pose.
(define (map-record-reading!)
  (let ((d (sonar-read-cm)))
    (set! floor-map
      (cons (list car-x car-y d car-heading) floor-map))
    d))

;;; Clear the map and reset pose.
(define (map-reset!)
  (set! floor-map '())
  (set! car-x 0.0)
  (set! car-y 0.0)
  (set! car-heading 0.0))

;;; Return the map as a list (for LLIP retrieval).
(define (map-get) floor-map)

;;; Update pose after a forward move of dist-cm.
(define (pose-advance! dist-cm)
  (let ((rad (* car-heading (/ (* 4 (atan 1)) 180))))
    (set! car-x (+ car-x (* dist-cm (cos rad))))
    (set! car-y (+ car-y (* dist-cm (sin rad)))))

  )

;;; Update pose after a turn of delta-deg (CCW positive).
(define (pose-turn! delta-deg)
  (set! car-heading (modulo (+ car-heading delta-deg) 360)))
```

3.2 Motion Primitives

The motion wrappers combine hardware commands with pose tracking and sonar sampling. A sonar reading is taken before every forward move; if an obstacle is closer than map-obstacle-cm the move is cancelled and 'obstacle is returned.

```
;;; map-obstacle-cm -- tunable threshold; default 20 cm.
(define map-obstacle-cm 20)

;;; Move forward step-cm, recording a sonar reading first.
;;; Returns 'moved or 'obstacle.
(define (map-step-forward! step-cm)
  (let ((d (map-record-reading!)))
    (if (< d map-obstacle-cm)
        'obstacle
        (begin
          (robot-move! 'forward (round (* step-cm 10))) ; ms ~= cm * 10
          (pose-advance! step-cm)
          'moved))))

;;; Turn in place, updating heading.
(define (map-turn! delta-deg)
  (robot-turn! delta-deg)
  (pose-turn! delta-deg))
```


3.3 Autonomous Fill Behaviour

When no waypoint is pending the car enters **autonomous fill** mode: it continues taking sonar readings and nudging forward in small steps, building map density around its current position. If it hits an obstacle it rotates in 10° increments until it finds a clear path, or waits if no path exists within a 360° sweep.

This behaviour runs inside the car's `loop` procedure and yields to the Jetson whenever a new waypoint arrives.

```
;;; Autonomous fill step -- called from loop when no waypoint is pending.
;;; Returns 'ok, 'obstacle-all-around, or 'at-waypoint.
(define (map-fill-step!)
  (let ((result (map-step-forward! 5))) ; 5 cm per fill step
    (if (eq? result 'moved)
        'ok
        ;; Try rotating to find a clear direction.
        (let search ((tries 36))
          (if (= tries 0)
              'obstacle-all-around
              (begin
                 (map-turn! 10)
                 (if (eq? (map-step-forward! 5) 'moved)
                     'ok
                     (search (- tries 1))))))))))

;;; Waypoint queue -- list of (x y) targets, set by the Jetson via LLIP.
(define waypoint-queue '())

(define (waypoint-pending?) (not (null? waypoint-queue)))

(define (waypoint-next!)
  (let ((wp (car waypoint-queue)))
    (set! waypoint-queue (cdr waypoint-queue))
    wp))

;;; Add a waypoint (called remotely by the Jetson).
(define (waypoint-push! x y)
  (set! waypoint-queue
        (append waypoint-queue (list (list x y)))))
```

3.4 Waypoint Navigation

Navigating to a waypoint uses a simple proportional heading controller: compute the bearing to the target, turn to face it, then drive forward in steps until the car is within `map-arrival-cm` of the target.

```
(define map-arrival-cm 10) ; close-enough radius

(define (bearing-to x y)
  (let* ((dx (- x car-x))
        (dy (- y car-y))
        (b (* (atan dy dx) (/ 180 (* 4 (atan 1))))))
    (modulo b 360)))

(define (distance-to x y)
  (let ((dx (- x car-x))
        (dy (- y car-y)))
    (sqrt (+ (* dx dx) (* dy dy)))))

;;; Navigate to (tx ty). Returns 'arrived or 'blocked.
(define (map-navigate-to! tx ty)
  (let loop ()
    (let ((d (distance-to tx ty)))
      (if (< d map-arrival-cm)
          'arrived
          (let* ((bearing (bearing-to tx ty))
                 (delta (let ((raw (- bearing car-heading)))
                          (cond ((> raw 180) (- raw 360))
                                ((< raw -180) (+ raw 360))
                                (else raw))))
                 (when (> (abs delta) 5)
                     (map-turn! delta))))
            loop))))
```

```

      (map-turn! delta))
    (let ((result (map-step-forward! 10)))
      (if (eq? result 'obstacle)
          'blocked
          (loop)))))))))

;;; Main loop procedure -- replaces the application's (loop).
(define (loop)
  (if (waypoint-pending?)
      (let* ((wp (waypoint-next!))
             (tx (car wp))
             (ty (cadr wp))
             (result (map-navigate-to! tx ty)))
        (when (eq? result 'blocked)
          (set! waypoint-queue (cons wp waypoint-queue))) ; re-queue
        (map-fill-step!)))
      (loop)))

```

3.5 Bee Strategy Module

The **bee strategy** is a reactive alternative to the pre-planned boustrophedon path. The car drives forward until it hits an obstacle, then asks the cloud strategy server for the next heading. It records every unblocked cell it traverses — a separate list from the sonar-obstacle map — so the cloud can choose the next heading that biases toward unexplored territory.

```

;;; bee.scm -- reactive coverage for the 4WD car
;;; Copyright 2026 by Frobenius Norm LLC

;;; Traversed (unblocked) cells -- separate from the sonar obstacle map.
(define bee-path '())

(define (bee-path-record!)
  (set! bee-path (cons (list car-x car-y) bee-path)))

(define (bee-path-get) bee-path)

(define (bee-path-reset!)
  (set! bee-path '()))

;;; One forward step in bee mode.
;;; Returns 'moved or 'blocked.
(define (bee-forward! step-cm)
  (let ((d (map-record-reading!)))
    (if (< d map-obstacle-cm)
        'blocked
        (begin
         (robot-move! 'forward (round (* step-cm 10)))
         (pose-advance! step-cm)
         (bee-path-record!)
         'moved))))))

;;; Run the bee strategy for n-steps total forward moves.
;;; After each obstacle the car asks the cloud for a new heading.
(define (bee-run! cloud n-steps step-cm)
  (let loop ((remaining n-steps))
    (when (> remaining 0)
      (when (eq? (bee-forward! step-cm) 'blocked)
        (let* ((new-h (llip-eval cloud
                                `(bee-next-heading ,car-heading ',bee-path)))
               (delta (let ((raw (- new-h car-heading)))
                        (cond ((> raw 180) (- raw 360))
                              ((< raw -180) (+ raw 360))
                              (else raw))))))
          (map-turn! delta)))
      (loop (- remaining 1)))))

```

7.1.4 4. Jetson Supervisor

The Jetson runs `data/floor-map-supervisor.scm`. It connects to both the car and the cloud server, requests a coverage plan, and feeds it to the car as a waypoint stream. It also runs the digital twin sync loop in parallel using LambLisp's **timed-queue** deferred task facility.

4.1 Connections

```
;;; floor-map-supervisor.scm -- Jetson side
;;; Copyright 2026 by Frobenius Norm LLC

(load "llip-client.scm" 0)

;;; Connect to the car and the cloud strategy server.
(define car (llip-connect "192.168.1.100" 5555))
(define cloud (llip-connect "llip.example.com" 5555))

;;; Upload the on-board modules and reset the map.
(llip-upload car "floor-map.scm" "/floor-map.scm")
(llip-eval car '(load "/floor-map.scm"))
(llip-eval car '(map-reset!))
(display "Car ready\n")
```

4.2 Requesting a Coverage Plan

The cloud server runs a boustrophedon planner that accepts the room outline as a list of (x y) corner coordinates and returns a list of (x y) waypoints.

```
;;; Describe the room outline (clockwise, centimetres from origin).
(define room-outline
  '((0 0) (400 0) (400 300) (0 300)))

;;; Ask the cloud for a complete coverage plan.
;;; Returns a list of (x y) waypoints.
(define coverage-plan
  (llip-eval cloud `(boustrophedon-plan ',room-outline 40)))
  ;; 40 cm strip width

(display (string-append "Plan: " (number->string (length coverage-plan))
  " waypoints\n"))
```

4.3 Feeding Waypoints to the Car

Waypoints are pushed to the car's queue via `waypoint-push!`. The Jetson sends them in batches: a new batch is sent whenever the queue drains below a low-water mark, allowing the car's autonomous fill behaviour to run between waypoints while the Jetson is computing the next batch.

```
(define wp-low-water 3) ;; refill when fewer than this many waypoints remain

(define (queue-length!)
  (llip-eval car '(length waypoint-queue)))

;;; Push one waypoint to the car.
(define (push-waypoint! wp)
  (llip-eval car `(waypoint-push! ,(car wp) ,(cadr wp))))

;;; Feed the full plan, refilling the queue as it drains.
(define (run-mission! plan)
  (let loop ((remaining plan))
    (unless (null? remaining)
      (let ((qlen (queue-length!)))
        (when (< qlen wp-low-water)
          (push-waypoint! (car remaining))
          (loop (cdr remaining)))))))
```

```
(let refill ((wps remaining) (sent 0))
  (if (or (null? wps) (= sent 5)) ; send up to 5 at a time
      (set! remaining wps)
      (begin
        (push-waypoint! (car wps))
        (refill (cdr wps) (+ sent 1))))))
(loop remaining))))

(run-mission! coverage-plan)
```

4.4 Digital Twin Sync

A timed task runs every 500 ms on the Jetson, pulling live telemetry from the car and updating a local twin dict. The twin is available for query at any time without blocking the mission loop.

```
;;; Local twin state -- updated by sync-twin! every 500 ms.
(define twin
  (list->dict '((x . 0.0) (y . 0.0) (heading . 0.0)
               (sonar . 0.0) (battery . 100.0) (loop-ms . 0))))

(define (sync-twin!)
  (dict-rebind-bang twin 'x      (llip-eval car 'car-x))
  (dict-rebind-bang twin 'y      (llip-eval car 'car-y))
  (dict-rebind-bang twin 'heading (llip-eval car 'car-heading))
  (dict-rebind-bang twin 'sonar   (llip-eval car '(sonar-read-cm)))
  (dict-rebind-bang twin 'battery (llip-eval car '(get-battery))))

;;; Schedule twin sync every 500 ms using the timed queue.
(tq-repeat! 500 sync-twin!)
```

4.5 Collecting the Map

After the mission completes the Jetson retrieves the full map from the car and relays it to the cloud for storage and visualisation.

```
;;; Pull the completed map from the car.
(define raw-map (llip-eval car '(map-get)))

(display (string-append "Map: " (number->string (length raw-map))
                        " readings\n"))

;;; Send the map to the cloud for storage and rendering.
(llip-eval cloud `(map-store! 'floor-run-1 ',raw-map))

;;; Request an ASCII render for quick inspection.
(display (llip-eval cloud '(map-render-ascii 'floor-run-1 80 24)))
```

7.1.5 5. Cloud Strategy Server

The cloud server is a Linux process running a LambLisp LLIP server. It provides the boustrophedon planner and the map store. Any number of Jetson supervisors can connect simultaneously; maps from different runs accumulate in a global dict keyed by run name.

5.1 Boustrophedon Coverage Planner




Figure 7.2 Boustrophedon Coverage Path

A boustrophedon path covers a rectangular area in parallel horizontal strips, reversing direction at the end of each strip. The planner takes the room outline and strip width and returns a list of (x y) waypoints.

For this example the room outline is assumed to be axis-aligned; a full polygon planner would decompose non-convex rooms into monotone cells first.

```
;;; coverage.scm -- boustrophedon coverage planner (cloud side)
;;; Copyright 2026 by Frobenius Norm LLC

;;; Compute axis-aligned bounding box of a polygon.
(define (bbox corners)
  (let ((xs (map car corners))
        (ys (map cadr corners)))
    (list (apply min xs) (apply min ys)
          (apply max xs) (apply max ys))))

;;; Generate boustrophedon waypoints for the bounding box.
;;; strip-cm: strip width in centimetres.
;;; Returns a flat list of (x y) waypoints.
(define (boustrophedon-plan corners strip-cm)
```

```
(let* ((box (bbox corners))
      (x0 (+ (car box) (/ strip-cm 2)))
      (y0 (+ (cadr box) (/ strip-cm 2)))
      (x1 (- (caddr box) (/ strip-cm 2)))
      (y1 (- (caddr box) (/ strip-cm 2)))
      (wps '()))
  (let strip-loop ((y y0) (left->right #t))
    (when (<= y y1)
      (let row-loop ((x (if left->right x0 x1)))
        (when (if left->right (<= x x1) (>= x x0))
          (set! wps (cons (list x y) wps))
          (row-loop (if left->right (+ x strip-cm) (- x strip-cm)))))
      (strip-loop (+ y strip-cm) (not left->right))))
  (reverse wps)))
```

5.2 Map Store and ASCII Render

Maps are stored as association lists keyed by run name. The ASCII renderer projects the (x y dist-cm) readings onto a character grid, marking obstacle hits with # and free space with ..

```
;;; Global map store: dict of run-name -> readings list.
(define map-store (list->dict '()))

(define (map-store! run-name readings)
  (dict-bind-bang map-store (string->symbol (symbol->string run-name))
    readings map-store))

(define (map-get-run run-name)
  (dict-ref map-store (string->symbol (symbol->string run-name))))

;;; ASCII render: cols x rows character grid.
(define (map-render-ascii run-name cols rows)
  (let* ((readings (map-get-run run-name))
        (xs (map car readings))
        (ys (map cadr readings))
        (x0 (apply min xs)) (x1 (apply max xs))
        (y0 (apply min ys)) (y1 (apply max ys))
        (xrng (max 1 (- x1 x0)))
        (yrng (max 1 (- y1 y0)))
        (grid (make-vector (* cols rows) #\.)))
    (for-each
      (lambda (r)
        (let* ((rx (car r)) (ry (cadr r)) (rd (caddr r))
              (col (inexact->exact (round (* (/ (- rx x0) xrng) (- cols 1)))))
              (row (inexact->exact (round (* (/ (- ry y0) yrng) (- rows 1)))))
              (idx (+ (* (- rows 1 row) cols) col)))
          (vector-set! grid idx (if (< rd 20) #\# #\.)
            readings))
        (let build ((i 0) (out ""))
          (if (= i (* cols rows))
            out
            (build (+ i 1)
              (string-append out
                (string (vector-ref grid i))
                (if (= (modulo (+ i 1) cols) 0) "\n" "")))))))
```



```
        cone-deg))
      path)))

;;; Select the next heading.
;;; Tries golden-angle increments 1..12; returns the first candidate
;;; whose 30-degree forward cone contains fewer than 4 visited cells.
;;; Falls back to the 1x golden-angle rotation if all cones are dense.
(define (bee-next-heading current-heading path)
  (let try ((n 1))
    (let ((candidate (modulo (+ current-heading (* n bee-golden-angle)) 360)))
      (if (or (> n 12)
              (< (cells-in-cone candidate 30 path) 4))
          candidate
          (try (+ n 1))))))
```

The 12-step limit means the algorithm tries a full 1650 degrees of rotation (12 x 137.5) before giving up, which reliably finds an unvisited sector in any room that is not completely surrounded by walls.

7.1.6 6. Running the Example

6.1 Start the Cloud Server

On any Linux host with LambLisp built for `linux_x86_64`:

```
# Upload the cloud-side files and start the LLIP server.
pio run -e linux_x86_64 -t exec -- \
  --load coverage.scm --load bee-strategy.scm \
  --load llip-server.scm --server 5555
```

The server accepts connections on port 5555 and exposes `boustrophedon-plan`, `bee-next-heading`, `map-store!`, and `map-render-ascii` to any LLIP client. Both planners are loaded simultaneously; the supervisor chooses which to use.

6.2 Flash and Start the Car

Flash the ESP32 with the standard Freenove 4WD build environment:

```
./w3 make any Freenove
./w3 push Freenove
```

On boot, `setup.scm` loads `robot.scm`, `sonar.scm`, and `llip-server.scm` and starts the control loop. The car is now reachable at its WiFi IP address on port 5555.

6.3 Run the Supervisor on the Jetson

```
ssh jetson
cd ~/Lamb-v01
pio run -e linux_aarch64 -t exec -- --load floor-map-supervisor.scm
```

The supervisor prints progress as it connects, retrieves the plan, and feeds waypoints to the car:

[illegible]

6.4 Tuning Parameters

All key parameters are top-level Scheme variables and can be adjusted live from the Jetson without reflashing the car:

```
;;; From the Jetson REPL or any LLIP client:

;;; Widen the obstacle clearance zone.
(llip-eval car '(set! map-obstacle-cm 30))

;;; Slow down the coverage drive speed (ms per cm).
(llip-eval car '(set! map-step-ms-per-cm 15))

;;; Increase sonar averaging.
(llip-eval car '(set! sonar-samples 5))

;;; Inject a manual waypoint to investigate an area of interest.
(llip-eval car '(waypoint-push! 200 150))

;;; Switch to bee strategy mid-mission (no reflash needed).
;;; Upload bee.scm, reset bee state, then run 200 steps at 5 cm each.
(llip-upload car "bee.scm" "/bee.scm")
(llip-eval car '(load "/bee.scm"))
(llip-eval car '(bee-path-reset!))
(llip-eval car `(bee-run! ,cloud 200 5))

;;; Tune the golden-angle cone sensitivity on the cloud side.
(llip-eval cloud '(set! bee-golden-angle 120.0)) ; denser spiral
```

7.1.7 7. Architecture Notes

7.1 Why Three Tiers?

The three-tier split follows naturally from where compute, latency, and connectivity requirements differ:

Concern	Car (ESP32)	Jetson	Cloud
Real-time loop	yes (hard deadline)	no	no
Sensor/actuator I/O	yes	no	no
Mission orchestration	no	yes	no
Map aggregation	partial	yes	yes
Path planning	no	capable	yes (preferred)
Persistent storage	no	local	yes (global)
Multi-fleet visibility	no	no	yes

The car's 240 MHz Xtensa LX6 core is powerful enough to run the full LambLisp stack and real-time sensor loops simultaneously, but its SPIFFS filesystem is not suited to persistent large-map storage. The Jetson bridges the gap: it has a full Linux filesystem and can buffer map data indefinitely even if the cloud link is intermittent.

7.2 LLIP as the Universal Glue

Every inter-tier call in this example is a plain `llip-eval`. There is no custom protocol, no serialisation library, and no schema. A list is a list whether it lives on the ESP32, the Jetson, or the cloud server. Updating the map format — adding a timestamp or a confidence value to each reading — requires changing one line in `floor-map.scm` and nothing else.

7.3 Extending the Example

Natural extensions to this worked example include:

- **Multi-robot coordination** — two or more cars upload maps to the same cloud run; the planner partitions the coverage path between them and prevents redundant work.
- **Obstacle classification** — replace the binary obstacle/free reading with a machine-learning classifier running on the Jetson that labels obstacle types (wall, furniture, person) based on sonar return profiles.
- **Replanning on the fly** — if the car reports `'obstacle-all-around`, the Jetson requests a new partial plan from the cloud starting at the car's current pose, circumventing the blocked area.
- **Persistent map merging** — subsequent runs load the previous cloud map and bias the coverage planner toward unexplored or low-density areas, building up a progressively more detailed map over multiple sessions.

See also: [LLIP — LambLisp Interaction Protocol](#) · [Digital Twin](#) · [Freenove 4WD Car Kit — Peripheral Summary](#) · [Ultrasonic Sonar \(HC-SR04\)](#)

Chapter 8

LambLisp Compatibility Matrix

8.1 Legend

Symbol	Meaning
	Fully supported
	Partial – implemented but incomplete, or stubbed to return #<undef>
	Not supported
	Not defined by this specification
	LambLisp extension (not in any Scheme standard)

8.2 Coverage Summary

One row per section.  and  count as covered;  as partial;  as missing.

Section	R5RS defined	R7RS adds	LambLisp covers	Gaps
Language (syntax, binding, macros)	core	extensions incl. case-lambda	 most, incl. case-lambda	make-parameter, cond-expand, import/library

Section	R5RS defined	R7RS adds	LambLisp covers	Gaps
Type predicates	full	null? bytevector?	 all	–
Numbers	integer, real, rational, complex	exact-integer? finite? infinite? nan?	 integer, real, rational, complex (native types), all predicates	–
Pairs & Lists	full	make-list list-copy	 full + extras	–
Characters	full	char-ci digit-value char-foldcase	 all	–
Strings	full	string-ci string-copy! string-map string-foldcase	 all	–
Vectors	full	vector->string vector-append vector-map	 full + extras	–
Bytevectors	–	full	 full	–
Control features	apply map for-each values call/cc dynamic-wind	string/vector map/for-each	 all except call/cc dynamic-wind	call/cc dynamic-wind
Environments & eval	full	–	 full	–
Ports	partial	full	 full	–
I/O	partial	full	 most	write-shared/write-simple lack shared/↔ circular structure detection
System interface	load	file-exists? command-line exit	 load, time functions, file ops	command-line exit
Exception handling	–	full	 full	raise-continuable behaves as raise; full continuable semantics require call/cc, which is not supported
LambLisp extensions	–	–		OTA updates

8.3 Detailed Compatibility

8.3.1 Core Language

Feature	R5RS	R7RS	LambLisp	Notes
Procedures				
define lambda	Y	Y	Y	
nlambda	--	--	ext	LambLisp/InterLisp extension
macro	--	--	ext	LambLisp extension
Conditionals				
if cond case and or not else	Y	Y	Y	
when unless	--	Y	Y	
cond-expand	--	Y	N	
case-lambda	--	Y	Y	Generates dispatch lambda at expansion time; rest-arg formals supported
Binding forms				
set! let let* letrec letrec* =>	Y	Y	Y	
Named let	Y	Y	Y	
do	Y	Y	Y	
Syntax definition				
define-syntax	Y	Y	Y	Evaluates transformer spec; binds in current scope only
let-syntax letrec-syntax	Y	Y	Y	let-syntax evals transformers in outer env; letrec-syntax in new env
syntax-rules	Y	Y	Y	All tests pass
syntax-error	--	Y	Y	Raises as a runtime error; syntax context not tracked

Feature	R5RS	R7RS	LambLisp	Notes
Multiple values				
values call-with-values	Y	Y	Y	
let-values let*-values define-values	--	Y	Y	
Quotation				
quote quasiquote unquote unquote-splicing	Y	Y	Y	
Reader abbreviations ' ` , ,@	Y	Y	Y	
Sequencing				
begin	Y	Y	Y	
Delayed evaluation				
delay force	Y	Y	Y	Implemented in Scheme
delay-force promise? make-promise	--	Y	Y	Implemented in Scheme
Dynamic binding				
make-parameter parameterize	--	Y	N	
File inclusion				
include	--	Y	Y	Via load
include-ci	--	Y	N	
Libraries / modules				
define-library import export	--	Y	N	
Records				
define-record-type	--	Y	Y	Implemented in Scheme; dict-backed records, O(1) field access
Continuations				
call-with-current-continuation	Y	Y	N	By design; C++ contiguous stack makes capture impractical
dynamic-wind	Y	Y	N	Requires call/cc

Feature	R5RS	R7RS	LambLisp	Notes
Type predicates				
boolean? char? number? symbol?				
pair? vector? procedure? string?				
port? eof-object?				
null?				
bytevector?				
atom?				
Equivalence				
eq? eqv? equal?				
Datum labels				
#0= #0#				Partially supported by reader

8.3.3 The Numeric Tower

Scheme defines a numeric tower with four levels, each a superset of the level below:

`integer -> rational -> real -> complex` (each a subset of the next)

In LambLisp all four levels are **native VM cell types**:

Type	Cell	Storage	Exactness
Integer	T_INT	Int_t (32-bit on ESP32, 64-bit on x86-64/AArch64) in car	always exact
Real	T_REAL	float (32-bit) in car — all platforms	always inexact
Rational	T_RATIONAL	(numerator . denominator) as Int_t pair, GCD-reduced	always exact
Complex	T_COMPLEX	(real . imaginary) as float pair	always inexact

The standard predicates `integer?`, `rational?`, `real?`, `complex?`, `exact?`, and `inexact?` all dispatch on these native types. Coercion follows the tower upward automatically: integer arithmetic stays exact; mixing exact and inexact promotes to inexact; `exact->inexact` and `inexact->exact` cross the boundary explicitly.

Higher-level arithmetic for rational and complex numbers is provided in the loadable libraries:

- `rxrs_rational.scm` — rationalize, arithmetic operators extended to handle `T_RATIONAL` values
- `rxrs_complex.scm` — complex arithmetic, make-rectangular, make-polar, magnitude, angle

For embedded applications where rational or complex numbers are not required, these libraries need not be loaded.

8.3.4 Numbers

Feature	R5RS	R7RS	LambLisp	Notes
Types				
Integer	Y	Y	Y	
Real (inexact float)	Y	Y	Y	
Rational (exact fractions)	Y	Y	Y	Native <code>T_RATIONAL</code> cell type; GCD-reduced (numerator . denominator) pair. Extended arithmetic in loadable <code>rxrs_rational.scm</code> (<code>rationalize</code> , <code>gcd</code> , <code>lcm</code>).
Complex	Y	Y	Y	Native <code>T_COMPLEX</code> cell type; single-precision (real . imaginary) float pair. Extended arithmetic in loadable <code>rxrs_complex.scm</code> ; rectangular and polar construction, magnitude, angle, <code>sqrt</code> for negative reals.
Exactness				
<code>exact? inexact?</code>	Y	Y	Y	
<code>exact->inexact inexact->exact</code>	Y	--	Y	R5RS names
<code>exact inexact</code>	--	Y	Y	R7RS names (aliases)
<code>exact-integer?</code>	--	Y	Y	
Numeric predicates				
<code>number? integer? real?</code>	Y	Y	Y	
<code>rational? complex?</code>	Y	Y	Y	

Feature	R5RS	R7RS	LambLisp	Notes
zero? positive? negative? odd? even?	Y	Y	Y	
finite? nan?	--	Y	Y	
infinite?	--	Y	Y	
Arithmetic				
+ - * / abs max min	Y	Y	Y	
< <= = >= >	Y	Y	Y	
quotient remainder modulo	Y	Y	Y	
expt log square sqrt	Y	Y	Y	
sin cos tan asin acos atan	Y	Y	Y	
floor (single-arg)	Y	Y	Y	Returns exact integer from inexact input (R5RS compliant)
ceiling	Y	Y	Y	Returns exact integer
truncate	Y	Y	Y	Returns exact integer
round	Y	Y	Y	Banker's rounding (half-to-even) per R5RS
floor/ floor-quotient floor-remainder	--	Y	Y	Integer args; floor/ returns two values via values
truncate/ truncate-quotient truncate-remainder	--	Y	Y	Integer args; identical to quotient/remainder
gcd lcm	Y	Y	Y	
numerator denominator	Y	Y	Y	
rationalize	Y	Y	Y	Stern-Brocot simplest-in-interval algorithm
exact-integer-sqrt	--	Y	Y	Returns two values via values
real-part imag-part magnitude angle	Y	Y	Y	

Feature	R5RS	R7RS	LambLisp	Notes
make-rectangular make-polar				
number->string string->number				

8.3.5 Symbols

Feature	R5RS	R7RS	LambLisp	Notes
symbol? symbol->string string->symbol				
symbol=?				
gensym				Used internally by syntax-rules hygiene

8.3.6 Pairs and Lists

Feature	R5RS	R7RS	LambLisp	Notes
pair? cons car cdr set-car! set-cdr!				
null? list?				
atom?				
list				
make-list				
caar ... cddddr (all compositions)				
length				
append reverse				

Feature	R5RS	R7RS	LambLisp	Notes
reverse!	--	--	ext	Destructive in-place reversal
list-tail list-ref list-set!	Y	Y	Y	
list-copy	--	Y	Y	Preserves improper list tail
memq memv member	Y	Y	Y	
assq assv assoc	Y	Y	Y	
vector->alist alist->vector	--	--	ext	

8.3.7 Characters

Feature	R5RS	R7RS	LambLisp	Notes
char?	Y	Y	Y	
char=? char<? char>? char<=? char>=?	Y	Y	Y	
char-ci=? char-ci<? char-ci>? char-ci<=? char-ci>=?	Y	Y	Y	
char-alphabetic? char-numeric? char-whitespace?	Y	Y	Y	
char-upper-case? char-lower-case?	Y	Y	Y	
char-upcase char-downcase	Y	Y	Y	
char-foldcase	--	Y	Y	tolower() (ASCII); in ll_vm_ai_↵ rxrs.cpp
char->integer integer->char	Y	Y	Y	
digit-value	--	Y	Y	

8.3.8 Strings

Feature	R5RS	R7RS	LambLisp	Notes
string? make-string string	Y	Y	Y	
string-length string-ref string-set!	Y	Y	Y	
string=? string<? string>? string<=? string>=?	Y	Y	Y	
string-ci=? string-ci<? string-ci>? string-ci<=? string-ci>=?	Y	Y	Y	
string-foldcase	--	Y	Y	ASCII downcase loop; in ll_vm↔ ai_rxrs.cpp
substring string-append	Y	Y	Y	
string->list list->string	Y	Y	Y	
string-copy	Y	Y	Y	
string-copy!	--	Y	Y	
string-fill!	Y	Y	Y	

8.3.9 Vectors

Feature	R5RS	R7RS	LambLisp	Notes
vector? make-vector vector	Y	Y	Y	
vector-length vector-ref vector-set!	Y	Y	Y	
vector-fill!	Y	Y	Y	
vector->list list->vector	Y	Y	Y	
vector-copy	--	Y	Y	
vector->string string->vector	--	Y	Y	

Feature	R5RS	R7RS	LambLisp	Notes
vector-append	--	Y	Y	
vector->alist alist->vector	--	--	ext	

8.3.10 Bytevectors

Feature	R5RS	R7RS	LambLisp	Notes
bytevector? make-bytevector bytevector	--	Y	Y	
bytevector-length	--	Y	Y	
bytevector-u8-ref bytevector-u8-set!	--	Y	Y	
bytevector-copy bytevector-copy! bytevector-append	--	Y	Y	
utf8->string string->utf8	--	Y	Y	ASCII-only; multi-byte UTF-8 sequences not supported
#u8(...) literal	--	Y	Y	

8.3.11 Control Features

Feature	R5RS	R7RS	LambLisp	Notes
procedure? apply	Y	Y	Y	
map for-each	Y	Y	Y	Multi-list variants supported
string-map string-for-each	--	Y	Y	Multi-string variants supported
vector-map vector-for-each	--	Y	Y	Multi-vector variants supported

Feature	R5RS	R7RS	LambLisp	Notes
values call-with-values	Y	Y	Y	
delay force	Y	Y	Y	Implemented in Scheme
delay-force promise? make-promise	--	Y	Y	Implemented in Scheme
call-with-current-continuation	Y	Y	N	By design; see FAQ
dynamic-wind	Y	Y	N	Requires call/cc

8.3.12 Environments and Evaluation

Feature	R5RS	R7RS	LambLisp	Notes
eval	Y	Y	Y	
scheme-report-environment	Y	Y	Y	
null-environment	Y	Y	Y	
interaction-environment	Y	Y	Y	
environment	--	Y	Y	

8.3.13 Ports

Feature	R5RS	R7RS	LambLisp	Notes
port? input-port? output-port?	Y	Y	Y	
textual-port? binary-port?	--	Y	Y	
input-port-open? output-port-open?	--	Y	Y	

Feature	R5RS	R7RS	LambLisp	Notes
current-input-port current-output-port current-error-port	Y	Y	Y	
close-port close-input-port close-output-port	Y	Y	Y	
open-input-file open-output-file	Y	Y	Y	
open-input-binary-file open-output-binary-file	--	Y	Y	
open-input-string	--	Y	Y	
open-output-string get-output-string	--	Y	Y	
open-input-bytevector open-output-bytevector	--	Y	Y	
with-input-from-file with-output-to-file	Y	Y	Y	
call-with-input-file call-with-output-file	Y	Y	Y	
call-with-port	--	Y	Y	

8.3.14 Input and Output

Feature	R5RS	R7RS	LambLisp	Notes
read read-char peek-char	Y	Y	Y	
eof-object?	Y	Y	Y	
char-ready?	Y	Y	Y	
read-line	--	Y	Y	
read-string	--	Y	N	Stubbed
read-u8 peek-u8 u8-ready?	--	Y	Y	

Feature	R5RS	R7RS	LambLisp	Notes
read-bytevector read-bytevector!	--	Y	N	Stubbed
write display newline	Y	Y	Y	
write-char	Y	Y	Y	
write-u8	--	Y	Y	
write-bytevector flush-output-port	--	Y	~	
write-shared write-simple	--	Y	N	

8.3.15 System Interface

Feature	R5RS	R7RS	LambLisp	Notes
load	Y	Y	Y	(load "file.scm" verbosity) – second arg: 1=verbose, 0=quiet
file-exists? delete-file	--	Y	Y	
current-jiffy jiffies-per-second	Y	Y	Y	
current-second	Y	Y	Y	Returns millis()/1000.0 as inexact real
command-line	--	Y	N	Not applicable on embedded targets
exit emergency-exit	--	Y	N	Not applicable on embedded targets
get-environment-variables	--	Y	N	Not applicable on embedded targets
features	--	Y	N	

8.3.16 Exception Handling

Feature	R5RS	R7RS	LambLisp	Notes
raise	--	Y	Y	
raise-continuable	--	Y	Y	Aliased to raise; full continuable behavior requires call/cc
with-exception-handler	--	Y	Y	
guard	--	Y	Y	Implemented as a syntax-rules macro
error	--	Y	Y	
error-object? error-object-message error-object-irritants	--	Y	Y	
read-error? file-error?	--	Y	Y	

Chapter 9

LambLisp Frequently Asked Questions

9.1 Getting Started

9.1.1 How do I try LambLisp right now?

The Linux build is the fastest way to explore LambLisp — no hardware required. Build and run with a single PlatformIO command:

```
pio run -e linux_x86_64 -t exec
```

The REPL starts immediately with the `LL=>` prompt. Load any `.scm` file with `(load "file.scm" 1)`. The same source files run unchanged on ESP32 once deployed.

See: [Quick Start: Try LambLisp on Linux · Deploy to ESP32](#)

9.1.2 What hardware does LambLisp run on?

The primary target is the **ESP32-S3** (240 MHz dual-core Xtensa LX7, 8 MB PSRAM, 4 MB flash filesystem (SPIFFS or LittleFS)). The Freenove 4WD car kit uses the original **ESP32** (single-core), which is also supported. A secondary target is the **NVIDIA Jetson Orin Nano** (ARM64, Linux). LambLisp also builds and runs on **Linux x86-64**, which is the primary development and testing platform.

See: [Supported Platforms and Boards](#)

9.1.3 Is LambLisp open source?

LambLisp is commercial software. It may be used for non-commercial purposes at no cost and with no support or warranty. Commercial use — generating revenue, promoting products or services, or building production systems — requires a license. Licensing includes customization, hardware application development assistance, post-deployment support, and warranty.

See: [License](#)

9.1.4 How do I update code without reflashing?

Upload a Scheme file to the running ESP32 over WiFi using LLIP, then evaluate `(load)` in the running loop — no reboot, no flash erase, no control gap:

```
;; From a supervisory host (Linux or Jetson):  
(define board (llip-connect "10.0.0.11" 5555))  
(llip-write board "/new-policy.scm" source-string)  
(llip-eval board '(load "/new-policy.scm"))
```

The new definitions replace the old ones in the interaction environment immediately. The GC reclaims the previous code on the next sweep.

See: [Code Lifecycle: Incremental Hot-Reload · LLIP](#)

9.2 Concepts

9.2.1 What is a real-time control system?

A real-time system guarantees that outputs are updated within a bounded time after inputs change — the *deadline*. LambLisp uses **loop-based control**: the application's `loop` procedure runs repeatedly, reading sensors and driving actuators each cycle. The GC runs incrementally in quanta so no single allocation event blocks the loop for more than a configurable slice. Deadlines may be *hard* (missing one requires a fallback) or *soft* (graceful degradation); most control problems can be factored to eliminate the hard aspect entirely.

See: [Real-time Guarantees · Memory and GC Tuning](#)

9.2.2 What is LLIP?

LLIP (LambLisp Interaction Protocol) is a line-oriented S-expression protocol over TCP that connects any two LambLisp nodes — ESP32 to Linux, ESP32 to Jetson, or any combination. It provides two services over one connection: remote filesystem operations (upload, download, directory listing) and remote evaluation (`llip-eval`). A supervisory controller on a Jetson or Linux host can query live sensor state, tune PID parameters, and hot-load new algorithms into a fleet of ESP32 boards without interrupting their control loops.

See: [LLIP — LambLisp Interaction Protocol](#)

9.2.3 How fast is LambLisp?

Speed depends on which execution tier is active. On an ESP32 LX6 at 240 MHz (D0WD-V3, 2026-05-20):

Benchmark	AST interpreter	Bytecode	NCG native	NCG/BC	NCG/Interp
tak(7,4,1) — mutual recursion	1,223 ms	28 ms	3.8 ms	7×	324×

Benchmark	AST interpreter	Bytecode	NCG native	NCG/BC	NCG/Interp
sieve(200) — loop-heavy	558 ms	4.8 ms	1.6 ms	3×	341×
fib-memo(25) — memoization	1,346 ms	4.9 ms	1.8 ms	3×	742×
tree(8) — allocation-heavy	299 ms	92 ms	14 ms	7×	22×
ackermann(3,1) — deep recursion	36 ms	9.5 ms	1.5 ms	6×	24×
qsort(200) — vector in-place	2,293 ms	49 ms	10 ms	5×	219×
fib-iter(1000) — tail recursion	720 ms	18 ms	9.3 ms	2×	77×
closure(500) — closure alloc	582 ms	113 ms	26 ms	4×	22×

The NCG compiles recursive and arithmetic-heavy procedures to native Xtensa instructions with direct call dispatch, achieving 2–10× over bytecode and 22–742× over the interpreter on recursion-dominated workloads.

See: [Execution Tiers: AST → Bytecode → Native Code](#)

9.2.4 What is bytecode / NCG compilation?

LambLisp has three execution tiers that mix at the procedure level:

Tier	Invocation	Typical speedup vs AST
AST interpreter	Default; no compile step	1×
Bytecode (BC)	(procedure->bytecode fn)	2–10×
Native Code Generator (NCG)	(ncg-compile fn)	2–10× over BC, up to 742× over Interp on ESP32

Each procedure can be compiled independently; the rest of the program runs interpreted. Recursive calls between compiled procedures bypass the interpreter entirely, which is why deeply recursive benchmarks show the largest NCG speedups. On ESP32, NCG→NCG direct calls skip argument-list allocation entirely.

See: [Execution Tiers: AST → Bytecode → Native Code · Bytecode Virtual Machine · Native Code Generator \(NCG\)](#)

9.2.5 Does the same Scheme code run on all platforms?

Yes, with a small exception for hardware I/O calls. All language semantics — closures, tail calls, the macro system, the exception system, the GC, the REPL — behave identically on ESP32, Jetson, and Linux. Any .scm file that does not call hardware-specific functions (digitalRead, Wire.begin, CUDA bindings) runs without modification on all three targets.

For the small amount of code that must branch on platform, use the boolean predicates `linux?`, `esp32?`, and `jetson?` defined in `setup.scm`, or inspect the `lamb-board` string directly.

See: [Multi-Platform Support · Deploy to ESP32](#)

9.3 Language and Conformance

9.3.1 Why is it called LambLisp instead of Scheme?

Scheme emerged in the 1970s as a minimal, principled dialect of Lisp that unified lexical scoping, tail recursion, and continuations. LambLisp implements R5RS/R7RS-small and adds practical extensions (`nlambda`, `macro`, the R7RS exception system), but deliberately omits `call-with-current-continuation`. For many Scheme enthusiasts `call/cc` is the *sine qua non* of the language, so the *Scheme* label would be misleading. See *Why no call-with-current-continuation?* below for the design rationale.

See: [Brief history of Lisp and Scheme](#)

9.3.2 Why no `call-with-current-continuation`?

Full `call/cc` requires capturing the entire C call stack as a first-class object — either by copying the stack, by converting all code to continuation-passing style (CPS), or by limiting continuations to escape-only uses. On a 240 MHz microcontroller with tight memory, the overhead of full CPS transformation is prohibitive.

Escape-only continuations cover the vast majority of real-world uses: structured loops, early return, and exception handling. LambLisp provides these through the R7RS exception system (`with-exception-handler`, `raise`, `guard`) — the practical use cases without the overhead of full continuations.

The Scheme community has itself debated whether full `call/cc` is worth the implementation cost; it is possible that a future R8RS specification will optionalize or remove it.

See: [Language Extensions](#)

9.3.3 Where does LambLisp sit in the Lisp landscape?

LambLisp targets the gap between tiny MCU Lisps and full desktop Schemes:

Feature	LispBM / uLisp &c	LambLisp	ChezScheme / Racket
Runs entirely on-board	no	yes	yes
Real-time incremental GC	no	yes	no
R5RS / R7RS conformance	no	yes (few exceptions)	yes (few exceptions)
File system, WiFi, BLE	no	yes	n/↔ a
<code>call/cc</code>	no	no	yes
Target RAM	< 64 KB	2–8 MB PSRAM	GB-class
Bytecode + NCG option	bytecode only	yes	yes

LambLisp was designed around these constraints: substantially conform to R5RS/R7RS-small; large enough to exploit 32-bit PSRAM but not a *Big Lisp*; real-time GC with no stop-the-world pauses; scalable execution from AST to NCG; easy reuse of Arduino-style drivers.

See: [Other Scheme Implementations](#)

9.3.4 Why not use embedded Python derivatives?

MicroPython and CircuitPython are useful for simple scripting but do not thrive as embedded control platforms. Their garbage collectors are tuned for throughput, not latency, and provide no real-time guarantees. A Python derivative called Serpent comes closest, with a real incremental collector, but it is not Python, requires off-board processing, has no generalized hardware interface, and does not fill the same solution space as LambLisp.

Python's indentation-sensitive block syntax makes tokenizing and parsing substantially heavier than Lisp's uniform prefix-parenthesis grammar. A Lisp reader is a few dozen lines of code; a Python parser requires tracking indent levels, dedent tokens, and a grammar with many special cases. That overhead compounds inside a microcontroller.

Part of Python's strength is its large library ecosystem, but the embedded Python variants cannot take full advantage of those libraries. LambLisp draws on lessons from Python where they apply — dictionaries are a first-class type in both languages — and extends them with hierarchical parent/child frames.

See: [Other Scheme Implementations](#)

9.3.5 Why not use LLVM for native code generation?

The NCG JIT-compiler Scheme lambdas **on the device at runtime**. The ESP32-S3 has 240 MHz and roughly 2 MB of PSRAM total. LLVM is a multi-megabyte library that requires a hosted environment with far more memory than an ESP32 provides. LLVM on the device is not a viable option regardless of output quality.

For the Linux and Jetson targets LLVM would fit in memory, but using a different code generator per platform would split the runtime semantics and make the codebase much harder to maintain. The hand-written NCG runs identically on all three targets, compiles a function in microseconds, and has no external dependencies. What it sacrifices in optimizer sophistication it gains in predictability — compilation latency is bounded and does not vary with loop complexity.

9.4 Integration and Operation

9.4.1 What are the advantages of *LambLisp* in combination with C/C++ on micro-controllers?

- **Interactive development** — the REPL lets you inspect state, redefine procedures, and test changes without reflashing.
- **Live code updates** — new definitions can be loaded over WiFi while the control loop runs.
- **Real-time safe GC** — incremental tricolor collector with configurable quanta keeps loop times bounded.
- **C++ interop** — any C++ function registered as a `mop3` appears as a native Scheme procedure; existing Arduino drivers need no changes.
- **Scalable execution** — bytecode and NCG compilation deliver 2–10× over BC and up to 742× over interpreter on hot recursive paths without touching the rest of the program.
- **AI-friendly** — uniform S-expression syntax makes LambLisp code easy to generate, modify, and verify programmatically.

See: [Why *LambLisp*? · C++ Integration · Scalability](#)

9.4.2 Does LambLisp interoperate with C/C++ code?

Yes. Any C++ function with the `mop3` signature appears as a native Scheme procedure in the interaction environment:

```
Sexpr_t my_fn(Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec);
```

Evaluated-argument procedures (`T_MOP3_PROC`) and syntax/unevaluated-argument procedures (`T_MOP3_↔NPROC`) are both supported. An installer function registers them in the interaction environment at startup. Existing Arduino libraries, hardware drivers, and C++ classes require no modification.

See: [C++ Integration](#)

9.4.3 Can I use LambLisp without an RTOS?

Yes. LambLisp runs directly on top of the Arduino `loop()` / `setup()` framework with no operating system. There is no dependency on FreeRTOS, Zephyr, or any other RTOS. The GC and the control loop share a single thread; incremental GC quanta are interleaved with Scheme evaluation so that loop deadlines are met without cooperative multitasking.

See: [Real-time Guarantees](#)

9.4.4 What is the memory footprint?

On the ESP32-S3 with 8 MB PSRAM:

- **Cell pool** — configurable; by default, cells are allocated from PSRAM until the pool is full.
- **GC heap** — the same cell pool; there is no separate allocation arena.
- **Scheme stack** — held as a linked list of cells on the GC heap; no fixed C stack limit.
- **C++ stack** — 96 KB in the current build (overridden from the 8 KB ESP-IDF default via `LL_STACK_SIZE`). This larger size is needed for deep recursion in the conformance test suite; most production applications can safely reduce it to 16–32 KB.
- **Flash filesystem** — `.scm` source files and `setup.scm`; stored on SPIFFS or LittleFS, selected via `LL_↔_SPIFFS=1 / LL_LITTLEFS=1`.

A fully loaded `setup.scm` with `syntax-rules`, the R7RS exception system, and standard library definitions uses a few thousand cells; the cell pool is large enough that this is not a concern in practice.

See: [Memory and GC Tuning](#)

Chapter 10

Other Scheme Implementations

This project was started after a review of available options for run-time languages to augment C++ embedded applications, and the likely candidates were all on the Lisp/Scheme spectrum. Some of the existing implementations that were examined include: TinyScheme, ulisp, microlisp, picolisp, T, L, Pre-Scheme, LispBM, chicken, bigloo, racket, chibi and others. Plug those into your favorite search engine to find even more.

General objections to existing Lisp/Scheme implementations
Required off-board compilation.
Out of scale for a microprocessor.
Non-portable hardware tricks (e.g., using the bottom bits of addresses as flag bits).
Non-Scheme lisps (e.g., no lexical scoping or tail recursion).
Bad fit on the interpret<->compile spectrum (e.g., low-level interpreter or huge compiler).
Lisp-on-Python, Lisp-on-JavaScript, and other such proofs.
Garbage collection: Real-time applications must avoid pauses, so stop-the-world, reference counting and generational GC cannot be used.

Chapter 11

Topic Documentation

11.1 Hardware & Platform Builtins (xmop3)

Topics

- [GPIO & Basic I/O](#)
- [ESP32 System](#)
- [Camera \(ESP32-S3-EYE\)](#)
- [LCD1602 Character Display](#)
- [1-Wire & DS18B20](#)
- [Over-the-Air Firmware Updates](#)
- [PCA9685 PWM / Servo Driver](#)
- [Platform \(generic\)](#)
- [Ultrasonic Sonar \(HC-SR04\)](#)
- [SPI Bus](#)
- [WiFi](#)
- [I2C \(Wire\)](#)
- [WS2812 / NeoPixel LEDs](#)
- [CUDA GPU](#)
- [HIP / ROCm GPU](#)

11.1.1 Detailed Description

The C++ `11_xmop3_*` builtins that expose hardware and platform features to Scheme – GPIO, buses (I2C/SPI/1-Wire), sensors, displays, LEDs, WiFi, camera, GPU, and platform control.

11.1.2 GPIO & Basic I/O

Functions

- `Sexpr_t mop3_cio_digitalRead` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_digitalWrite` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_pinMode` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_analogRead` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_analogWrite` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_delay_ms` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_delay_us` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_analogReadResolution` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_analogReference` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_analogWriteResolution` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_noTone` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_pulseIn` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_pulseInLong` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_shiftIn` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_shiftOut` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_tone` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_bit` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_bitClear` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_bitRead` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_bitSet` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_bitWrite` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_highByte` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_lowByte` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_attachInterrupt` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_detachInterrupt` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_digitalPinToInterrupt` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_interrupts` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_cio_noInterrupts` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t CommonIO_install_mop3` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Detailed Description

LambLisp GPIO & Basic I/O builtins.

Function Documentation

`mop3_cio_digitalRead()`

`Sexpr_t mop3_cio_digitalRead` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Read digital pin state; returns #t if HIGH, #f if LOW.

mop3_cio_digitalWrite()

```
Sexpr_t mop3_cio_digitalWrite (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Write digital pin HIGH/LOW. Level is Scheme-truthy (only #f is LOW) with one pragmatic extension: a NUMERIC ZERO also drives LOW. So (digitalWrite p #f) and (digitalWrite p 0) both give LOW; (digitalWrite p #t) and (digitalWrite p 1) give HIGH.

mop3_cio_pinMode()

```
Sexpr_t mop3_cio_pinMode (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Configure a pin as INPUT, INPUT_PULLUP, or OUTPUT.

mop3_cio_analogRead()

```
Sexpr_t mop3_cio_analogRead (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Read analog pin; returns raw ADC integer value.

mop3_cio_analogWrite()

```
Sexpr_t mop3_cio_analogWrite (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Write analog PWM value to pin.

mop3_cio_delay_ms()

```
Sexpr_t mop3_cio_delay_ms (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Delay execution by the given number of milliseconds.

mop3_cio_delay_us()

```
Sexpr_t mop3_cio_delay_us (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Delay execution by the given number of microseconds.

mop3_cio_analogReadResolution()

```
Sexpr_t mop3_cio_analogReadResolution (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set ADC resolution in bits (stub – platform-dependent).

mop3_cio_analogReference()

```
Sexpr_t mop3_cio_analogReference (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set ADC voltage reference (stub – platform-dependent).

mop3_cio_analogWriteResolution()

```
Sexpr_t mop3_cio_analogWriteResolution (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set DAC/PWM resolution in bits (stub – platform-dependent).

mop3_cio_noTone()

```
Sexpr_t mop3_cio_noTone (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Stop tone generation on a pin (stub).

mop3_cio_pulseIn()

```
Sexpr_t mop3_cio_pulseIn (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Measure a pulse on a pin; returns duration in microseconds (stub).

mop3_cio_pulseInLong()

```
Sexpr_t mop3_cio_pulseInLong (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Like pulseIn but supports longer pulses (stub).

mop3_cio_shiftIn()

```
Sexpr_t mop3_cio_shiftIn (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Shift in a byte of data one bit at a time (stub).

mop3_cio_shiftOut()

```
Sexpr_t mop3_cio_shiftOut (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Shift out a byte of data one bit at a time (stub).

mop3_cio_tone()

```
Sexpr_t mop3_cio_tone (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Generate a square wave tone on a pin at the given frequency and optional duration (ms).

mop3_cio_bit()

```
Sexpr_t mop3_cio_bit (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return a bitmask with only bit n set (stub).

mop3_cio_bitClear()

```
Sexpr_t mop3_cio_bitClear (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Clear a specific bit in an integer (stub).

mop3_cio_bitRead()

```
Sexpr_t mop3_cio_bitRead (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Read a specific bit from an integer (stub).

mop3_cio_bitSet()

```
Sexpr_t mop3_cio_bitSet (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set a specific bit in an integer (stub).

mop3_cio_bitWrite()

```
Sexpr_t mop3_cio_bitWrite (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Write a specific bit in an integer (stub).

mop3_cio_highByte()

```
Sexpr_t mop3_cio_highByte (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the high byte of a two-byte integer (stub).

mop3_cio_lowByte()

```
Sexpr_t mop3_cio_lowByte (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the low byte of a two-byte integer (stub).

mop3_cio_attachInterrupt()

```
Sexpr_t mop3_cio_attachInterrupt (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Attach an ISR to a digital pin interrupt (stub).

mop3_cio_detachInterrupt()

```
Sexpr_t mop3_cio_detachInterrupt (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Detach the ISR from a digital pin interrupt (stub).

mop3_cio_digitalPinToInterrupt()

```
Sexpr_t mop3_cio_digitalPinToInterrupt (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the interrupt number for a given digital pin (stub).

mop3_cio_interrupts()

```
Sexpr_t mop3_cio_interrupts (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Re-enable interrupts globally (stub).

mop3_cio_noInterrupts()

```
Sexpr_t mop3_cio_noInterrupts (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Disable interrupts globally (stub).

11.1.3 ESP32 System

Functions

- Sexpr_t mop3_esp32_heapinfo (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t mop3_esp32_gcmem (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t mop3_esp32_flash_probe (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t mop3_esp32_efuse_mac_get_default (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t mop3_esp32_efuse_pkg_ver (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t mop3_esp32_efuse_read_block (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t mop3_esp32_efuse_write_block (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **ESP32_install_mop3** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Detailed Description

LambLisp ESP32 System builtins.

Function Documentation

mop3_esp32_heapinfo()

`Sexpr_t mop3_esp32_heapinfo (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Report heap free/largest-block for the three caps that matter to flash reads. Returns (internal-free internal-largest dma-free dma-largest spiram-free spiram-largest) and also logs it. Diagnostic: littlefs reads ENOMEM (esp_partition_read rc 257).

mop3_esp32_gcmem()

`Sexpr_t mop3_esp32_gcmem (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

B117 diagnostic: print GC-buffer / cell-block memory placement (INTERNAL vs PSRAM) + sizes.

mop3_esp32_flash_probe()

`Sexpr_t mop3_esp32_flash_probe (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Reproduce a raw esp_partition_read under our own logging, to isolate the ENOMEM. (esp32-flash-probe "label" offset len where) where: 0=internal DMA buffer, 1=PSRAM buffer. Logs heap + destination region + rc; returns the esp_partition_read rc as an integer.

mop3_esp32_efuse_mac_get_default()

`Sexpr_t mop3_esp32_efuse_mac_get_default (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Read the default factory MAC address from eFuse; returns a 6-byte bytevector.

mop3_esp32_efuse_pkg_ver()

`Sexpr_t mop3_esp32_efuse_pkg_ver (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Return the chip package version from eFuse.

mop3_esp32_efuse_read_block()

`Sexpr_t mop3_esp32_efuse_read_block (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Read Nbits bits from an eFuse block at a bit offset; returns a bytevector or error code.

mop3_esp32_efuse_write_block()

```
Sexpr_t mop3_esp32_efuse_write_block (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Write Nbits bits from a bytevector or integer into an eFuse block; returns esp_err_t.

11.1.4 Camera (ESP32-S3-EYE)**Macros**

- #define **EYE_LCD_HOST** SPI3_HOST
- #define **EYE_LCD_SCLK** 21
- #define **EYE_LCD_MOSI** 47
- #define **EYE_LCD_CS** 44
- #define **EYE_LCD_DC** 43
- #define **EYE_LCD_RST** (-1)
- #define **EYE_LCD_BL** 48
- #define **EYE_BL_TIMER** LEDC_TIMER_1
- #define **EYE_BL_CHANNEL** LEDC_CHANNEL_1
- #define **EYE_LCD_W** 240
- #define **EYE_LCD_H** 240
- #define **EYE_LCD_HZ** (40 * 1000 * 1000)
- #define **EYE_LCD_BAND_ROWS** 16
- #define **EYE_CAM_EARLY_DMA_RESERVE** 0
- #define **EYE_CAM_XCLK** 15
- #define **EYE_CAM_SIOD** 4
- #define **EYE_CAM_SIOC** 5
- #define **EYE_CAM_VSYNC** 6
- #define **EYE_CAM_HREF** 7
- #define **EYE_CAM_PCLK** 13
- #define **EYE_CAM_D0** 11
- #define **EYE_CAM_D1** 9
- #define **EYE_CAM_D2** 8
- #define **EYE_CAM_D3** 10
- #define **EYE_CAM_D4** 12
- #define **EYE_CAM_D5** 18
- #define **EYE_CAM_D6** 17
- #define **EYE_CAM_D7** 16

Functions

- `Sexpr_t mop3_lcd_init` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_lcd_clear` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_lcd_box` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_lcd_rgb565` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_camera_init` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_camera_to_lcd` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_camera_preview` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_camera_thumb` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_camera_set` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_camera_capture` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_frame_brightness` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_lcd_backlight` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t mop3_lcd_testpat` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t EyeCam_install_mop3` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Detailed Description

LambLisp Camera (ESP32-S3-EYE) builtins.

Macro Definition Documentation

EYE_LCD_RST

```
#define EYE_LCD_RST (-1)
```

reset not wired (tied high on the board)

EYE_LCD_BL

```
#define EYE_LCD_BL 48
```

backlight – ACTIVE-LOW, LEDC/PWM (per esp-bsp)

EYE_BL_TIMER

```
#define EYE_BL_TIMER LEDC_TIMER_1
```

dedicated timer (camera XCLK uses LEDC_TIMER_0)

EYE_BL_CHANNEL

```
#define EYE_BL_CHANNEL LEDC_CHANNEL_1
```

dedicated channel (camera XCLK uses LEDC_CHANNEL_0)

EYE_LCD_HZ

```
#define EYE_LCD_HZ (40 * 1000 * 1000)
```

40 MHz pclk (10 MHz test ruled out SPI signal integrity)

EYE_LCD_BAND_ROWS

```
#define EYE_LCD_BAND_ROWS 16
```

16-row band -> ~15 draw_bitmap calls/frame (was 240, one per row -> a visible top-to-bottom paint = "wiggle"). ~7.5 KB internal-DMA; falls back to 1 row if the alloc is tight. The camera's 23 KB line buffer is reserved internal-DMA line buffer barely fits, so keep the LCD minimal

EYE_CAM_SIOD

```
#define EYE_CAM_SIOD 4
```

SCCB SDA.

EYE_CAM_SIOC

```
#define EYE_CAM_SIOC 5
```

SCCB SCL.

Function Documentation

mop3_lcd_init()

```
Sexpr_t mop3_lcd_init (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(lcd-init) -> #t. Bring up SPI3, the ST7789 panel and the backlight.

mop3_lcd_clear()

```
Sexpr_t mop3_lcd_clear (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(lcd-clear [rgb565]) -> void. Clear the whole screen (default black).

mop3_lcd_box()

```
Sexpr_t mop3_lcd_box (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(lcd-box x y w h rgb565) -> void. Draw a filled rectangle.

mop3_lcd_rgb565()

```
Sexpr_t mop3_lcd_rgb565 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(lcd-rgb565! bvec x y w h) -> void. Blit a big-endian RGB565 region to the panel.

mop3_camera_init()

```
Sexpr_t mop3_camera_init (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(camera-init [mode]) -> #t/#f. mode 'rgb565 (default, drives the LCD) or 'jpeg (for capture).

mop3_camera_to_lcd()

```
Sexpr_t mop3_camera_to_lcd (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(camera->lcd) -> #t/#f. Grab one RGB565 frame and blit it to the LCD (the hot path).

PER-FRAME DEBUG, OFF BY DEFAULT. This logged unconditionally, and the eye's app-loop calls camera->lcd on EVERY tick once the live preview auto-starts – ~20 lines/second onto the same serial line the REPL and every test harness read. On 2026-08-29 it buried the gcpause output and the --- done --- sentinel under a wall of frame dumps, so a leg that had ALREADY failed for a different reason (the bench file was missing from the manifest) looked busy for 146 s instead of obviously broken. A log inside a frame-rate path is a debugging aid, never shipped behaviour.

mop3_camera_preview()

```
Sexpr_t mop3_camera_preview (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(camera-preview [n]) -> n. Live preview: blit n frames to the LCD (default 300).

mop3_camera_thumb()

`Sexpr_t mop3_camera_thumb (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

(camera-thumb [n]) -> #t/#f. Grab one RGB565 frame, nearest-neighbour downsample to n*n (default 24, clamped 4..48), and print it over serial as: THUMB <n> <srcW> <srcH> <n rows of n 4-hex RGB565 values, byte-swapped to native> For OFF-DEVICE visual verification (reconstruct the image on the host) without the LCD – isolates a camera/frame fault from an LCD/blit fault. Output is gated by the quiet REPL, so call (verbose) first.

mop3_camera_set()

`Sexpr_t mop3_camera_set (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

(camera-set key val) -> void. Tune the OV2640 via its sensor_t control block. keys: framesize brightness contrast hmirror vflip special-effect.

mop3_camera_capture()

`Sexpr_t mop3_camera_capture (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

(camera-capture) -> bytevector. Return one JPEG frame (needs (camera-init 'jpeg)); returns #f if no frame is available.

mop3_frame_brightness()

`Sexpr_t mop3_frame_brightness (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

(frame-brightness) -> integer 0..255. Cheap mean luma over a downsampled RGB565 frame – a scalar for trigger-policy demos ("brighten past N -> fire").

mop3_lcd_backlight()

`Sexpr_t mop3_lcd_backlight (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

(lcd-backlight [percent]) -> void. Set backlight brightness 0..100 (default 100). Handy for demos: (lcd-backlight 0) blanks the screen without tearing down the panel, (lcd-backlight 100) restores it. Independent of panel content, so it also isolates a backlight fault from a draw fault when bringing up new hardware.

mop3_lcd_testpat()

`Sexpr_t mop3_lcd_testpat (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

(lcd-testpat) -> #t. Blit a KNOWN 240x240 pattern through the SAME path as camera->lcd (byte-swap + banded draw) – nothing to do with the camera. Four horizontal zones: rows 0- 59 : vertical colour bars (R G B W x2) – X-varying solid data rows 60-119 : horizontal colour bars (R G B W) – Y-varying solid data rows 120-179 : smooth grayscale ramp (black->white L->R) rows 180-239 : fine 2px vertical black/white stripes – high-frequency detail Whatever the panel shows per zone pinpoints the display-path corruption (I control every input pixel).

EyeCam_install_mop3()

`Sexpr_t EyeCam_install_mop3 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Installer – always compiles; registers the eye mops only under LL_CAMERA.

11.1.5 LCD1602 Character Display

Typedefs

- typedef LiquidCrystal_I2C **LCD1602**

Functions

- Sexpr_t **LCD1602_mop3_begin** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_clear** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_home** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_display** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_cursor** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_backlight** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_write** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_print** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_mop3_setCursor** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **LCD1602_install_mop3** (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Variables

- LCD1602 * **lcd1602** = 0

Detailed Description

LambLisp LCD1602 Character Display builtins.

Function Documentation

LCD1602_mop3_begin()

```
Sexpr_t LCD1602_mop3_begin (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Initialize the LCD1602 at the given I2C address with cols x rows dimensions.

LCD1602_mop3_clear()

```
Sexpr_t LCD1602_mop3_clear (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Clear the LCD display and return cursor to home.

LCD1602_mop3_home()

```
Sexpr_t LCD1602_mop3_home (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Move cursor to the home position (0, 0).

LCD1602_mop3_display()

```
Sexpr_t LCD1602_mop3_display (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Turn the display on or off; optional arg defaults to #t (on).

LCD1602_mop3_cursor()

```
Sexpr_t LCD1602_mop3_cursor (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Show or hide the cursor; optional arg defaults to #t (show).

LCD1602_mop3_backlight()

`Sexpr_t LCD1602_mop3_backlight (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Turn backlight on or off; optional arg defaults to #t (on).

LCD1602_mop3_write()

`Sexpr_t LCD1602_mop3_write (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Write a string at the current cursor position.

LCD1602_mop3_print()

`Sexpr_t LCD1602_mop3_print (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Move cursor to (0,0) then write a string – convenience for full-line refresh.

LCD1602_mop3_setCursor()

`Sexpr_t LCD1602_mop3_setCursor (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Move cursor to column col and row row (zero-based).

11.1.6 1-Wire & DS18B20

Functions

- `Sexpr_t OneWire_mop3_reset (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t OneWire_mop3_write_byte (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t OneWire_mop3_read_byte (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t OneWire_mop3_write_bit (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t OneWire_mop3_read_bit (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t OneWire_mop3_search (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t OneWire_install_mop3 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Detailed Description

LambLisp 1-Wire & DS18B20 builtins.

Function Documentation

OneWire_mop3_reset()

```
Sexpr_t OneWire_mop3_reset (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(onewire-reset pin) -> #t/#f

OneWire_mop3_write_byte()

```
Sexpr_t OneWire_mop3_write_byte (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(onewire-write-byte pin b)

OneWire_mop3_read_byte()

```
Sexpr_t OneWire_mop3_read_byte (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(onewire-read-byte pin) -> byte

OneWire_mop3_write_bit()

```
Sexpr_t OneWire_mop3_write_bit (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(onewire-write-bit pin b)

OneWire_mop3_read_bit()

```
Sexpr_t OneWire_mop3_read_bit (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(onewire-read-bit pin) -> 0/1

OneWire_mop3_search()

```
Sexpr_t OneWire_mop3_search (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(onewire-search pin) -> (list rom-bytevector ...), each 8 bytes.

OneWire_install_mop3()

```
Sexpr_t OneWire_install_mop3 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Install all OneWire symbols in base environment.

11.1.7 Over-the-Air Firmware Updates

Macros

- `#define OTA_NS "ota-p125"`

Functions

- `Sexpr_t ota_mop3_request_install (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t ota_mop3_request_recover (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t ota_mop3_request_console (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t ota_mop3_set_network (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t ota_mop3_reboot_to_loader (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t ota_mop3_boot_ok (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`
- `Sexpr_t ota_install_mop3 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)`

Detailed Description

LambLisp Boot-Loader Handshake (OTA) builtins.

Function Documentation

ota_mop3_request_install()

```
Sexpr_t ota_mop3_request_install (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(ota-request-install!) — ask the loader to INSTALL (pending_update=1: fetch + write a new ota_0). Returns #t on success, #f off-device.

ota_mop3_request_recover()

```
Sexpr_t ota_mop3_request_recover (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(ota-request-recover!) — ask the loader to RECOVER (pending_update=2: re-fetch a known-good image). Returns #t on success, #f off-device.

ota_mop3_request_console()

```
Sexpr_t ota_mop3_request_console (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(ota-request-console!) — ask the loader for the CONSOLE fallback (pending_update=3: operator BLE/↔UART). Returns #t on success, #f off-device.

ota_mop3_set_network()

```
Sexpr_t ota_mop3_set_network (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(ota-set-network! ssid pass url) — write the network-rung config keys wifi_ssid / wifi_pass / img↔_url (all strings) to NVS. Returns #t on success, #f off-device.

ota_mop3_reboot_to_loader()

```
Sexpr_t ota_mop3_reboot_to_loader (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(ota-reboot-to-loader) — point otadata at the factory loader partition and esp_restart; never returns on device. Returns #f if there is no factory partition (or off-device).

ota_mop3_boot_ok()

```
Sexpr_t ota_mop3_boot_ok (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(ota-boot-ok!) — confirm THIS image is healthy: cancel the native OTA rollback and clear pending_update + install_tries (the loader's anti-loop counter). The loader's install-retry guard waits for this. Returns #t (vacuously #t off-device).

ota_install_mop3()

```
Sexpr_t ota_install_mop3 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Installer: bind the OTA loader-handshake builtins (ota-request-install! / -recover! / -console!, ota-set-network!, ota-reboot-to-loader, ota-boot-ok!) into env_target.

11.1.8 PCA9685 PWM / Servo Driver

Functions

- Sexpr_t `PCA9685_mop3_begin` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)
- Sexpr_t `PCA9685_mop3_setupSingleDevice` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)
- Sexpr_t `PCA9685_mop3_setSingleDeviceToFrequency` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_↔_exec`)
- Sexpr_t `PCA9685_mop3_getSingleDeviceFrequency` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)
- Sexpr_t `PCA9685_mop3_setChannelDutyCycle` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)
- Sexpr_t `PCA9685_mop3_setChannelPulseWidth` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)
- Sexpr_t `PCA9685_mop3_setChannelServoPulseDuration` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_↔_exec`)
- Sexpr_t `PCA9685_mop3_setChannelOnAndOffTime` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)
- Sexpr_t `PCA9685_install_mop3` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)

Variables

- PCA9685 * `pca9685` = 0

Detailed Description

LambLisp PCA9685 PWM / Servo Driver builtins.

Function Documentation

PCA9685_mop3_begin()

Sexpr_t `PCA9685_mop3_begin` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)

Initialize PCA9685 at I2C addr via Wire; sets 50 Hz PWM frequency.

PCA9685_mop3_setupSingleDevice()

Sexpr_t `PCA9685_mop3_setupSingleDevice` (`Lamb & lamb`, Sexpr_t `sexpr`, Sexpr_t `env_exec`)

Configure a single PCA9685 device on the Wire bus at the given I2C address.

PCA9685_mop3_setSingleDeviceToFrequency()

```
Sexpr_t PCA9685_mop3_setSingleDeviceToFrequency (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set PWM oscillator frequency (Hz) for a single device at the given I2C address.

PCA9685_mop3_getSingleDeviceFrequency()

```
Sexpr_t PCA9685_mop3_getSingleDeviceFrequency (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Read back the PWM oscillator frequency (Hz) for a single device at the given address.

PCA9685_mop3_setChannelDutyCycle()

```
Sexpr_t PCA9685_mop3_setChannelDutyCycle (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set channel duty cycle as a fraction [0.0, 1.0] with optional phase shift fraction.

PCA9685_mop3_setChannelPulseWidth()

```
Sexpr_t PCA9685_mop3_setChannelPulseWidth (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set channel pulse width in raw 12-bit ticks (0–4095).

PCA9685_mop3_setChannelServoPulseDuration()

```
Sexpr_t PCA9685_mop3_setChannelServoPulseDuration (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set channel servo pulse duration in microseconds.

PCA9685_mop3_setChannelOnAndOffTime()

```
Sexpr_t PCA9685_mop3_setChannelOnAndOffTime (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set channel on-time and off-time in raw 12-bit ticks; on-time anchored at 0.

11.1.9 Platform (generic)

Functions

- Sexpr_t `Platform_mop3_free_heap` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_ncg_pool` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_free_stack` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_micros` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_millis` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_loop_elapsed_ms` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_loop_elapsed_us` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_lamb_reboot` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_exit` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_delay_ms` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_gc_idle_task` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_gc_collect` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_random_bytevector` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_random_integer` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_random_real` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_cell_block_size` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_n_cell_blocks` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_total_cells` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_nfree` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_max_cell_blocks` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_set_max_cell_blocks` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_extension_block_size` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_set_extension_block_size` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_expand_to_n_blocks` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_gc_diag` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_gc_budget_ns` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_set_gc_budget_ns` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_gcloud_target_pct` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_set_gcloud_target_pct` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_verbosity` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_set_verbosity` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_ncg_frame_pool_n` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_ncg_frame_pool_init` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Platform_mop3_repl_quiet` (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `__lamb_platform_generic_install_mop3` (Lamb &lamb, Sexpr_t list_containing_env_target, Sexpr_t env_exec)

Variables

- LambPlatform **lambPlatform**

Detailed Description

LambLisp Platform (generic) builtins.

Function Documentation

Platform_mop3_free_heap()

```
Sexpr_t Platform_mop3_free_heap (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return bytes of free heap memory on the platform.

Platform_mop3_ncg_pool()

```
Sexpr_t Platform_mop3_ncg_pool (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

B245: return (used . size) of the NCG executable pool, in bytes.

Platform_mop3_free_stack()

```
Sexpr_t Platform_mop3_free_stack (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return bytes of free stack space on the platform.

Platform_mop3_micros()

```
Sexpr_t Platform_mop3_micros (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return microseconds since boot as an integer.

Platform_mop3_millis()

```
Sexpr_t Platform_mop3_millis (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return milliseconds since boot as an integer.

Platform_mop3_loop_elapsed_ms()

```
Sexpr_t Platform_mop3_loop_elapsed_ms (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return elapsed milliseconds since the start of the current Arduino loop() call.

Platform_mop3_loop_elapsed_us()

```
Sexpr_t Platform_mop3_loop_elapsed_us (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return elapsed microseconds since the start of the current Arduino loop() call.

Platform_mop3_lamb_reboot()

```
Sexpr_t Platform_mop3_lamb_reboot (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Reboot the platform (embedded) or do nothing (POSIX).

Platform_mop3_exit()

```
Sexpr_t Platform_mop3_exit (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Flush stdio and exit (POSIX) or reboot (embedded).

< release platform resources

< embedded: reboot after cleanup

< unreachable

Platform_mop3_delay_ms()

```
Sexpr_t Platform_mop3_delay_ms (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Delay for the given number of milliseconds using platform delay().

Platform_mop3_gc_idle_task()

```
Sexpr_t Platform_mop3_gc_idle_task (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Run the GC idle task for up to the given number of microseconds.

Platform_mop3_gc_collect()

```
Sexpr_t Platform_mop3_gc_collect (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Run one complete GC cycle (mark + sweep); use before benchmarks to reset GC phase.

Platform_mop3_random_bytevector()

```
Sexpr_t Platform_mop3_random_bytevector (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return a bytevector of n cryptographically random bytes.

Platform_mop3_random_integer()

```
Sexpr_t Platform_mop3_random_integer (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return a cryptographically random integer.

Platform_mop3_random_real()

```
Sexpr_t Platform_mop3_random_real (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return a cryptographically random float32 in [0, 1).

$< [0, 2^{64}-1] * 2^{-64} \rightarrow [0,1)$ exact

Platform_mop3_cell_block_size()

```
Sexpr_t Platform_mop3_cell_block_size (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the number of cells in the first (fixed) memory block.

Platform_mop3_n_cell_blocks()

```
Sexpr_t Platform_mop3_n_cell_blocks (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the current number of allocated memory blocks.

Platform_mop3_total_cells()

```
Sexpr_t Platform_mop3_total_cells (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

B229: THE CELL POOL, computed ONCE in the VM. Callers used to reconstruct it as `nblocks * cell_block_size` which is right only if every block is the same size. With a 65,536-cell base block and 8,192-cell expansions that overstated the pool by 7.2x (4,194,304 reported against a true 581,632), and every occupancy figure derived from it was wrong by the same factor. `total_cells()` is `baseBlockSize + NexansionBlocks * expansionBlockSize` – ask the VM, do not re-derive it.

Platform_mop3_nfree()

```
Sexpr_t Platform_mop3_nfree (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the current number of free cells in the GC heap.

Platform_mop3_max_cell_blocks()

```
Sexpr_t Platform_mop3_max_cell_blocks (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the maximum number of GC blocks allowed before the memory_cap_reached warning fires.

Platform_mop3_set_max_cell_blocks()

```
Sexpr_t Platform_mop3_set_max_cell_blocks (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the maximum number of GC blocks before memory_cap_reached warning fires.

Platform_mop3_set_extension_block_size()

```
Sexpr_t Platform_mop3_set_extension_block_size (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the number of cells in each future extension memory block.

Platform_mop3_expand_to_n_blocks()

```
Sexpr_t Platform_mop3_expand_to_n_blocks (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Expand the heap to at least n blocks; allocates extension blocks as needed.

Platform_mop3_gc_diag()

```
Sexpr_t Platform_mop3_gc_diag (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

B147: (Platform.gc-diag) -> (pass-max-us urgent-max-us urgent-calls). Reported through the RESULT of a mop3 so the bench can display it – warn() output never survives the embedded path (W3_REPL_← QUIET at the source, plus serial_run.py dropping any line naming LambMemory or gc_pass), so a warn-based probe silently reports nothing and looks like a clean result. Index 10: monotonic total us inside gc_← pass. APPENDED at the tail on purpose – the field order here is positional and gc-pause-bench.scm reads it by list-ref, and the labels went stale once before when the list was reordered. Add at the end; never insert. Index 11/12: yuasa_M and yuasa_N. APPENDED, never inserted – gc-pause-bench.scm reads this list POSITIONALLY by list-ref and the labels went stale once before on a reorder.

Platform_mop3_gc_budget_ns()

```
Sexpr_t Platform_mop3_gc_budget_ns (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the GC time budget per loop iteration in nanoseconds.

Platform_mop3_set_gc_budget_ns()

```
Sexpr_t Platform_mop3_set_gc_budget_ns (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the GC time budget per loop iteration in nanoseconds.

Platform_mop3_gcloud_target_pct()

```
Sexpr_t Platform_mop3_gcloud_target_pct (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the target GC load as a percentage of loop time.

Platform_mop3_set_gcloud_target_pct()

```
Sexpr_t Platform_mop3_set_gcloud_target_pct (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the target GC load percentage (0–100).

Platform_mop3_verbosity()

```
Sexpr_t Platform_mop3_verbosity (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the current verbosity level (0 = silent, higher = more logging).

Platform_mop3_set_verbosity()

```
Sexpr_t Platform_mop3_set_verbosity (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the verbosity level.

Platform_mop3_ncg_frame_pool_n()

```
Sexpr_t Platform_mop3_ncg_frame_pool_n (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the number of NcgFrames currently in the free-list pool.

Platform_mop3_ncg_frame_pool_init()

```
Sexpr_t Platform_mop3_ncg_frame_pool_init (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Add n NcgFrames to the free-list pool (useful after pool exhaustion).

Platform_mop3_repl_quiet()

```
Sexpr_t Platform_mop3_repl_quiet (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Quiet REPL: enable/disable line-mode quiet console (suppress async log() output and the per-char redisplay() prompt-reprint). The LLIP tunnel uses this – (llt) sets it on, (en) clears it – so a tool's frames aren't garbled by the prompt-redraw flood.

11.1.10 Ultrasonic Sonar (HC-SR04)**Functions**

- void IRAM_ATTR `sonar_echo_isr` (void *arg)
- Sexpr_t `Sonar_mop3_begin` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_mop3_start` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_mop3_result` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_mop3_ping` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_mop3_poke` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_mop3_us_to_distance` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_mop3_set_timeout_us` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_mop3_set_min_interval_us` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t `Sonar_install_mop3` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Variables

- Sonar * `sonar` = 0

Detailed Description

Copyright 2026 by Frobenius Norm LLC 2026-04-23 00:00:00.

LambLisp Ultrasonic Sonar (HC-SR04) builtins.

Function Documentation

sonar_echo_isr()

```
void IRAM_ATTR sonar_echo_isr (void * arg)
```

FALLING-edge ISR: ECHO pin went LOW – echo received (or timeout on sensor side).

Sonar_mop3_begin()

```
Sexpr_t Sonar_mop3_begin (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Initialize the Sonar singleton with trigger and echo pin numbers.

Sonar_mop3_start()

```
Sexpr_t Sonar_mop3_start (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Start a ping. On ESP32: non-blocking. On other platforms: blocks on pulseIn.

Sonar_mop3_result()

```
Sexpr_t Sonar_mop3_result (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Returns: #f – echo not yet received (still within timeout window) 0 – timed out, no echo; caller should re-arm with Sonar.start integer – elapsed microseconds; caller should convert and re-arm

Sonar_mop3_ping()

```
Sexpr_t Sonar_mop3_ping (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Blocking ping: fire trigger, wait for echo, return elapsed microseconds or #f on timeout.

Sonar_mop3_poke()

```
Sexpr_t Sonar_mop3_poke (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Blocking ping converted to distance in meters; returns #f on timeout.

Sonar_mop3_us_to_distance()

```
Sexpr_t Sonar_mop3_us_to_distance (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Convert elapsed microseconds to distance in meters; optional Celsius temperature arg.

Sonar_mop3_set_timeout_us()

```
Sexpr_t Sonar_mop3_set_timeout_us (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(Sonar.timeout-us! <us>) – set the no-echo timeout at runtime; returns the effective (clamped-to-range-floor) value so the caller can confirm it.

Sonar_mop3_set_min_interval_us()

```
Sexpr_t Sonar_mop3_set_min_interval_us (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(Sonar.min-interval-us! <us>) – set the HC-SR04 inter-ping settle gap at runtime; returns the effective value. Larger = steadier readings, lower max ping rate.

Sonar_install_mop3()

```
Sexpr_t Sonar_install_mop3 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Install all Sonar symbols in base environment.

Variable Documentation

sonar

```
Sonar* sonar = 0
```

Sonar singleton.

11.1.11 SPI Bus

Functions

- `Sexpr_t SPI_mop3_begin` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t SPI_mop3_beginTransaction` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t SPI_mop3_endTransaction` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t SPI_mop3_transfer` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t SPI_mop3_transfer_bytevector` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t SPI_mop3_write_bytevector` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t SPI_mop3_end` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t SPI_install_mop3` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Detailed Description

LambLisp SPI Bus builtins.

Function Documentation

SPI_mop3_begin()

`Sexpr_t SPI_mop3_begin` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

(SPI.begin [sck miso mosi]) – start the HW SPI peripheral on the given pins (or defaults).

SPI_mop3_beginTransaction()

`Sexpr_t SPI_mop3_beginTransaction` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

(SPI.beginTransaction clock-hz [mode]) – open a transaction; mode 0..3 (default 0).

SPI_mop3_endTransaction()

`Sexpr_t SPI_mop3_endTransaction` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

(SPI.endTransaction) – close the current transaction.

SPI_mop3_transfer()

```
Sexpr_t SPI_mop3_transfer (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(SPI.transfer byte) -> byte – full-duplex single-byte exchange.

SPI_mop3_transfer_bytevector()

```
Sexpr_t SPI_mop3_transfer_bytevector (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(SPI.transfer-bytevector! bv) – in-place full-duplex block transfer; returns bv.

SPI_mop3_write_bytevector()

```
Sexpr_t SPI_mop3_write_bytevector (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(SPI.write-bytevector bv) – write-only (MOSI), MISO discarded; returns byte count.

SPI_mop3_end()

```
Sexpr_t SPI_mop3_end (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(SPI.end) – release the SPI peripheral.

SPI_install_mop3()

```
Sexpr_t SPI_install_mop3 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Install all SPI symbols in base environment.

11.1.12 WiFi

Functions

- `Sexpr_t WiFi_mop3_begin` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_disconnect` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_reconnect` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_config` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_setDNS` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_waitForConnectResult` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_isConnected` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_setAutoReconnect` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_SSID` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_BSSID` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_RSSI` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_localIP` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_subnetMask` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_gatewayIP` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_broadcastIP` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_networkID` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_dnsIP` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_subnetCIDR` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_psk` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_macAddress` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_encryptionType` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_status` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_mode` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_getHostname` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_setHostname` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_setSleep` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_channel` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_scanNetworks` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_hostByName` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAP` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPConfig` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPdisconnect` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPgetStationNum` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPIP` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPBroadcastIP` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPSSID` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPmacAddress` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPgetHostname` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_mop3_softAPsetHostname` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WiFi_install_mop3` (`Lamb &lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Variables

- `WiFiClass * LL_WiFi`

Detailed Description

LambLisp WiFi builtins.

Function Documentation

WiFi_mop3_begin()

```
Sexpr_t WiFi_mop3_begin (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Connect to an access point with SSID and optional password; returns wl_status_t.

WiFi_mop3_disconnect()

```
Sexpr_t WiFi_mop3_disconnect (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Disconnect from WiFi; optional args wifioff and eraseap.

WiFi_mop3_reconnect()

```
Sexpr_t WiFi_mop3_reconnect (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Reconnect to the last access point; returns #t on success.

WiFi_mop3_config()

```
Sexpr_t WiFi_mop3_config (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set static IP configuration; all args are IP strings or #f for INADDR_ANY.

WiFi_mop3_setDNS()

```
Sexpr_t WiFi_mop3_setDNS (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Reconfigure DNS servers without changing IP/GW/subnet.

WiFi_mop3_waitForConnectResult()

```
Sexpr_t WiFi_mop3_waitForConnectResult (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Block until connected or timeout (default 60000 ms); returns wl_status_t integer.

WiFi_mop3_isConnected()

```
Sexpr_t WiFi_mop3_isConnected (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return #t if currently connected to an access point.

WiFi_mop3_setAutoReconnect()

```
Sexpr_t WiFi_mop3_setAutoReconnect (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Enable or disable automatic reconnection; optional arg defaults to #t (enable).

WiFi_mop3_SSID()

```
Sexpr_t WiFi_mop3_SSID (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the SSID of the connected access point as a string.

WiFi_mop3_BSSID()

```
Sexpr_t WiFi_mop3_BSSID (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the BSSID (AP MAC address) as a colon-separated hex string.

WiFi_mop3_RSSI()

```
Sexpr_t WiFi_mop3_RSSI (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the received signal strength in dBm.

WiFi_mop3_localIP()

```
Sexpr_t WiFi_mop3_localIP (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the local IP address as a dotted-decimal string.

WiFi_mop3_subnetMask()

```
Sexpr_t WiFi_mop3_subnetMask (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the subnet mask as a dotted-decimal string.

WiFi_mop3_gatewayIP()

```
Sexpr_t WiFi_mop3_gatewayIP (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the gateway IP address as a dotted-decimal string.

WiFi_mop3_broadcastIP()

```
Sexpr_t WiFi_mop3_broadcastIP (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the broadcast IP address as a dotted-decimal string.

WiFi_mop3_networkID()

```
Sexpr_t WiFi_mop3_networkID (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the network base address as a dotted-decimal string.

WiFi_mop3_dnsIP()

```
Sexpr_t WiFi_mop3_dnsIP (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the DNS server IP string; optional arg n selects primary (0) or secondary (1).

WiFi_mop3_subnetCIDR()

```
Sexpr_t WiFi_mop3_subnetCIDR (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the subnet prefix length as an integer (CIDR notation).

WiFi_mop3_psk()

```
Sexpr_t WiFi_mop3_psk (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the WPA/WPA2 pre-shared key as a string.

WiFi_mop3_macAddress()

```
Sexpr_t WiFi_mop3_macAddress (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the station MAC address as a colon-separated hex string.

WiFi_mop3_encryptionType()

```
Sexpr_t WiFi_mop3_encryptionType (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the encryption type of the connected network as an integer.

WiFi_mop3_status()

```
Sexpr_t WiFi_mop3_status (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return a list (status-int status-string) describing the current WiFi connection state.

WiFi_mop3_mode()

```
Sexpr_t WiFi_mop3_mode (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set WiFi mode: 0=NULL 1=STA 2=AP 3=APSTA; returns #t on success.

WiFi_mop3_getHostname()

```
Sexpr_t WiFi_mop3_getHostname (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the mDNS hostname as a string.

WiFi_mop3_setHostname()

```
Sexpr_t WiFi_mop3_setHostname (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the mDNS hostname; returns #t on success.

WiFi_mop3_setSleep()

```
Sexpr_t WiFi_mop3_setSleep (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set power-save sleep mode: #t/#f or integer 0=none 1=modem 2=light.

WiFi_mop3_channel()

```
Sexpr_t WiFi_mop3_channel (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the current WiFi channel number.

WiFi_mop3_scanNetworks()

```
Sexpr_t WiFi_mop3_scanNetworks (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Scan for nearby networks and return the count found.

WiFi_mop3_hostByName()

```
Sexpr_t WiFi_mop3_hostByName (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Resolve a hostname to an IP string; returns #f on failure.

WiFi_mop3_softAP()

```
Sexpr_t WiFi_mop3_softAP (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Start a soft access point with SSID and optional password/channel/hidden/max-conn.

WiFi_mop3_softAPConfig()

```
Sexpr_t WiFi_mop3_softAPConfig (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set soft-AP IP configuration; ip, gateway, subnet are dotted-decimal strings.

WiFi_mop3_softAPdisconnect()

```
Sexpr_t WiFi_mop3_softAPdisconnect (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Shut down the soft-AP; optional arg wifioff also disables the radio.

WiFi_mop3_softAPgetStationNum()

```
Sexpr_t WiFi_mop3_softAPgetStationNum (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the number of stations currently connected to the soft-AP.

WiFi_mop3_softAPIP()

```
Sexpr_t WiFi_mop3_softAPIP (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the soft-AP IP address as a dotted-decimal string.

WiFi_mop3_softAPBroadcastIP()

```
Sexpr_t WiFi_mop3_softAPBroadcastIP (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the soft-AP broadcast IP address as a dotted-decimal string.

WiFi_mop3_softAPSSID()

```
Sexpr_t WiFi_mop3_softAPSSID (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the soft-AP SSID as a string.

WiFi_mop3_softAPmacAddress()

```
Sexpr_t WiFi_mop3_softAPmacAddress (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the soft-AP MAC address as a colon-separated hex string.

WiFi_mop3_softAPgetHostname()

```
Sexpr_t WiFi_mop3_softAPgetHostname (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the soft-AP mDNS hostname as a string.

WiFi_mop3_softAPsetHostname()

```
Sexpr_t WiFi_mop3_softAPsetHostname (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the soft-AP mDNS hostname; returns #t on success.

11.1.13 I2C (Wire)

Functions

- Sexpr_t [Wire_mop3_setTimeout](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_setPins](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_begin](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_requestFrom](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_setClock](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_beginTransmission](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_available](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_read](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_end](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_onReceive](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_onRequest](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_clearWireTimeoutFlag](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_getWireTimeoutFlag](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_endTransmission](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_mop3_write](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [Wire_install_mop3](#) (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Variables

- TwoWire * **LL_Wire** = 0

WiFi and Wire, in case we need to rationalize conflicting implementations.

- WiFiClass * **LL_WiFi**

Detailed Description

LambLisp I2C (Wire) builtins.

Function Documentation

Wire_mop3_setTimeout()

```
Sexpr_t Wire_mop3_setTimeout (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set Wire timeout; no args = 1000 ms default; (timeout reset-on-timeout-flag) form also accepted.

(Wire.setTimeout) OR (Wire.setTimeout timeout reset-on-timeout-flag)

With no parameters, some generally useful (but as yet unknown) values are used.

Wire_mop3_setPins()

```
Sexpr_t Wire_mop3_setPins (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Assign SDA and SCL pins before calling Wire.begin.

Wire_mop3_begin()

```
Sexpr_t Wire_mop3_begin (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(begin [addr]) If address is provided , it is our slave bus address. If no address is provided, then we are the bus controller.

Wire_mop3_requestFrom()

```
Sexpr_t Wire_mop3_requestFrom (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

(Wire.requestFrom addr quantity [release-bus-after]) OR (Wire.requestFrom addr quantity)

If not provided, release-bus-after defaults to #t.

Wire_mop3_setClock()

```
Sexpr_t Wire_mop3_setClock (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Set the I2C clock frequency in Hz.

Wire_mop3_beginTransmission()

```
Sexpr_t Wire_mop3_beginTransmission (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Begin a transmission to the I2C device at the given address.

Wire_mop3_available()

```
Sexpr_t Wire_mop3_available (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the number of bytes available to read from the I2C buffer.

Wire_mop3_read()

```
Sexpr_t Wire_mop3_read (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Read one byte from the I2C receive buffer; returns -1 if none available.

Wire_mop3_end()

```
Sexpr_t Wire_mop3_end (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Release the Wire bus and free resources; returns #t on success.

Wire_mop3_onReceive()

```
Sexpr_t Wire_mop3_onReceive (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Register a receive callback for slave mode (stub).

Wire_mop3_onRequest()

```
Sexpr_t Wire_mop3_onRequest (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Register a request callback for slave mode (stub).

Wire_mop3_clearWireTimeoutFlag()

```
Sexpr_t Wire_mop3_clearWireTimeoutFlag (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Clear the Wire timeout flag (stub).

Wire_mop3_getWireTimeoutFlag()

```
Sexpr_t Wire_mop3_getWireTimeoutFlag (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the Wire timeout flag (stub).

Wire_mop3_endTransmission()

```
Sexpr_t Wire_mop3_endTransmission (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

End the current transmission; optional stop arg; returns #f on success or error string.

Wire_mop3_write()

```
Sexpr_t Wire_mop3_write (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Write a byte, character, string, or bytevector to the I2C transmit buffer; returns bytes written.

11.1.14 WS2812 / NeoPixel LEDs

Functions

- `Sexpr_t mop3_neopixelWrite` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WS2812_mop3_begin` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WS2812_mop3_setBrightness` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WS2812_mop3_setLedColorData` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WS2812_mop3_Wheel` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WS2812_mop3_show` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)
- `Sexpr_t WS2812_install_mop3` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Variables

- `WS2812 * ws2812 = 0`

Detailed Description

LambLisp WS2812 / NeoPixel LEDs builtins.

Function Documentation

`mop3_neopixelWrite()`

`Sexpr_t mop3_neopixelWrite` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Write a single NeoPixel directly: (`neopixelWrite pin r g b`).

`WS2812_mop3_begin()`

`Sexpr_t WS2812_mop3_begin` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Initialize WS2812 strip: (`WS2812.begin num-leds pin channel led-type`).

`WS2812_mop3_setBrightness()`

`Sexpr_t WS2812_mop3_setBrightness` (`Lamb & lamb`, `Sexpr_t sexpr`, `Sexpr_t env_exec`)

Set global brightness scale (0–255) applied to all subsequent `show()` calls.

WS2812_mop3_setLedColorData()

Sexpr_t WS2812_mop3_setLedColorData (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Set LED color data: (WS2812.setLedColorData index rgb) or (... index r g b).

WS2812_mop3_Wheel()

Sexpr_t WS2812_mop3_Wheel (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Map a wheel position (0–255) to an RGB packed integer cycling through the color wheel.

WS2812_mop3_show()

Sexpr_t WS2812_mop3_show (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Push the color buffer to the physical LED strip; returns 0 on success.

11.1.15 CUDA GPU

Functions

- `__global__ void cudakernel_vectorAdd` (const LL_float32 *A, const LL_float32 *B, LL_float32 *C, int n)
- `__global__ void kernel1` (LL_float32 *data, int n)
- `cudaError_t launch_kernel1` (int nblocks, int thrPerBlk, int shared, cudaStream_t stream, LL_float32 *data, int n)
- `__global__ void kernel2` (LL_float32 *data, int n)
- `void cudaDevicePtrDestroyer` (Ptr_t p)
- `void cudaHostPtrDestroyer` (Ptr_t p)
- `int array_matmul_cublas_f32` (const LL_float32 *A, const LL_float32 *B, LL_float32 *C, int m, int k, int n)
- `int array_matmul_cublas_f64` (const LL_float64 *A, const LL_float64 *B, LL_float64 *C, int m, int k, int n)
- `Sexpr_t mop3_cpp_vector_add` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaLaunch_kernel1` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudakernel_vectorAdd` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMalloc` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMallocHost` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaFree` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMemcpy` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMemset` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `void cudaStreamDestroyer` (Ptr_t p)
- `Sexpr_t mop3_cudaStreamCreate` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaStreamSynchronize` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaDeviceSynchronize` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaGetDeviceCount` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaGetDeviceProp` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t Cuda_install_mop3` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Detailed Description

LambLisp CUDA GPU builtins.

Function Documentation

cudaDevicePtrDestroyer()

```
void cudaDevicePtrDestroyer (Ptr_t p)
```

GC destructor for device memory.

cudaHostPtrDestroyer()

```
void cudaHostPtrDestroyer (Ptr_t p)
```

GC destructor for pinned host memory.

array_matmul_cublas_f32()

```
int array_matmul_cublas_f32 (const LL_float32 * A, const LL_float32 * B, LL_float32 * C, int m, int k, int n)
```

Row-major $C(m \times n) = A(m \times k) * B(k \times n)$ via cuBLAS, on HOST buffers. Our arrays are host-resident (shape . data) pairs, so we round-trip through the device: cudaMalloc + H2D + gemm + D2H + cudaFree.

cuBLAS is COLUMN-major but our buffers are ROW-major. Rather than copy/transpose, use the identity $(A*B)^T = B^T * A^T$. A row-major $(m \times k)$ buffer, reinterpreted column-major, already IS $A^T(k \times m)$. So computing, in cuBLAS column-major terms, $C(n \times m) = B^T(n \times k) * A^T(k \times m)$ — i.e. `sgemm(N,N, n, m, k, dB, n, dA, k, dC, n)` — writes exactly our desired row-major $C(m \times n)$ into `dC` with no transposes. Returns 0 on success, nonzero on any CUDA/cuBLAS error.

array_matmul_cublas_f64()

```
int array_matmul_cublas_f64 (const LL_float64 * A, const LL_float64 * B, LL_float64 * C, int m, int k, int n)
```

Double-precision counterpart of `array_matmul_cublas_f32` (uses `cublasDgemm`).

mop3_cpp_vector_add()

```
Sexpr_t mop3_cpp_vector_add (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

CPU-side vector add: $C[i] = A[i] + B[i]$ for i in $[0, n)$. Args: (A B C n) – A, B, C are realvectors; n is element count. Available with or without CUDA; used in three-way correctness tests.

mop3_cudaLaunch_kernel1()

```
Sexpr_t mop3_cudaLaunch_kernel1 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Launch kernel1 (element $\ast = 2$) on a device buffer; args: (stream data [n]).

mop3_cudakernel_vectorAdd()

```
Sexpr_t mop3_cudakernel_vectorAdd (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

GPU vector add $C[i] = A[i] + B[i]$; args: (A B C n) as device/pinned pointers or bytevectors.

mop3_cudaMalloc()

```
Sexpr_t mop3_cudaMalloc (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Allocate n bytes of device memory; returns a T_CPP_HEAP object with cudaFree destructor.

mop3_cudaMallocHost()

```
Sexpr_t mop3_cudaMallocHost (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Allocate n bytes of page-locked host memory; returns a T_CPP_HEAP object with cudaFreeHost destructor.

mop3_cudaFree()

```
Sexpr_t mop3_cudaFree (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Free one or more device memory objects returned by cudaMalloc.

mop3_cudaMemcpy()

```
Sexpr_t mop3_cudaMemcpy (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Combines cudaMemcpy() and cudaMemcpyAsync(). If the stream argument is provided then the asynchronous copy will be used. count is in bytes. dst and src may be device pointers (cudaMalloc), pinned host pointers (cudaMallocHost), or bytevectors.

mop3_cudaMemset()

```
Sexpr_t mop3_cudaMemset (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Fill count bytes of device memory with a byte value: (cudaMemset ptr val count).

mop3_cudaStreamCreate()

```
Sexpr_t mop3_cudaStreamCreate (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Create a CUDA stream; returns a T_CPP_HEAP object with cudaStreamDestroy destructor.

mop3_cudaStreamSynchronize()

```
Sexpr_t mop3_cudaStreamSynchronize (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Block until all operations in the given CUDA stream have completed.

mop3_cudaDeviceSynchronize()

```
Sexpr_t mop3_cudaDeviceSynchronize (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Block until all CUDA work on the current device has completed.

mop3_cudaGetDeviceCount()

```
Sexpr_t mop3_cudaGetDeviceCount (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return the number of CUDA-capable devices available on the host.

mop3_cudaGetDeviceProp()

```
Sexpr_t mop3_cudaGetDeviceProp (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Return an association list of device properties for the given device index (default 0).

11.1.16 HIP / ROCm GPU

Functions

- `__global__ void cudakernel_vectorAdd` (const LL_float32 *A, const LL_float32 *B, LL_float32 *C, int n)
- `__global__ void kernel1` (LL_float32 *data, int n)
- `hipError_t launch_kernel1` (int nblocks, int thrPerBlk, int shared, hipStream_t stream, LL_float32 *data, int n)
- `__global__ void kernel2` (LL_float32 *data, int n)
- `void cudaDevicePtrDestroyer` (Ptr_t p)
- `void cudaHostPtrDestroyer` (Ptr_t p)
- `Sexpr_t mop3_cpp_vector_add` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaLaunch_kernel1` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudakernel_vectorAdd` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMalloc` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMallocHost` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaFree` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMemcpy` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaMemset` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `void cudaStreamDestroyer` (Ptr_t p)
- `Sexpr_t mop3_cudaStreamCreate` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaStreamSynchronize` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaDeviceSynchronize` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaGetDeviceCount` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t mop3_cudaGetDeviceProp` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- `Sexpr_t Hip_install_mop3` (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)

Detailed Description

LambLisp HIP / ROCm GPU builtins.

Function Documentation

cudaDevicePtrDestroyer()

```
void cudaDevicePtrDestroyer (Ptr_t p)
```

GC destructor for device memory.

cudaHostPtrDestroyer()

```
void cudaHostPtrDestroyer (Ptr_t p)
```

GC destructor for pinned host memory.

mop3_cpp_vector_add()

```
Sexpr_t mop3_cpp_vector_add (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

CPU-side vector add: $C[i] = A[i] + B[i]$ for i in $[0, n)$. Args: (A B C n) – A, B, C are realvectors; n is element count. Available with or without CUDA; used in three-way correctness tests.

mop3_cudaLaunch_kernel1()

```
Sexpr_t mop3_cudaLaunch_kernel1 (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Launch kernel1 (element $\ast = 2$) on a device buffer; args: (stream data [n]).

mop3_cudakernel_vectorAdd()

```
Sexpr_t mop3_cudakernel_vectorAdd (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

GPU vector add $C[i] = A[i] + B[i]$; args: (A B C n) as device/pinned pointers or bytevectors.

mop3_cudaMalloc()

```
Sexpr_t mop3_cudaMalloc (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Allocate n bytes of device memory; returns a T_CPP_HEAP object with hipFree destructor.

mop3_cudaMallocHost()

```
Sexpr_t mop3_cudaMallocHost (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Allocate n bytes of page-locked host memory; returns a T_CPP_HEAP object with hipHostFree destructor.

mop3_cudaFree()

```
Sexpr_t mop3_cudaFree (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Free one or more device memory objects returned by hipMalloc.

mop3_cudaMemcpy()

```
Sexpr_t mop3_cudaMemcpy (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Combines hipMemcpy() and hipMemcpyAsync(). If the stream argument is provided then the asynchronous copy will be used. count is in bytes. dst and src may be device pointers (hipMalloc), pinned host pointers (hipHostMalloc), or bytevectors.

mop3_cudaMemset()

```
Sexpr_t mop3_cudaMemset (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Fill count bytes of device memory with a byte value: (hipMemset ptr val count).

mop3_cudaStreamCreate()

```
Sexpr_t mop3_cudaStreamCreate (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Create a CUDA stream; returns a T_CPP_HEAP object with hipStreamDestroy destructor.

mop3_cudaStreamSynchronize()

```
Sexpr_t mop3_cudaStreamSynchronize (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Block until all operations in the given CUDA stream have completed.

mop3_cudaDeviceSynchronize()

```
Sexpr_t mop3_cudaDeviceSynchronize (Lamb & lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

Block until all CUDA work on the current device has completed.

mop3_cudaGetDeviceCount()

```
Sexpr_t mop3_cudaGetDeviceCount (Lamb & lamb, Sexpr_t seexpr, Sexpr_t env_exec)
```

Return the number of CUDA-capable devices available on the host.

mop3_cudaGetDeviceProp()

```
Sexpr_t mop3_cudaGetDeviceProp (Lamb & lamb, Sexpr_t seexpr, Sexpr_t env_exec)
```

Return an association list of device properties for the given device index (default 0).

Hip_install_mop3()

```
Sexpr_t Hip_install_mop3 (Lamb & lamb, Sexpr_t seexpr, Sexpr_t env_exec)
```

Installer: bind the HIP/ROCm GPU builtins into `env_target` (AMD-GPU counterpart of the CUDA bindings); host builds register safe no-ops.

Chapter 12

Class Documentation

12.1 Cell Class Reference

Public Types

Cell type enumeration

<i>Note the following properties of this enumeration:</i>
<i>NIL is a singleton, in the middle of the enumeration.</i>
<i>Types $\leq$ NIL are atoms.</i>
<i>Types $\geq$ NIL are lists.</i>
<i>Types $>$ NIL are pairs.</i>
<i>Types $\leq$ T_PORT_HEAP are heap-allocated objects, requiring specialized finalizing but not specialized marking.</i>
<i>Types $\leq$ T_ANY_HEAP_SVEC are heap allocated vectors requiring specialized marking in addition to finalizing.</i>
<i>For Scheme purposes, there is one pair type, but there are several other specialized pair types that facilitate execution.</i>
<i>These additional pair types are atoms, but list operations such as car, cdr, and append can operate on them.</i>

```

• enum {
  T_SVEC_HEAP = 0 , T_SVEC2N_HEAP , T_ANY_HEAP_SVEC = T_SVEC2N_HEAP ,
  T_SYM_HEAP ,
  T_BVEC_HEAP , T_BIGNUM , T_STR_HEAP , T_CPP_HEAP ,
  T_PORT_HEAP , T_NEEDS_FINALIZING = T_PORT_HEAP , T_BVEC_EXT , T_STR_EXT ,
  T_BVEC_IMM , T_STR_IMM , T_GENSYM , T_BOOL ,
  T_CHAR , T_INT32 , T_INT64 , T_FLOAT32 ,
  T_FLOAT64 , T_COMPLEX , T_RATIONAL , T_MOP3_PROC ,
  T_MOP3_NPROC , T_VOID , T_UNDEF , T_NIL ,
  T_PAIR , T_SVEC_IMM , T_PROC , T_NPROC ,
  T_MACRO , T_DICT , T_THUNK_SEXPR , T_THUNK_BODY ,
  T_ERROR , T_VALUES , T_BYTECODE , T_BYTECODE_N ,
  T_IDENT , Ntypes }

```

Public Member Functions

- `Sexpr_t mk_error` (const char *fmt,...) CHECKPRINTF_pos2
- `Sexpr_t mk_error` (Sexpr_t irritants, const char *fmt,...) CHECKPRINTF_pos3

Static Cell constructors, used only for "well-known" atomic cells (nil, true, false etc).

In general, it is best not to do anything complex at static construction time. There is no guarantee that dependencies will be ready. In particular, use of the terminal may cause the system to crash. Most applications should leave any construction activities to more purposeful code at runtime, using only cells obtained through cons.

- `Cell` (Word_t typ, Word_t w1, Word_t w2)
- `Cell` (Word_t typ, Sexpr_t w1, Sexpr_t w2)

Low-level cell field manipulation.

Note that the type field occupies all of byte 0, although not all bits are required. This speeds up the very common type-checking operation, because no mask is required to extract the type bits.

The flags are in byte 1, and a mask operation is required to access the individual flag fields.

- void `zero` ()
- LL_int32 `type` (void)
- void `type` (LL_int32 t)
- LL_int32 `flags` (void)
- void `flags_set` (LL_int32 f)
- void `flags_clr` (LL_int32 f)
- void `rplaca` (Word_t p)
- void `rplacd` (Word_t p)

word-type accessors — meaningful for T_PAIR cells only.

- int `car_word_type` () const
- int `cdr_word_type` () const
- bool `car_is_ptr` () const
- bool `cdr_is_ptr` () const
- bool `car_is_embedded` () const
- bool `cdr_is_embedded` () const
- bool `is_any_embedded` () const
- LL_int32 `car_int32` () const
- LL_float32 `car_float32` () const
- LL_int32 `cdr_int32` () const
- LL_float32 `cdr_float32` () const

Cell value extractors, dependent on Cell type.

- Ptr_t `get_car_addr` ()
- Word_t `get_car` (void)
- Word_t `get_cdr` (void)
- Bool_t `as_Bool_t` ()
- Char_t `as_Char_t` ()
- Il_codepoint_t `as_codepoint` ()
- LL_int32 `as_int32` () const
- LL_int64 `as_int64` () const
- LL_float32 `as_float32` () const

- `LL_float64 as_float64 ()` const
- `Ptr_t as_Ptr_t ()`
- `Charst_t as_Charst_t ()`
- `Bytest_t as_Bytest_t ()`
- `CharVec_t as_CharVec_t ()`
- `ByteVec_t as_ByteVec_t ()`
- `Portst_t as_Portst_t ()`

Cell hash value

Hashing is used extensively throughout LambLisp. Each *Cell* has a hash value, calculated as follows:

- If the *Cell* is a symbol, the *Cell*'s hash value is the hash value of the symbol.
 - Otherwise, the address of the *Cell* is hashed, and that is the result.
- `LL_int32 hash_sexpr (void)`
 - `LL_int32 hash_contents (void)`
 - `LL_int32 hash (void)`

Cell setters converting the C types into S-expressions.

- `Sexpr_t set (LL_int32 typ, Word_t w1, Word_t w2)`
- `Sexpr_t set (LL_int32 typ, LL_int32 a, Sexpr_t b)`
- `Sexpr_t set (LL_int32 typ, Word_t a, Sexpr_t b)`
- `Sexpr_t set (LL_int32 typ, Sexpr_t a, Sexpr_t b)`
- `Sexpr_t set (Bool_t b)`
- `Sexpr_t set (Char_t c)`
- `Sexpr_t set (LL_int32 i)`
- `Sexpr_t set (LL_int64 v)`
- `Sexpr_t set (LL_float32 r)`
- `Sexpr_t set (LL_float64 d)`
- `Sexpr_t set (LL_float32 r, LL_float32 i)`
- `LL_float32 cpx_real ()` const
- `LL_float32 cpx_imag ()` const
- `Sexpr_t set (LL_int32 num, LL_int32 den)`
- `void rat_set_car_i32 (LL_int32 v)`
- `void rat_set_car_ptr (Sexpr_t p)`
- `void rat_set_cdr_i32 (LL_int32 v)`
- `void rat_set_cdr_ptr (Sexpr_t p)`
- `LL_int64 int_as_i64 ()` const
- `Sexpr_t rat_car_ptr ()` const
- `Sexpr_t rat_cdr_ptr ()` const
- `bool rat_car_big ()` const
- `bool rat_cdr_big ()` const
- `bool rat_is_big ()` const
- `bool exact_is_big ()` const
- `LL_int64 as_numerator ()` const
- `LL_int64 as_denominator ()` const
- `int exact_sign ()`
- `Sexpr_t set (Port_t &p)`
- `Sexpr_t set (LL_int32 typ, LL_int32 a, Charst_t b)`
- `Sexpr_t set (LL_int32 typ, LL_int32 a, Bytest_t b)`
- `Sexpr_t set (LL_int32 typ, LL_int32 a, CharVec_t b)`
- `Sexpr_t set (LL_int32 typ, LL_int32 a, ByteVec_t b)`

Accessors for use when the type is known.

If the cell type is already known, then these accessors can be used to efficiently access the cell contents.

- Sexpr_t **prechecked_anypair_get_car** ()
- Sexpr_t **prechecked_anypair_get_cdr** ()
- LL_int32 **prechecked_sym_heap_get_hash** ()
- Charst_t **prechecked_sym_heap_get_chars** ()
- void **prechecked_sym_heap_get_info** (LL_int32 &hsh, Charst_t &chars)
- CharVec_t **prechecked_str_heap_get_chars** ()
- CharVec_t **prechecked_str_ext_get_chars** ()
- CharVec_t **prechecked_str_imm_get_chars** ()
- void **prechecked_bvec_imm_set_length** (LL_int32 l)
- Charst_t **prechecked_gensym_get_chars** ()
- void **prechecked_gensym_get_info** (LL_int32 &hsh, Charst_t &chars)
- Sexpr_t **prechecked_error_get_irritants** ()
- Sexpr_t **prechecked_error_get_str** ()
- Charst_t **prechecked_error_get_chars** ()

Accessors for use when the cell type is unverified.

The **any** and **mustbe** accessors will perform type checking and throw an error if an improper access is attempted. The **coerce** operators will perform C coercion on its operand if possible, otherwise throw an error.

- Bool_t **mustbe_Bool_t** ()
- Char_t **mustbe_Char_t** ()
- ll_codepoint_t **mustbe_codepoint** ()
- LL_int32 **mustbe_int32** ()
- LL_float32 **mustbe_float32** ()
- Sexpr_t **mustbe_any_str_t** ()
- Sexpr_t **mustbe_cppobj_t** ()
- CPPDeleterPtr **prechecked_cppobj_get_deleter** ()
- Ptr_t **prechecked_cppobj_get_ptr** ()
- CPPDeleterPtr **any_cppobj_get_deleter** ()
- Ptr_t **any_cppobj_get_ptr** ()
- void **any_cppobj_get_info** (CPPDeleterPtr &d, Ptr_t &p)
- LL_float32 **coerce_float32** ()
- LL_float64 **coerce_float64** ()
- LL_int32 **coerce_int32** ()
- LL_int64 **coerce_int64** ()
- int **exact_rank** () const
- int **inexact_rank** () const
- bool **is_exact** () const
- bool **is_inexact** () const
- bool **is_complex** () const
- int **num_rank** () const
- LL_int64 **numerator64** () const
- LL_int64 **denominator64** () const
- LL_float64 **coerce_cpx_real** ()
- LL_float64 **coerce_cpx_imag** ()
- Sexpr_t **error_get_str** ()
- Sexpr_t **error_get_irritants** ()
- Charst_t **error_get_chars** ()
- Sexpr_t **any_error_get_str** ()
- Sexpr_t **any_error_get_irritants** ()
- LL_int32 **any_sym_get_hash** ()
- Charst_t **any_sym_get_chars** ()
- void **any_sym_get_info** (LL_int32 &hsh, Charst_t &chars)

Operations on strings

*LambLisp supports several subtypes of **strings**. At the time of writing, there are heap-allocated strings (read-write), strings whose characters outside of LambLisp's managed memory, (aka external or EXT strings), and short **immediate** strings that are contained completely within a single cell. Additional subtypes (such as load-on-demand strings) may be added in future.*

There are functions of the form `any_xxx()` that can be used with any subtype, and there are type-specific functions of the form `any_xxx_yyy()` for use where the type is already known, with `yyy` being a code hint.

- `CharVec_t any_str_get_chars ()`
- `LL_int32 any_str_get_length ()`
- `void any_str_get_info (LL_int32 &len, CharVec_t &chars)`

Operations on vectors and sub-types of vectors

*Within the LambLisp virtual machine, a Lisp vector is referred to as **svec**. This is an array of S-expressions having a fixed dimension. There are also bytevectors; these are an array of bytes, also of fixed dimension.*

As with strings, there are several subtypes of S-expression vectors. There is a heap-allocated vector, which may be of any size. There is a second type of heap-allocated vector, that is always sized to be a power of 2. These are provided to support efficient hash tables. There are immediate vectors, which may be of 0, 1, or 2 elements. The 2-element immediate vector can also be used as a hash table. This can reduce search time by half without requiring any heap allocation, at the cost of 1 extra cell allocation.

Bytevectors are also diverse, having heap, external, and immediate variants. Heap bytevectors are allocated, obviously, on the system heap, and the heap space is freed when the bytevector is garbage-collected.

External bytevectors operate on bytes provided externally to LambLisp's memory manager. This space may have been dynamically allocated from outside LambLisp, or may be located in read-only memory. When a C++ object is injected into LambLisp, it can optionally be provided with a garbage collector callback; in that case the external object can be automatically garbage collected when no it's longer used in the Lisp program.

*Immediate bytevectors are contained within a LambLisp **Cell**. The maximum size is limited by the word size of the underlying platform.*

*Within the LambLisp virtual machine, there are generic functions of the form `any_xvec_xxx()` that can operate on any *xvec* (*svec* or *bvec*) subtype, as well as type-specific functions of the form `xvec_yyy_xxx()`, where `yyy` is a code hint for the cell storage type (heap, immediate, ROM).*

- `void any_svec_get_info (LL_int32 &Nelems, Sexpr_t *&elems)`
- `Sexpr_t * any_svec_get_elems ()`
- `LL_int32 any_bvec_get_length ()`
- `ByteVec_t any_bvec_get_elems ()`
- `void any_bvec_get_info (LL_int32 &Nelems, ByteVec_t &elems)`
- `LL_int32 bvec_elem_type (void) const`
- `void bvec_set_elem_type (LL_int32 t)`

Cell conversions to printable representation

*These functions convert a **Cell** (or parts of a **Cell**) to a printable representation of the S-expression contents of the **Cell**. Because environments are often included in the descendants of the **Cell** being printed, the depth of environment recursiveness is limited.*

- `String cell_name (void)`
- `Charst_t type_name (LL_int32 typ)`
- `Charst_t type_name (void)`
- `Charst_t gcstate_name (void)`

- String `dump` ()
- String `str` (Sexpr_t sx, Bool_t as_write_or_display, DatumLabels *lbl, LL_int32 max_depth)
- String `str` (Sexpr_t sx, Bool_t as_write_or_display, LL_int32 env_depth, LL_int32 max_depth)
- String `str` (Bool_t as_write_or_display, LL_int32 env_depth, LL_int32 max_depth)
- String `str` (Bool_t as_write_or_display, LL_int32 env_depth)
- String `str` (Bool_t as_write_or_display)
- String `str` (LL_int32 env_depth)
- String `str` (void)

Cell flags, including gc states for multi-pass incremental gc, tail marker, and spares.

- enum `WordType` { `Wt_ptr` =0 , `Wt_int32` =1 , `Wt_float32` =2 , `Wt_reserved` =3 }
- static const int `F_GC01` = 0x01
- static const int `F_GC02` = 0x02
- static const int `F_GC04` = 0x04
- static const int `F_TAIL` = 0x08
- static const int `F_CAR_WT0` = 0x10
- static const int `F_CAR_WT1` = 0x20
- static const int `F_CDR_WT0` = 0x40
- static const int `F_CDR_WT1` = 0x80
- static const int `GC_STATE_MASK` = `F_GC01` | `F_GC02` | `F_GC04`

Flag testing & setting for garbage collection and tail recursion.

The garbage collection algorithm is based on the *tricolor abstraction* described in *Dijkstra 1978*. The original set of 3 colors was enlarged to 4 with *Kung and Song 1977*, with their correctness proof of the incremental GC algorithm. In LambLisp, it is useful to have a fifth state, and to think of the state or color as a stage in the `Cell` life cycle. When GC is in progress, Cells may advance in their life cycle, but may also be moved back to an earlier gc state as the result of an assignment. This enumeration allows some tests to be combined in an inequality rather than a sequence of equality tests or a C `switch`.

LambLisp* uses a **trampoline technique** to implement tail recursion. The **tail** of a series of expressions is the last one evaluated; this is the expression that will return the value of the series.

In the *trampoline*, instead of evaluating the last expression in the series and returning the value (as in C/C++), the expression is returned unevaluated, with the **tail flag** set. The evaluator checks every result to see if it is a *tail* that needs additional evaluation, or is a final result to be returned.

This removes the need for an additional stack frame during recursion.

- enum {
 `gcst_idle` , `gcst_issued` , `gcst_stacked` , `gcst_marked` ,
 `gcst_free` , `Ngcstates` }
- static const int `F_BC_VARIADIC` = 0x01
- LL_int32 `gc_state` (void)
- Sexpr_t `gc_state` (LL_int32 st)
- LL_int32 `tail_state` (void)
- Sexpr_t `tail_state_set` (void)
- Sexpr_t `tail_state_clr` (void)
- bool `bytecode_is_variadic` () const
- void `bytecode_set_variadic` ()

Cell type testing

The cell type enumeration has been designed to group together cell types according to their most common operations.

The main groups are those needing special sweep and finalizing during garbage collection, those that are treated as pairs, and simple atoms. There is a subgroup of those needing special sweep, that also need specialised marking during the garbage collection mark phase.

There are some types that have optimized subtypes (such as **immediate** types), and they differ in their processing at GC time. Type-testing functions below have `is_any_x()` predicates that group all subtypes together (e.g., `is_any_str_atom()` returns **true** for any type of string (heap, immediate, external). There are also functions further down that will return the contents of complex cells, and throw an error if called with incorrect type. Therefore it is often not necessary to use the predicates before accessing the atom internals with `any_x_get_info()`.

Note also: it is sometimes faster to check the types directly, rather than check the cell feature table. Preliminary testing indicates more than 2 type tests should use the table instead.

- static const CellFeatures **features** [Ntypes]
- void **init_static_data** ()
- Bool_t **is_atom** (void)
- Bool_t **is_pair** (void)
- Bool_t **is_ident** (void)
- Bool_t **is_any_pair** (void)
- Bool_t **is_any_svec_atom** ()
- Bool_t **is_any_svec2n_atom** ()
- Bool_t **is_any_str_atom** (void)
- Bool_t **is_any_sym_atom** (void)
- Bool_t **is_any_bvec_atom** (void)
- Bool_t **is_any_int** (void)
- Bool_t **is_bignum** (void)
- Bool_t **is_any_real** (void)

12.1.1 Detailed Description

The `Cell` class is the foundational class for the LambLisp Virtual Machine.

Nearly all the `Cell` methods are *inline* for performance. A `Cell` has only getters and setters; there are no other side effects. `Cell` fields are changed only through explicit requests, and not as the result of any other mutations. This means that Cells do not participate in garbage collection, which is managed from outside the `Cell` class. Indeed, there is no concept within the `Cell` that there might be a plurality of them; only the behavior of a single `Cell` is defined.

A `Cell` may be one of several types. Historically, the number of types varied with the variant of Lisp. In LambLisp, there are `Cell` types that correspond directly to the types described in the *Scheme RxRS* specifications (integers, procedures etc), and there are additional types that implement underlying behavior to support the higher-level language (thunks, dictionaries).

For example, LambLisp supports several types of *strings* internally, depending on whether the character are stored on the heap, in externally-provided memory, or immediately within the cell (for fast operations on short strings). Each of these *string* types is compatible with the string type in the *Scheme RxRS* specifications. Likewise, bytevectors have external and immediate variants.

LambLisp also supports specialized *pair* types, such as the LambLisp *dictionary*. These act as *Lisp* pairs for purposes of allocation and garbage collection, but their specialized operators are executed by the underlying LambLisp virtual machine and therefore operate at C++ speed.

The [Cell](#) type enumeration is purposefully ordered in such a way as to allow efficient integer comparisons instead of a C switch in most cases, enhancing runtime performance as well as garbage collection.

Cell enumeration characteristics:
Cells requiring submarking
Cells requiring finalizing
Simple atoms: immediate cell types like bool char int real etc
External (non-GC) objects and pointers to native C++ code
Cells that point to C++-allocated objects, paired with optional GC finalizers.
"well-known singleton atoms" such as NIL, undefined, and void types.
Pair - the Cell type that responds to the Lisp <i>pair?</i> predicate.
Other pairs - specialized pair types

With this ordering of [Cell](#) types, the solution to many common cases during expression evaluation can be obtained with an inequality, rather than a C++ *switch* statement. This provides a significant performance boost, because the most common cases are checked first, and the total number of cases is reduced.

Cell optimized type tests
(Types < "simple atoms") need specialized GC marking and heap reclamation.
(Types <= "cells requiring finalizing") need specialized GC heap reclamation.
(Types >= T_NIL) are lists.
(Types > T_NIL) are pairs (but not only <i>Lisp</i> pairs, extended pairs too).
(Types != T_PAIR) are atoms, including numbers, strings, vectors, and singletons such as NIL.
(Types > T_PAIR) are extended pairs, used for specialized lists known to LambVM, but receiving regular <i>pair</i> GC processing.

There are cases where the ordering does not help so much. For example the LambLisp type system allows for several types of strings. Because the type system is ordered according to GC requirements, these are not adjacent in the enumeration and less amenable to inequality tests. For these cases, there is a cell feature matrix available to be queried by type and feature. This allows the determination to be made with an array lookup, which is faster than several sequential tests or a C *switch*, putting a permanent cap on the cost of a type check. It also allows other details such as *is-immediate?* to be implemented inexpensively.

The garbage collector states are ordered for similar reasons. See the garbage collector chapter for details on the [Cell](#) life cycle.

12.1.2 Member Enumeration Documentation

anonymous enum

anonymous enum

Enumerator

T_SVEC_HEAP	(LL_int32 Sexpr_t[]) A vector of Sexprs is a (len Sexpr_t*) pair.
T_SVEC2N_HEAP	(LL_int32 Sexpr_t[]) Same as vector, but length must be a power of 2; useful for hash tables.
T_ANY_HEAP_SVEC	Any types less than or equal to this are heap-stored vectors of S-expressions. They need specialized marking at GC time.
T_SYM_HEAP	(LL_int32 Charst_t) Symbol is a (hash char*) pair.
T_BVEC_HEAP	(LL_int32 Bytest_t) Bytevector is a (len byte*) pair.
T_BIGNUM	(LL_int32 Bytest_t) Arbitrary-precision integer; same layout as T_BVEC_↵ HEAP: car=byte-length, cdr=ptr to payload [sign_nlimbs int32][limb0 LL_↵ int32]...
T_STR_HEAP	(reserved Charst_t) The reserved field may be used in future to store the string length (not possible with immediate strings, so some details to be worked).
T_CPP_HEAP	(ptr-to-cpp-deleter ptr-to-cpp-object) Deleter is a function void f(void *cpp_obj) of the appropriate type T of the C++ object, and performs delete (T *) cpp_obj;
T_PORT_HEAP	(reserved Ptr_t) car is unused, cdr is ptr to underlying C++ port instance
T_NEEDS_FINALIZING	Any types less than or equal to this have data stored in the heap, and need specialized finalizing at GC time.
T_BVEC_EXT	(LL_int32 Bytest_t) Same as T_BYTEVEC but externally allocated; the byte array is not freed at GC time.
T_STR_EXT	(reserved Charst_t) Same as T_STR but externally allocated; the character array is not freed at GC time.
T_BVEC_IMM	Special format: type and flag bytes as usual, byte 2 is vector length, remaining bytes are vector elements.
T_STR_IMM	Special format: type and flag bytes as usual, remaining bytes are 0-terminated string embedded in the cell.
T_GENSYM	Runtime symbol generation with no heap operations.
T_BOOL	(Bool_t reserved) Boolean atom
T_CHAR	(Char_t reserved) Character atom
T_INT32	(LL_int32 reserved) Exact 32-bit integer; value in car via memcpy(4).
T_INT64	(LL_int64 wide) Exact 64-bit integer; 8-byte value at &car (spans car+cdr on ESP32). GC-safe: type < T_NIL.
T_FLOAT32	(LL_float32 reserved) Inexact 32-bit LL_float32; value in car via memcpy(4).
T_FLOAT64	(LL_float64 wide) Inexact 64-bit LL_float32; 8-byte value at &car (spans car+cdr on ESP32). GC-safe: type < T_NIL.

T_COMPLEX	(LL_float32 LL_float32) Complex number atom; car=real part (LL_float32), cdr=imaginary part (LL_float32). GC atom: car/cdr not traversed as pointers.
T_RATIONAL	(LL_int32 LL_int32) Exact rational atom; car=numerator, cdr=denominator (always ≥ 2 , always GCD-reduced). GC atom.
T_MOP3_PROC	(reserved *Mop3st_t) pointer to native function - args are evaluated before calling
T_MOP3_NPROC	(reserved *Mop3st_t) pointer to native macro processor - args are not evaluated before calling
T_VOID	(don't care) VOID has its own type and a singleton instance OBJ_VOID
T_UNDEF	(ERROR ERROR) UNDEF has its own type and a singleton instance OBJ_UNDEF
T_NIL	(ERROR ERROR) NIL has its own type and a singleton instance. • The NIL singleton is both a list and an atom, but not a pair. • Types $\leq$ T_NIL are atoms, types $\geq$ T_NIL are lists, types $>$ T_NIL are pairs, but only T_PAIR responds to the Scheme <i>pair?</i> predicate.
T_PAIR	(Sexpr_t Sexpr_t) Normal untyped cons cell. All C++ types $<$ T_PAIR are atoms. • Pair types. Types $>$ NIL and types $\geq$ T_PAIR are pairs, but only T_PAIR responds to Scheme (pair?). • All pair types can be garbage collected without type-specific GC marking or finalizing.
T_SVEC_IMM	(Sexpr_t Sexpr_t) Vector of 0, 1 or 2 elements.
T_PROC	(Sexpr_t Sexpr_t) Procedure pair; car is lambda (formals + body containing free variables), cdr is environment.
T_NPROC	(Sexpr_t Sexpr_t) Non-evaluating procedure pair; car is nlambda (formals + body containing free variables), cdr is environment.
T_MACRO	(Sexpr_t Sexpr_t) Macro; car is transformer (proc of 1 arg) and cdr is env.
T_DICT	(Sexpr_t Sexpr_t) Dictionary pair; car is local frame, cdr is list of parent frames, used for environments & object instances.
T_THUNK_SEXPR	(Sexpr_t Sexpr_t) S-expression thunk pair; car is sexpr, cdr is environment with all variable bindings.
T_THUNK_BODY	(Sexpr_t Sexpr_t) Code body thunk pair; car is body (i.e., list of sexprs), cdr is environment with all variable bindings.
T_ERROR	(Sexpr_t Sexpr_t) An error cell; car is a T_STRING, cdr is a pointer to <i>irritants</i> .
T_VALUES	(Sexpr_t Sexpr_t) Multiple return values; car is first value, cdr is list of remaining values.
T_BYTECODE	(Sexpr_t Sexpr_t) Compiled bytecode procedure; car is T_SVEC_HEAP [bvec,arity,maxloc,maxstk,const...], cdr is closure env.
T_BYTECODE_N	(Sexpr_t Sexpr_t) Compiled bytecode nlambda (non-evaluating); same layout as T_BYTECODE, but caller does NOT evaluate arguments.

T_IDENT	(Sexpr_t Sexpr_t) Explicit-renaming identifier (syntactic closure over ONE id): car=symbol, cdr=definition-env. Pair-region type (GC-traversed as a pair). Made by <code>rename</code> ; free refs resolve in cdr (referential transparency), bound occurrences alpha-rename via gensym. No wraps/marks/substs.
Ntypes	Number of LambLisp virtual machine types.

12.1.3 Member Function Documentation

zero()

```
void Cell::zero () [inline]
```

Set all cell bits to zero.

type() [1/2]

```
LL_int32 Cell::type (void ) [inline]
```

Return the type of the cell as a small integer.

type() [2/2]

```
void Cell::type (LL_int32 t) [inline]
```

Set the type of this cell.

flags()

```
LL_int32 Cell::flags (void ) [inline]
```

Return the entire set of cell flags.

flags_set()

```
void Cell::flags_set (LL_int32 f) [inline]
```

Set the selected flags.

flags_clr()

```
void Cell::flags_clr (LL_int32 f) [inline]
```

Clear the selected flags.

rplaca()

```
void Cell::rplaca (Word_t p) [inline]
```

Replace the cell car field. Called rplaca for historical reasons, and to distinguish it from set-car!, which must respect the GC flags.

rplacd()

```
void Cell::rplacd (Word_t p) [inline]
```

Replace the cell cdr field. Called rplacd for historical reasons, and to distinguish it from set-cdr!, which must respect the GC flags.

is_atom()

```
Bool_t Cell::is_atom (void ) [inline]
```

Return true if the cell is an atom.

is_pair()

```
Bool_t Cell::is_pair (void ) [inline]
```

Return true if the cell is a cons pair.

is_ident()

```
Bool_t Cell::is_ident (void ) [inline]
```

Return true if the cell is an ER renamed identifier (T_IDENT).

is_any_pair()

```
Bool_t Cell::is_any_pair (void ) [inline]
```

Return true if the cell is any pair type.

is_any_svec_atom()

```
Bool_t Cell::is_any_svec_atom () [inline]
```

Return true if the cell is any kind of Sexpr_t vector.

is_any_svec2n_atom()

```
Bool_t Cell::is_any_svec2n_atom () [inline]
```

Return true if the cell is Sexpr_t vector of size 2^n .

is_any_str_atom()

```
Bool_t Cell::is_any_str_atom (void ) [inline]
```

Return true if the cell is any kind of string.

is_any_sym_atom()

```
Bool_t Cell::is_any_sym_atom (void ) [inline]
```

Return true if the cell is any kind of symbol.

is_any_bvec_atom()

```
Bool_t Cell::is_any_bvec_atom (void ) [inline]
```

Return true if the cell is any kind of bytevector.

is_any_int()

```
Bool_t Cell::is_any_int (void ) [inline]
```

Return true if the cell is any exact integer type (T_INT32, T_INT64, or T_BIGNUM).

is_bignum()

```
Bool_t Cell::is_bignum (void ) [inline]
```

Return true if the cell is T_BIGNUM.

is_any_real()

```
Bool_t Cell::is_any_real (void ) [inline]
```

Return true if the cell is any inexact real type (T_FLOAT32 or T_FLOAT64).

car_word_type()

```
int Cell::car_word_type () const [inline]
```

Car word-type: 0=ptr 1=int32 2=float32 3=reserved.

cdr_word_type()

```
int Cell::cdr_word_type () const [inline]
```

Cdr word-type: 0=ptr 1=int32 2=float32 3=reserved.

car_is_ptr()

```
bool Cell::car_is_ptr () const [inline]
```

True when car word holds a [Cell](#) pointer.

cdr_is_ptr()

```
bool Cell::cdr_is_ptr () const [inline]
```

True when cdr word holds a [Cell](#) pointer.

car_is_embedded()

```
bool Cell::car_is_embedded () const [inline]
```

True when car holds an embedded int32 or float32 (single bitwise op).

cdr_is_embedded()

```
bool Cell::cdr_is_embedded () const [inline]
```

True when cdr holds an embedded int32 or float32 (single bitwise op).

is_any_embedded()

```
bool Cell::is_any_embedded () const [inline]
```

True when car or cdr holds an embedded value; single bitwise op on flags byte.

car_int32()

```
LL_int32 Cell::car_int32 () const [inline]
```

Car word as signed int32 (valid when [car_word_type\(\)](#)==Wt_int32).

car_float32()

```
LL_float32 Cell::car_float32 () const [inline]
```

Car word as float32 (valid when [car_word_type\(\)](#)==Wt_float32).

cdr_int32()

```
LL_int32 Cell::cdr_int32 () const [inline]
```

Cdr word as signed int32 (valid when `cdr_word_type()==Wt_int32`).

cdr_float32()

```
LL_float32 Cell::cdr_float32 () const [inline]
```

Cdr word as float32 (valid when `cdr_word_type()==Wt_float32`).

gc_state() [1/2]

```
LL_int32 Cell::gc_state (void ) [inline]
```

Return the garbage collection state of this cell.

gc_state() [2/2]

```
Sexpr_t Cell::gc_state (LL_int32 st) [inline]
```

Set the garbage collection state of this cell, and return the cell.

tail_state()

```
LL_int32 Cell::tail_state (void ) [inline]
```

Return the tail state of this cell.

tail_state_set()

```
Sexpr_t Cell::tail_state_set (void ) [inline]
```

Set the tail state flag on this cell, and return the cell.

tail_state_clr()

```
Sexpr_t Cell::tail_state_clr (void ) [inline]
```

Clear the tail state flag on this cell, and return the cell.

bytecode_is_variadic()

```
bool Cell::bytecode_is_variadic () const [inline]
```

Ph9: true if T_BYTECODE with rest-arg.

bytecode_set_variadic()

```
void Cell::bytecode_set_variadic () [inline]
```

Ph9: mark T_BYTECODE as variadic.

get_car_addr()

```
Ptr_t Cell::get_car_addr () [inline]
```

Return the address of the cell car.

get_car()

```
Word_t Cell::get_car (void ) [inline]
```

Return the value of the cell car.

get_cdr()

```
Word_t Cell::get_cdr (void ) [inline]
```

Return the value of the cell cdr.

as_Bool_t()

```
Bool_t Cell::as_Bool_t () [inline]
```

Return the value of this cell as a boolean.

as_Char_t()

```
Char_t Cell::as_Char_t () [inline]
```

Return the value of this cell as a character.

as_codepoint()

```
ll_codepoint_t Cell::as_codepoint () [inline]
```

Return full Unicode codepoint stored in T_CHAR cell.

as_int32()

```
LL_int32 Cell::as_int32 () const [inline]
```

T_INT32 cell value.

as_int64()

```
LL_int64 Cell::as_int64 () const [inline]
```

T_INT64 cell value; spans car+cdr on ESP32.

as_float32()

```
LL_float32 Cell::as_float32 () const [inline]
```

T_FLOAT32 cell value.

as_float64()

```
LL_float64 Cell::as_float64 () const [inline]
```

T_FLOAT64 cell value; spans car+cdr on ESP32.

as_Ptr_t()

```
Ptr_t Cell::as_Ptr_t () [inline]
```

Return the value of this cell as a generic pointer (i.e., void*).

as_Charst_t()

```
Charst_t Cell::as_Charst_t () [inline]
```

Return the value of this cell as a pointer to a zero-terminated character array.

as_Bytest_t()

```
Bytest_t Cell::as_Bytest_t () [inline]
```

Return the value of this cell as a pointer to an array of bytes.

as_CharVec_t()

```
CharVec_t Cell::as_CharVec_t () [inline]
```

Return the value of this cell as a pointer to a zero-terminated character array.

as_ByteVec_t()

```
ByteVec_t Cell::as_ByteVec_t () [inline]
```

Return the value of this cell as a pointer to an array of bytes.

as_Portst_t()

```
Portst_t Cell::as_Portst_t () [inline]
```

Return the value of this cell as a pointer to an instance of the system underlying "port" implementation.

hash_sexpr()

```
LL_int32 Cell::hash_sexpr (void )
```

For symbols, the hash of its characters; for numbers, the hash of the number; otherwise the hash of the S-expression itself (i.e., the address of a cell).

hash_contents()

```
LL_int32 Cell::hash_contents (void )
```

For symbols, the hash of its characters; for numbers, strings, vectors, and bytevectors, the hash of the contents; otherwise the hash of the S-expression itself.

hash()

```
LL_int32 Cell::hash (void ) [inline]
```

Return the hash value of this cell.

set() [1/17]

```
Sexpr_t Cell::set (LL_int32 typ, Word_t w1, Word_t w2) [inline]
```

This is the lowest-level generic "set" function.

set() [2/17]

```
Sexpr_t Cell::set (LL_int32 typ, LL_int32 a, Sexpr_t b) [inline]
```

Convenience: car is a raw (sign-extended) int32. ATOMS only – never a pair type (a pair car must be a full pointer; see the Word_t/Sexpr_t overload below, added for B182).

set() [3/17]

```
Sexpr_t Cell::set (LL_int32 typ, Word_t a, Sexpr_t b) [inline]
```

B182: a (Word_t,Sexpr_t) call formerly resolved to the (LL_int32,Sexpr_t) overload and SIGN-TRUNCATED the car (e.g. `expand()`'s `set(T_PAIR,(Word_t)NIL,free_list_head)` stored (LL_int32)NIL). This full-width overload wins that resolution.

set() [4/17]

```
Sexpr_t Cell::set (LL_int32 typ, Sexpr_t a, Sexpr_t b) [inline]
```

This is a convenience function for common cases.

set() [5/17]

```
Sexpr_t Cell::set (Bool_t b) [inline]
```

Set cell as boolean.

set() [6/17]

```
Sexpr_t Cell::set (Char_t c) [inline]
```

Set cell as character.

set() [7/17]

```
Sexpr_t Cell::set (LL_int32 i) [inline]
```

Set cell as T_INT32.

set() [8/17]

```
Sexpr_t Cell::set (LL_int64 v) [inline]
```

Set cell as T_INT64 (8 bytes at &car; spans cdr on ESP32).

set() [9/17]

```
Sexpr_t Cell::set (LL_float32 r) [inline]
```

Set cell as T_FLOAT32.

set() [10/17]

```
Sexpr_t Cell::set (LL_float64 d) [inline]
```

Set cell as T_FLOAT64 (8 bytes at &car; spans cdr on ESP32).

set() [11/17]

```
Sexpr_t Cell::set (LL_float32 r, LL_float32 i) [inline]
```

Set cell as T_COMPLEX (single precision).

cpx_real()

```
LL_float32 Cell::cpx_real () const [inline]
```

Real part of a T_COMPLEX cell.

cpx_imag()

```
LL_float32 Cell::cpx_imag () const [inline]
```

Imaginary part of a T_COMPLEX cell.

set() [12/17]

```
Sexpr_t Cell::set (LL_int32 num, LL_int32 den) [inline]
```

< both slots Wt_int32 (preserve gc_state/F_TAIL)

rat_set_car_i32()

```
void Cell::rat_set_car_i32 (LL_int32 v) [inline]
```

embedded int32 numerator (Wt_int32)

rat_set_car_ptr()

```
void Cell::rat_set_car_ptr (Sexpr_t p) [inline]
```

boxed-int numerator ptr (Wt_ptr)

rat_set_cdr_i32()

```
void Cell::rat_set_cdr_i32 (LL_int32 v) [inline]
```

embedded int32 denominator (Wt_int32)

rat_set_cdr_ptr()

```
void Cell::rat_set_cdr_ptr (Sexpr_t p) [inline]
```

boxed-int denominator ptr (Wt_ptr)

int_as_i64()

```
LL_int64 Cell::int_as_i64 () const [inline]
```

value of a boxed T_INT32/T_INT64 (NOT T_BIGNUM)

rat_car_ptr()

```
Sexpr_t Cell::rat_car_ptr () const [inline]
```

boxed numerator cell (valid when [car_word_type\(\)](#)==Wt_ptr)

rat_cdr_ptr()

```
Sexpr_t Cell::rat_cdr_ptr () const [inline]
```

boxed denominator cell (valid when `cdr_word_type()==Wt_ptr`)

rat_is_big()

```
bool Cell::rat_is_big () const [inline]
```

a component is a T_BIGNUM

exact_is_big()

```
bool Cell::exact_is_big () const [inline]
```

does not fit the int64 fast path

as_numerator()

```
LL_int64 Cell::as_numerator () const [inline]
```

Numerator as int64 (guard `rat_is_big()`).

as_denominator()

```
LL_int64 Cell::as_denominator () const [inline]
```

Denominator as int64 (guard `rat_is_big()`).

exact_sign()

```
int Cell::exact_sign () [inline]
```

Sign (-1/0/+1) of any exact number; for a rational, the numerator's sign (denominator is positive).

set() [13/17]

```
Sexpr_t Cell::set (Port_t & p) [inline]
```

Set cell as port.

set() [14/17]

```
Sexpr_t Cell::set (LL_int32 typ, LL_int32 a, Charst_t b) [inline]
```

Set cell as a immutable string.

set() [15/17]

```
Sexpr_t Cell::set (LL_int32 typ, LL_int32 a, Bytest_t b) [inline]
```

Set cell as a immutable bytevector.

set() [16/17]

```
Sexpr_t Cell::set (LL_int32 typ, LL_int32 a, CharVec_t b) [inline]
```

Set cell as a mutable string.

set() [17/17]

```
Sexpr_t Cell::set (LL_int32 typ, LL_int32 a, ByteVec_t b) [inline]
```

Set cell as a mutable bytevector.

mk_error() [1/2]

```
Sexpr_t Cell::mk_error (const char * fmt, ...)
```

Fills and returns the single Cell-level T_ERROR object.

mk_error() [2/2]

```
Sexpr_t Cell::mk_error (Sexpr_t irritants, const char * fmt, ...)
```

Fills and returns the single Cell-level T_ERROR object.

prechecked_str_heap_get_chars()

```
CharVec_t Cell::prechecked_str_heap_get_chars () [inline]
```

Return a pointer to the character array in the heap.

prechecked_str_ext_get_chars()

```
CharVec_t Cell::prechecked_str_ext_get_chars () [inline]
```

Return a pointer to the character array located outside the heap.

prechecked_str_imm_get_chars()

```
CharVec_t Cell::prechecked_str_imm_get_chars () [inline]
```

Return a pointer to the character array embedded in this cell.

mustbe_Bool_t()

```
Bool_t Cell::mustbe_Bool_t () [inline]
```

Return the value of this cell as a boolean.

mustbe_Char_t()

```
Char_t Cell::mustbe_Char_t () [inline]
```

Return the value of this cell as a character.

mustbe_codepoint()

```
ll_codepoint_t Cell::mustbe_codepoint () [inline]
```

Return full codepoint; throw if not T_CHAR.

mustbe_int32()

```
LL_int32 Cell::mustbe_int32 () [inline]
```

Return the value of this T_INT32 cell as int32.

mustbe_float32()

```
LL_float32 Cell::mustbe_float32 () [inline]
```

Return the value of this T_FLOAT32 cell as a real.

mustbe_any_str_t()

```
Sexpr_t Cell::mustbe_any_str_t () [inline]
```

Return this cell if it is a string.

mustbe_cppobj_t()

```
Sexpr_t Cell::mustbe_cppobj_t () [inline]
```

Return this cell if it is a CPP object.

prechecked_cppobj_get_deleter()

```
CPPDeleterPtr Cell::prechecked_cppobj_get_deleter () [inline]
```

Return the function to be called at garbage collection time to recycle the C++ object.

prechecked_cppobj_get_ptr()

```
Ptr_t Cell::prechecked_cppobj_get_ptr () [inline]
```

Return a pointer to a C++ object obtained earlier.

coerce_float32()

```
LL_float32 Cell::coerce_float32 () [inline]
```

< P70: a component may be a boxed bignum -> coerce it via the boxed cell

coerce_float64()

```
LL_float64 Cell::coerce_float64 () [inline]
```

< P70: a component may be a boxed bignum -> coerce it via the boxed cell

exact_rank()

```
int Cell::exact_rank () const [inline]
```

Exact numeric rank: 0=T_INT32 1=T_INT64 2=T_RATIONAL 3=T_BIGNUM -1=not exact.

inexact_rank()

```
int Cell::inexact_rank () const [inline]
```

Inexact numeric rank: 0=T_FLOAT32 1=T_FLOAT64 2=T_COMPLEX -1=not inexact.

is_exact()

```
bool Cell::is_exact () const [inline]
```

True for T_INT32, T_INT64, T_RATIONAL, T_BIGNUM.

is_inexact()

```
bool Cell::is_inexact () const [inline]
```

True for T_FLOAT32, T_FLOAT64, T_COMPLEX.

is_complex()

```
bool Cell::is_complex () const [inline]
```

True iff T_COMPLEX.

num_rank()

```
int Cell::num_rank () const [inline]
```

Legacy scalar rank: 0=int 1=float 2=rational 3=complex -1=non-numeric.

numerator64()

```
LL_int64 Cell::numerator64 () const [inline]
```

Numerator of any exact number as LL_int64.

denominator64()

```
LL_int64 Cell::denominator64 () const [inline]
```

Denominator of any exact number as LL_int64 (1 for integer types).

coerce_cpx_real()

```
LL_float64 Cell::coerce_cpx_real () [inline]
```

Real part: T_COMPLEX → cpx_real; others → coerce_float64.

coerce_cpx_imag()

```
LL_float64 Cell::coerce_cpx_imag () [inline]
```

Imaginary part: T_COMPLEX $\rightarrow$ cpx_imag; others $\rightarrow$ 0.

any_error_get_str()

```
Sexpr_t Cell::any_error_get_str () [inline]
```

`any_error_get_*`: like `error_get_*` but the T_ERROR typecheck is folded in and a non-error yields a benign sentinel (NIL) instead of throwing – lets a predicate write `o->any_error_get_irritants()` without a separate `type()` guard (less source, no missed-check bug).

any_str_get_chars()

```
CharVec_t Cell::any_str_get_chars () [inline]
```

If the cell is any kind of string, return a pointer to the zero-terminated character array.

any_svec_get_info()

```
void Cell::any_svec_get_info (LL_int32 & Nelems, Sexpr_t *& elems) [inline]
```

<B196: NOT OBJ_VOID – that is a real value

bvec_elem_type()

```
LL_int32 Cell::bvec_elem_type (void ) const [inline]
```

Typed-vector element type tag stored in [Cell](#) byte 3 (valid for all bvec subtypes).

bvec_set_elem_type()

```
void Cell::bvec_set_elem_type (LL_int32 t) [inline]
```

Set the typed-vector element type tag in [Cell](#) byte 3.

cell_name()

```
String Cell::cell_name (void )
```

A convenience feature to produce a string name for cells which are "well known" like NIL. Otherwise the name is the hex representation of the cell address.

type_name()

```
Charst_t Cell::type_name (LL_int32 typ)
```

Return a pointer to the C string corresponding to the cell type.

gcstate_name()

```
Charst_t Cell::gcstate_name (void )
```

Return a pointer to a C string corresponding to the given GC state.

dump()

```
String Cell::dump ()
```

Return a printable representation of the [Cell](#) internals.

12.1.4 Member Data Documentation

F_GC01

```
const int Cell::F_GC01 = 0x01 [static]
```

gc flag

F_GC02

```
const int Cell::F_GC02 = 0x02 [static]
```

gc flag

F_GC04

```
const int Cell::F_GC04 = 0x04 [static]
```

gc flag

F_TAIL

```
const int Cell::F_TAIL = 0x08 [static]
```

trampoline tail marker

F_CAR_WT0

```
const int Cell::F_CAR_WT0 = 0x10 [static]
```

T_PAIR car word-type bit 0 (see WordType enum; bits[5:4] = car word-type).

F_CAR_WT1

```
const int Cell::F_CAR_WT1 = 0x20 [static]
```

T_PAIR car word-type bit 1 (Wt_float32 has bit 1 set, bit 0 clear).

F_CDR_WT0

```
const int Cell::F_CDR_WT0 = 0x40 [static]
```

T_PAIR cdr word-type bit 0 (bits[7:6] = cdr word-type).

F_CDR_WT1

```
const int Cell::F_CDR_WT1 = 0x80 [static]
```

T_PAIR cdr word-type bit 1 (Wt_float32 has bit 1 set, bit 0 clear).

GC_STATE_MASK

```
const int Cell::GC_STATE_MASK = F_GC01 | F_GC02 | F_GC04 [static]
```

Mask for obtaining garbage collection state bits from [Cell](#) flag byte.

F_BC_VARIADIC

```
const int Cell::F_BC_VARIADIC = 0x01 [static]
```

byte 2 bit 0: final locals slot holds rest-arg list (T_BYTECODE only).

12.2 Lamb Class Reference

Public Types

- typedef Sexpr_t(* [Mop3st_t](#)) ([Lamb](#) &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
- typedef Sexpr_t(* [Mop3st_argv_t](#)) ([Lamb](#) &lamb, Sexpr_t *argv, int nargs, Sexpr_t env_exec)

Public Member Functions

- Sexpr_t [load](#) (Charst_t name, Sexpr_t env_exec, LL_int32 verbosity=0)
- Sexpr_t [mk_input_string_port](#) (Charst_t inp, Sexpr_t env_exec)
- Sexpr_t [mk_output_string_port](#) (Sexpr_t env_exec)
- Sexpr_t [mk_input_bytevector_port](#) (ByteVec_t src, LL_int32 n, Sexpr_t env_exec)
- Sexpr_t [mk_output_bytevector_port](#) (Sexpr_t env_exec)
- void [diag_memory_report](#) ()

Arduino-compatible loop-based control interface.

- Sexpr_t [setup](#) (void)
- Sexpr_t [loop](#) (void)
- void [end](#) (void)
- ~[Lamb](#) ()

A few foundational functions for embedded debugging.

- void [log](#) (const char *fmt,...) CHECKPRINTF_pos2
- void [printf](#) (const char *fmt,...) CHECKPRINTF_pos2
- bool [debug](#) (void)
- void [debug](#) (bool onoff)

Information about the current build.

- bool `build_isDebug` ()
- unsigned long `build_version` ()
- unsigned long `build_UTC` ()
- unsigned long `build_pushUTC` ()
- const char * `build_buildRelease` ()
- const char * `build_buildDate` ()
- const char * `build_pushDate` ()

Cell constructors

Garbage collection may happen at any time during the execution of the *Cell* constructors. The cell constructors that accept S-expression arguments will protect those arguments from GC during the operation, using the `gc root push/pop` functions.

- void `expand` ()
- Sexpr_t `tcons` (LL_int32 typ, Word_t a, Word_t b, Sexpr_t env_exec)
- Sexpr_t `tcons` (LL_int32 typ, Sexpr_t a, Sexpr_t b, Sexpr_t env_exec)
- Sexpr_t `cons` (Sexpr_t a, Sexpr_t b, Sexpr_t env_exec)
- void `gc_root_push` (Sexpr_t p)
- void `gc_root_pop` (LL_int32 n=1)
- LL_int32 `gc_root_depth` ()
- void `gc_root_setdepth` (LL_int32 n)
- void `gc_idle_task_us` (LL_int32 us, Sexpr_t env_exec)
- void `gc_collect` (Sexpr_t env_exec)
- LL_int32 `gc_Qm` () const
- void `gc_pass_n` (LL_int32 nsteps)
- LL_int32 `mem_cell_block_size` ()
- LL_int32 `mem_n_cell_blocks` ()
- LL_int32 `mem_total_cells` ()
- LL_int32 `mem_ngc` ()
- LL_int32 `mem_markstack_grows` ()
- LL_int32 `mem_urgent_calls` ()
- LL_int32 `mem_urgent_max_us` ()
- void `mem_diag_quiet` (Bool_t q)
- LL_int32 `mem_startmark_max_us` ()
- LL_int32 `mem_startmark_nroots` ()
- LL_int32 `mem_extraroots_max_us` ()
- LL_int32 `mem_extraroots_pushes` ()
- LL_int32 `mem_nfree` ()
- LL_int32 `mem_extension_block_size` ()
- void `mem_set_extension_block_size` (LL_int32 sz)
- void `mem_expand_to_n_blocks` (LL_int32 n, Sexpr_t env_exec)
- LL_int32 `mem_max_cell_blocks` ()
- void `mem_set_max_cell_blocks` (LL_int32 n)
- LL_int32 `mem_gc_budget_ns` ()
- LL_int32 `mem_dbg_urgent_max_us` ()
- LL_int32 `mem_dbg_urgent_calls` ()
- unsigned long `mem_dbg_t_gc_total_us` ()
- LL_int32 `mem_dbg_amax` ()
- LL_int32 `mem_dbg_live_marked` ()
- unsigned long `mem_dbg_n_idle` ()
- unsigned long `mem_dbg_n_work` ()
- LL_int32 `mem_dbg_t_mark_ns` ()
- LL_int32 `mem_dbg_t_sweep_ns` ()
- LL_int32 `mem_dbg_mark_q` ()
- LL_int32 `mem_dbg_sweep_q` ()
- LL_int32 `mem_dbg_yuasa_M` ()
- LL_int32 `mem_dbg_yuasa_N` ()
- void `mem_set_gc_budget_ns` (LL_int32 ns)

- LL_int32 **mem_gcload_target_pct** ()
- void **mem_set_gcload_target_pct** (LL_int32 pct)
- LL_int32 **verbosity** ()
- void **set_verbosity** (LL_int32 v)
- void **set_repl_quiet** (bool q)
- bool **repl_quiet** (void) const
- LL_int32 **ncg_frame_pool_n** ()
- void **ncg_frame_pool_init** (LL_int32 n)

A subset of the car/cdr accessors for list cells. This is the subset used within LambLisp.

- Sexpr_t **car** (Sexpr_t l)
- Sexpr_t **cdr** (Sexpr_t l)
- Sexpr_t **prechecked_anypair_get_car** (Sexpr_t c)
- Sexpr_t **prechecked_anypair_get_cdr** (Sexpr_t c)
- Sexpr_t **caar** (Sexpr_t l)
- Sexpr_t **cadr** (Sexpr_t l)
- Sexpr_t **cdar** (Sexpr_t l)
- Sexpr_t **cddr** (Sexpr_t l)
- Sexpr_t **caddr** (Sexpr_t l)
- Sexpr_t **cdddr** (Sexpr_t l)
- Sexpr_t **cadddr** (Sexpr_t l)
- Sexpr_t **cddddr** (Sexpr_t l)

The "bang" functions are the "mutators" of GC literature.

- void **set_car_bang** (Sexpr_t c, Sexpr_t val)
- void **set_cdr_bang** (Sexpr_t c, Sexpr_t val)
- void **vector_set_bang** (Sexpr_t vec, LL_int32 k, Sexpr_t val)
- Sexpr_t **reverse_bang** (Sexpr_t l)
- Sexpr_t **bulk_alloc_pairs** (LL_int32 n, Sexpr_t env_exec)

Equivalence tests and sequential search

- Sexpr_t **eq_q** (Sexpr_t obj1, Sexpr_t obj2)
- Sexpr_t **eqv_q** (Sexpr_t obj1, Sexpr_t obj2)
- Sexpr_t **equal_q** (Sexpr_t obj1, Sexpr_t obj2)
- Sexpr_t **assq** (Sexpr_t obj, Sexpr_t alist)

Interned symbols

There are just a few operations on symbols:

- *Remember a new symbol.*
- *Check if symbol has already been seen.*
- *Compare 2 symbols for equality.*
- *At evaluation time, lookup the symbol in the current environment.*

For best performance, symbols in LambLisp are interned. This means a single copy of the symbol's characters are stored in a data structure containing the set of all interned symbols. Once interned, symbols can be tested for equality by address rather than character-by-character.

Traditionally, the data structure was called oblist if an association list was used, or called obarray if an array was used. Arrays allow for a hash table implementation, which greatly reduces search time, and can be further optimized if the array size is a power of 2 (2^n).

LambLisp uses the term oblist for the interned symbol table, implemented using a 2^n hash table. Symbol hashes are computed only once and are stored with the symbol, speeding runtime lookups.

- Sexpr_t [oblist_query](#) (Sexpr_t oblist, const char *identifier, Bool_t force, Sexpr_t env_exec)
- Sexpr_t [oblist_test](#) (Sexpr_t oblist, const char *identifier, Sexpr_t env_exec)
- Sexpr_t [oblist_intern](#) (Sexpr_t oblist, const char *identifier, Sexpr_t env_exec)
- Sexpr_t [oblist_analyze](#) (Sexpr_t oblist, LL_int32 verbosity, Sexpr_t env_exec)

Hierarchical Dictionary type

LambLisp's high-performance hierarchical dictionary implementation is used internally to represent the runtime environment. Dictionaries are also directly usable in the Lisp applications, and the dictionary data type is the basis for the LambLisp Object System (LOBS).

- Sexpr_t [dict_new](#) (LL_int32 framesize, Sexpr_t env_exec)
- Sexpr_t [dict_new](#) (Sexpr_t env_exec)
- Sexpr_t [dict_add_empty_frame](#) (Sexpr_t dict, LL_int32 framesize, Sexpr_t env_exec)
- Sexpr_t [dict_add_empty_frame](#) (Sexpr_t dict, Sexpr_t env_exec)
- Sexpr_t [dict_add_keyval_frame](#) (Sexpr_t dict, Sexpr_t keys, Sexpr_t vals, Sexpr_t env_exec)
- Sexpr_t [dict_add_keyval_frame_argv](#) (Sexpr_t dict, Sexpr_t keys, Sexpr_t *argv, int nargs, Sexpr_t env_exec)
- Sexpr_t [dict_add_vector_frame](#) (Sexpr_t parent, Sexpr_t names, Sexpr_t *vals, int ncaps, Sexpr_t env)
- Sexpr_t [dict_add_alist_frame](#) (Sexpr_t dict, Sexpr_t alist, Sexpr_t env_exec)
- Sexpr_t [dict_append_bang](#) (Sexpr_t d1, Sexpr_t d2, Sexpr_t env_exec)
- void [dict_bind_bang](#) (Sexpr_t dict, Sexpr_t key, Sexpr_t value, Sexpr_t env_exec)
- void [dict_bind_local_bang](#) (Sexpr_t dict, Sexpr_t key, Sexpr_t value, Sexpr_t env_exec)
- void [dict_rebind_bang](#) (Sexpr_t dict, Sexpr_t key, Sexpr_t value, Sexpr_t env_exec)
- void [dict_bind_alist_bang](#) (Sexpr_t dict, Sexpr_t alist, Sexpr_t env_exec)
- void [dict_rebind_alist_bang](#) (Sexpr_t dict, Sexpr_t alist, Sexpr_t env_exec)
- Sexpr_t [dict_ref_q](#) (Sexpr_t dict, Sexpr_t key)
- Sexpr_t [dict_ref](#) (Sexpr_t dict, Sexpr_t key)
- Sexpr_t [dict_keys](#) (Sexpr_t dict, Sexpr_t env_exec)
- Sexpr_t [dict_values](#) (Sexpr_t dict, Sexpr_t env_exec)
- Sexpr_t [dict_to_alist](#) (Sexpr_t dict, Sexpr_t env_exec)
- Sexpr_t [dict_to_2list](#) (Sexpr_t dict, Sexpr_t env_exec)
- Sexpr_t [alist_to_dict](#) (Sexpr_t alist, Sexpr_t env_exec)
- Sexpr_t [twolist_to_dict](#) (Sexpr_t twolist, Sexpr_t env_exec)
- Sexpr_t [dict_analyze](#) (Sexpr_t dict, LL_int32 verbosity=0)
- Sexpr_t [integrity_check](#) (Sexpr_t env_exec, LL_int32 verbosity=1)

Reading and writing

- Sexpr_t [read](#) (LL_Port &src, Sexpr_t env_exec)
- Sexpr_t [write_or_display](#) (Sexpr_t sexpr, Bool_t do_write, WriteMode mode=WriteMode::Cyclic)
- Sexpr_t [write_simple](#) (Sexpr_t sexpr)
- String [sprintf](#) (Charst_t fmt, Sexpr_t args, Sexpr_t env_exec)
- Sexpr_t [printf](#) (Sexpr_t args, LL_Port &outp)
- Sexpr_t [printf](#) (Sexpr_t args)

Evaluation and function application

- Sexpr_t [eval](#) (Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [eval_list](#) (Sexpr_t args, Sexpr_t env_exec)
- Sexpr_t [apply_proc_partial](#) (Sexpr_t proc, Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t [map_proc](#) (Sexpr_t proc, Sexpr_t lists, Sexpr_t env_exec)

Querying underlying system features defined by Scheme.

- Sexpr_t **r5_base_environment** ()
- Sexpr_t **r5_interaction_environment** ()
- Sexpr_t **lamb_oblist** ()
- Sexpr_t **current_input_port** ()
- Sexpr_t **current_output_port** ()
- Sexpr_t **current_error_port** ()
- void **set_current_input_port** (Sexpr_t p)
- void **set_current_output_port** (Sexpr_t p)
- void **set_current_error_port** (Sexpr_t p)
- bool **reader_fold_case** ()
- void **reader_fold_case_set** (bool v)

Useful list processing used internally by the LambLisp virtual machine.

- Sexpr_t **append** (Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **append2** (Sexpr_t lis, Sexpr_t obj, Sexpr_t env_exec)
- Sexpr_t **list_copy** (Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **list_analyze** (Sexpr_t sexpr, Sexpr_t env_exec)
- Sexpr_t **list_to_vector** (Sexpr_t l, Sexpr_t env_exec)
- Sexpr_t **vector_to_list** (Sexpr_t v, Sexpr_t env_exec)
- Sexpr_t **vector_copy** (Sexpr_t from, Sexpr_t env_exec)

Sparse vectors

The sparse vector representation is a sorted association list, in which the car of each pair is the vector index and the cdr is the vector value at that index.

- Sexpr_t **vector_to_sparsevec** (Sexpr_t vec, Sexpr_t skip, Sexpr_t env_exec)
- Sexpr_t **sparsevec_to_vector** (Sexpr_t alist, Sexpr_t fill, Sexpr_t env_exec)

Makers for error types

These functions will `sprintf` a message into a buffer, and return an error object containing the message. They differ in that one (`mk_error`) creates a new error object, and the other (`mk_syserror`) avoids use of the memory manager by reusing the system error singleton.

- Sexpr_t **mk_error** (Sexpr_t env_exec, Sexpr_t irritants, const char *fmt,...) CHECKPRINTF_pos4
- Sexpr_t **mk_error** (Sexpr_t env_exec, const char *fmt,...) CHECKPRINTF_pos3
- Sexpr_t **mk_syserror** (const char *fmt,...) CHECKPRINTF_pos2
- Sexpr_t **tag_error** (Sexpr_t err, const char *kind)
- Sexpr_t **mk_syserror** (Sexpr_t irritants, const char *fmt,...) CHECKPRINTF_pos3

Makers for simple compact types with embedded values, without external storage or dependencies

These types all have their values embedded in the `Cell`. Note that other types may also hold immediate values, but are "less simple" because they are C++ subtypes that are variations on the Lisp-declared types.

- Sexpr_t **gensym** (Sexpr_t env_exec)
- Sexpr_t **mk_bool** (Bool_t b, Sexpr_t env_exec)
- Sexpr_t **mk_character** (ll_codepoint_t cp, Sexpr_t env_exec)
- Sexpr_t **mk_integer** (LL_int64 i, Sexpr_t env_exec)
- Sexpr_t **mk_int32** (LL_int32 i, Sexpr_t env_exec)
- Sexpr_t **mk_int64** (LL_int64 i, Sexpr_t env_exec)
- Sexpr_t **mk_float32** (LL_float32 r, Sexpr_t env_exec)

- `Sexpr_t mk_float64` (LL_float64 d, Sexpr_t env_exec)
- `Sexpr_t mk_complex` (LL_float32 r, LL_float32 i, Sexpr_t env_exec)
- `void rat_store_i64` (Sexpr_t c, bool is_num, LL_int64 v, Sexpr_t env_exec)
- `Sexpr_t mk_rational` (LL_int64 num, LL_int64 den, Sexpr_t env_exec)
- `Sexpr_t mk_number` (Charst_t str, Sexpr_t env_exec)
- `Sexpr_t mk_sharp_const` (Charst_t name, Sexpr_t env_exec)

Makers for heap storage types

In principle, these types require space from the system heap. At GC time, they may also require specialized marking, and also specialized sweeping.

*As an optimization, LambLisp also implements **immediate** types. When the size required for the object is small, the internal bytes of the [Cell](#) can be used to hold the contents, avoiding heap operations. Using immediate types in this way has been common practice for integers and real numbers. LambLisp makes some additional use of immediate types. Immediate types are a LambLisp internal optimization, and they are all subtypes of the main Lisp type.*

The types that employ the immediate optimization are:

- *strings, providing up to 9 content bytes plus trailing 0 (32 bit word) or 22 content bytes (64 bit).*
- *bytevectors a count plus up to 9 or 22 content bytes.*
- *vectors of length 0, 1 or 2.*

*There are also **EXTERNAL** versions of those data types. These data types can be used to access data that should not be garbage collected, such as data contained in the application image, or C++ objects provided by another application.*

A variant of the EXTERNAL types is also provided, so that objects created in C++ can be passed to LambLisp, along with an optional deleter function. In C++, the deleter is best implemented as a one-line C++ lambda function accepting 1 argument (a void pointer to a C++ object). The deleter function casts the void* pointer to match the C++ object type, and applies the C++ delete or delete[] operators to free the object space and any dependent resources.*

By passing the C++ objects to LambLisp, they can be automatically garbage collected when LambLisp has finished with them.

- `Sexpr_t mk_string` (Sexpr_t env_exec, const char *fmt,...) CHECKPRINTF_pos3
- `void *payload_alloc` (size_t nbytes, const char *who, Sexpr_t env_exec)
- `Sexpr_t mk_string` (LL_int32 k, Charst_t src, Sexpr_t env_exec)
- `Sexpr_t mk_string` (String s, Sexpr_t env_exec)
- `Sexpr_t mk_symbol_or_number` (Charst_t str, Sexpr_t env_exec)
- `Sexpr_t mk_symbol` (Charst_t str, Sexpr_t env_exec)
- `Sexpr_t mk_bytevector` (LL_int32 k, Sexpr_t env_exec)
- `Sexpr_t mk_bytevector` (LL_int32 k, Bytest_t src, Sexpr_t env_exec)
- `Sexpr_t mk_bytevector` (LL_int32 k, LL_int32 fill, Sexpr_t env_exec)
- `Sexpr_t mk_bytevector_ext` (LL_int32 k, Bytest_t ext, Sexpr_t env_exec)
- `Sexpr_t mk_intvector` (LL_int32 k, Sexpr_t env_exec)
- `Sexpr_t mk_realvector` (LL_int32 k, Sexpr_t env_exec)
- `Sexpr_t mk_intvector` (LL_int32 k, LL_int32 fill, Sexpr_t env_exec)
- `Sexpr_t mk_realvector` (LL_int32 k, LL_float32 fill, Sexpr_t env_exec)
- `Sexpr_t mk_vector` (LL_int32 len, Sexpr_t fill, Sexpr_t env_exec)
- `Sexpr_t mk_hashtbl` (LL_int32 len, Sexpr_t fill, Sexpr_t env_exec)
- `Sexpr_t mk_serial_port` (Sexpr_t env_exec)
- `Sexpr_t mk_input_file_port` (Charst_t name, Sexpr_t env_exec)
- `Sexpr_t mk_output_file_port` (Charst_t name, Sexpr_t env_exec)

Network port makers (TCP and TLS)

These makers are available when `LL_WIFI` (ESP32 WiFi) or `LL_POSIX` (Linux/macOS) is defined. They are exposed in the Scheme interaction environment as `open-tcp-client-port`, `open-tcp-server-port`, `server-accept`, `open-tls-client-port`, and `open-tls-server-port`.

All network ports are bidirectional `T_PORT_HEAP` objects and work with the standard R7RS I/O procedures (`read`, `write`, `read-char`, `write-string`, `close-port`, etc.).

The LLIP (LambLisp Interaction Protocol) layer (`llip-server.scm` / `llip-client.scm`) is built entirely on top of these primitives. It speaks a line-oriented S-expression protocol over plain TCP (development) or TLS (production).

Wire protocol summary:

- Server sends: (`hello "llip/1.0"`) on connect.
- Client sends: (`auth "token"`). Server replies (`ok`) or closes.
- Each request is one S-expression per line; response is (`ok ...`) or (`error "msg"`).
- Supported ops: `ping`, `bye`, `ls`, `stat`, `read`, `write`, `append`, `delete`, `rename`, `load`, `eval`.
- (`bye`) triggers a graceful close; EOF from the client also closes the session.

See also

`data/llip-server.scm` – Scheme server (load on the target board).

`data/llip-client.scm` – Scheme client (load on the host or peer board).

`ll_tests/llip-tests.scm` – integration test suite.

- `Sexpr_t mk_tcp_client_port` (`Charst_t host`, unsigned port, `Sexpr_t env_exec`)
- `Sexpr_t mk_tcp_server_port` (unsigned port, `Sexpr_t env_exec`)
- `Sexpr_t server_accept_port` (`LL_Port *server`, `Sexpr_t env_exec`)
- `Sexpr_t mk_tls_client_port` (`Charst_t host`, unsigned port, `Charst_t ca_cert_pem`, `Sexpr_t env_exec`)
- `Sexpr_t mk_tls_server_port` (unsigned port, `Charst_t cert_pem`, `Charst_t key_pem`, `Sexpr_t env_↔exec`)

Makers for interface to native procedures and C++ objects.

- `Sexpr_t mk_Mop3_procst_t` (`Mop3st_t f`, `Sexpr_t env_exec`)
- `Sexpr_t mk_Mop3_procst_argv_t` (`Mop3st_t f`, `Mop3st_argv_t fa`, `Sexpr_t env_exec`)
- `Sexpr_t mk_Mop3_nprocst_t` (`Mop3st_t f`, `Sexpr_t env_exec`)
- `Sexpr_t mk_cppobj` (`void *obj`, `CPPDeleterPtr deleter`, `Sexpr_t env_exec`)

Makers for pair types.

These **pair** types all have the same structure for construction and garbage collection purposes, but are not Scheme pair types. LambLisp*'s scalable type system presents opportunities for efficiency when implementing common operations on lists. These are executed by the LambLisp virtual machine at C++ speed, rather than at the slower speed of the Lisp evaluation loop.

- *procedures ... the evaluator must be able to identify these directly.*
- *dictionaries ... often containing cycles when used as environments, so a depth limit is imposed when traversing.*
- *thunks ... used in combination with the trampoline technique to implement tail recursion.*

Note that there is no need for `mk_pair()`, it is just `cons()`.

- `Sexpr_t mk_macro` (`Sexpr_t nlam`, `Sexpr_t env_nlam`, `Sexpr_t env_exec`)

- `Sexpr_t mk_ident` (`Sexpr_t sym`, `Sexpr_t def_env`, `Sexpr_t env_exec`)
- `Sexpr_t mk_procedure` (`Sexpr_t formals`, `Sexpr_t body`, `Sexpr_t env_proc`, `Sexpr_t env_exec`)
- `Sexpr_t mk_nprocedure` (`Sexpr_t formals`, `Sexpr_t body`, `Sexpr_t env_nproc`, `Sexpr_t env_exec`)
- `Sexpr_t mk_thunk_sexpr` (`Sexpr_t sexpr`, `Sexpr_t env_thunk`, `Sexpr_t env_exec`)
- `Sexpr_t mk_thunk_body` (`Sexpr_t body`, `Sexpr_t env_thunk`, `Sexpr_t env_exec`)
- `Sexpr_t mk_dict` (`Sexpr_t frame`, `Sexpr_t base`, `Sexpr_t env_exec`)
- `Sexpr_t mk_dict` (`LL_int32 framesize`, `Sexpr_t env_exec`)

Macro support

These macro primitives behave as described in Steele's Common Lisp.

- `Sexpr_t macroexpand` (`Sexpr_t form`, `Sexpr_t env_exec`)
- `Sexpr_t macroexpand1` (`Sexpr_t form`, `Sexpr_t env_exec`)
- `Sexpr_t macroexpand` (`Sexpr_t proc`, `Sexpr_t args`, `Sexpr_t env_exec`)
- `Sexpr_t macroexpand1` (`Sexpr_t proc`, `Sexpr_t args`, `Sexpr_t env_exec`)

Native Code Generator (NCG)

- `NcgFrame * ncg_frame_top` = `nullptr`
- `NcgFrame * ncg_frame_pool_head` = `nullptr`
- `int ncg_frame_pool_count` = 0
- `NcgErrJmp * ncg_err_ctx` = `nullptr`
- `void ncg_mark_frames` ()
- `Sexpr_t ncg_eval` (`Sexpr_t bc`, `Sexpr_t args`, `Sexpr_t env_exec`)
- `Sexpr_t ncg_eval_argv` (`Sexpr_t bc`, `Sexpr_t *argv`, `int nargs`, `Sexpr_t env_exec`)
- `Sexpr_t ncg_compile` (`Sexpr_t bc`, `Sexpr_t env_exec`, `Sexpr_t captured_names=nullptr`)
- `void ncg_disassemble` (`Sexpr_t bc`, `Sexpr_t env_exec`)
- `void ncg_vector_set_ptr` (`Sexpr_t *data_ptr`, `LL_int32 k`, `Sexpr_t val`)
- `Sexpr_t ncg_vref_fast` (`Sexpr_t vec`, `Sexpr_t idx`)
- `Sexpr_t ncg_vset_fast` (`Sexpr_t vec`, `Sexpr_t idx`, `Sexpr_t val`)
- `Sexpr_t ncg_cons_fast` (`Sexpr_t a`, `Sexpr_t d`)
- `Sexpr_t ncg_box_fast` (`LL_int64 v`)

native-recursion guard (catch C-stack overflow before it SIGSEGVs)

Both the interpreter (`eval`) and the compiled path (`ncg_eval/ncg_eval_argv`) recurse on the native C stack; unbounded/non-tail Scheme recursion overflows it into the OS guard page = an uncatchable SIGSEGV. Stack grows down, so the outermost frame has the highest address; we track that high-water mark as the base and, once a recursive call has grown more than `eval_stack_budget_bytes` below it, raise a normal (catchable) "recursion too deep" error. `eval_recursion_check()` is inlined into the hot recursion entry points; the throw lives in the cold out-of-line `eval_recursion_overflow()`.

- `const char * eval_stack_base_` = `nullptr`
- `size_t eval_stack_budget_` = `LL_EVAL_STACK_BUDGET`
- `LL_int32 eval_root_depth_max_` = `LL_EVAL_ROOT_DEPTH_MAX`
- `void eval_recursion_overflow` (`const char *where`)
- `void eval_recursion_check` (`const char *where`)

12.2.1 Detailed Description

The `Lamb` class represents a single Lisp virtual machine. To closely match the Arduino-style of loop-based control, LambLisp provides `begin()`, `loop()`, and `end()`.

To initialize the LambLisp VM, `begin()` should be called once before any other call. Generally LambLisp `loop()` will be called one time for every main `loop()`, but it is not necessary. If `end()` is called, all LambLisp resources are freed; `Lamb::setup()` can be called again to restart that LambLisp VM.

The LambLisp VM functions are grouped this way:

- Logging functions
- Build and version informational functions.
- `Cell` constructors, getters & setters
- Base data structures for the interpreter: alist, hash tables, and environments, along with their getters & setters.
- Printers
- Partial evaluation and application
- Port access from *Lisp*
- Vector, list, and dictionary utilities, used internally and also available in *Lisp*.
- GC interface functions to protect critical sections.
- Bindings
- Makers

12.2.2 Member Typedef Documentation

Mop3st_t

```
typedef Sexpr_t(* Lamb::Mop3st_t) (Lamb &lamb, Sexpr_t sexpr, Sexpr_t env_exec)
```

This is a pointer to a C++ native function that interacts directly with the S-expression in the given environment. Every LambLisp native function shares this signature. New external functions that conform to this signature can be used from within LambLisp and will run at full speed.

Mop3st_argv_t

```
typedef Sexpr_t(* Lamb::Mop3st_argv_t) (Lamb &lamb, Sexpr_t *argv, int nargs, Sexpr_t env_exec)
```

direct-argv calling convention; `argv[0..nargs-1]` are already-evaluated args. GC CONTRACT for argv variants: the caller (`apply_fast1/2_argv`, `ncg`) passes a C-array `argv` that is NOT on the root stack, so a variant MUST read every `argv[i]` it needs into a local BEFORE its first allocation (arg cells are otherwise unrooted and can be swept mid-call – B43). To build a Scheme list from argv, use `argv_to_list()`, which roots all elements for the build.

12.2.3 Constructor & Destructor Documentation

~Lamb()

```
Lamb::~~Lamb ()
```

Real teardown: finalize every cell, free the blocks and the memory manager. Pairs with `new Lamb + setup()`. Written 2026-08-24; before that the class had no destructor at all and nothing ever tore an instance down, so the leaks below had never been exercised.

12.2.4 Member Function Documentation

setup()

```
Sexpr_t Lamb::setup (void )
```

Run once after base platform has started. In particular, *Serial* should be initialized before calling `setup()`.

loop()

```
Sexpr_t Lamb::loop (void )
```

Run often to maintain control. The main responsibility of C++ `Lamb::loop()` is to call the LambLisp function of the same name (`loop`).

end()

```
void Lamb::end (void )
```

Reset roots and force a full collection. NOTE: despite the name this does NOT release anything – the pool, the blocks and the payloads all stay. Use the destructor for real teardown.

log()

```
void Lamb::log (const char * fmt, ...)
```

C-level printf feature with a limit on the length of strings produced. Takes care of log prompt.

printf() [1/3]

```
void Lamb::printf (const char * fmt, ...)
```

C-level printf feature with a limit on the length of strings produced.

debug() [1/2]

```
bool Lamb::debug (void )
```

Return the state of the [Lamb](#) internal debug flag.

debug() [2/2]

```
void Lamb::debug (bool onoff)
```

Set the state of the [Lamb](#) debug flag.

build_isDebug()

```
bool Lamb::build_isDebug ()
```

Return true if this build is not checked in. This is unrelated to the runtime [debug](#) flag.

build_version()

```
unsigned long Lamb::build_version ()
```

Return the build version as a long integer (implemented as a UTC time stamp).

build_UTC()

```
unsigned long Lamb::build_UTC ()
```

Return the UTC time of this build.

build_pushUTC()

```
unsigned long Lamb::build_pushUTC ()
```

Return the UTC time this repo was last pushed.

build_buildRelease()

```
const char * Lamb::build_buildRelease ()
```

Return a pointer to a character array containing the release description.

build_buildDate()

```
const char * Lamb::build_buildDate ()
```

Return a pointer to a character array containing the build date.

build_pushDate()

```
const char * Lamb::build_pushDate ()
```

Return a pointer to a character array containing the date this repo was last pushed.

expand()

```
void Lamb::expand ()
```

Add another block of cells to the population.

tcons() [1/2]

```
Sexpr_t Lamb::tcons (LL_int32 typ, Word_t a, Word_t b, Sexpr_t env_exec)
```

Generic cell constructor with no GC protection for its arguments.

tcons() [2/2]

```
Sexpr_t Lamb::tcons (LL_int32 typ, Sexpr_t a, Sexpr_t b, Sexpr_t env_exec)
```

Constructor for any pair type; protects both S-expression arguments.

cons()

```
Sexpr_t Lamb::cons (Sexpr_t a, Sexpr_t b, Sexpr_t env_exec) [inline]
```

[Cell](#) constructor for normal pairs.

gc_root_push()

```
void Lamb::gc_root_push (Sexpr_t p)
```

Preserve the cell given from GC, until popped.

gc_root_pop()

```
void Lamb::gc_root_pop (LL_int32 n = 1)
```

Release the preserved cells for normal GC processing.

gc_root_depth()

```
LL_int32 Lamb::gc_root_depth ()
```

Return current rootstack depth (save before a try block).

gc_root_setdepth()

```
void Lamb::gc_root_setdepth (LL_int32 n)
```

Restore rootstack to saved depth, discarding entries above n (call in catch to drop orphaned roots without disturbing the caller's entries).

gc_idle_task_us()

```
void Lamb::gc_idle_task_us (LL_int32 us, Sexpr_t env_exec)
```

Run GC for up to us microseconds of spare time; bank extra quanta as credit to skip during later allocations.

gc_collect()

```
void Lamb::gc_collect (Sexpr_t env_exec)
```

Run one complete GC cycle (mark + sweep); leaves heap in idle phase with all unreachable cells freed.

gc_Qm()

```
LL_int32 Lamb::gc_Qm () const
```

return current mark quantum size.

gc_pass_n()

```
void Lamb::gc_pass_n (LL_int32 nsteps)
```

run nsteps mark steps; no-op if not in marking phase.

mem_dbg_t_gc_total_us()

```
unsigned long Lamb::mem_dbg_t_gc_total_us ()
```

monotonic total us in gc_pass – measured GC load

set_repl_quiet()

```
void Lamb::set_repl_quiet (bool q)
```

Quiet REPL mode: suppress the per-char redisplay prompt reprint and the "Input:" echo. NOT platform-specific despite its old home – the terminal reaches the platform only through LambStdio (LL_Term::get↵_input_serial/get_output_serial both return it), and the flag's readers are the REPL echo in ll_vm_lamb.cpp and the (quiet)/(verbose) procs. Declared HERE, defined in ll_vm_term.cpp beside the ll_term singleton that stores it: a shipped ll_xmop3_*.cpp must be able to set it WITHOUT including a stripped ll_vm_* header, which is what made every customer package unbuildable from source (2026-08-26).

ncg_frame_pool_n()

```
LL_int32 Lamb::ncg_frame_pool_n () [inline]
```

Frames currently in the free-list pool.

ncg_frame_pool_init()

```
void Lamb::ncg_frame_pool_init (LL_int32 n)
```

Allocate n frames and push them onto the pool free list.

car()

```
Sexpr_t Lamb::car (Sexpr_t c) [inline]
```

Return the first element (*car*) of a pair cell; throws if 1 is not a pair type.

gc_protect idiom for the inline [Lamb](#) members below. `ll_gc_protect` (`ll_vm_mem.h`) is mem->rootpush-based and is NOT in scope here – [LambLisp.h](#) is included before `ll_vm_mem.h` in every TU – so this Lamb-method variant gives the same scoped push/code/pop shape for `car()/cdr()`. Body must be multi-line if it contains a top-level comma (same preprocessor rule as `ll_gc_protect`).

cdr()

```
Sexpr_t Lamb::cdr (Sexpr_t l) [inline]
```

Return the second element (*cdr*) of a pair cell; throws if 1 is not a pair type.

prechecked_anypair_get_car()

```
Sexpr_t Lamb::prechecked_anypair_get_car (Sexpr_t c) [inline]
```

Car of a known pair; handles embedded int32/float32; no type check.

prechecked_anypair_get_cdr()

```
Sexpr_t Lamb::prechecked_anypair_get_cdr (Sexpr_t c) [inline]
```

Cdr of a known pair; handles embedded int32/float32; no type check.

caar()

```
Sexpr_t Lamb::caar (Sexpr_t l) [inline]
```

Composition: `car(car(l))`.

cadr()

```
Sexpr_t Lamb::cadr (Sexpr_t l) [inline]
```

Composition: `car(cdr(l))`; returns the second element of a list.

cdar()

```
Sexpr_t Lamb::cdar (Sexpr_t l) [inline]
```

Composition: `cdr(car(l))`.

cddr()

```
Sexpr_t Lamb::cddr (Sexpr_t l) [inline]
```

Composition: `cdr(cdr(l))`.

caddr()

```
Sexpr_t Lamb::caddr (Sexpr_t l) [inline]
```

Composition: `car(cdr(cdr(l)))`; returns the third element of a list.

cdddr()

```
Sexpr_t Lamb::cdddr (Sexpr_t l) [inline]
```

Composition: `cdr(cdr(cdr(l)))`.

caddr()

```
Sexpr_t Lamb::caddr (Sexpr_t l) [inline]
```

Composition: `car(cdr(cdr(cdr(l))))`; returns the fourth element of a list.

cddddr()

```
Sexpr_t Lamb::cddddr (Sexpr_t l) [inline]
```

Composition: `cdr(cdr(cdr(cdr(l))))`.

set_car_bang()

```
void Lamb::set_car_bang (Sexpr_t c, Sexpr_t val)
```

Replace the car field in the cell with *val*. GC flags will be maintained as required.

set_cdr_bang()

```
void Lamb::set_cdr_bang (Sexpr_t c, Sexpr_t val)
```

Replace the cdr field in the cell with *val*. GC flags will be maintained as required.

vector_set_bang()

```
void Lamb::vector_set_bang (Sexpr_t vec, LL_int32 k, Sexpr_t val)
```

Replace the specified vector element with *val*. GC flags will be maintained as required.

reverse_bang()

```
Sexpr_t Lamb::reverse_bang (Sexpr_t l)
```

Reverse the list in-place and return the new list head (the former list tail).

bulk_alloc_pairs()

```
Sexpr_t Lamb::bulk_alloc_pairs (LL_int32 n, Sexpr_t env_exec)
```

Grab n free T_PAIR cells atomically; mark gcst_issued during mark/sweep; return NIL-terminated chain.

eq_q()

```
Sexpr_t Lamb::eq_q (Sexpr_t obj1, Sexpr_t obj2)
```

Return true if 2 cells are the same cell, or are atoms with the same value.

eqv_q()

```
Sexpr_t Lamb::eqv_q (Sexpr_t obj1, Sexpr_t obj2)
```

Return true if 1 cells are the same cell, or are atoms with the same value.

equal_q()

```
Sexpr_t Lamb::equal_q (Sexpr_t obj1, Sexpr_t obj2)
```

Returns true when obj1 and obj2 are eqv?, and also all their descendants.

assq()

```
Sexpr_t Lamb::assq (Sexpr_t obj, Sexpr_t alist)
```

Search an association list for a matching key, and return the (key . value) pair, or **false** if not found.

oblist_query()

```
Sexpr_t Lamb::oblist_query (Sexpr_t oblist, const char * identifier, Bool_t force, Sexpr_t env_exec)
```

Look up identifier in oblist. If force is true, intern the symbol (creating it if absent) and return it. If force is false, return the existing symbol or HASHF if not found.

oblist_test()

```
Sexpr_t Lamb::oblist_test (Sexpr_t oblist, const char * identifier, Sexpr_t env_exec) [inline]
```

Return the interned symbol for identifier, or HASHF if not present.

oblist_intern()

```
Sexpr_t Lamb::oblist_intern (Sexpr_t oblist, const char * identifier, Sexpr_t env_exec) [inline]
```

Return the interned symbol for identifier, creating it if necessary.

oblist_analyze()

```
Sexpr_t Lamb::oblist_analyze (Sexpr_t oblist, LL_int32 verbosity, Sexpr_t env_exec)
```

Report statistics on the oblist hash table (load factor, chain lengths, etc.) at the given verbosity level. Returns NIL.

dict_new() [1/2]

```
Sexpr_t Lamb::dict_new (LL_int32 framesize, Sexpr_t env_exec) [inline]
```

Return a new empty dictionary with the given top frame size.

dict_new() [2/2]

```
Sexpr_t Lamb::dict_new (Sexpr_t env_exec) [inline]
```

Return a new empty dictionary with alist top frame.

dict_add_empty_frame() [1/2]

```
Sexpr_t Lamb::dict_add_empty_frame (Sexpr_t dict, LL_int32 framesize, Sexpr_t env_exec)
```

Prepend a new empty frame to dict and return the extended dictionary. Frame storage is chosen by framesize: < 4 = alist; 4–7 = 2-slot hash table; >= 8 = heap-allocated hash table of framesize slots.

dict_add_empty_frame() [2/2]

```
Sexpr_t Lamb::dict_add_empty_frame (Sexpr_t dict, Sexpr_t env_exec) [inline]
```

Returns a new dictionary with an empty alist top frame and dict as parent.

dict_add_keyval_frame()

```
Sexpr_t Lamb::dict_add_keyval_frame (Sexpr_t dict, Sexpr_t keys, Sexpr_t vals, Sexpr_t env_exec)
```

Returns a new dictionary with a new top frame containing the keys bound to the values.

dict_add_keyval_frame_argv()

```
Sexpr_t Lamb::dict_add_keyval_frame_argv (Sexpr_t dict, Sexpr_t keys, Sexpr_t * argv, int nargs, Sexpr_t env_exec)
```

same as dict_add_keyval_frame but values come from argv[] instead of a cons list.

dict_add_vector_frame()

```
Sexpr_t Lamb::dict_add_vector_frame (Sexpr_t parent, Sexpr_t names, Sexpr_t * vals, int ncaps, Sexpr_t env)
```

compact kv-vec frame (T_SVEC_IMM for ncaps==1); alist fallback for ncaps>1.

dict_add_alist_frame()

```
Sexpr_t Lamb::dict_add_alist_frame (Sexpr_t dict, Sexpr_t alist, Sexpr_t env_exec) [inline]
```

Returns a new dictionary with the alist bindings added in a new frame on top of the base dictionary.

dict_append_bang()

```
Sexpr_t Lamb::dict_append_bang (Sexpr_t d1, Sexpr_t d2, Sexpr_t env_exec)
```

Append d2 as a parwent dictionary of d1.

dict_bind_bang()

```
void Lamb::dict_bind_bang (Sexpr_t dict, Sexpr_t key, Sexpr_t value, Sexpr_t env_exec)
```

Modify the target dictionary; value is re-assigned to key wherever first found in chain, else created in the top frame.

dict_bind_local_bang()

```
void Lamb::dict_bind_local_bang (Sexpr_t dict, Sexpr_t key, Sexpr_t value, Sexpr_t env_exec)
```

Modify only the top frame of dict; update key if found there, else create it there. Never walks parent frames. Use for define/define-syntax per R7RS 5.3.2.

dict_rebind_bang()

```
void Lamb::dict_rebind_bang (Sexpr_t dict, Sexpr_t key, Sexpr_t value, Sexpr_t env_exec)
```

Modify the target dictionary; value is assigned to key wherever first found in env, else error if not found.

dict_bind_alist_bang()

```
void Lamb::dict_bind_alist_bang (Sexpr_t dict, Sexpr_t alist, Sexpr_t env_exec)
```

Modify the target dictionary; binding keys in the alist to their corresponding values.

dict_rebind_alist_bang()

```
void Lamb::dict_rebind_alist_bang (Sexpr_t dict, Sexpr_t alist, Sexpr_t env_exec)
```

Modify the target dictionary; rebinding keys in the alist to their corresponding values.

dict_ref_q()

```
Sexpr_t Lamb::dict_ref_q (Sexpr_t dict, Sexpr_t key)
```

Returns (key value) pair, or **false** if key is unbound.

dict_ref()

```
Sexpr_t Lamb::dict_ref (Sexpr_t dict, Sexpr_t key)
```

Returns the value associated with the key in the given dictionary; throws error if key is unbound.

dict_keys()

```
Sexpr_t Lamb::dict_keys (Sexpr_t dict, Sexpr_t env_exec)
```

Note that keys and values are guaranteed to return in the corresponding order in dict_keys and dict_values, duplicate keys may occur.

Return a list of all the keys in this dictionary. If the dictionary has parents, keys may appear multiple times.

dict_values()

```
Sexpr_t Lamb::dict_values (Sexpr_t dict, Sexpr_t env_exec)
```

Return a list of all the values in this dictionary. If the dictionary has parents, values may appear multiple times for each key.

dict_to_alist()

```
Sexpr_t Lamb::dict_to_alist (Sexpr_t dict, Sexpr_t env_exec)
```

Convert the dictionary to an alist, retaining only the top-level (key . value) pairs.

dict_to_2list()

```
Sexpr_t Lamb::dict_to_2list (Sexpr_t dict, Sexpr_t env_exec)
```

Convert the dictionary to a list of 2-element lists, retaining only the top-level (key value) sublists.

alist_to_dict()

```
Sexpr_t Lamb::alist_to_dict (Sexpr_t alist, Sexpr_t env_exec)
```

Convert an alist into a dictionary. The resulting dictionary has no parent.

twolist_to_dict()

```
Sexpr_t Lamb::twolist_to_dict (Sexpr_t twolist, Sexpr_t env_exec)
```

Convert a list of 2-element lists into a dictionary. The resulting dictionary has no parent.

dict_analyze()

```
Sexpr_t Lamb::dict_analyze (Sexpr_t dict, LL_int32 verbosity = 0)
```

Internal integrity check.

integrity_check()

```
Sexpr_t Lamb::integrity_check (Sexpr_t env_exec, LL_int32 verbosity = 1)
```

GC heap integrity check; callable from mop3.

read()

```
Sexpr_t Lamb::read (LL_Port & src, Sexpr_t env_exec) [inline]
```

Read one S-expression from *src*; returns OBJ_EOF at end of input.

write_or_display()

```
Sexpr_t Lamb::write_or_display (Sexpr_t sexpr, Bool_t do_write, WriteMode mode = WriteMode::Cyclic)
```

Write *sexpr* to the current output port. *do_write* true => *write* semantics (strings quoted, chars as #\\x); false => *display* (strings/chars as-is). *mode* selects datum-label emission: Cyclic labels nodes on a cycle (default, terminates on cycles), Shared labels all shared nodes, Simple emits no labels. Returns OBJ_VOID.

write_simple()

```
Sexpr_t Lamb::write_simple (Sexpr_t sexpr)
```

Write *sexpr* using *write* semantics without following shared structure. Returns OBJ_VOID.

sprintf()

```
String Lamb::sprintf (Charst_t fmt, Sexpr_t args, Sexpr_t env_exec)
```

Format a string from a template and a list of S-expression arguments. Format specifiers: `~a` / `~A` inserts the next argument using *display* formatting; `~n` or `~%` inserts a newline; `~~` inserts a literal `~`. Backslash escapes `\\n` (newline) and `\\\\` (literal backslash) are also recognized. Throws if the argument count does not match the format string.

printf() [2/3]

```
Sexpr_t Lamb::printf (Sexpr_t args, LL_Port & outp)
```

Format and write to *outp*; *args* is (fmt-string val ...). See [sprintf\(\)](#) for format specifier syntax. Returns OBJ_VOID.

printf() [3/3]

```
Sexpr_t Lamb::printf (Sexpr_t args)
```

Format and write to the current output port; *args* is (fmt-string val ...). See [sprintf\(\)](#) for format specifier syntax. Returns OBJ_VOID.

eval()

```
Sexpr_t Lamb::eval (Sexpr_t sexpr, Sexpr_t env_exec)
```

Evaluate the S-expression in the environment provided.

eval_list()

```
Sexpr_t Lamb::eval_list (Sexpr_t args, Sexpr_t env_exec)
```

Evaluate the list of S-expressions in the environment provided, and return a llist of results.

apply_proc_partial()

```
Sexpr_t Lamb::apply_proc_partial (Sexpr_t proc, Sexpr_t sexpr, Sexpr_t env_exec)
```

Evaluate the function arguments and then apply the function to them. The result may be a tail requiring further evaluation.

map_proc()

```
Sexpr_t Lamb::map_proc (Sexpr_t proc, Sexpr_t lists, Sexpr_t env_exec)
```

Apply the procedure to the lists per R5RS.

lamb_oblist()

```
Sexpr_t Lamb::lamb_oblist () [inline]
```

B149: the current in/out/error ports are ordinary global bindings in the base env (which is permanently rooted), so the ports are traced like any other binding – NOT special GC roots. These read/write the cached binding cell's cdr (the port value) for O(1) access.

set_current_input_port()

```
void Lamb::set_current_input_port (Sexpr_t p) [inline]
```

for with-input-from-file

set_current_output_port()

```
void Lamb::set_current_output_port (Sexpr_t p) [inline]
```

for with-output-to-file

load()

```
Sexpr_t Lamb::load (Charst_t name, Sexpr_t env_exec, LL_int32 verbosity = 0)
```

Load and evaluate S-expressions from the named file.

append()

```
Sexpr_t Lamb::append (Sexpr_t sexpr, Sexpr_t env_exec)
```

Concatenate a list of lists per R5RS [append](#); the last element may be a non-list, producing an improper list. Returns the concatenated result.

append2()

```
Sexpr_t Lamb::append2 (Sexpr_t lis, Sexpr_t obj, Sexpr_t env_exec)
```

Destructively attach *obj* to the end of *lis*, which must be a proper list.

list_copy()

```
Sexpr_t Lamb::list_copy (Sexpr_t sexpr, Sexpr_t env_exec)
```

Return a shallow copy of the list: new pair cells with the same elements in the same order.

list_analyze()

```
Sexpr_t Lamb::list_analyze (Sexpr_t sexpr, Sexpr_t env_exec)
```

Traverse *sexpr* using Floyd's cycle-detection algorithm and return a structural description:

- proper list of length $N \rightarrow$ integer N
- improper list (non-NIL terminator) after N pairs $\rightarrow$ (improper . N)
- circular structure $\rightarrow$ #f

list_to_vector()

```
Sexpr_t Lamb::list_to_vector (Sexpr_t l, Sexpr_t env_exec)
```

Convert a proper list to a heap-allocated S-expression vector (T_SVEC_HEAP); list order is preserved.

vector_to_list()

```
Sexpr_t Lamb::vector_to_list (Sexpr_t v, Sexpr_t env_exec)
```

Convert an S-expression vector to a proper list; element order is preserved.

vector_copy()

```
Sexpr_t Lamb::vector_copy (Sexpr_t from, Sexpr_t env_exec)
```

Return a shallow copy of the S-expression vector *from*.

vector_to_sparsevec()

```
Sexpr_t Lamb::vector_to_sparsevec (Sexpr_t vec, Sexpr_t skip, Sexpr_t env_exec)
```

Convert *vec* to a sorted (index . value) alist, omitting elements equal to *skip*.

sparsevec_to_vector()

```
Sexpr_t Lamb::sparsevec_to_vector (Sexpr_t alist, Sexpr_t fill, Sexpr_t env_exec)
```

Reconstruct a dense vector from a sparse (index . value) alist; positions not present in the alist are initialized to *fill*.

mk_error() [1/2]

```
Sexpr_t Lamb::mk_error (Sexpr_t env_exec, Sexpr_t irritants, const char * fmt, ...)
```

Fill a new error object with the information given, and return it.

mk_error() [2/2]

```
Sexpr_t Lamb::mk_error (Sexpr_t env_exec, const char * fmt, ...)
```

Fill a new error object with the information given, and return it. The *irritants* field is left NIL.

mk_syserror() [1/2]

```
Sexpr_t Lamb::mk_syserror (const char * fmt, ...)
```

Fill the system error object with the info given, and return it.

tag_error()

```
Sexpr_t Lamb::tag_error (Sexpr_t err, const char * kind)
```

Tag a fresh T_ERROR with a kind marker in its irritants (e.g. "%file-error"/"%read-error") so file-error?/read-error? can classify it. Returns *err*.

mk_syserror() [2/2]

```
Sexpr_t Lamb::mk_syserror (Sexpr_t irritants, const char * fmt, ...)
```

Fill the system error object with the info given, and return it.

gensym()

```
Sexpr_t Lamb::gensym (Sexpr_t env_exec)
```

Produce a new unique symbol.

Perf: these makers allocate a blank cell then embed the value with `->set()`, so the tcons a/b args are throw-aways – pass (Word_t)0,(Word_t)0 to hit the NON-pushing Word_t tcons overload. Do NOT change back to NIL, NIL: that selects the Sexpr_t overload, whose `ll_gc_protect2(a,b)` pushes NIL twice on EVERY `mk_↔int32/float32/...` (i.e. every materializing car/cdr) for nothing.

mk_bool()

```
Sexpr_t Lamb::mk_bool (Bool_t b, Sexpr_t env_exec) [inline]
```

Return HASHT or HASHF.

mk_character()

```
Sexpr_t Lamb::mk_character (ll_codepoint_t cp, Sexpr_t env_exec) [inline]
```

Allocate a new T_CHAR cell holding Unicode codepoint *cp*.

mk_integer()

```
Sexpr_t Lamb::mk_integer (LL_int64 i, Sexpr_t env_exec) [inline]
```

`mk_integer`: T_INT32 if fits, else T_INT64. Returns cached cell for [-2048..4096].

THE DOUBLE WRITE BELOW LOOKS LIKE WASTE. REMOVING IT DID NOT PAY – THREE VARIANTS MEASURED. `tcons(T_PAIR,0,0,env)` initialises the cell as an empty pair and `->set(v)` immediately overwrites it. Tried 2026-09-01 with an `alloc_cell()` returning a zeroed UNTYPED cell so the caller writes it once. All three were CORRECT (rxrs 443/0, extend 629/0, conform 1337/0, and numeric spot-checks across `int32/int64/float/complex/bignum`). S3 devkit, fib-iter(45) AST, 3 s window, fresh boot; baseline 17,784 / 17,924 / 17,994 us:

1. shared allocator, out-of-line 18,384 (+3.4%)
2. shared allocator, in-class 18,139-18,209 (+2.3%)
3. DUPLICATED pop, `tcons_unguarded` BYTE-IDENTICAL 18,135-18,348 (+1.5-2.4%) Variant 3 is the control: it cannot have touched the cons path, and it lost the same amount. So the "sharing grew a hot function" story I first wrote is WRONG. What is left is that the effect sits at or below the level at which code LAYOUT alone moves this benchmark on this device – adding a function perturbs placement, and the baseline itself spans 1.2%. CONCLUSION: no demonstrable benefit; do not re-attempt without a profile that can separate layout from mechanism, and measure cons-heavy work on the DEVICE, not on x86.

mk_int32()

```
Sexpr_t Lamb::mk_int32 (LL_int32 i, Sexpr_t env_exec) [inline]
```

<T_INT32 cell; returns cached cell for small values.

mk_int64()

```
Sexpr_t Lamb::mk_int64 (LL_int64 i, Sexpr_t env_exec) [inline]
```

T_INT64 cell (always allocates).

mk_float32()

```
Sexpr_t Lamb::mk_float32 (LL_float32 r, Sexpr_t env_exec) [inline]
```

T_FLOAT32 cell.

mk_float64()

```
Sexpr_t Lamb::mk_float64 (LL_float64 d, Sexpr_t env_exec) [inline]
```

T_FLOAT64 cell.

mk_complex()

```
Sexpr_t Lamb::mk_complex (LL_float32 r, LL_float32 i, Sexpr_t env_exec) [inline]
```

T_COMPLEX cell.

rat_store_i64()

```
void Lamb::rat_store_i64 (Sexpr_t c, bool is_num, LL_int64 v, Sexpr_t env_exec) [inline]
```

Store an int64 exact value into a rational slot: embedded int32 if it fits, else a boxed T_INT64.

< boxed (issued -> survives; stored immediately)

mk_rational()

```
Sexpr_t Lamb::mk_rational (LL_int64 num, LL_int64 den, Sexpr_t env_exec) [inline]
```

Canonical rational from int64 components: GCD-reduces, normalises sign, returns an integer when den==1, else a T_RATIONAL whose components box to T_INT64 when they overflow int32 (P70 – no float64 fallback, so exactness is preserved). When LL_RATIONAL=0, always returns inexact float64.

mk_number()

```
Sexpr_t Lamb::mk_number (Charst_t str, Sexpr_t env_exec)
```

Parse str as a number; returns T_INT, T_REAL, T_RATIONAL, or T_COMPLEX as appropriate, or OBJ_UNDEF if str is not a valid numeric literal.

mk_sharp_const()

```
Sexpr_t Lamb::mk_sharp_const (Charst_t name, Sexpr_t env_exec)
```

Parse a #-prefixed constant name (e.g., "t", "f", "true", "false") and return the corresponding singleton; returns OBJ_UNDEF if unrecognized.

payload_alloc()

```
void * Lamb::payload_alloc (size_t nbytes, const char * who, Sexpr_t env_exec)
```

Allocate a bulk PAYLOAD (string/bytevector/vector/bignum/bytecode body) with a DEFINED out-of-memory policy, mirroring what gc_urgent() does for cells.

B129: ll_payload_alloc_bytes() falls back to malloc() when PSRAM is full, and malloc() returns NULL – which every call site then wrote through (memset/memcpy on address 0), producing Guru Meditation Error: StoreProhibited and a BOOT LOOP that only a reflash cleared. A real-time node must not die that way, and must not die silently.

Policy, in order:

1. try the allocation;
2. on failure run the collector to completion – payload memory is reclaimed by the SWEEP (finalize_one↵_cell frees the payloads of dead cells), so exhaustion is often transient – and retry ONCE;
3. if it still fails, throw a normal LambLisp error. It is catchable with guard, so an application can shed load, and the REPL survives to report it. Never returns NULL.

mk_symbol_or_number()

```
Sexpr_t Lamb::mk_symbol_or_number (Charst_t str, Sexpr_t env_exec)
```

Try to parse `str` as a number; if that fails, intern it as a symbol. Used by the reader for unquoted tokens.

mk_symbol()

```
Sexpr_t Lamb::mk_symbol (Charst_t str, Sexpr_t env_exec) [inline]
```

Intern `str` in the oblist and return the symbol S-expression. Two calls with identical strings always return the same cell pointer.

mk_bytevector() [1/3]

```
Sexpr_t Lamb::mk_bytevector (LL_int32 k, Sexpr_t env_exec)
```

Simplest heap allocation with no initialization.

mk_bytevector() [2/3]

```
Sexpr_t Lamb::mk_bytevector (LL_int32 k, Bytest_t src, Sexpr_t env_exec)
```

Heap allocation with initialization.

mk_bytevector() [3/3]

```
Sexpr_t Lamb::mk_bytevector (LL_int32 k, LL_int32 fill, Sexpr_t env_exec)
```

Heap allocation with initialization.

mk_bytevector_ext()

```
Sexpr_t Lamb::mk_bytevector_ext (LL_int32 k, Bytest_t ext, Sexpr_t env_exec) [inline]
```

Injection of externally allocated memory, which will not be freed at GC time.

mk_intvector()

```
Sexpr_t Lamb::mk_intvector (LL_int32 k, Sexpr_t env_exec)
```

Note that integer vectors and real vectors are just bytevectors underneath.

Allocates a bytevector from the heap to ensure proper alignment (no IMM type).

mk_realvector()

```
Sexpr_t Lamb::mk_realvector (LL_int32 k, Sexpr_t env_exec)
```

Allocates a bytevector from the heap to ensure proper alignment (no IMM type).

mk_vector()

```
Sexpr_t Lamb::mk_vector (LL_int32 len, Sexpr_t fill, Sexpr_t env_exec)
```

Allocate a T_SVEC_HEAP vector of *len* elements, all initialized to *fill*.

mk_hashtbl()

```
Sexpr_t Lamb::mk_hashtbl (LL_int32 len, Sexpr_t fill, Sexpr_t env_exec)
```

Allocate a T_SVEC2N_HEAP hash-table vector; actual size is rounded up to the next power of 2, minimum *len* slots, all initialized to *fill*.

mk_serial_port()

```
Sexpr_t Lamb::mk_serial_port (Sexpr_t env_exec)
```

Wrap the platform serial port (Arduino `Serial` or POSIX `stdin/stdout`) as a LambLisp port object.

mk_input_file_port()

```
Sexpr_t Lamb::mk_input_file_port (Charst_t name, Sexpr_t env_exec)
```

Open the named filesystem file for reading; return a port. Throws if the file cannot be opened.

mk_output_file_port()

```
Sexpr_t Lamb::mk_output_file_port (Charst_t name, Sexpr_t env_exec)
```

Open the named filesystem file for writing; return a port. Throws if the file cannot be opened.

mk_tcp_client_port()

```
Sexpr_t Lamb::mk_tcp_client_port (Charst_t host, unsigned port, Sexpr_t env_exec)
```

Open a TCP connection to `host:port`; return a bidirectional port. Scheme: `(open-tcp-client-port host port)`. Throws on connection failure.

mk_tcp_server_port()

```
Sexpr_t Lamb::mk_tcp_server_port (unsigned port, Sexpr_t env_exec)
```

Create a TCP listener on `port`; return a server port. Scheme: `(open-tcp-server-port port)`. Use `server_accept_port()` to accept incoming connections.

server_accept_port()

```
Sexpr_t Lamb::server_accept_port (LL_Port * server, Sexpr_t env_exec)
```

Accept one pending connection from `server`; returns a bidirectional client port, or `HASHF(#f)` if no connection is waiting (non-blocking). Scheme: `(server-accept server-port)`.

mk_tls_client_port()

```
Sexpr_t Lamb::mk_tls_client_port (Charst_t host, unsigned port, Charst_t ca_cert_pem, Sexpr_t env_exec)
```

Open a TLS client connection to `host:port`. Pass a PEM CA certificate string in `ca_cert_pem` to verify the server; pass `nullptr` to skip verification (insecure, development only). Scheme: `(open-tls-client-port host port)` or `(open-tls-client-port host port ca-cert-pem)`.

mk_tls_server_port()

```
Sexpr_t Lamb::mk_tls_server_port (unsigned port, Charst_t cert_pem, Charst_t key_pem, Sexpr_t env_exec)
```

Create a TLS server on `port` using PEM-encoded certificate `cert_pem` and private key `key_pem`. Scheme: `(open-tls-server-port port cert-pem key-pem)`.

mk_input_string_port()

```
Sexpr_t Lamb::mk_input_string_port (Charst_t inp, Sexpr_t env_exec)
```

Wrap the C string *inp* as an input port; reading consumes characters left to right.

mk_output_string_port()

```
Sexpr_t Lamb::mk_output_string_port (Sexpr_t env_exec)
```

Create an output port that accumulates written characters; retrieve the result with Scheme `get-output-string`.

mk_input_bytevector_port()

```
Sexpr_t Lamb::mk_input_bytevector_port (ByteVec_t src, LL_int32 n, Sexpr_t env_exec)
```

Wrap *n* bytes of *src* as a binary input port.

mk_output_bytevector_port()

```
Sexpr_t Lamb::mk_output_bytevector_port (Sexpr_t env_exec)
```

Create a binary output port that accumulates written bytes; retrieve the result with Scheme `get-output-bytevector`.

mk_Mop3_procst_t()

```
Sexpr_t Lamb::mk_Mop3_procst_t (Mop3st_t f, Sexpr_t env_exec) [inline]
```

car=0 (no argv variant)

mk_Mop3_procst_argv_t()

```
Sexpr_t Lamb::mk_Mop3_procst_argv_t (Mop3st_t f, Mop3st_argv_t fa, Sexpr_t env_exec) [inline]
```

car=argv fn, *cdr*=standard fn

mk_macro()

```
Sexpr_t Lamb::mk_macro (Sexpr_t nlam, Sexpr_t env_nlam, Sexpr_t env_exec) [inline]
```

Allocate a T_MACRO cell: *nlam* is the nlambda transformer, *env_nlam* is its definition environment. At expansion time the transformer receives the unevaluated argument list; its result is evaluated in the call-site environment.

mk_ident()

```
Sexpr_t Lamb::mk_ident (Sexpr_t sym, Sexpr_t def_env, Sexpr_t env_exec) [inline]
```

Allocate a T_IDENT cell (ER renamed identifier): *sym* is the original symbol, *def_env* its definition environment. *car*=*sym*, *cdr*=*def_env*.

mk_procedure()

```
Sexpr_t Lamb::mk_procedure (Sexpr_t formals, Sexpr_t body, Sexpr_t env_proc, Sexpr_t env_exec) [inline]
```

Allocate a T_PROC closure: *formals* is the parameter list, *body* is the list of body expressions, *env_proc* is the captured lexical environment.

mk_nprocedure()

```
Sexpr_t Lamb::mk_nprocedure (Sexpr_t formals, Sexpr_t body, Sexpr_t env_nproc, Sexpr_t env_exec) [inline]
```

As [mk_procedure\(\)](#) but returns a T_NPROC cell; arguments are passed unevaluated at call time.

mk_thunk_sexpr()

```
Sexpr_t Lamb::mk_thunk_sexpr (Sexpr_t sexpr, Sexpr_t env_thunk, Sexpr_t env_exec) [inline]
```

Allocate a T_THUNK_SEXPR trampoline cell: *sexpr* will be evaluated in *env_thunk* on the next trampoline iteration.

mk_thunk_body()

```
Sexpr_t Lamb::mk_thunk_body (Sexpr_t body, Sexpr_t env_thunk, Sexpr_t env_exec) [inline]
```

Allocate a T_THUNK_BODY trampoline cell: *body* is a list of expressions evaluated in sequence in *env_thunk*; the last expression's value is the result.

macroexpand() [1/2]

```
Sexpr_t Lamb::macroexpand (Sexpr_t form, Sexpr_t env_exec)
```

Fully expand all macros in `form`, applying transformers repeatedly until no further expansion is possible. Returns the fully expanded form.

macroexpand1() [1/2]

```
Sexpr_t Lamb::macroexpand1 (Sexpr_t form, Sexpr_t env_exec)
```

Apply at most one macro expansion step to `form`. Returns the expansion if `form` is a macro call, otherwise returns `form` unchanged.

macroexpand() [2/2]

```
Sexpr_t Lamb::macroexpand (Sexpr_t proc, Sexpr_t args, Sexpr_t env_exec)
```

Fully expand the macro `proc` applied to `args`.

macroexpand1() [2/2]

```
Sexpr_t Lamb::macroexpand1 (Sexpr_t proc, Sexpr_t args, Sexpr_t env_exec)
```

Apply one expansion step of macro `proc` to `args`.

ncg_mark_frames()

```
void Lamb::ncg_mark_frames ()
```

Push all NCG live-frame values onto GC markstack; called from `start_marking()`.

ncg_eval()

```
Sexpr_t Lamb::ncg_eval (Sexpr_t bc, Sexpr_t args, Sexpr_t env_exec)
```

Execute a T_BYTECODE cell that has been NCG-compiled.

ncg_eval_argv()

```
Sexpr_t Lamb::ncg_eval_argv (Sexpr_t bc, Sexpr_t * argv, int nargs, Sexpr_t env_exec)
```

Phase 6: direct NCG→NCG call; takes argv[] instead of list, skips apply_proc_partial.

ncg_compile()

```
Sexpr_t Lamb::ncg_compile (Sexpr_t bc, Sexpr_t env_exec, Sexpr_t captured_names = nullptr)
```

Compile a T_BYTECODE cell to native code; sets slot BC_SLOT_NCG_CODE in its cvec.

ncg_disassemble()

```
void Lamb::ncg_disassemble (Sexpr_t bc, Sexpr_t env_exec)
```

Log basic info about an NCG-compiled proc.

ncg_vector_set_ptr()

```
void Lamb::ncg_vector_set_ptr (Sexpr_t * data_ptr, LL_int32 k, Sexpr_t val)
```

Write-barrier store for inlined NCG_STORE_LOCAL (Phase 1).

ncg_vref_fast()

```
Sexpr_t Lamb::ncg_vref_fast (Sexpr_t vec, Sexpr_t idx)
```

fast vector-ref (T_SVEC + T_INT32 fast path).

ncg_vset_fast()

```
Sexpr_t Lamb::ncg_vset_fast (Sexpr_t vec, Sexpr_t idx, Sexpr_t val)
```

fast vector-set (write-barrier inline).

ncg_cons_fast()

```
Sexpr_t Lamb::ncg_cons_fast (Sexpr_t a, Sexpr_t d)
```

fast cons; returns nullptr if Nfree<=yuasa_M or sweeping.

ncg_box_fast()

```
Sexpr_t Lamb::ncg_box_fast (LL_int64 v)
```

fast integer box; nullptr if Nfree<=yuasa_M or sweeping.

eval_recursion_overflow()

```
void Lamb::eval_recursion_overflow (const char * where)
```

cold path: throws the catchable "recursion too deep" error

eval_recursion_check()

```
void Lamb::eval_recursion_check (const char * where) [inline]
```

< hot path: cheap C-stack + root-stack depth check at a recursion entry

diag_memory_report()

```
void Lamb::diag_memory_report ()
```

B117: dump GC-buffer / cell-block placement (INTERNAL vs PSRAM) + sizes.

12.2.5 Member Data Documentation

ncg_frame_top

```
NcgFrame* Lamb::ncg_frame_top = nullptr
```

linked list of live frames; head threaded through frame->prev

ncg_frame_pool_head

```
NcgFrame* Lamb::ncg_frame_pool_head = nullptr
```

free-list head; recycled frames linked via frame->prev

ncg_frame_pool_count

```
int Lamb::ncg_frame_pool_count = 0
```

frames currently available in the pool

ncg_err_ctx

```
NcgErrJump* Lamb::ncg_err_ctx = nullptr
```

innermost escape context; set by ncg_eval before JIT call

eval_stack_base_

```
const char* Lamb::eval_stack_base_ = nullptr
```

highest stack address seen at a recursion entry (== outermost frame)

eval_stack_budget_

```
size_t Lamb::eval_stack_budget_ = LL_EVAL_STACK_BUDGET
```

trip threshold: bytes of C-stack growth below base; tunable at runtime

eval_root_depth_max_

```
LL_int32 Lamb::eval_root_depth_max_ = LL_EVAL_ROOT_DEPTH_MAX
```

trip threshold: rootstack depth (cap is 65536); tunable at runtime

12.3 LambPreSettings Struct Reference

Public Member Functions

- void `load` ()

Public Attributes

- LL_int32 `cell_block_size` = 8 * 1024
- LL_int32 `extension_block_size` = 4 * 1024
- LL_int32 `ncg_frame_pool_size` = 16

12.3.1 Detailed Description

Read from LAMB_SETTINGS_PATH before any cells exist. Default matches the hardcoded initial cell block size. Silent if file is missing (first-boot safe).

12.3.2 Member Function Documentation

`load()`

```
void LambPreSettings::load ()
```

Reads LAMB_SETTINGS_PATH_RAW via fopen; silent if missing.

12.3.3 Member Data Documentation

`cell_block_size`

```
LL_int32 LambPreSettings::cell_block_size = 8 * 1024
```

Cells in the initial GC block (8K default; perf pushes 16K to match comps).

`extension_block_size`

```
LL_int32 LambPreSettings::extension_block_size = 4 * 1024
```

Cells per on-demand expansion block (4K).

ncg_frame_pool_size

```
LL_int32 LambPreSettings::ncg_frame_pool_size = 16
```

NcgFrames pre-allocated into the free-list pool at startup.

12.4 LambStdioClass Class Reference

Public Member Functions

- int **setTxBufferSize** (int n)
- int **setRxBufferSize** (int n)
- void **begin** (void)
- void **begin** (unsigned long baudrate)
- void **end** ()
- int **available** (void)
- int **availableForWrite** (void)
- int **read** (void)
- int **write** (uint8_t c)
- int **write** (char c)
- void **flush** (void)
- int **read** (byte *buf, int max)
- int **read** (char *s, int max)
- int **write** (const char *s, int n)
- int **write** (const byte *b, size_t n)
- int **write** (const char *s)
- **operator bool** ()

12.4.1 Detailed Description

A wrapper around the underlying stdin/stdout implementation. On an embedded system, this class will use the primary serial in/out (`Serial` on Arduino-compatibles). On Linux, this class will set the terminal to byte-at-a-time mode (called "non-CANONICAL" mode).

12.5 LL_File Class Reference

Public Member Functions

- bool **isOpen** ()
- int **read** (void)
- int **write** (byte b)
- int **seek** (unsigned long target, int whence=SEEK_SET)
- int **tell** ()
- int **size** ()
- int **close** ()
- int **read** (byte *b, int n)
- int **read** (char *s, int n)
- int **write** (const byte *b, int n)
- int **write** (const char *s, int n)
- int **peek** ()

Public Attributes

- File_Native **_theFile**
- String **_path**
- String **_mode**

12.5.1 Detailed Description

The *file* type is ultimately provided by the underlying operating system, not by LambLisp. This class elides the differences between file types on different platforms, providing a POSIX-like interface.

Note that there is no *open* operation on files. A file is opened by the *file system* and then a *file* is returned. After a *file* is closed, the same file object cannot be opened again; instead a new file must be requested from the *file system*.

12.6 LL_File_System Class Reference

Public Member Functions

- **LL_File * open** (const char *path, const char *mode)
- bool **exists** (const char *path)
- int **rm** (const char *path)
- int **mv** (const char *from, const char *to)
- int **mkdir** (const char *path)

12.6.1 Detailed Description

The "file system" type elides the differences between different underlying platforms. For example, it will deal with the leading '/' required by LittleFS. This minimal file system interface is platform-independent.

12.7 NcgErrJump Struct Reference

Public Attributes

- jmp_buf **jmpbuf**
- Sexpr_t **error**
- LL_int32 **gc_depth**

12.7.1 Detailed Description

Forward decl; defined in ll_vm_ncg.h.

Exception-escape context set up by ncg_eval/ncg_eval_argv so that C++ shims called FROM JIT code can longjmp rather than throw through the JIT frame. JIT frames have no DWARF unwind info; uncaught C++ exceptions propagating through them call std::terminate(). The longjmp skips the JIT frame entirely and returns control to the setjmp site in ncg_eval, which then re-throws.

12.7.2 Member Data Documentation

jmpbuf

`jmp_buf NcgErrJump::jmpbuf`

longjmp target; set by ncg_eval before calling fn(lamb,fp)

error

`Sexpr_t NcgErrJump::error`

exception value stored by ncg_err_longjmp before longjmp

gc_depth

`LL_int32 NcgErrJump::gc_depth`

gc_root_depth() saved before fn(); restored in setjmp handler

Chapter 13

Acknowledgements

I reviewed many *Lisp* and *Scheme* implementations, but the one I studied in detail was *TinyScheme 1.42*, written in 5000+ lines of high-quality circa-1992 C code. It was too big for the processor I was working on at the time (in 2020). Ultimately I returned to the embedded Lisp problem, but made different design choices than the TinyScheme authors, especially on the compile-vs-interpret spectrum and garbage collection techniques. The study was hugely informative nonetheless.

Chez Scheme is the ultimate large implementation. It is the work of a lifetime, primarily of one man, and a standard for performance and completeness in the Scheme world. During the development of *LambLisp*, *Chez Scheme* was used to provide benchmark behavior in cases where the spec is not clear or offers a choice of outcomes.

I have provided an extensive bibliography that also informed my efforts. If the reader would like to better understand Lisp design, begin your journey there. No doubt you will discover a different destination.

Some images in this document are used under Creative Commons license:

- Lambda symbol image: Harke, CC BY-SA 3.0, via Wikimedia Commons. <https://commons.wikimedia.org/w/index.php?curid=127589172>
- Greek letter image: user:rogerb (VectorVoyager), own work, CC BY-SA 3.0. <https://creativecommons.org/licenses/by-sa/3.0>

Bill White 2026

Chapter 14

Glossary

abstract machine

The terms *abstract machine* and *virtual machine* are often used interchangeably, but there is a subtle difference. An *abstract machine* is a mathematics exercise; it is a set of elements, relations among the elements, and transformations between them. A *virtual machine* is an implementation of an abstract machine. When a VM runs, it behaves as defined by the abstract machine, although it runs on top of some other lower-level machine. That lower-level machine may itself be a software or hardware implementation of another abstract machine.

association list, or alist

A list of pairs. The first part of each pair is treated as a key, and the second part is treated as a value. The operations *set* and *get* are provided, to install new pairs or to retrieve existing pairs. Operations to *set* might either replace an existing pair, or add it if not already existing, or simply add the new pair to the front of the list, shadowing any matching pairs later in the list. Alists require linear search and so are used where the lists are expected to be short, as they are in hash tables.

AST, or Abstract Syntax Tree

A data structure created as output of analyzing source code. The AST is an internal representation of the source code elements and their relationships. The AST can be used to re-create the original source code it was created from (not counting white space and comments). Creating the AST is the "front half" of a compiler.

atom

A data item that is not a *pair*. Atoms are the leaf nodes of the S-expression tree; they have no *car* or *cdr*. Most atoms are simple scalar values: integers, real numbers, booleans, characters, and strings. Compound atoms such as vectors and bytevectors carry internal structure but expose it only through their own dedicated operators (e.g., *vector-ref*), not through *car* and *cdr*.

In *LambLisp*, the tag word of a [Cell](#) encodes the atom type: `T_INT`, `T_REAL`, `T_BOOL`, `T_SYM_HEAP`, `T_SYM_ATOM`, and so on. The distinction between atom and pair is the termination condition of every list traversal: a list walk continues while the current cell is a pair, and stops when it reaches an atom (typically `NIL`).

bytecode

A compact, machine-independent instruction set produced by a compiler and executed by a virtual machine. Rather than interpreting the source AST node by node, a bytecode compiler makes a single pass to emit a sequence of encoded instructions stored as a vector; the bytecode executor then runs those instructions in a tight loop with lower dispatch overhead than tree-walking interpretation.

LambLisp includes a bytecode compiler and executor (T_BYTECODE). Bytecode compilation is performed at load time; the resulting instruction vector lives in the *LambLisp* heap and is subject to garbage collection like any other object. Benchmarks show a 1.7–3.3x speedup over the pure-interpreter baseline for compute-bound code.

*car** and **cdr*

Historically, *car* stands for *Contents of the Address Register* and *cdr* for *Contents of the Decrement Register* – fields in the memory word of the IBM 704 computer on which Lisp was first implemented in 1958.

Refers to the first and second data word in a [Cell](#), respectively. In *LambLisp*, a [Cell](#) consists of 3 sequential memory *words* (1 tag + 2 data), where a single *word* is large enough to hold a pointer or an integer, and 2 consecutive *words* are large enough to hold a real number.

In the 3-word [Cell](#), the first word is the *tag* word, the second word is the *car*, and the third word is the *cdr*.

The *tag* word contains a [Cell](#) type indicator, and flags representing the garbage collection state of each cell.

Integers are stored in the *car* word. Real numbers are stored in the *car* word, and may possibly extend into the *cdr* word, depending on the underlying microprocessor. Some complex atoms, such as vectors, store a count in the *car* and a pointer in the *cdr*. Symbols store a hash in the *car*, and a pointer to the symbol identifier in the *cdr*.

Cell

A triplet of information: tag, first word (aka *car*), second word (aka *cdr*). In *LambLisp*, the three elements of the triplet are consecutive in memory, and are all the same size, large enough to hold a pointer or integer. The tag indicates how the first and second words are to be interpreted. When the tag indicates a [Cell](#) is a *pair*, the *car* and *cdr* both point to other Cells (or possibly the same [Cell](#)).

A *list* is formed when the *car* of a pair [Cell](#) points to an S-expression (which is an item in the list), and the *cdr* of that [Cell](#) points to another *pair* [Cell](#). That *pair* [Cell](#) may also have a list item in its *car*. By convention, lists terminate with a special S-expression called *the empty list*, denoted '()', or NIL, or nil. NIL is considered a list (*the empty list*), but is not a pair, so it has no *car* or *cdr*.

Cells that are not pairs are called *atoms*. Atoms do not have any directly accessible internal structure, but may have operations that work on them, such as vector subscripting. Many atoms are simple integers, strings, or floating-point numbers, but there are compound atoms such as vectors and bytevectors.

In *LambLisp*, there are several additional *pair* types that facilitate high performance in conjunction with the underlying *LambLisp* virtual machine. In standard Lisp or Scheme code, these would be considered atoms, but in *LambLisp* they are a bit of a hybrid, because *car*, *cdr*, *append*, and a few other list operations will work as expected.

closure

A procedure combined with the *lexical environment* in which it was created. When a `lambda` expression is evaluated, the result is a closure: the procedure body paired with a captured reference to the enclosing environment frame. This captured environment gives the closure access to variables from its defining scope even when called later in a different context.

Closures are the mechanism behind higher-order functions, callbacks, and object-based programming patterns common in Scheme. In *LambLisp*, a closure is represented as a `T_PROC Cell` whose *cdr* points to the captured environment chain.

compiler, interpreter

Both are language translators; in addition, an interpreter is also the language executor. There is a spectrum between compilers and interpreters, but we can identify a key difference:

Translator type	Description
<i>compiler</i>	a software process that reads source code and produces new code to be executed later, possibly by a different machine.
<i>interpreter</i>	a software process that reads source code, performs some amount of processing, and executes the result immediately on the same machine.

In a compiler, and in any nontrivial interpreter, the software process first produces an Abstract Syntax Tree (AST). Creating the AST is the first half of the compilation process; in the second half, code is generated for later execution. An interpreter will create the AST and then traverse it, performing the operations as they are encountered.

In *Lisp*, code compiling is commonly done through the use of a *macro* facility. Macro expansion allows code to be transformed as it is read in, factoring out invariants, and producing new code with the invariants removed or calculated once. For example, replacing a symbol with its value "freezes" the symbol; subsequent reads of the symbol will be replaced with the frozen value whenever they are encountered. When that replacement is done at read time, the value is located and replaces the symbol in the executable program. The symbol-value lookup operation is therefore no longer required at run time. In this way symbols become replaced by their read-time values, effectively compiling a block of code into the execution stream.

This method of compilation, called *incremental compilation* or *just-in-time compilation*, blurs the line between compiler and interpreter.

dictionary

A *dictionary* is a data structure containing (key value) pairs. Pairs in the dictionary can be retrieved by matching a target key with the keys of each pair. Dictionaries also allow creating, updating, or deleting (key value) pairs.

A simple dictionary might be a list of (key value) pairs, requiring linear search. This is appropriate for small dictionaries, or short-lived dictionaries, where the construction overhead or runtime overhead of a faster solution outweighs any potential speed improvement.

For larger or long-lived dictionaries, intelligent key creation and caching when combined with hash tables makes this a very fast data structure upon which to operate.

What's described above is a *flat dictionary*, having a single layer of (key value) pairs.

environments & hierarchical dictionaries

Often called the "symbol table" in other programming languages, in LambLisp the environment is a *hierarchical dictionary*. The *flat dictionary* described above provides a single *dictionary frame* for a *hierarchical dictionary*. A *hierarchical dictionary* is simply a list of flat dictionaries. The frame at the *car* of the list is a *child frame*, while frames in the *cdr* are called *parent frames* or *ancestor frames*.

The *hierarchical dictionary* provides a high-performing foundation for the LambLisp execution environment. It also provides a simple and natural foundation for the LambLisp Object System. Inheritance (including multiple inheritance) and shadowing are inherent behaviors of the *hierarchical dictionary*.

evaluation

Lisp operates by reading symbolic expressions, evaluating them, and returning the result. The *eval/apply* cycle is the core of a Lisp interpreter:

Step	Action
<i>eval(atom)</i>	Numbers and strings evaluate to themselves; symbols look up their binding in the current environment.
<i>eval(list)</i>	Evaluate the <i>car</i> to obtain a procedure; evaluate each element of the <i>cdr</i> to obtain arguments; <i>apply</i> the procedure to the arguments.
<i>apply(proc, args)</i>	Bind the formal parameters to the argument values in a new environment frame; evaluate the procedure body in that frame.

Special forms — *if*, *define*, *lambda*, *quote*, and others — short-circuit normal *eval/apply*. Their sub-expressions are not pre-evaluated; the special form decides which expressions to evaluate, when, and how many times.

In *LambLisp*, tail calls are handled by returning a *thunk* to the trampoline rather than calling the next procedure directly, keeping the C stack at constant depth.

hash function

A function that takes a series of bytes as input, and produces a fixed-length number called a *hash*. When given a series of inputs, the hash function should give a series of outputs passing relevant tests for randomness. It is not necessary that the series be truly random, only that the results be widely dispersed. For any given input, the output *hash* must always be the same. It can then be used as an index, replacing a search with an index lookup.

hash table

A data structure used for efficiently storing (key value) pairs. Commonly implemented as a set of operations employing a vector and a *hash function*. To search for an item by its key:

- Obtain the hash of the key *H*, which may have been computed once and stored.
- Obtain the size of the hash table *N*.
- Compute the hash table index, which is $(H \% N)$, or $(H \& (N-1))$ if *N* is power of 2.
- Search the (very short) alist at that index of the hash table.

Form	Arguments at call	Return value
<code>lambda</code>	evaluated before body runs	body result returned to caller
<code>nlambda</code>	passed unevaluated (as source)	body result returned to caller
<code>macro</code>	passed unevaluated (as source)	body result is evaluated; that result returned to caller

`nlambda` is a LambLisp extension inherited from *InterLisp*; it is the foundation used to implement all *Scheme* syntax forms. `syntax-rules` is available as the hygienic alternative for portable macro libraries. See the *Advanced topics* section for a full treatment including examples and hygiene discussion.

lambda calculus

A branch of mathematics developed during the 1930's, where the elements of interest are functions rather than data values.

LISP

A symbolic programming language developed about 1960. LISP (always uppercase in those days) offered a platform for reasoning, while then-existing languages focused more on numeric applications. The objects of interest in *LISP* are called *symbolic expressions* or *S-expressions*.

list

A **list** is data structure consisting of a backbone of Cells linked by their *cdrs*. The last Cell in the list points to a special NIL Cell to mark the end of the list. The *cars* of each Cell in the list backbone point to the item that is in the list at that position. Note that every item in the list requires 2 entities: the item itself, and a *pair* Cell that connects the item to the list backbone.

In *Lisp* source code, lists are represented between parenthesis, as in `(this is a list)`.

NIL

All *Lisps* have some version of NIL, `nil`, `()` – otherwise called "the empty list". This is used to mark the end of a list. In some *Lisps*, it may also be considered to be equivalent to "false", or equivalent to zero, but not in *Scheme* or *LambLisp*.

In LambLisp, NIL is implemented as a pointer to a special statically-allocated Cell. The space for the NIL Cell is a singleton, outside the blocks of Cells managed by the *LambLisp* garbage collector. This allows NIL to be constant at compile time, speeding up *LambLisp*.

There are a few other singleton Cells implemented this way, e.g., for undefined and void values.

nprocedure

The value type returned by `nlambda`. An *nprocedure* has the same internal structure as a lambda procedure – a parameter list, a body, and a captured lexical environment – but its arguments are passed unevaluated (as source S-expressions) rather than being evaluated before the call. See **nlambda** and the *Advanced topics* section for usage and examples.

oblist

Lisp is a symbolic computing language and *symbol* is a first-class type. For best performance, LambLisp collects all the symbols encountered into an *interned symbol table*, for traditional reasons called the *obarray* or *oblist*. Symbol hashes are calculated on the symbol's identifier and then stored with the symbol in *oblist*. By this means symbols can be compared with a pointer comparison operation; 2 symbols are the same if they point to the same entry in the oblist. In LambLisp, the oblist is neither an array nor a list; instead it is a hash table, but the *oblist* name is retained.

pair

A *Cell* whose purpose is to connect 2 other cells. In a typical *Lisp* tree structure, *pair* Cells will make up half of the Cells used, while the leaves of the tree will be atoms.

prefix, infix, postfix

In typical algebra notation, to compose two functions *f* and *g* that each require a single parameter *x* we write:
`f(g(x))`

The functions are composed using *prefix* notation. Algebra and C/C++ use a mix of parenthesis, prefix, infix, and postfix notation like this:

```
--x + abs(x) + 2 + (x++ - y[z])
```

Lambda calculus uses prefix notation and makes frequent use of the greek symbol lambda (λ), which is not conducive to keyboard use. Instead *Lisp* spells out *lambda* when required.

Lisp (like lambda calculus) uses prefix notation, with parenthesis and symbols being the fundamental lexical elements. To compose functions *f* and *g*, we write:

```
(f (g x))
```

and to compare 2 elements *x* and *y*, there is the *eq?* operator:

```
(eq? x y)
```

To summarize prefix/infix/postfix, using C++ examples:

Notation		
Prefix	<code>++x</code>	<code>f(x)</code>
Infix	<code>a * b</code>	<code>a + b & c</code>
Postfix	<code>y++</code>	<code>y->foo</code>

REPL

Read-Eval-Print Loop. The interactive command cycle of a Lisp system: *read* an S-expression from the user, *evaluate* it, *print* the result, and *loop* back to read again. *LambLisp* presents the prompt `LL=>` over the serial connection.

The REPL makes Lisp development highly interactive: expressions can be tested immediately, definitions added or redefined without restarting the system, and the live environment inspected at any time. On embedded targets this is particularly valuable, since it allows in-circuit development without a reflash cycle.

Scheme

A dialect of *Lisp* standardized in the R5RS (1998) and R7RS-small (2013) specifications. Scheme is distinguished from other Lisps by its minimalist design: a small, orthogonal core with well-defined semantics, lexical scoping, proper tail recursion, and first-class continuations (`call/cc`).

LambLisp implements the R5RS core with substantial R7RS-small coverage. The primary deliberate departure from standard Scheme is the omission of `call/cc`, which is incompatible with the C-stack trampoline used to achieve real-time performance.

symbolic expression

In its simplest implementations, a *Lisp* **symbol** is any sequence of characters not containing spaces or parenthesis. Some of these symbols sequences are defined by the language (`if`, `lambda`, and so on), along with a value for those symbols, but others are just symbols, to which a value may be attached. These are all symbolic expressions, or **S-expressions**, and combinations of S-expressions and balanced sets of parenthesis are also valid S-expressions.

The basic *Lisp* language can only test for equality of symbols, but real-world Lisps define useful data types like numbers and strings, along with operations on those.

The LambLisp implementation depends on a [Cell](#) structure to contain the details of an **atomic** S-expression (such as an integer), or the links between S-expressions (as in a list). When programming LambLisp at the C++ level, an *S-expression* is therefore a pointer to a [Cell](#).

tail call, tail recursion

A *tail call* is a procedure call that is the last action performed by the calling procedure – the caller has no further use for the return value beyond passing it to its own caller. Because the caller's stack frame is no longer needed, it may be discarded before the call is made, keeping stack depth constant. *Tail recursion* is the specific case where a procedure calls itself in tail position.

The R5RS and R7RS specifications *require* proper tail calls. *LambLisp* implements them via a *trampoline*: a tail-position call packages itself as a *thunk* and returns it to the trampoline loop rather than calling directly, so the C stack never grows beyond a fixed depth regardless of recursion depth.

trampoline

A control-flow technique for implementing proper *tail calls* without growing the underlying C stack. Instead of calling a tail-position procedure directly, *LambLisp* packages the call as a *thunk* (a zero-argument closure) and returns it to an outer loop. That loop iteratively invokes thunks until it receives a non-thunk result, which it then returns.

Because each iteration begins with a fresh C stack frame, stack depth remains constant regardless of how deeply tail-recursive the Scheme code is. The trampoline is also the mechanism that makes *LambLisp's* real-time scheduling possible: the outer loop can perform GC work and service hardware timers between thunk invocations.

Chapter 15

List of Figures

1	Eponymous Lamb	1
1.1	LambLisp REPL Cycle	4
1.2	ESP32-S3 Hardware Context	8
4.1	Dictionary: Hash Table Layout (djb2)	19
4.2	Tail Recursion: Stack Growth vs. Constant Stack	20
4.3	Trampoline: O(1) Tail Calls	21
5.1	Environment Chain and Closure Capture	26
5.2	Performance Characteristics of Several Implementation Strategies	29
5.3	Runtime RAM by Implementation Strategy	30
5.4	Bytecode Compilation Pipeline	31
5.5	A CPU shown as a stack of Virtual Machines	32
5.6	Typical Control Application Organization	33
5.7	LambLisp Block Diagram	34
5.8	LambLisp VM Cell Structure	36
5.9	Typical List Structure	37
5.10	LambLisp VM Cell Life Cycle	39
5.11	Tricolor Marking and Yuasa Write Barrier	39
5.12	Incremental GC Timeline	40
5.13	Argument Evaluation: lambda vs. nlambda vs. macro	51
5.14	Adding a C++ Primitive to the Scheme Environment (mop3)	60
5.15	T_CPP_HEAP cell layout and GC lifecycle	60
5.16	Hardware Device Access Patterns	63
6.1	LLIP Connection Lifecycle	83
6.2	LambLisp Digital Twin	92
6.3	Unified Distributed Environment	93
6.4	Typical Modbus Field Architecture	106
6.5	Non-blocking request/poll flow	108
6.6	Modbus register types	111
6.7	PROFIBUS DP: LambLisp as master or as slave	118
6.8	The state machine the master drives, per slave	120
6.9	PROFINET runs as two processes	125
7.1	Three-Tier Architecture	136
7.2	Boustrophedon Coverage Path	143

7.3 Bee Strategy Path 145

Chapter 16

Index

Index

~Lamb
 Lamb, [268](#)

1-Wire & DS18B20, [190](#)
 OneWire_install_mop3, [191](#)
 OneWire_mop3_read_bit, [191](#)
 OneWire_mop3_read_byte, [191](#)
 OneWire_mop3_reset, [191](#)
 OneWire_mop3_search, [191](#)
 OneWire_mop3_write_bit, [191](#)
 OneWire_mop3_write_byte, [191](#)

alist_to_dict
 Lamb, [280](#)

any_error_get_str
 Cell, [256](#)

any_str_get_chars
 Cell, [256](#)

any_svec_get_info
 Cell, [256](#)

append
 Lamb, [283](#)

append2
 Lamb, [283](#)

apply_proc_partial
 Lamb, [282](#)

array_matmul_cublas_f32
 CUDA GPU, [220](#)

array_matmul_cublas_f64
 CUDA GPU, [220](#)

as_Bool_t
 Cell, [243](#)

as_Bytest_t
 Cell, [245](#)

as_ByteVec_t
 Cell, [245](#)

as_Char_t
 Cell, [244](#)

as_Charst_t
 Cell, [245](#)

as_CharVec_t
 Cell, [245](#)

as_codepoint
 Cell, [244](#)

as_denominator
 Cell, [250](#)

as_float32
 Cell, [244](#)

as_float64
 Cell, [244](#)

as_int32
 Cell, [244](#)

as_int64
 Cell, [244](#)

as_numerator
 Cell, [250](#)

as_Portst_t
 Cell, [245](#)

as_Ptr_t
 Cell, [245](#)

assq
 Lamb, [276](#)

build_buildDate
 Lamb, [270](#)

build_buildRelease
 Lamb, [270](#)

build_isDebug
 Lamb, [269](#)

build_pushDate
 Lamb, [270](#)

build_pushUTC
 Lamb, [269](#)

build_UTC
 Lamb, [269](#)

build_version
 Lamb, [269](#)

bulk_alloc_pairs
 Lamb, [275](#)

bvec_elem_type
 Cell, [256](#)

bvec_set_elem_type
 Cell, [256](#)

bytecode_is_variadic
 Cell, [243](#)

bytecode_set_variadic
 Cell, [243](#)

caar
 Lamb, [273](#)

caddr

Lamb, [274](#)
caddr
 Lamb, [274](#)
cadr
 Lamb, [274](#)
Camera (ESP32-S3-EYE), [183](#)
 EYE_BL_CHANNEL, [184](#)
 EYE_BL_TIMER, [184](#)
 EYE_CAM_SIOC, [185](#)
 EYE_CAM_SIOD, [185](#)
 EYE_LCD_BAND_ROWS, [185](#)
 EYE_LCD_BL, [184](#)
 EYE_LCD_HZ, [185](#)
 EYE_LCD_RST, [184](#)
 EyeCam_install_mop3, [188](#)
 mop3_camera_capture, [187](#)
 mop3_camera_init, [186](#)
 mop3_camera_preview, [186](#)
 mop3_camera_set, [187](#)
 mop3_camera_thumb, [186](#)
 mop3_camera_to_lcd, [186](#)
 mop3_frame_brightness, [187](#)
 mop3_lcd_backlight, [187](#)
 mop3_lcd_box, [186](#)
 mop3_lcd_clear, [185](#)
 mop3_lcd_init, [185](#)
 mop3_lcd_rgb565, [186](#)
 mop3_lcd_testpat, [187](#)
car
 Lamb, [273](#)
car_float32
 Cell, [241](#)
car_int32
 Cell, [241](#)
car_is_embedded
 Cell, [241](#)
car_is_ptr
 Cell, [240](#)
car_word_type
 Cell, [240](#)
cdar
 Lamb, [274](#)
cddddr
 Lamb, [275](#)
cdddr
 Lamb, [274](#)
cddr
 Lamb, [274](#)
cdr
 Lamb, [273](#)
cdr_float32
 Cell, [242](#)
cdr_int32
 Cell, [241](#)
cdr_is_embedded
 Cell, [241](#)
cdr_is_ptr
 Cell, [240](#)
cdr_word_type
 Cell, [240](#)
Cell, [227](#)
 any_error_get_str, [256](#)
 any_str_get_chars, [256](#)
 any_svec_get_info, [256](#)
 as_Bool_t, [243](#)
 as_Bytest_t, [245](#)
 as_ByteVec_t, [245](#)
 as_Char_t, [244](#)
 as_Charst_t, [245](#)
 as_CharVec_t, [245](#)
 as_codepoint, [244](#)
 as_denominator, [250](#)
 as_float32, [244](#)
 as_float64, [244](#)
 as_int32, [244](#)
 as_int64, [244](#)
 as_numerator, [250](#)
 as_Portst_t, [245](#)
 as_Ptr_t, [245](#)
 bvec_elem_type, [256](#)
 bvec_set_elem_type, [256](#)
 bytecode_is_variadic, [243](#)
 bytecode_set_variadic, [243](#)
 car_float32, [241](#)
 car_int32, [241](#)
 car_is_embedded, [241](#)
 car_is_ptr, [240](#)
 car_word_type, [240](#)
 cdr_float32, [242](#)
 cdr_int32, [241](#)
 cdr_is_embedded, [241](#)
 cdr_is_ptr, [240](#)
 cdr_word_type, [240](#)
 cell_name, [256](#)
 coerce_cpx_imag, [255](#)
 coerce_cpx_real, [255](#)
 coerce_float32, [254](#)
 coerce_float64, [254](#)
 cpx_imag, [248](#)
 cpx_real, [248](#)
 denominator64, [255](#)
 dump, [257](#)
 exact_is_big, [250](#)
 exact_rank, [254](#)
 exact_sign, [250](#)
 F_BC_VARIADIC, [259](#)
 F_CAR_WT0, [258](#)
 F_CAR_WT1, [258](#)

F_CDR_WT0, 258
F_CDR_WT1, 258
F_GC01, 257
F_GC02, 257
F_GC04, 257
F_TAIL, 258
flags, 237
flags_clr, 237
flags_set, 237
gc_state, 242
GC_STATE_MASK, 258
gcstate_name, 257
get_car, 243
get_car_addr, 243
get_cdr, 243
hash, 246
hash_contents, 246
hash_sexpr, 246
inexact_rank, 254
int_as_i64, 249
is_any_bvec_atom, 239
is_any_embedded, 241
is_any_int, 239
is_any_pair, 238
is_any_real, 240
is_any_str_atom, 239
is_any_svec2n_atom, 239
is_any_svec_atom, 239
is_any_sym_atom, 239
is_atom, 238
is_bignum, 240
is_complex, 255
is_exact, 254
is_ident, 238
is_inexact, 254
is_pair, 238
mk_error, 251
mustbe_any_str_t, 253
mustbe_Bool_t, 252
mustbe_Char_t, 252
mustbe_codepoint, 252
mustbe_cppobj_t, 253
mustbe_float32, 253
mustbe_int32, 253
Ntypes, 237
num_rank, 255
numerator64, 255
prechecked_cppobj_get_deleter, 253
prechecked_cppobj_get_ptr, 253
prechecked_str_ext_get_chars, 252
prechecked_str_heap_get_chars, 252
prechecked_str_imm_get_chars, 252
rat_car_ptr, 249
rat_cdr_ptr, 249
rat_is_big, 250
rat_set_car_i32, 248
rat_set_car_ptr, 249
rat_set_cdr_i32, 249
rat_set_cdr_ptr, 249
rplaca, 238
rplacd, 238
set, 246–248, 250, 251
T_ANY_HEAP_SVEC, 235
T_BIGNUM, 235
T_BOOL, 235
T_BVEC_EXT, 235
T_BVEC_HEAP, 235
T_BVEC_IMM, 235
T_BYTECODE, 236
T_BYTECODE_N, 236
T_CHAR, 235
T_COMPLEX, 236
T_CPP_HEAP, 235
T_DICT, 236
T_ERROR, 236
T_FLOAT32, 235
T_FLOAT64, 235
T_GENSYM, 235
T_IDENT, 237
T_INT32, 235
T_INT64, 235
T_MACRO, 236
T_MOP3_NPROC, 236
T_MOP3_PROC, 236
T_NEEDS_FINALIZING, 235
T_NIL, 236
T_NPROC, 236
T_PAIR, 236
T_PORT_HEAP, 235
T_PROC, 236
T_RATIONAL, 236
T_STR_EXT, 235
T_STR_HEAP, 235
T_STR_IMM, 235
T_SVEC2N_HEAP, 235
T_SVEC_HEAP, 235
T_SVEC_IMM, 236
T_SYM_HEAP, 235
T_THUNK_BODY, 236
T_THUNK_SEXPR, 236
T_UNDEF, 236
T_VALUES, 236
T_VOID, 236
tail_state, 242
tail_state_clr, 242
tail_state_set, 242
type, 237
type_name, 257

- zero, [237](#)
- cell_block_size
 - LambPreSettings, [298](#)
- cell_name
 - Cell, [256](#)
- coerce_cpx_imag
 - Cell, [255](#)
- coerce_cpx_real
 - Cell, [255](#)
- coerce_float32
 - Cell, [254](#)
- coerce_float64
 - Cell, [254](#)
- cons
 - Lamb, [271](#)
- cpx_imag
 - Cell, [248](#)
- cpx_real
 - Cell, [248](#)
- CUDA GPU, [219](#)
 - array_matmul_cublas_f32, [220](#)
 - array_matmul_cublas_f64, [220](#)
 - cudaDevicePtrDestroyer, [220](#)
 - cudaHostPtrDestroyer, [220](#)
 - mop3_cpp_vector_add, [220](#)
 - mop3_cudaDeviceSynchronize, [222](#)
 - mop3_cudaFree, [221](#)
 - mop3_cudaGetDeviceCount, [222](#)
 - mop3_cudaGetDeviceProp, [222](#)
 - mop3_cudakernel_vectorAdd, [221](#)
 - mop3_cudaLaunch_kernel1, [221](#)
 - mop3_cudaMalloc, [221](#)
 - mop3_cudaMallocHost, [221](#)
 - mop3_cudaMemcpy, [221](#)
 - mop3_cudaMemset, [222](#)
 - mop3_cudaStreamCreate, [222](#)
 - mop3_cudaStreamSynchronize, [222](#)
- cudaDevicePtrDestroyer
 - CUDA GPU, [220](#)
 - HIP / ROCm GPU, [223](#)
- cudaHostPtrDestroyer
 - CUDA GPU, [220](#)
 - HIP / ROCm GPU, [223](#)
- debug
 - Lamb, [269](#)
- denominator64
 - Cell, [255](#)
- diag_memory_report
 - Lamb, [296](#)
- dict_add_alist_frame
 - Lamb, [278](#)
- dict_add_empty_frame
 - Lamb, [277](#)
- dict_add_keyval_frame
 - Lamb, [278](#)
- dict_add_keyval_frame_argv
 - Lamb, [278](#)
- dict_add_vector_frame
 - Lamb, [278](#)
- dict_analyze
 - Lamb, [281](#)
- dict_append_bang
 - Lamb, [278](#)
- dict_bind_alist_bang
 - Lamb, [279](#)
- dict_bind_bang
 - Lamb, [278](#)
- dict_bind_local_bang
 - Lamb, [279](#)
- dict_keys
 - Lamb, [280](#)
- dict_new
 - Lamb, [277](#)
- dict_rebind_alist_bang
 - Lamb, [279](#)
- dict_rebind_bang
 - Lamb, [279](#)
- dict_ref
 - Lamb, [279](#)
- dict_ref_q
 - Lamb, [279](#)
- dict_to_2list
 - Lamb, [280](#)
- dict_to_alist
 - Lamb, [280](#)
- dict_values
 - Lamb, [280](#)
- dump
 - Cell, [257](#)
- end
 - Lamb, [268](#)
- eq_q
 - Lamb, [276](#)
- equal_q
 - Lamb, [276](#)
- eqv_q
 - Lamb, [276](#)
- error
 - NcgErrJmp, [301](#)
- ESP32 System, [181](#)
 - mop3_esp32_efuse_mac_get_default, [182](#)
 - mop3_esp32_efuse_pkg_ver, [182](#)
 - mop3_esp32_efuse_read_block, [182](#)
 - mop3_esp32_efuse_write_block, [182](#)
 - mop3_esp32_flash_probe, [182](#)
 - mop3_esp32_gcmem, [182](#)

mop3_esp32_heapinfo, 182

eval
 Lamb, 282

eval_list
 Lamb, 282

eval_recursion_check
 Lamb, 296

eval_recursion_overflow
 Lamb, 296

eval_root_depth_max_
 Lamb, 297

eval_stack_base_
 Lamb, 297

eval_stack_budget_
 Lamb, 297

exact_is_big
 Cell, 250

exact_rank
 Cell, 254

exact_sign
 Cell, 250

expand
 Lamb, 270

extension_block_size
 LambPreSettings, 298

EYE_BL_CHANNEL
 Camera (ESP32-S3-EYE), 184

EYE_BL_TIMER
 Camera (ESP32-S3-EYE), 184

EYE_CAM_SIOC
 Camera (ESP32-S3-EYE), 185

EYE_CAM_SIOD
 Camera (ESP32-S3-EYE), 185

EYE_LCD_BAND_ROWS
 Camera (ESP32-S3-EYE), 185

EYE_LCD_BL
 Camera (ESP32-S3-EYE), 184

EYE_LCD_HZ
 Camera (ESP32-S3-EYE), 185

EYE_LCD_RST
 Camera (ESP32-S3-EYE), 184

EyeCam_install_mop3
 Camera (ESP32-S3-EYE), 188

F_BC_VARIADIC
 Cell, 259

F_CAR_WT0
 Cell, 258

F_CAR_WT1
 Cell, 258

F_CDR_WT0
 Cell, 258

F_CDR_WT1
 Cell, 258

F_GC01
 Cell, 257

F_GC02
 Cell, 257

F_GC04
 Cell, 257

F_TAIL
 Cell, 258

flags
 Cell, 237

flags_clr
 Cell, 237

flags_set
 Cell, 237

gc_collect
 Lamb, 272

gc_depth
 NcgErrJump, 301

gc_idle_task_us
 Lamb, 271

gc_pass_n
 Lamb, 272

gc_Qm
 Lamb, 272

gc_root_depth
 Lamb, 271

gc_root_pop
 Lamb, 271

gc_root_push
 Lamb, 271

gc_root_setdepth
 Lamb, 271

gc_state
 Cell, 242

GC_STATE_MASK
 Cell, 258

gcstate_name
 Cell, 257

gensym
 Lamb, 286

get_car
 Cell, 243

get_car_addr
 Cell, 243

get_cdr
 Cell, 243

GPIO & Basic I/O, 176

 mop3_cio_analogRead, 177

 mop3_cio_analogReadResolution, 177

 mop3_cio_analogReference, 178

 mop3_cio_analogWrite, 177

 mop3_cio_analogWriteResolution, 178

 mop3_cio_attachInterrupt, 180

- mop3_cio_bit, 179
- mop3_cio_bitClear, 179
- mop3_cio_bitRead, 179
- mop3_cio_bitSet, 179
- mop3_cio_bitWrite, 180
- mop3_cio_delay_ms, 177
- mop3_cio_delay_us, 177
- mop3_cio_detachInterrupt, 180
- mop3_cio_digitalPinToInterrupt, 180
- mop3_cio_digitalRead, 176
- mop3_cio_digitalWrite, 176
- mop3_cio_highByte, 180
- mop3_cio_interrupts, 181
- mop3_cio_lowByte, 180
- mop3_cio_noInterrupts, 181
- mop3_cio_noTone, 178
- mop3_cio_pinMode, 177
- mop3_cio_pulseIn, 178
- mop3_cio_pulseInLong, 178
- mop3_cio_shiftIn, 178
- mop3_cio_shiftOut, 179
- mop3_cio_tone, 179
- Hardware & Platform Builtins (xmop3), 175
- hash
 - Cell, 246
- hash_contents
 - Cell, 246
- hash_sexpr
 - Cell, 246
- HIP / ROCm GPU, 223
 - cudaDevicePtrDestroyer, 223
 - cudaHostPtrDestroyer, 223
 - Hip_install_mop3, 226
 - mop3_cpp_vector_add, 224
 - mop3_cudaDeviceSynchronize, 225
 - mop3_cudaFree, 224
 - mop3_cudaGetDeviceCount, 225
 - mop3_cudaGetDeviceProp, 226
 - mop3_cudakernel_vectorAdd, 224
 - mop3_cudaLaunch_kernel1, 224
 - mop3_cudaMalloc, 224
 - mop3_cudaMallocHost, 224
 - mop3_cudaMemcpy, 225
 - mop3_cudaMemset, 225
 - mop3_cudaStreamCreate, 225
 - mop3_cudaStreamSynchronize, 225
- Hip_install_mop3
 - HIP / ROCm GPU, 226
- I2C (Wire), 214
 - Wire_mop3_available, 216
 - Wire_mop3_begin, 215
 - Wire_mop3_beginTransmission, 216
 - Wire_mop3_clearWireTimeoutFlag, 217
 - Wire_mop3_end, 216
 - Wire_mop3_endTransmission, 217
 - Wire_mop3_getWireTimeoutFlag, 217
 - Wire_mop3_onReceive, 216
 - Wire_mop3_onRequest, 217
 - Wire_mop3_read, 216
 - Wire_mop3_requestFrom, 215
 - Wire_mop3_setClock, 216
 - Wire_mop3_setPins, 215
 - Wire_mop3_setTimeout, 215
 - Wire_mop3_write, 217
- inexact_rank
 - Cell, 254
- int_as_i64
 - Cell, 249
- integrity_check
 - Lamb, 281
- is_any_bvec_atom
 - Cell, 239
- is_any_embedded
 - Cell, 241
- is_any_int
 - Cell, 239
- is_any_pair
 - Cell, 238
- is_any_real
 - Cell, 240
- is_any_str_atom
 - Cell, 239
- is_any_svec2n_atom
 - Cell, 239
- is_any_svec_atom
 - Cell, 239
- is_any_sym_atom
 - Cell, 239
- is_atom
 - Cell, 238
- is_bignum
 - Cell, 240
- is_complex
 - Cell, 255
- is_exact
 - Cell, 254
- is_ident
 - Cell, 238
- is_inexact
 - Cell, 254
- is_pair
 - Cell, 238
- jmpbuf
 - NcgErrJump, 301
- Lamb, 259

~Lamb, 268
alist_to_dict, 280
append, 283
append2, 283
apply_proc_partial, 282
assq, 276
build_buildDate, 270
build_buildRelease, 270
build_isDebug, 269
build_pushDate, 270
build_pushUTC, 269
build_UTC, 269
build_version, 269
bulk_alloc_pairs, 275
caar, 273
caddr, 274
caddr, 274
cadr, 274
car, 273
cdar, 274
cddddr, 275
cdddr, 274
cddr, 274
cdr, 273
cons, 271
debug, 269
diag_memory_report, 296
dict_add_alist_frame, 278
dict_add_empty_frame, 277
dict_add_keyval_frame, 278
dict_add_keyval_frame_argv, 278
dict_add_vector_frame, 278
dict_analyze, 281
dict_append_bang, 278
dict_bind_alist_bang, 279
dict_bind_bang, 278
dict_bind_local_bang, 279
dict_keys, 280
dict_new, 277
dict_rebind_alist_bang, 279
dict_rebind_bang, 279
dict_ref, 279
dict_ref_q, 279
dict_to_2list, 280
dict_to_alist, 280
dict_values, 280
end, 268
eq_q, 276
equal_q, 276
eqv_q, 276
eval, 282
eval_list, 282
eval_recursion_check, 296
eval_recursion_overflow, 296
eval_root_depth_max_, 297
eval_stack_base_, 297
eval_stack_budget_, 297
expand, 270
gc_collect, 272
gc_idle_task_us, 271
gc_pass_n, 272
gc_Qm, 272
gc_root_depth, 271
gc_root_pop, 271
gc_root_push, 271
gc_root_setdepth, 271
gensym, 286
integrity_check, 281
lamb_oblist, 283
list_analyze, 284
list_copy, 284
list_to_vector, 284
load, 283
log, 268
loop, 268
macroexpand, 293, 294
macroexpand1, 294
map_proc, 282
mem_dbg_t_gc_total_us, 272
mk_bool, 286
mk_bytevector, 289
mk_bytevector_ext, 289
mk_character, 286
mk_complex, 287
mk_error, 285
mk_float32, 287
mk_float64, 287
mk_hashtbl, 290
mk_ident, 293
mk_input_bytevector_port, 292
mk_input_file_port, 290
mk_input_string_port, 291
mk_int32, 286
mk_int64, 287
mk_integer, 286
mk_intvector, 289
mk_macro, 292
mk_Mop3_procst_argv_t, 292
mk_Mop3_procst_t, 292
mk_nprocedure, 293
mk_number, 288
mk_output_bytevector_port, 292
mk_output_file_port, 290
mk_output_string_port, 292
mk_procedure, 293
mk_rational, 287
mk_realvector, 290
mk_serial_port, 290

mk_sharp_const, 288
mk_symbol, 289
mk_symbol_or_number, 288
mk_syserror, 285
mk_tcp_client_port, 291
mk_tcp_server_port, 291
mk_thunk_body, 293
mk_thunk_sexpr, 293
mk_tls_client_port, 291
mk_tls_server_port, 291
mk_vector, 290
Mop3st_argv_t, 267
Mop3st_t, 267
ncg_box_fast, 296
ncg_compile, 295
ncg_cons_fast, 295
ncg_disassemble, 295
ncg_err_ctx, 297
ncg_eval, 294
ncg_eval_argv, 294
ncg_frame_pool_count, 297
ncg_frame_pool_head, 296
ncg_frame_pool_init, 273
ncg_frame_pool_n, 272
ncg_frame_top, 296
ncg_mark_frames, 294
ncg_vector_set_ptr, 295
ncg_vref_fast, 295
ncg_vset_fast, 295
oblist_analyze, 277
oblist_intern, 277
oblist_query, 276
oblist_test, 276
payload_alloc, 288
prechecked_anypair_get_car, 273
prechecked_anypair_get_cdr, 273
printf, 268, 282
rat_store_i64, 287
read, 281
reverse_bang, 275
server_accept_port, 291
set_car_bang, 275
set_cdr_bang, 275
set_current_input_port, 283
set_current_output_port, 283
set_repl_quiet, 272
setup, 268
sparsevec_to_vector, 285
sprintf, 281
tag_error, 285
tcons, 270
twolist_to_dict, 280
vector_copy, 284
vector_set_bang, 275
vector_to_list, 284
vector_to_sparsevec, 284
write_or_display, 281
write_simple, 281
lamb_oblist
 Lamb, 283
LambPreSettings, 298
 cell_block_size, 298
 extension_block_size, 298
 load, 298
 ncg_frame_pool_size, 298
LambStdioClass, 299
LCD1602_Character_Display, 188
 LCD1602_mop3_backlight, 189
 LCD1602_mop3_begin, 189
 LCD1602_mop3_clear, 189
 LCD1602_mop3_cursor, 189
 LCD1602_mop3_display, 189
 LCD1602_mop3_home, 189
 LCD1602_mop3_print, 190
 LCD1602_mop3_setCursor, 190
 LCD1602_mop3_write, 190
LCD1602_mop3_backlight
 LCD1602_Character_Display, 189
LCD1602_mop3_begin
 LCD1602_Character_Display, 189
LCD1602_mop3_clear
 LCD1602_Character_Display, 189
LCD1602_mop3_cursor
 LCD1602_Character_Display, 189
LCD1602_mop3_display
 LCD1602_Character_Display, 189
LCD1602_mop3_home
 LCD1602_Character_Display, 189
LCD1602_mop3_print
 LCD1602_Character_Display, 190
LCD1602_mop3_setCursor
 LCD1602_Character_Display, 190
LCD1602_mop3_write
 LCD1602_Character_Display, 190
list_analyze
 Lamb, 284
list_copy
 Lamb, 284
list_to_vector
 Lamb, 284
LL_File, 299
LL_File_System, 300
load
 Lamb, 283
 LambPreSettings, 298
log
 Lamb, 268
loop

- Lamb, [268](#)
- macroexpand
 - Lamb, [293](#), [294](#)
- macroexpand1
 - Lamb, [294](#)
- map_proc
 - Lamb, [282](#)
- mem_dbg_t_gc_total_us
 - Lamb, [272](#)
- mk_bool
 - Lamb, [286](#)
- mk_bytevector
 - Lamb, [289](#)
- mk_bytevector_ext
 - Lamb, [289](#)
- mk_character
 - Lamb, [286](#)
- mk_complex
 - Lamb, [287](#)
- mk_error
 - Cell, [251](#)
 - Lamb, [285](#)
- mk_float32
 - Lamb, [287](#)
- mk_float64
 - Lamb, [287](#)
- mk_hashtbl
 - Lamb, [290](#)
- mk_ident
 - Lamb, [293](#)
- mk_input_bytevector_port
 - Lamb, [292](#)
- mk_input_file_port
 - Lamb, [290](#)
- mk_input_string_port
 - Lamb, [291](#)
- mk_int32
 - Lamb, [286](#)
- mk_int64
 - Lamb, [287](#)
- mk_integer
 - Lamb, [286](#)
- mk_intvector
 - Lamb, [289](#)
- mk_macro
 - Lamb, [292](#)
- mk_Mop3_procst_argv_t
 - Lamb, [292](#)
- mk_Mop3_procst_t
 - Lamb, [292](#)
- mk_nprocedure
 - Lamb, [293](#)
- mk_number
 - Lamb, [288](#)
- mk_output_bytevector_port
 - Lamb, [292](#)
- mk_output_file_port
 - Lamb, [290](#)
- mk_output_string_port
 - Lamb, [292](#)
- mk_procedure
 - Lamb, [293](#)
- mk_rational
 - Lamb, [287](#)
- mk_realvector
 - Lamb, [290](#)
- mk_serial_port
 - Lamb, [290](#)
- mk_sharp_const
 - Lamb, [288](#)
- mk_symbol
 - Lamb, [289](#)
- mk_symbol_or_number
 - Lamb, [288](#)
- mk_syserror
 - Lamb, [285](#)
- mk_tcp_client_port
 - Lamb, [291](#)
- mk_tcp_server_port
 - Lamb, [291](#)
- mk_thunk_body
 - Lamb, [293](#)
- mk_thunk_sexpr
 - Lamb, [293](#)
- mk_tls_client_port
 - Lamb, [291](#)
- mk_tls_server_port
 - Lamb, [291](#)
- mk_vector
 - Lamb, [290](#)
- mop3_camera_capture
 - Camera (ESP32-S3-EYE), [187](#)
- mop3_camera_init
 - Camera (ESP32-S3-EYE), [186](#)
- mop3_camera_preview
 - Camera (ESP32-S3-EYE), [186](#)
- mop3_camera_set
 - Camera (ESP32-S3-EYE), [187](#)
- mop3_camera_thumb
 - Camera (ESP32-S3-EYE), [186](#)
- mop3_camera_to_lcd
 - Camera (ESP32-S3-EYE), [186](#)
- mop3_cio_analogRead
 - GPIO & Basic I/O, [177](#)
- mop3_cio_analogReadResolution
 - GPIO & Basic I/O, [177](#)
- mop3_cio_analogReference

- GPIO & Basic I/O, [178](#)
- mop3_cio_analogWrite
 - GPIO & Basic I/O, [177](#)
- mop3_cio_analogWriteResolution
 - GPIO & Basic I/O, [178](#)
- mop3_cio_attachInterrupt
 - GPIO & Basic I/O, [180](#)
- mop3_cio_bit
 - GPIO & Basic I/O, [179](#)
- mop3_cio_bitClear
 - GPIO & Basic I/O, [179](#)
- mop3_cio_bitRead
 - GPIO & Basic I/O, [179](#)
- mop3_cio_bitSet
 - GPIO & Basic I/O, [179](#)
- mop3_cio_bitWrite
 - GPIO & Basic I/O, [180](#)
- mop3_cio_delay_ms
 - GPIO & Basic I/O, [177](#)
- mop3_cio_delay_us
 - GPIO & Basic I/O, [177](#)
- mop3_cio_detachInterrupt
 - GPIO & Basic I/O, [180](#)
- mop3_cio_digitalPinToInterrupt
 - GPIO & Basic I/O, [180](#)
- mop3_cio_digitalRead
 - GPIO & Basic I/O, [176](#)
- mop3_cio_digitalWrite
 - GPIO & Basic I/O, [176](#)
- mop3_cio_highByte
 - GPIO & Basic I/O, [180](#)
- mop3_cio_interrupts
 - GPIO & Basic I/O, [181](#)
- mop3_cio_lowByte
 - GPIO & Basic I/O, [180](#)
- mop3_cio_noInterrupts
 - GPIO & Basic I/O, [181](#)
- mop3_cio_noTone
 - GPIO & Basic I/O, [178](#)
- mop3_cio_pinMode
 - GPIO & Basic I/O, [177](#)
- mop3_cio_pulseIn
 - GPIO & Basic I/O, [178](#)
- mop3_cio_pulseInLong
 - GPIO & Basic I/O, [178](#)
- mop3_cio_shiftIn
 - GPIO & Basic I/O, [178](#)
- mop3_cio_shiftOut
 - GPIO & Basic I/O, [179](#)
- mop3_cio_tone
 - GPIO & Basic I/O, [179](#)
- mop3_cpp_vector_add
 - CUDA GPU, [220](#)
 - HIP / ROCm GPU, [224](#)
- mop3_cudaDeviceSynchronize
 - CUDA GPU, [222](#)
 - HIP / ROCm GPU, [225](#)
- mop3_cudaFree
 - CUDA GPU, [221](#)
 - HIP / ROCm GPU, [224](#)
- mop3_cudaGetDeviceCount
 - CUDA GPU, [222](#)
 - HIP / ROCm GPU, [225](#)
- mop3_cudaGetDeviceProp
 - CUDA GPU, [222](#)
 - HIP / ROCm GPU, [226](#)
- mop3_cudakernel_vectorAdd
 - CUDA GPU, [221](#)
 - HIP / ROCm GPU, [224](#)
- mop3_cudaLaunch_kernel1
 - CUDA GPU, [221](#)
 - HIP / ROCm GPU, [224](#)
- mop3_cudaMalloc
 - CUDA GPU, [221](#)
 - HIP / ROCm GPU, [224](#)
- mop3_cudaMallocHost
 - CUDA GPU, [221](#)
 - HIP / ROCm GPU, [224](#)
- mop3_cudaMemcpy
 - CUDA GPU, [221](#)
 - HIP / ROCm GPU, [225](#)
- mop3_cudaMemset
 - CUDA GPU, [222](#)
 - HIP / ROCm GPU, [225](#)
- mop3_cudaStreamCreate
 - CUDA GPU, [222](#)
 - HIP / ROCm GPU, [225](#)
- mop3_cudaStreamSynchronize
 - CUDA GPU, [222](#)
 - HIP / ROCm GPU, [225](#)
- mop3_esp32_efuse_mac_get_default
 - ESP32 System, [182](#)
- mop3_esp32_efuse_pkg_ver
 - ESP32 System, [182](#)
- mop3_esp32_efuse_read_block
 - ESP32 System, [182](#)
- mop3_esp32_efuse_write_block
 - ESP32 System, [182](#)
- mop3_esp32_flash_probe
 - ESP32 System, [182](#)
- mop3_esp32_gcmem
 - ESP32 System, [182](#)
- mop3_esp32_heapinfo
 - ESP32 System, [182](#)
- mop3_frame_brightness
 - Camera (ESP32-S3-EYE), [187](#)
- mop3_lcd_backlight
 - Camera (ESP32-S3-EYE), [187](#)

mop3_lcd_box
 Camera (ESP32-S3-EYE), [186](#)
mop3_lcd_clear
 Camera (ESP32-S3-EYE), [185](#)
mop3_lcd_init
 Camera (ESP32-S3-EYE), [185](#)
mop3_lcd_rgb565
 Camera (ESP32-S3-EYE), [186](#)
mop3_lcd_testpat
 Camera (ESP32-S3-EYE), [187](#)
mop3_neopixelWrite
 WS2812 / NeoPixel LEDs, [218](#)
Mop3st_argv_t
 Lamb, [267](#)
Mop3st_t
 Lamb, [267](#)
mustbe_any_str_t
 Cell, [253](#)
mustbe_Bool_t
 Cell, [252](#)
mustbe_Char_t
 Cell, [252](#)
mustbe_codepoint
 Cell, [252](#)
mustbe_cppobj_t
 Cell, [253](#)
mustbe_float32
 Cell, [253](#)
mustbe_int32
 Cell, [253](#)

ncg_box_fast
 Lamb, [296](#)
ncg_compile
 Lamb, [295](#)
ncg_cons_fast
 Lamb, [295](#)
ncg_disassemble
 Lamb, [295](#)
ncg_err_ctx
 Lamb, [297](#)
ncg_eval
 Lamb, [294](#)
ncg_eval_argv
 Lamb, [294](#)
ncg_frame_pool_count
 Lamb, [297](#)
ncg_frame_pool_head
 Lamb, [296](#)
ncg_frame_pool_init
 Lamb, [273](#)
ncg_frame_pool_n
 Lamb, [272](#)
ncg_frame_pool_size
 LambPreSettings, [298](#)
ncg_frame_top
 Lamb, [296](#)
ncg_mark_frames
 Lamb, [294](#)
ncg_vector_set_ptr
 Lamb, [295](#)
ncg_vref_fast
 Lamb, [295](#)
ncg_vset_fast
 Lamb, [295](#)
NcgErrJump, [300](#)
 error, [301](#)
 gc_depth, [301](#)
 jmpbuf, [301](#)
Ntypes
 Cell, [237](#)
num_rank
 Cell, [255](#)
numerator64
 Cell, [255](#)

oblist_analyze
 Lamb, [277](#)
oblist_intern
 Lamb, [277](#)
oblist_query
 Lamb, [276](#)
oblist_test
 Lamb, [276](#)
OneWire_install_mop3
 1-Wire & DS18B20, [191](#)
OneWire_mop3_read_bit
 1-Wire & DS18B20, [191](#)
OneWire_mop3_read_byte
 1-Wire & DS18B20, [191](#)
OneWire_mop3_reset
 1-Wire & DS18B20, [191](#)
OneWire_mop3_search
 1-Wire & DS18B20, [191](#)
OneWire_mop3_write_bit
 1-Wire & DS18B20, [191](#)
OneWire_mop3_write_byte
 1-Wire & DS18B20, [191](#)
ota_install_mop3
 Over-the-Air Firmware Updates, [193](#)
ota_mop3_boot_ok
 Over-the-Air Firmware Updates, [193](#)
ota_mop3_reboot_to_loader
 Over-the-Air Firmware Updates, [193](#)
ota_mop3_request_console
 Over-the-Air Firmware Updates, [192](#)
ota_mop3_request_install
 Over-the-Air Firmware Updates, [192](#)

- ota_mop3_request_recover
 - Over-the-Air Firmware Updates, 192
- ota_mop3_set_network
 - Over-the-Air Firmware Updates, 193
- Over-the-Air Firmware Updates, 192
 - ota_install_mop3, 193
 - ota_mop3_boot_ok, 193
 - ota_mop3_reboot_to_loader, 193
 - ota_mop3_request_console, 192
 - ota_mop3_request_install, 192
 - ota_mop3_request_recover, 192
 - ota_mop3_set_network, 193
- payload_alloc
 - Lamb, 288
- PCA9685 PWM / Servo Driver, 194
 - PCA9685_mop3_begin, 194
 - PCA9685_mop3_getSingleDeviceFrequency, 195
 - PCA9685_mop3_setChannelDutyCycle, 195
 - PCA9685_mop3_setChannelOnAndOffTime, 195
 - PCA9685_mop3_setChannelPulseWidth, 195
 - PCA9685_mop3_setChannelServoPulseDuration, 195
 - PCA9685_mop3_setSingleDeviceToFrequency, 194
 - PCA9685_mop3_setupSingleDevice, 194
- PCA9685_mop3_begin
 - PCA9685 PWM / Servo Driver, 194
- PCA9685_mop3_getSingleDeviceFrequency
 - PCA9685 PWM / Servo Driver, 195
- PCA9685_mop3_setChannelDutyCycle
 - PCA9685 PWM / Servo Driver, 195
- PCA9685_mop3_setChannelOnAndOffTime
 - PCA9685 PWM / Servo Driver, 195
- PCA9685_mop3_setChannelPulseWidth
 - PCA9685 PWM / Servo Driver, 195
- PCA9685_mop3_setChannelServoPulseDuration
 - PCA9685 PWM / Servo Driver, 195
- PCA9685_mop3_setSingleDeviceToFrequency
 - PCA9685 PWM / Servo Driver, 194
- PCA9685_mop3_setupSingleDevice
 - PCA9685 PWM / Servo Driver, 194
- Platform (generic), 196
 - Platform_mop3_cell_block_size, 199
 - Platform_mop3_delay_ms, 198
 - Platform_mop3_exit, 198
 - Platform_mop3_expand_to_n_blocks, 200
 - Platform_mop3_free_heap, 197
 - Platform_mop3_free_stack, 197
 - Platform_mop3_gc_budget_ns, 200
 - Platform_mop3_gc_collect, 198
 - Platform_mop3_gc_diag, 200
 - Platform_mop3_gc_idle_task, 198
 - Platform_mop3_gcloud_target_pct, 201
 - Platform_mop3_lamb_reboot, 198
 - Platform_mop3_loop_elapsed_ms, 197
 - Platform_mop3_loop_elapsed_us, 197
 - Platform_mop3_max_cell_blocks, 200
 - Platform_mop3_micros, 197
 - Platform_mop3_millis, 197
 - Platform_mop3_n_cell_blocks, 199
 - Platform_mop3_ncg_frame_pool_init, 202
 - Platform_mop3_ncg_frame_pool_n, 201
 - Platform_mop3_ncg_pool, 197
 - Platform_mop3_nfree, 199
 - Platform_mop3_random_bytevector, 198
 - Platform_mop3_random_integer, 199
 - Platform_mop3_random_real, 199
 - Platform_mop3_repl_quiet, 202
 - Platform_mop3_set_extension_block_size, 200
 - Platform_mop3_set_gc_budget_ns, 201
 - Platform_mop3_set_gcloud_target_pct, 201
 - Platform_mop3_set_max_cell_blocks, 200
 - Platform_mop3_set_verbosity, 201
 - Platform_mop3_total_cells, 199
 - Platform_mop3_verbosity, 201
 - Platform_mop3_cell_block_size
 - Platform (generic), 199
 - Platform_mop3_delay_ms
 - Platform (generic), 198
 - Platform_mop3_exit
 - Platform (generic), 198
 - Platform_mop3_expand_to_n_blocks
 - Platform (generic), 200
 - Platform_mop3_free_heap
 - Platform (generic), 197
 - Platform_mop3_free_stack
 - Platform (generic), 197
 - Platform_mop3_gc_budget_ns
 - Platform (generic), 200
 - Platform_mop3_gc_collect
 - Platform (generic), 198
 - Platform_mop3_gc_diag
 - Platform (generic), 200
 - Platform_mop3_gc_idle_task
 - Platform (generic), 198
 - Platform_mop3_gcloud_target_pct
 - Platform (generic), 201
 - Platform_mop3_lamb_reboot
 - Platform (generic), 198
 - Platform_mop3_loop_elapsed_ms
 - Platform (generic), 197
 - Platform_mop3_loop_elapsed_us
 - Platform (generic), 197
 - Platform_mop3_max_cell_blocks
 - Platform (generic), 200

Platform_mop3_micros
 Platform (generic), 197
Platform_mop3_millis
 Platform (generic), 197
Platform_mop3_n_cell_blocks
 Platform (generic), 199
Platform_mop3_ncg_frame_pool_init
 Platform (generic), 202
Platform_mop3_ncg_frame_pool_n
 Platform (generic), 201
Platform_mop3_ncg_pool
 Platform (generic), 197
Platform_mop3_nfree
 Platform (generic), 199
Platform_mop3_random_bytevector
 Platform (generic), 198
Platform_mop3_random_integer
 Platform (generic), 199
Platform_mop3_random_real
 Platform (generic), 199
Platform_mop3_repl_quiet
 Platform (generic), 202
Platform_mop3_set_extension_block_size
 Platform (generic), 200
Platform_mop3_set_gc_budget_ns
 Platform (generic), 201
Platform_mop3_set_gpload_target_pct
 Platform (generic), 201
Platform_mop3_set_max_cell_blocks
 Platform (generic), 200
Platform_mop3_set_verbosity
 Platform (generic), 201
Platform_mop3_total_cells
 Platform (generic), 199
Platform_mop3_verbosity
 Platform (generic), 201
prechecked_anypair_get_car
 Lamb, 273
prechecked_anypair_get_cdr
 Lamb, 273
prechecked_cppobj_get_deleter
 Cell, 253
prechecked_cppobj_get_ptr
 Cell, 253
prechecked_str_ext_get_chars
 Cell, 252
prechecked_str_heap_get_chars
 Cell, 252
prechecked_str_imm_get_chars
 Cell, 252
printf
 Lamb, 268, 282
rat_car_ptr
 Cell, 249
rat_cdr_ptr
 Cell, 249
rat_is_big
 Cell, 250
rat_set_car_i32
 Cell, 248
rat_set_car_ptr
 Cell, 249
rat_set_cdr_i32
 Cell, 249
rat_set_cdr_ptr
 Cell, 249
rat_store_i64
 Lamb, 287
read
 Lamb, 281
Real-time Lisp for Microprocessors, 1
reverse_bang
 Lamb, 275
rplaca
 Cell, 238
rplacd
 Cell, 238
server_accept_port
 Lamb, 291
set
 Cell, 246–248, 250, 251
set_car_bang
 Lamb, 275
set_cdr_bang
 Lamb, 275
set_current_input_port
 Lamb, 283
set_current_output_port
 Lamb, 283
set_repl_quiet
 Lamb, 272
setup
 Lamb, 268
sonar
 Ultrasonic Sonar (HC-SR04), 204
sonar_echo_isr
 Ultrasonic Sonar (HC-SR04), 203
Sonar_install_mop3
 Ultrasonic Sonar (HC-SR04), 204
Sonar_mop3_begin
 Ultrasonic Sonar (HC-SR04), 203
Sonar_mop3_ping
 Ultrasonic Sonar (HC-SR04), 203
Sonar_mop3_poke
 Ultrasonic Sonar (HC-SR04), 203
Sonar_mop3_result

- Ultrasonic Sonar (HC-SR04), [203](#)
- Sonar_mop3_set_min_interval_us
 - Ultrasonic Sonar (HC-SR04), [204](#)
- Sonar_mop3_set_timeout_us
 - Ultrasonic Sonar (HC-SR04), [204](#)
- Sonar_mop3_start
 - Ultrasonic Sonar (HC-SR04), [203](#)
- Sonar_mop3_us_to_distance
 - Ultrasonic Sonar (HC-SR04), [204](#)
- sparsevec_to_vector
 - Lamb, [285](#)
- SPI Bus, [205](#)
 - SPI_install_mop3, [206](#)
 - SPI_mop3_begin, [205](#)
 - SPI_mop3_beginTransaction, [205](#)
 - SPI_mop3_end, [206](#)
 - SPI_mop3_endTransaction, [205](#)
 - SPI_mop3_transfer, [205](#)
 - SPI_mop3_transfer_bytevector, [206](#)
 - SPI_mop3_write_bytevector, [206](#)
- SPI_install_mop3
 - SPI Bus, [206](#)
- SPI_mop3_begin
 - SPI Bus, [205](#)
- SPI_mop3_beginTransaction
 - SPI Bus, [205](#)
- SPI_mop3_end
 - SPI Bus, [206](#)
- SPI_mop3_endTransaction
 - SPI Bus, [205](#)
- SPI_mop3_transfer
 - SPI Bus, [205](#)
- SPI_mop3_transfer_bytevector
 - SPI Bus, [206](#)
- SPI_mop3_write_bytevector
 - SPI Bus, [206](#)
- sprintf
 - Lamb, [281](#)
- T_ANY_HEAP_SVEC
 - Cell, [235](#)
- T_BIGNUM
 - Cell, [235](#)
- T_BOOL
 - Cell, [235](#)
- T_BVEC_EXT
 - Cell, [235](#)
- T_BVEC_HEAP
 - Cell, [235](#)
- T_BVEC_IMM
 - Cell, [235](#)
- T_BYTECODE
 - Cell, [236](#)
- T_BYTECODE_N
 - Cell, [236](#)
- T_CHAR
 - Cell, [235](#)
- T_COMPLEX
 - Cell, [236](#)
- T_CPP_HEAP
 - Cell, [235](#)
- T_DICT
 - Cell, [236](#)
- T_ERROR
 - Cell, [236](#)
- T_FLOAT32
 - Cell, [235](#)
- T_FLOAT64
 - Cell, [235](#)
- T_GENSYM
 - Cell, [235](#)
- T_IDENT
 - Cell, [237](#)
- T_INT32
 - Cell, [235](#)
- T_INT64
 - Cell, [235](#)
- T_MACRO
 - Cell, [236](#)
- T_MOP3_NPROC
 - Cell, [236](#)
- T_MOP3_PROC
 - Cell, [236](#)
- T_NEEDS_FINALIZING
 - Cell, [235](#)
- T_NIL
 - Cell, [236](#)
- T_NPROC
 - Cell, [236](#)
- T_PAIR
 - Cell, [236](#)
- T_PORT_HEAP
 - Cell, [235](#)
- T_PROC
 - Cell, [236](#)
- T_RATIONAL
 - Cell, [236](#)
- T_STR_EXT
 - Cell, [235](#)
- T_STR_HEAP
 - Cell, [235](#)
- T_STR_IMM
 - Cell, [235](#)
- T_SVEC2N_HEAP
 - Cell, [235](#)
- T_SVEC_HEAP
 - Cell, [235](#)
- T_SVEC_IMM
 - Cell, [236](#)

- Cell, [236](#)
- T_SYM_HEAP
 - Cell, [235](#)
- T_THUNK_BODY
 - Cell, [236](#)
- T_THUNK_SEXPR
 - Cell, [236](#)
- T_UNDEF
 - Cell, [236](#)
- T_VALUES
 - Cell, [236](#)
- T_VOID
 - Cell, [236](#)
- tag_error
 - Lamb, [285](#)
- tail_state
 - Cell, [242](#)
- tail_state_clr
 - Cell, [242](#)
- tail_state_set
 - Cell, [242](#)
- tcons
 - Lamb, [270](#)
- twolist_to_dict
 - Lamb, [280](#)
- type
 - Cell, [237](#)
- type_name
 - Cell, [257](#)
- Ultrasonic Sonar (HC-SR04), [202](#)
 - sonar, [204](#)
 - sonar_echo_isr, [203](#)
 - Sonar_install_mop3, [204](#)
 - Sonar_mop3_begin, [203](#)
 - Sonar_mop3_ping, [203](#)
 - Sonar_mop3_poke, [203](#)
 - Sonar_mop3_result, [203](#)
 - Sonar_mop3_set_min_interval_us, [204](#)
 - Sonar_mop3_set_timeout_us, [204](#)
 - Sonar_mop3_start, [203](#)
 - Sonar_mop3_us_to_distance, [204](#)
- vector_copy
 - Lamb, [284](#)
- vector_set_bang
 - Lamb, [275](#)
- vector_to_list
 - Lamb, [284](#)
- vector_to_sparsevec
 - Lamb, [284](#)
- WiFi, [207](#)
 - WiFi_mop3_begin, [208](#)
 - WiFi_mop3_broadcastIP, [210](#)
 - WiFi_mop3_BSSID, [209](#)
 - WiFi_mop3_channel, [212](#)
 - WiFi_mop3_config, [208](#)
 - WiFi_mop3_disconnect, [208](#)
 - WiFi_mop3_dnsIP, [210](#)
 - WiFi_mop3_encryptionType, [211](#)
 - WiFi_mop3_gatewayIP, [210](#)
 - WiFi_mop3_getHostname, [211](#)
 - WiFi_mop3_hostByName, [212](#)
 - WiFi_mop3_isConnected, [209](#)
 - WiFi_mop3_localIP, [209](#)
 - WiFi_mop3_macAddress, [211](#)
 - WiFi_mop3_mode, [211](#)
 - WiFi_mop3_networkID, [210](#)
 - WiFi_mop3_psk, [211](#)
 - WiFi_mop3_reconnect, [208](#)
 - WiFi_mop3_RSSI, [209](#)
 - WiFi_mop3_scanNetworks, [212](#)
 - WiFi_mop3_setAutoReconnect, [209](#)
 - WiFi_mop3_setDNS, [208](#)
 - WiFi_mop3_setHostname, [212](#)
 - WiFi_mop3_setSleep, [212](#)
 - WiFi_mop3_softAP, [212](#)
 - WiFi_mop3_softAPBroadcastIP, [213](#)
 - WiFi_mop3_softAPConfig, [213](#)
 - WiFi_mop3_softAPdisconnect, [213](#)
 - WiFi_mop3_softAPgetHostname, [214](#)
 - WiFi_mop3_softAPgetStationNum, [213](#)
 - WiFi_mop3_softAPIP, [213](#)
 - WiFi_mop3_softAPmacAddress, [214](#)
 - WiFi_mop3_softAPsetHostname, [214](#)
 - WiFi_mop3_softAPSSID, [213](#)
 - WiFi_mop3_SSID, [209](#)
 - WiFi_mop3_status, [211](#)
 - WiFi_mop3_subnetCIDR, [210](#)
 - WiFi_mop3_subnetMask, [210](#)
 - WiFi_mop3_waitForConnectResult, [208](#)

WiFi, [210](#)
WiFi_mop3_getHostname
WiFi, [211](#)
WiFi_mop3_hostByName
WiFi, [212](#)
WiFi_mop3_isConnected
WiFi, [209](#)
WiFi_mop3_localIP
WiFi, [209](#)
WiFi_mop3_macAddress
WiFi, [211](#)
WiFi_mop3_mode
WiFi, [211](#)
WiFi_mop3_networkID
WiFi, [210](#)
WiFi_mop3_psk
WiFi, [211](#)
WiFi_mop3_reconnect
WiFi, [208](#)
WiFi_mop3_RSSI
WiFi, [209](#)
WiFi_mop3_scanNetworks
WiFi, [212](#)
WiFi_mop3_setAutoReconnect
WiFi, [209](#)
WiFi_mop3_setDNS
WiFi, [208](#)
WiFi_mop3_setHostname
WiFi, [212](#)
WiFi_mop3_setSleep
WiFi, [212](#)
WiFi_mop3_softAP
WiFi, [212](#)
WiFi_mop3_softAPBroadcastIP
WiFi, [213](#)
WiFi_mop3_softAPConfig
WiFi, [213](#)
WiFi_mop3_softAPdisconnect
WiFi, [213](#)
WiFi_mop3_softAPgetHostname
WiFi, [214](#)
WiFi_mop3_softAPgetStationNum
WiFi, [213](#)
WiFi_mop3_softAPIP
WiFi, [213](#)
WiFi_mop3_softAPmacAddress
WiFi, [214](#)
WiFi_mop3_softAPsetHostname
WiFi, [214](#)
WiFi_mop3_softAPSSID
WiFi, [213](#)
WiFi_mop3_SSID
WiFi, [209](#)
WiFi_mop3_status
WiFi, [211](#)
WiFi_mop3_subnetCIDR
WiFi, [210](#)
WiFi_mop3_subnetMask
WiFi, [210](#)
WiFi_mop3_waitForConnectResult
WiFi, [208](#)
Wire_mop3_available
I2C (Wire), [216](#)
Wire_mop3_begin
I2C (Wire), [215](#)
Wire_mop3_beginTransmission
I2C (Wire), [216](#)
Wire_mop3_clearWireTimeoutFlag
I2C (Wire), [217](#)
Wire_mop3_end
I2C (Wire), [216](#)
Wire_mop3_endTransmission
I2C (Wire), [217](#)
Wire_mop3_getWireTimeoutFlag
I2C (Wire), [217](#)
Wire_mop3_onReceive
I2C (Wire), [216](#)
Wire_mop3_onRequest
I2C (Wire), [217](#)
Wire_mop3_read
I2C (Wire), [216](#)
Wire_mop3_requestFrom
I2C (Wire), [215](#)
Wire_mop3_setClock
I2C (Wire), [216](#)
Wire_mop3_setPins
I2C (Wire), [215](#)
Wire_mop3_setTimeout
I2C (Wire), [215](#)
Wire_mop3_write
I2C (Wire), [217](#)
write_or_display
Lamb, [281](#)
write_simple
Lamb, [281](#)
WS2812 / NeoPixel LEDs, [218](#)
mop3_neopixelWrite, [218](#)
WS2812_mop3_begin, [218](#)
WS2812_mop3_setBrightness, [218](#)
WS2812_mop3_setLedColorData, [218](#)
WS2812_mop3_show, [219](#)
WS2812_mop3_Wheel, [219](#)
WS2812_mop3_begin
WS2812 / NeoPixel LEDs, [218](#)
WS2812_mop3_setBrightness
WS2812 / NeoPixel LEDs, [218](#)
WS2812_mop3_setLedColorData
WS2812 / NeoPixel LEDs, [218](#)

WS2812_mop3_show
 WS2812 / NeoPixel LEDs, [219](#)
WS2812_mop3_Wheel
 WS2812 / NeoPixel LEDs, [219](#)

zero
 Cell, [237](#)